



LEARN

# JAVA WITH PROJECTS

---

A concise practical guide to learning everything  
a Java professional really needs to know

**Dr. Seán Kennedy | Maaike van Putten**

# Learn Java with Projects

A concise practical guide to learning everything a Java professional really needs to know

**Dr. Seán Kennedy**

**Maaike van Putten**



BIRMINGHAM—MUMBAI

# Learn Java with Projects

Copyright © 2023 Packt Publishing

*All rights reserved.* No part of this book may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, without the prior written permission of the publisher, except in the case of brief quotations embedded in critical articles or reviews.

Every effort has been made in the preparation of this book to ensure the accuracy of the information presented. However, the information contained in this book is sold without warranty, either express or implied. Neither the author(s), nor Packt Publishing or its dealers and distributors, will be held liable for any damages caused or alleged to have been caused directly or indirectly by this book.

Packt Publishing has endeavored to provide trademark information about all of the companies and products mentioned in this book by the appropriate use of capitals. However, Packt Publishing cannot guarantee the accuracy of this information.

**Associate Group Product Manager:** Kunal Sawant

**Book Project Manager:** Deeksha Thakkar

**Senior Editor:** Rounak Kulkarni

**Technical Editor:** Jubit Pincy

**Copy Editor:** Safis Editing

**Indexer:** Hemangini Bari

**Production Designer:** Vijay Kamble

**DevRel Marketing Coordinator:** Shrinidhi Manoharan and Sonia Chauhan

**Business Development Executive:** Kriti Sharma

First published: September 2023

Production reference: 1231123

Published by Packt Publishing Ltd.

Grosvenor House

11 St Paul's Square

Birmingham

B3 1RB, UK

ISBN 978-1-83763-718-8

[www.packtpub.com](http://www.packtpub.com)



*To my wife Maria, and my daughters, Emily, Miriam, and Lilian.*

*– Dr. Seán Kennedy*

*To Adnane, for not distracting me.*

*– Maaike van Putten*

# Contributors

## About the authors

**Dr. Séan Kennedy** is a university lecturer with over 20 years of experience in teaching. He has a Ph.D. in IT and is Oracle-certified in Java at the Professional level (OCP). In his daily work, he has taught Java on Ericssons' bespoke Master's program for over a decade. He has several very popular Java courses on Udemy, one of which, *Java 21, Java 17, Java 11, Advanced Java 8'*, has been selected for inclusion in their Udemy Business program (where only the top 3% of courses qualify). He has a YouTube channel called *Let's Get Certified* that teaches Java at all levels and prepares candidates for Java certification. Outside of work, he enjoys tennis, walking, reading, TV, and nature.

*I want to thank those who have always supported me, especially my wife, Maria, and my late parents.*

**Maaike van Putten** is a software consultant and trainer with a passion for empowering others in their careers. Her love for Java shows in numerous software development projects she participated in and the 5 Oracle Java certifications she obtained. She has designed and delivered a broad spectrum of training courses catering to beginners and seasoned developers in Java and many other languages and frameworks. Next to that, she has authored multiple books and online courses through multiple platforms reaching over 600,000 learners.

*I want to thank all my students for inspiring and motivating me.*

## About the reviewers

**Alejandro Duarte** is a Software Engineer, published author, and award winner. He currently works for MariaDB as a Developer Relations Engineer. Starting his coding journey at 13 with BASIC on a rudimentary black screen, Alejandro quickly transitioned to C, C++, and Java during his academic years at the National University of Colombia. He relocated first to the UK and then to Finland to foster his career in the open-source software industry. He's a recognized figure in Java circles, credited with articles amassing millions of views, as well as technical videos and presentations at international events. You can contact him on X (Twitter) @alejandro\_du.

*Thanks to the global Java community for being part of such an amazing technology. Thanks to MariaDB plc for allowing me the flexibility to work on this project.*

**Thiago Braga Rodrigues**, a university lecturer, researcher, and Java specialist, has journeyed academically from Brazil to Ireland over the span of a decade. He holds a BEng in Biomedical Engineering and a Ph.D. in IT. Throughout his academic journey, Thiago's research endeavors have led him to globally recognized institutions, such as Harvard Medical School in the USA, and Athlone Institute of Technology in Ireland, where he delved into the intricacies of bioinformatics. His expertise encompasses programming, biomedical engineering, and software engineering, with a particular focus on developing hardware and software solutions for wearable systems. As both an author and reviewer for various academic journals and conferences, Thiago persistently contributes to the advancement of his field. Outside of academia and work, he is a classically trained violin teacher with experience playing in several orchestras. An avid gamer, Thiago seamlessly blends technology, arts, and leisure throughout his enriching journey.

**Giuseppe Pedullà** is a senior software developer and software architect, working since 2014. He has worked in some large companies and created some products for clients such as banks, insurance companies, and so on. He has developed complex solutions for his clients, making architectural decisions to improve software efficiency. He also teaches coding, conducts webinars, and so on.

*I thank my wife and my children who are very patient and push me to do things that realize me as a developer and so on. I thank them because they help me and they understand that I want to do something to help me improve in what I do.*



**Anurag Sharma** is a seasoned Integration Architect, boasting 11 years of rich experience in the IT realm, with deep-seated expertise in JAVA/J2EE and a broad suite of integration technologies, including a prominent presence in the Salesforce ecosystem as a MuleSoft Ambassador. As a JAVA-certified professional, Anurag has a track record of success in developing and designing solutions across diverse domains such as airlines, retail, and manufacturing, to name a few.

With a dual-level certification in JAVA and comprehensive certifications in MuleSoft, Anurag has been a vital asset to tech giants like Oracle and HCL, in both India and the UK. His thought leadership and insights have resonated on global stages such as Salesforce's London World Tour and Dreamforce, where he has been a sought-after speaker.

Anurag is known for his commitment to the Integration community, where he spearheads mentorship programs, nurturing both novices and seasoned professionals in the field. He solves problems on Tech helper forums. Beyond his mentorship roles, Anurag contributes to the vibrancy of the tech world through his writings, illuminating tech articles that decipher complex technical landscapes for a broader audience.

In the quieter moments, you will find Anurag immersed in literature, as reading and writing about technology are not just hobbies but gateways to staying ahead in the ever-evolving tech universe.



# Table of Contents

## Part 1: Java Fundamentals

**1**

<b>Getting Started with Java</b>		<b>3</b>
Technical requirements	3	<b>Compiling and running</b>
Exploring Java features	4	<b>Java programs</b>
OOP in Java	5	Understanding the compilation process
Working with OOP	5	Compiling the code with javac
Compiled language	6	on the command line
Write once, run anywhere	6	Running the compiled code
Automatic memory management	8	with Java on the command line
<b>Installing Java</b>	<b>9</b>	<b>Working with an IDE</b>
Checking whether Java is installed on your system	9	What is an IDE?
Installing Java on Windows	10	Choosing an IDE
Installing Java on macOS	14	
Installing Java on Linux	14	<b>Creating and running</b>
Running Java online	15	<b>a program with an IDE</b>
<b>Writing our first program</b>	<b>16</b>	21
Hello world	16	Creating a program in an IDE
Steps to create the program	16	Running your program
Understanding the program	17	Debugging a program
		25
		27
		<b>Exercises</b>
		<b>Project</b>
		30
		<b>Summary</b>
		30

**2****Variables and Primitive Data Types** **31**

<b>Technical requirements</b>	<b>31</b>	<b>Understanding Java's primitive data types</b>	<b>35</b>
<b>Understanding and declaring variables</b>	<b>31</b>	<b>Java's primitive data types</b>	<b>35</b>
What is a variable?	32	<b>Screen output</b>	<b>38</b>
Declaring a variable	32	<b>Exercises</b>	<b>39</b>
Naming a variable	33	<b>Project – dinosaur profile generator</b>	<b>40</b>
Accessing a variable	34	<b>Summary</b>	<b>41</b>
Accessing a variable that you have not declared	34		

**3****Operators and Casting** **43**

<b>Technical requirements</b>	<b>43</b>	Bitwise operators	<b>56</b>
<b>Learning how Java's operators cooperate</b>	<b>43</b>	Ternary operator	<b>58</b>
Order of precedence	44	Compound assignment operators	<b>59</b>
Associativity	44	<b>Explaining Java casting</b>	<b>61</b>
<b>Understanding Java's operators</b>	<b>46</b>	Widening	<b>61</b>
Unary operators	46	Narrowing	<b>63</b>
Arithmetic operators	48	<b>Exercises</b>	<b>65</b>
Relational operators	52	<b>Project – Dino meal planner</b>	<b>66</b>
Logical operators	53	<b>Summary</b>	<b>66</b>

**4****Conditional Statements** **69**

<b>Technical requirements</b>	<b>69</b>	else if statements	<b>73</b>
<b>Understanding scope</b>	<b>70</b>	else statements	<b>73</b>
What is a block?	70	<b>Mastering switch statements and expressions</b>	<b>78</b>
<b>Exploring if statements</b>	<b>71</b>	switch statements	<b>79</b>
The if statement itself	72		

switch expressions	83	<b>Project – Task allocation system</b>	89
<b>Exercises</b>	<b>88</b>	<b>Summary</b>	<b>89</b>

## 5

### Understanding Iteration 91

<b>Technical requirements</b>	<b>91</b>	Nested loops	102
<b>while loops</b>	<b>92</b>	<b>break and continue statements</b>	<b>104</b>
<b>do-while loops</b>	<b>94</b>	break statements	104
while versus do-while	95	continue statements	107
<b>for loops</b>	<b>98</b>	<b>Exercises</b>	<b>109</b>
Traditional for loop	98	<b>Project – Dino meal planner</b>	<b>109</b>
Enhanced for loop	101	<b>Summary</b>	<b>110</b>

## 6

### Working with Arrays 113

<b>Technical requirements</b>	<b>113</b>	Using the for loop	119
<b>Arrays – what, when, and why?</b>	<b>114</b>	Using the for each loop	121
Arrays explained	114	Choosing between the regular loop	
When to use arrays	115	and the enhanced for loop	121
<b>Declaring and initializing arrays</b>	<b>115</b>	<b>Handling multidimensional arrays</b>	<b>122</b>
Declaring arrays	115	Declaring and initializing	
Initializing arrays	115	multidimensional arrays	122
<b>Short syntax for array initialization</b>	<b>116</b>	Accessing and modifying elements	
<b>Accessing elements in an array</b>	<b>116</b>	of multidimensional arrays	123
Understanding indexing	117	Iterating over multidimensional arrays	124
Accessing array elements	117		
Modifying array elements	118	<b>Using Java's built-in</b>	
		<b>methods for arrays</b>	<b>125</b>
		Built-in Arrays class for working with arrays	125
<b>Working with length and bounds</b>	<b>118</b>	<b>Exercises</b>	<b>131</b>
Determining the length of an array	118	<b>Project – Dino tracker</b>	<b>131</b>
Dealing with the bounds of an array	118	<b>Summary</b>	<b>132</b>
<b>Iterating over arrays</b>	<b>119</b>		

**7**

<b>Methods</b>	<b>135</b>		
Technical requirements	135	Exploring method overloading	143
Explaining why methods are important	136	Method signature	144
Flow of control	136	Overloading a method	144
Abstraction	137	Explaining varargs	146
Code duplication	138	Mastering call by value	148
Understanding the difference between method definition and method execution	139	Primitives versus references in memory	148
Method definition	140	Exercises	152
Method execution	141	Project – Mesozoic Eden assistant	153
		Summary	155

**Part 2: Object-Oriented Programming****8**

<b>Classes, Objects, and Enums</b>	<b>159</b>		
Technical requirements	160	Applying access modifiers	176
Understanding the differences between classes and objects	160	private	176
Classes	160	Package-private	176
Objects	161	protected	177
Getting familiar with the new keyword	161	public	177
Contrasting instance with class members	163	packages	179
Instance members (methods/data)	163	Encapsulation	181
Class members (methods/data)	167	Achieving encapsulation	181
Exploring the “this” reference	170	Mastering advanced encapsulation	184
Associating an instance with the “this” reference	172	Call By value revisited	184
Shadowing or hiding an instance variable	173	The issue	184
		The solution	187
		Delving into the object life cycle	190
		Garbage collection	190

---

Object life cycle example	191	<b>Appreciating records</b>	201
<b>Explaining the instanceof keyword</b>	<b>195</b>	Record patterns	205
<b>Understanding enums</b>	<b>196</b>	<b>Exercises</b>	<b>206</b>
Simple enums	196	<b>Project – Mesozoic Eden park manager</b>	<b>207</b>
Complex enums	199	<b>Summary</b>	<b>209</b>

## 9

---

### Inheritance and Polymorphism 213

---

<b>Technical requirements</b>	<b>214</b>	The UML diagram	242
<b>Understanding inheritance</b>	<b>214</b>	The package with the protected member	244
Advantages of inheritance	215	The other package	245
Disadvantages of inheritance	215	Pattern matching for switch	247
Base class	215	<b>Explaining the abstract and final keywords</b>	<b>247</b>
Subclass	216	The abstract keyword	248
The “is-a” relationship	216	The final keyword	249
<b>Applying inheritance</b>	<b>218</b>	<b>Applying sealed classes</b>	<b>252</b>
extends	218	sealed and permits	253
implements	219	non-sealed	253
<b>Exploring polymorphism</b>	<b>221</b>	Example using sealed, permits, and non-sealed	253
Separating the reference type from the object type	221	<b>Understanding instance and static blocks</b>	<b>255</b>
Applying polymorphism	222	Instance blocks	256
JVM – object type versus reference type usage	228	static blocks	256
<b>Contrasting method overriding and method overloading</b>	<b>231</b>	<b>Mastering upcasting and downcasting</b>	<b>259</b>
Method overloading	231	Upcasting	260
Method overriding	234	Downcasting	261
<b>Exploring the super keyword</b>	<b>237</b>	<b>Exercises</b>	<b>263</b>
super()	237	<b>Project</b>	<b>264</b>
super.	238	<b>Summary</b>	<b>265</b>
An example of using super	238		
<b>Revisiting the protected access modifier</b>	<b>242</b>		



# 10

<b>Interfaces and Abstract Classes</b>	<b>269</b>		
Technical requirements	270	'static' interface methods	278
Understanding abstract classes	270	Multiple interface inheritance	280
Mastering interfaces	271	<b>Explaining 'private' interface methods</b>	281
Abstract methods in interfaces	272	<b>Exploring sealed interfaces</b>	284
Interface constants	272	<b>Exercises</b>	287
Multiple interface inheritance	274	<b>Project – unified park management system</b>	287
Examining default and static interface methods	276	<b>Summary</b>	289
'default' interface methods	277		

# 11

<b>Dealing with Exceptions</b>	<b>291</b>		
Technical requirements	292	Creating and throwing custom exceptions	301
Understanding exceptions	292	<b>The catch or declare principle</b>	302
What are exceptions?	292	Understanding the principle	302
Need for exception handling	293	Declaring exceptions using throws	302
Common situations that require exception handling	293	Handling exceptions with try-catch	302
Understanding the exception hierarchy	295	Handling exceptions with try-with-resources	306
Checked exceptions	295		
Unchecked exceptions	295		
<b>Working with basic I/O operations</b>	<b>297</b>	<b>Working with inheritance and exceptions</b>	<b>309</b>
Reading from a file using FileReader	297	Declaring exceptions in method signatures	309
Writing to a file using FileWriter	299	Overriding methods and exception handling	309
<b>Throwing exceptions</b>	<b>300</b>	<b>Exercises</b>	<b>311</b>
The throw keyword	300	<b>Project – dinosaur care system</b>	312
		<b>Summary</b>	312

# 12

<b>Java Core API</b>	<b>315</b>		
Technical requirements	315	The checklist	332
Understanding the Scanner class	315	<b>Examining List and ArrayList</b>	<b>339</b>
Using Scanner to read from the keyboard	317	List properties	339
Using Scanner to read from a file	319		
Using Scanner to read from a string	320	<b>Exploring the Date API</b>	<b>341</b>
		Dates and times	342
<b>Comparing String with StringBuilder</b>	<b>320</b>	Duration and Period	343
String class	321	Additional interesting types	343
StringBuilder class	327	Formatting dates and times	347
String versus StringBuilder example	329	<b>Exercises</b>	<b>353</b>
Designing a custom immutable type	331	<b>Project – dinosaur care system</b>	<b>353</b>
		<b>Summary</b>	<b>355</b>

## Part 3: Advanced Topics

# 13

<b>Generics and Collections</b>	<b>359</b>		
Technical requirements	359	HashMap	368
<b>Getting to know collections</b>	<b>359</b>	TreeMap	368
Overview of different collection types	360	LinkedHashMap	369
<b>List</b>	<b>362</b>	<b>Basic operations on maps</b>	<b>369</b>
ArrayList	362	<b>Queue</b>	<b>371</b>
LinkedList	362	Queue implementations	371
Exploring the basic operations for lists	363	Basic operations on the Queue interface	371
<b>Set</b>	<b>365</b>	<b>Sorting collections</b>	<b>374</b>
HashSet	365	Natural ordering	374
TreeSet	366	The Comparable and Comparator interfaces	374
LinkedHashSet	366	<b>Working with generics</b>	<b>380</b>
Performing basic operations on a set	367	Life before generics – objects	381
<b>Map</b>	<b>368</b>	Use case of generics	382



Syntax generics	384	Using hashCode() in custom generic types	388
Bounded generics	385	<b>Exercises</b>	<b>389</b>
<b>Hashing and overriding hashCode()</b>	<b>386</b>	<b>Project – advanced dinosaur care system</b>	<b>389</b>
Understanding basic hashing concepts	386	<b>Summary</b>	<b>391</b>
hashCode() and its role in collections	386		
Overriding hashCode() and best practices	387		

## 14

<b>Lambda Expressions</b>	<b>393</b>		
<b>Technical requirements</b>	<b>393</b>	UnaryOperator and BinaryOperator	412
<b>Understanding lambda expressions</b>	<b>393</b>	<b>Mastering method references</b>	<b>413</b>
Functional Interfaces	394	Bound method references	414
Lambda expressions	394	Unbound method references	416
final or effectively final	400	Static method references	418
<b>Exploring functional interfaces from the API</b>	<b>402</b>	Constructor method references	419
Predicate and BiPredicate	403	Method references and context	420
Supplier	405	<b>Exercises</b>	<b>423</b>
Consumer and BiConsumer	406	<b>Project – agile dinosaur care system</b>	<b>424</b>
Function and BiFunction	410	<b>Summary</b>	<b>425</b>

## 15

<b>Streams – Fundamentals</b>	<b>427</b>		
<b>Technical requirements</b>	<b>428</b>	<b>Mastering terminal operations</b>	<b>441</b>
<b>Understanding stream pipelines</b>	<b>428</b>	count()	443
Stream pipeline	428	min() and max()	443
<b>Exploring stream laziness</b>	<b>431</b>	findAny() and findFirst()	444
<b>Creating streams</b>	<b>434</b>	anyMatch(), allMatch, and noneMatch()	445
Streaming from an array	434	forEach()	447
Streaming from a collection	435	reduce()	448
Stream.of()	436	collect()	451
Streaming from a file	437	collect(Collector)	452
Infinite streams	438	<b>Exercises</b>	<b>462</b>

<b>Project – dynamic dinosaur care system</b>	<b>Summary</b>	<b>464</b>
	463	

## 16

<b>Streams: Advanced Concepts</b>	<b>465</b>
-----------------------------------	------------

<b>Technical requirements</b>	<b>466</b>	<b>Explaining Optionals</b>	<b>491</b>
<b>Examining intermediate operations</b>	<b>466</b>	Creating Optionals	492
filter(Predicate)	466	Using the Optional API	494
distinct()	466	Primitive Optionals	497
limit(long)	468	<b>Understanding parallel streams</b>	<b>499</b>
map(Function)	469	Creating parallel streams	499
flatMap(Function)	469	Parallel decomposition	501
sorted() and sorted(Comparator)	471	Parallel reductions using reduce()	503
<b>Delving into primitive streams</b>	<b>474</b>	Parallel reductions using collect()	505
Creating primitive streams	474	<b>Exercises</b>	<b>506</b>
Common primitive stream methods	477	<b>Project – dynamic dinosaur care system</b>	<b>506</b>
New primitive stream interfaces	482	<b>Summary</b>	<b>507</b>
<b>Mapping streams</b>	<b>484</b>		
Mapping from Object streams	485		
Mapping from primitive streams	488		

## 17

<b>Concurrency</b>	<b>509</b>
--------------------	------------

<b>Technical requirements</b>	<b>509</b>	Lambda expressions with Runnable	517
<b>Understanding concurrency</b>	<b>510</b>	<b>Thread management – sleep() and join()</b>	<b>517</b>
Multiprocessing	510	The Thread.sleep() method	517
Multitasking	510	Handling InterruptedException	518
Multithreading	511	Using the join() method	518
Importance of concurrency in modern applications	512	<b>Atomic classes</b>	<b>519</b>
Challenges in concurrent programming	512	AtomicInteger, AtomicLong, and AtomicReference	520
<b>Working with threads</b>	<b>513</b>	<b>The synchronized keyword</b>	<b>521</b>
The Thread class	514		
The Runnable interface	515		



Using synchronized methods	524	Future objects and their methods	537
Using synchronized blocks	525	Invoking multiple tasks and handling the results	538
Synchronized methods versus synchronized blocks	525	Thread pools and task execution	539
<b>The Lock interface</b>	<b>526</b>	ScheduledExecutorServices	542
ReentrantLock	526	<b>Data races</b>	<b>544</b>
Best practices for working with locks	529	<b>Threading problems</b>	<b>545</b>
<b>Concurrent collections</b>	<b>529</b>	Data races	546
Concurrent collection interfaces	530	Race conditions	546
Understanding SkipList collections	531	Deadlocks	548
Understanding CopyOnWrite collections	532	Livelocks	549
Synchronized collections	533	Starvation	552
<b>ExecutorService and thread pools</b>	<b>534</b>	<b>Exercises</b>	<b>553</b>
Executing tasks using SingleThreadExecutor	534	<b>Project – Park Operations System – the calm before the storm</b>	<b>554</b>
The Callable interface and Future	535	<b>Summary</b>	<b>557</b>
Submitting tasks and handling results	536		
<b>Index</b>			<b>559</b>
<b>Other Books You May Enjoy</b>			<b>574</b>

# Preface

Welcome to the world of Java programming! Java is one of the most versatile and widely used programming languages in the world. Its platform independence, object-oriented nature, and extensive library support make it an ideal choice for developing a wide range of applications, from desktop to mobile and enterprise solutions.

Whether you are a novice eager to learn the fundamentals or an experienced developer seeking to enhance your skills, this book is designed to be your comprehensive guide to mastering the Java programming language. We start with the basics: how to get set up with an editor; and primitive data types; and progress systematically to more advanced concepts such as lambdas, streams, and concurrency.

This book is more than just a guide; it's a companion on your journey to master Java. Our goal is to make this journey enjoyable and effective. This book adopts a hands-on approach, combining theoretical concepts with lots of practical exercises and a capstone project. Whether you are a self-learner or part of a formal educational program, the hands-on exercises and capstone project will help solidify your knowledge, making the learning experience engaging and practical.

*Learning by doing* is critical in mastering any programming language and we have taken that to heart in this book. At the end of each chapter, you'll have a few exercises and a project that will help you get some experience with Java. The exercises are typically smaller tasks, and the project is a bit bigger. In all of them, you'll have quite some freedom to choose how to implement it specifically. Why? Well, because that's what it's going to be like in real life as well! We will provide you with a sample solution, but if ours is a bit different, that doesn't mean yours is bad. If you're in doubt, ask some AI assistants such as ChatGPT what they think about your solution. Still unclear? We're always willing to help you as well!

Alright, back to the exercises and projects. We wanted to choose a common theme for our exercises and projects. Believe it or not, but one of your writers (hint: it's the female one) is surrounded by miniature gigantic historic reptile replicas during the writing of the book. In fact, currently, it's hard to type because this enormous battery-powered Tyrannosaurus Rex tries to destroy my laptop.

I couldn't help but draw quite some inspiration from that. So yes, all the exercises and projects will be dinosaur-themed, based on the collection and vivid fantasy play of my 5-year-old son. (I wake up around 5-6 AM. Not by my alarm clock, and definitely not because I'm part of the 5 AM club. No, I have this (most amazing) excited kid telling me fun facts about dinosaurs. Might as well put this knowledge to use!)

So, here's your context: congratulations, you're hired! You are now working for our special dinosaur zoo: Mesozoic Eden. It's a unique blend of prehistoric wilderness and modern comfort, where humans and dinosaurs coexist. People can visit for a day, or camp here for several weeks.

At Mesozoic Eden, we house a rich variety of dinosaur species, each with its distinct behavior and lifestyle, ranging from the colossal Brachiosaurus to the swift-footed Velociraptor, the regal Tyrannosaurus Rex, and many more. We even have a state-of-the-art laboratory where we continue to discover and study new dinosaur breeds.

As part of our team, your role is not only about taking care of these majestic creatures and ensuring their well-being but also about maintaining the safety and security of our guests. Our park employs cutting-edge technology and stringent protocols to ensure a safe environment for all.

In the exercises and projects, you'll take on various software development tasks as an employee of Mesozoic Eden, from coding software for feeding schedules to the app that handles emergency alerts, ensuring park operations run smoothly, and above all, creating an unforgettable experience for our visitors.

## Who this book is for

This book is suitable for beginners looking to take their first steps into programming as well as experienced developers transitioning to Java. Whether you are a student or professional, the content is structured to cater to a diverse audience.

If you are interested in Java 8 OCA Oracle certification, then this book is extremely helpful as it covers many important fundamental concepts by going “under the hood” to explain what is happening in memory. It is no coincidence that both authors are Oracle OCP qualified.

## What this book covers

*Chapter 1, Getting Started with Java*, starts by discussing the main Java features, such as OOP. How to install Java on various operating systems and how to write your first Java program with and without an IDE are also explored.

*Chapter 2, Variables and Primitive Data Types*, explains what a variable is and the fact that Java uses “strong-typing” (you must declare a variables type). This chapter also covers Java’s primitive data types, their sizes in bytes (needed to understand for later when discussing casting), and their ranges.

*Chapter 3, Operators and Casting*, explores how Java’s operators cooperate using precedence and associativity. We discuss Java’s various operators and explain both widening and narrowing when casting in Java.

*Chapter 4, Conditional Statements*, focuses on both scope and conditional statements. We initially examine Java’s use of block scope. We then explain the various forms of the `if` statement; and conclude the chapter with both `switch` statements and `switch` expressions.

---

*Chapter 5, Understanding Iteration*, discusses loops, including `while`, `do-while`, `for`, and enhanced `for`. This chapter also explores the `break` and `continue` statements.

*Chapter 6, Working with Arrays*, describes why one needs arrays. We show how to declare and initialize arrays of various primitive types, including using the shorthand syntax. We discuss how to loop through an array, processing each element. Multi-dimensional arrays are also covered; as is the `Arrays` class.

*Chapter 7, Methods*, discusses the importance of methods and the difference between the method definition and method execution. Method overloading is discussed and the `varargs` format is explained. Lastly, the important concept of call-by-value is explained.

*Chapter 8, Classes, Objects, and Enums*, is a significant OOP chapter and details the following: the difference between classes and objects; the `this` reference; access modifiers; basic and advanced encapsulation; the object life cycle; the `instanceof` keyword; enums and records.

*Chapter 9, Inheritance and Polymorphism*, explains inheritance and polymorphism. We detail what overriding means and discuss the `super`, `protected`, `abstract`, and `final` keywords. We also explore `sealed` classes and upcasting/downcasting.

*Chapter 10, Interfaces and Abstract Classes*, covers both `abstract` classes and interfaces. We explain `static`, `default`, and `private` interface methods and also `sealed` interfaces.

*Chapter 11, Dealing with Exceptions*, explains exceptions and their purpose. We explain the difference between checked and unchecked exceptions. We delve into throwing exceptions and how to create your own custom exceptions. The important catch or declare principle is discussed; as are the try-catch, try-catch-finally, and try-with-resources blocks.

*Chapter 12, Java Core API*, introduces important classes/interfaces from the API, such as `Scanner`. We compare and contrast `String` and `StringBuilder` and discuss how to create a custom immutable type. We example the `List` interface and its popular implementation `ArrayList`. Lastly, we explore the Date API.

*Chapter 13, Generics and Collections*, discusses the collections framework and its interfaces `List`, `Set`, `Map`, and `Queue`. We examine several implementations of each and basic operations. We explain sorting using both the `Comparable` and `Comparator` interfaces. We finish by examining generics and basic hashing concepts.

*Chapter 14, Lambda Expressions*, explains what lambda expressions are and their relationship to functional interfaces. Several functional interfaces from the API are examined. Lastly, method references and the role of context in understanding them are outlined.

*Chapter 15, Streams: Fundamentals*, is our first chapter on streams. In this chapter, we discuss what a stream pipeline is and what stream laziness means. We show different ways to create both finite and infinite streams. We examine the terminal operations that start off the streaming process - including reductions such as `collect()` which is very useful for extracting information out of a stream.



*Chapter 16, Streams: Advanced Concepts*, starts by examining intermediate operations such as `filter()`, `map()`, and `sorted()`. We explore the primitive streams followed by how to map from one stream to another, regardless of type. The `Optional` type is explained and we conclude with a discussion of parallel streams.

*Chapter 17, Concurrency*, starts by explaining what concurrency is. We examine working with threads and present issues with concurrent access. Mechanisms to resolve these issues are discussed; namely: atomic classes, synchronized blocks, and the `Lock` interface. Concurrent collections and the `ExecutorService` are explored next. We finish with a discussion on threading problems such as data races, deadlock, and livelock.

## To get the most out of this book

While much of the code will work with earlier Java versions, we would recommend installing or upgrading to Java 21 to avoid version-related compiler errors.

If you currently have nothing on your system, the following setup would be great:

- JDK 21 or later (Oracle's JDK or OpenJDK)
- IntelliJ IDEA (community edition is good enough) or Eclipse or Netbeans

Software/hardware covered in the book	Operating system requirements
Java 21+	Windows, macOS, or Linux

See *Chapter 1* for instructions on how to get both Java and an IDE installed on various operating systems.

If you are using the digital version of this book, we advise you to type the code yourself or access the code from the book's GitHub repository (a link is available in the next section). Doing so will help you avoid any potential errors related to the copying and pasting of code.

### Disclaimer

With the intention of the Publisher and Author, certain graphics included in this title are displaying large screen examples where the textual content is not relevant to the graphic example. We encourage our readers to download the digital copy included in their purchase for magnified and accessible content requirements.

## Download the example code files

You can download the example code files for this book from GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects>. If there's an update to the code, it will be updated in the GitHub repository.

We also have other code bundles from our rich catalog of books and videos available at <https://github.com/PacktPublishing/>. Check them out!

## Code in Action

The Code in Action videos for this book can be viewed at <https://bit.ly/3GdtYeC>

## Download the color images

We also provide a PDF file that has color images of the screenshots and diagrams used in this book. You can download it here: <https://packt.link/gbp/9781837637188>

## Conventions used

There are a number of text conventions used throughout this book.

**Code in text:** Indicates code words in text, database table names, folder names, filenames, file extensions, pathnames, dummy URLs, user input, and Twitter handles. Here is an example: “So, the `main()` method has now handed over control to the `simpleExample()` method, and control will not return to `main()` until `simpleExample()` exits”

A block of code is set as follows:

```
public class HelloWorld {  
    public static void main(String[] args) {  
        System.out.println("Hello world!");  
    }  
}
```

When we wish to draw your attention to a particular part of a code block, the relevant lines or items are set in bold:

```
//           int age = 25;  
           System.out.println(age);
```

Any command-line input or output is written as follows:

```
Enter a number (negative number to exit) -->  
1  
Enter a number (negative number to exit) -->  
2
```

**Tips or important notes**

Appear like this.

## Get in touch

Feedback from our readers is always welcome.

**General feedback:** If you have questions about any aspect of this book, email us at [customerservice@packtpub.com](mailto:customerservice@packtpub.com) and mention the book title in the subject of your message.

**Errata:** Although we have taken every care to ensure the accuracy of our content, mistakes do happen. If you have found a mistake in this book, we would be grateful if you would report this to us. Please visit [www.packtpub.com/support/errata](http://www.packtpub.com/support/errata) and fill in the form.

**Piracy:** If you come across any illegal copies of our works in any form on the internet, we would be grateful if you would provide us with the location address or website name. Please contact us at [copyright@packt.com](mailto:copyright@packt.com) with a link to the material.

**If you are interested in becoming an author:** If there is a topic that you have expertise in and you are interested in either writing or contributing to a book, please visit [authors.packtpub.com](http://authors.packtpub.com).

## Share Your Thoughts

Once you've read *Learn Java with Projects*, we'd love to hear your thoughts! Please click here to go straight to the Amazon review page for this book and share your feedback.

Your review is important to us and the tech community and will help us make sure we're delivering excellent quality content.

## Download a free PDF copy of this book

Thanks for purchasing this book!

Do you like to read on the go but are unable to carry your print books everywhere?

Is your eBook purchase not compatible with the device of your choice?

Don't worry, now with every Packt book you get a DRM-free PDF version of that book at no cost.

Read anywhere, any place, on any device. Search, copy, and paste code from your favorite technical books directly into your application.

The perks don't stop there, you can get exclusive access to discounts, newsletters, and great free content in your inbox daily

Follow these simple steps to get the benefits:

1. Scan the QR code or visit the link below



<https://packt.link/free-ebook/9781837637188>

2. Submit your proof of purchase
3. That's it! We'll send your free PDF and other benefits to your email directly

Trial Version



Wondershare  
**PDFelement**



# Part 1:

## Java Fundamentals

In this part, we will start by looking into the features of Java and how to get set up using Java and an IDE. We will examine variables and Java's eight primitive data types. Following that, we will discuss Java's operators and casting. We then move on to Java conditional statements and looping constructs. After that, we will look at using arrays, before finally finishing with methods.

This section has the following chapters:

- *Chapter 1, Getting Started with Java*
- *Chapter 2, Variables and Primitive Data Types*
- *Chapter 3, Operators and Casting*
- *Chapter 4, Conditional Statements*
- *Chapter 5, Understanding Iteration*
- *Chapter 6, Working with Arrays*
- *Chapter 7, Methods*

Trial Version



Wondershare  
**PDFelement**

## 1

# Getting Started with Java

Welcome to the exciting world of Java! Java is a very popular programming language. It is a multipurpose, powerful, and popular programming language that has been used by millions of developers worldwide to create a wide variety of applications. And yes, it really is multipurpose since it can be used to create all sorts of applications, from web and mobile apps to game development and beyond.

So, you've done a great job choosing a (new) language. We're going to take you on a (hopefully) fascinating journey that will provide you with valuable skills and open new opportunities in the ever-evolving field of technology.

What are we waiting for? In this chapter, we're going to cover the following main topics:

- Java features
- Installing Java
- Compiling and running Java programs
- Working with an **integrated development environment (IDE)**
- Creating and running a program with an IDE

## Technical requirements

Before diving into the magical world of Java programming, let's ensure you have the right hardware. If your hardware doesn't meet these requirements, don't worry; online alternatives are discussed later in this chapter. If you are using your work laptop, make sure that you have download rights. Here's a brief overview of the requirements:

- **Operating system:** Java can run on various operating systems, including Windows, macOS, and Linux. Ensure that you have a recent version of one of these operating systems installed on your computer.

- **Java Development Kit (JDK):** To compile and run Java programs, you'll need the JDK installed on your system. The JDK includes the **Java Runtime Environment (JRE)**, which contains the necessary libraries and components for running Java applications. We'll see how to install this later.
- **System resources:** More is always better, but Java isn't too demanding. It doesn't require high-end hardware, but it's still a good idea to have a system with sufficient resources for a smooth development experience. The following are the minimum and recommended system requirements:
  - **Minimum requirements:**
    - CPU: 1 GHz or faster processor
    - RAM: 2 GB
    - Disk space: 1 GB (for JDK installation and additional files)
  - **Recommended requirements:**
    - CPU: 2 GHz or faster multi-core processor
    - RAM: 4 GB or more
    - Disk space: 2 GB or more (for JDK installation, additional files, and projects)

Keep in mind that these requirements may change with future updates to the JDK and related tools. We have placed the files in a GitHub repository. You can clone the projects with the use of Git and import them to your computer this way. It's beyond the scope of explaining how to use Git here but it's recommended to look into it independently. You can access the files and examples used in this book here: <https://github.com/PacktPublishing/Learn-Java-with-Projects>.

## Exploring Java features

Java was developed by James Gosling at Sun Microsystems in the mid-1990s. When Java was created, it was originally designed as a language for consumer electronics. It attempted to support complex host architectures, focused on portability, and supported secure networking. However, Java outgrew its own ambitions. It quickly gained momentum as a versatile language for creating enterprise, web, and mobile applications. Today, Java no longer belongs to Sun Microsystems. Oracle Corporation acquired Sun Microsystems in 2010. And with that acquisition, Java became an integral part of Oracle's software ecosystem.

Java was very unique at the time it was created. The huge success of Java can be attributed to some of its core features. These features were very innovative at the time but are now found in many other (competing) languages. One of the core features is object-oriented programming. OOP allows us to structure our code in a neat way that helps with reusability and maintainability. We're going to start discussing the core features by having a look at **object-oriented programming (OOP)**.

## OOP in Java

Arguably the most important feature of Java is its support for OOP. If you ask any Java developer what Java is, the answer is often that it's an OOP language.

It's safe to say that OOP is a key feature. *What is this OOP thing?* you may wonder. OOP is a programming paradigm. It structures applications to model real-world objects and their interactions and behaviors. Let's go over the main concepts of OOP:

- **Objects:** This may be stating the obvious but, in OOP, **objects** are the main building blocks of your program. An object is a representation of a real-world entity, such as a user, an email, or a bank account. Each object has its own **attributes** (data fields) and behaviors (**methods**).
- **Classes:** Objects are created using their **class**. A class is a blueprint for creating objects. It defines the attributes and methods that objects of the class should have. For example, a `Car` class might define attributes such as color, make, and model, and methods such as start, accelerate, and brake.
- **Inheritance:** Another key feature is **inheritance**. Inheritance allows one class to inherit the attributes and methods of another class. For example, `Car` could inherit from a `Vehicle` class. We're not going to cover the details here, but inheritance helps to better structure the code. The code is more reusable, and the hierarchy of related classes opens doors in terms of what we can do with our types.
- **Encapsulation:** Encapsulation is giving a class control over its own data. This is done by bundling data (attributes) and methods that operate on that data. The attributes can only be accessed via these special methods from outside. Encapsulation helps to protect the internal state of an object and allows you to control how the object's data can be accessed or modified. Don't worry if this sounds tricky still, we'll deal with this in more detail later.
- **Polymorphism and Abstraction:** These are two key concepts of OOP that will be explained later when you're ready for them.

## Working with OOP

I can imagine this all sounds very abstract at this point, but before you know it, you'll be creating classes and instantiating objects yourself. OOP helps to make code more maintainable, better structured, and reusable. These things really help to be able to make changes to your application, solve problems, and scale up when needed.

OOP is just one key feature of Java. Another key feature is that it's a compiled language. Let's make sure you understand what is meant by that now.



## Compiled language

Java is a **compiled programming language**, which means that the source code you write must be transformed into a machine-readable format before it can be interpreted. This machine-readable format is called bytecode. This process is different from that of interpreted languages, where the source code is read, interpreted, and executed on the fly. During runtime, the computer interprets an interpreted language line by line. When a compiled language is running, the computer interprets the bytecode during runtime. We'll dive deeper into the compilation process in just a bit when we are going to compile our own code. For now, let's see what the benefits of compiled languages are.

### ***Benefits of Java being a compiled language***

Compiling code first requires an extra step, and it takes time in the beginning, but it brings advantages. First of all, the performance of compiled languages is typically better than interpreted languages. This is because the bytecode gets optimized for efficient execution on the target platform.

Another advantage of compilation is the early detection of syntax errors and certain other types of errors before the code is executed. This enables developers to identify and fix issues before deploying the application, reducing the likelihood of runtime errors.

Java code is turned into bytecode – a form of binary code - by the compiler. This bytecode is platform-independent. This means that it allows Java applications to run on different operating systems without modification. Platform independence is actually the key feature that we're going to be discussing next.

## **Write once, run anywhere**

Java's **Write Once, Run Anywhere (WORA)** principle is another key feature. This used to set Java apart from many other programming languages, but now, this is rather common, and many competing languages also implemented this feature. This principle ensures that Java code can run on different platforms without requiring different versions of the Java code for each platform. This means that a Java program is not tied to any specific operating system or hardware architecture.

When you have different versions of the code for each platform, this means that you have to maintain all these versions of the code as well. Let's say you have a code base for Linux, macOS, and Windows. When a new feature or a change is required, you need to add this to three places! You can imagine that WORA was a game-changer at the time Java came out. And it leads to an increased reach of your application – any device that can run Java applications can run yours.

### ***Understanding the WORA elements***

The WORA principle is made possible by bytecode and the **Java Virtual Machine (JVM)**. Bytecode is the compiled Java program. The compiler turns the Java code into this bytecode, and this bytecode is platform-independent. It can run on any device that can run the bytecode executer.

This bytecode executer is called the JVM. Each platform (Windows, macOS, Linux, and so on) has its own JVM implementation, which is specifically designed to translate bytecode into native machine code for that platform. Since the bytecode remains the same across platforms, the JVM handles the differences between operating systems and hardware architectures. The WORA principle is explained in *Figure 1.1*.

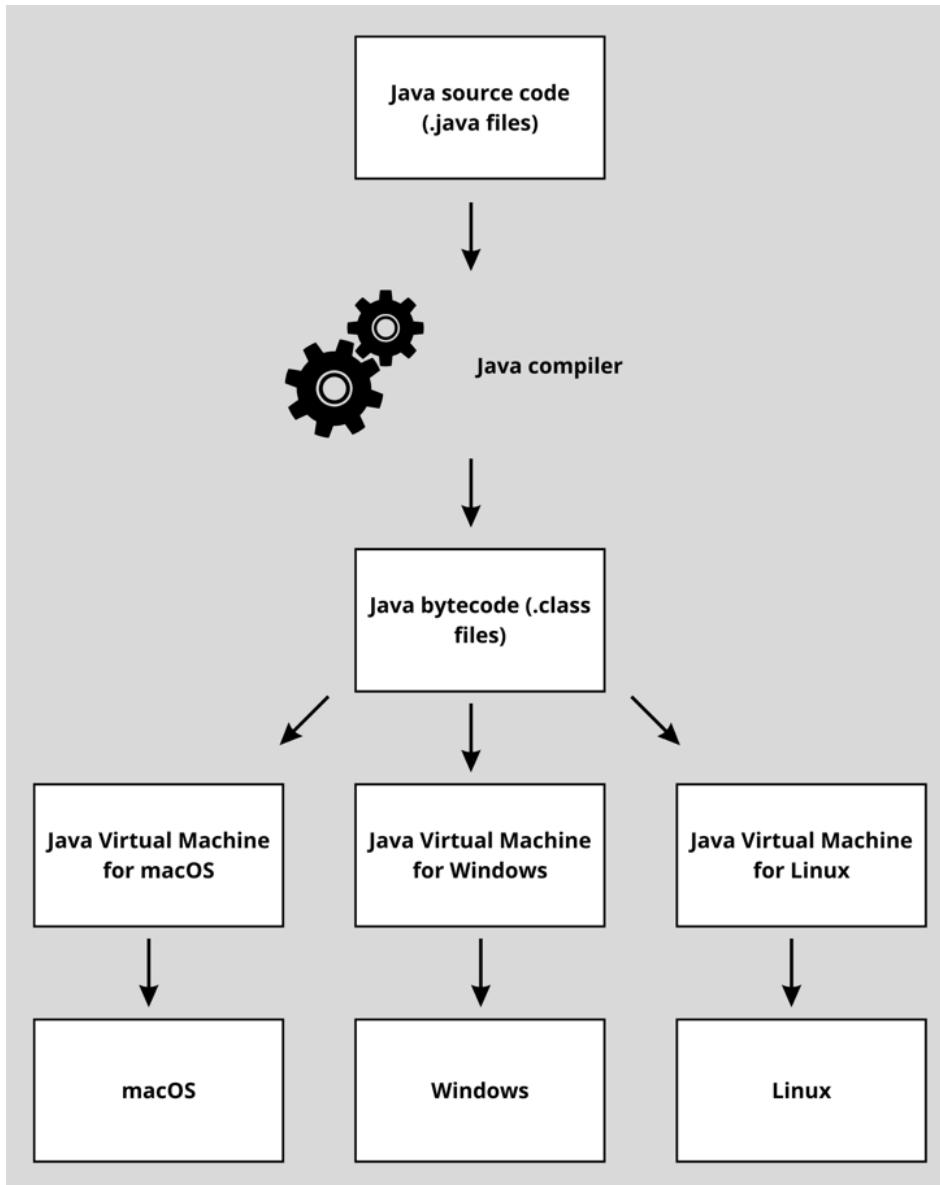


Figure 1.1 – The WORA principle in a diagram

You can see that the compiler creates bytecode and that this bytecode can be picked up by the JVM. The JVM is platform-specific and does the translation to the platform it's on. There's more that the JVM does for us, and that is automatic memory management. Let's explore this next.

## Automatic memory management

Another key feature that made Java great is its **automatic memory management**, which simplifies development and prevents common memory-related errors. Java handles memory allocation and garbage collection for you. The developer doesn't need to take care of manually managing the memory.

Nowadays, this is the rule and not the exception. Most other modern languages have automatic memory management as well. However, it is important to know what automatic memory management means. The memory allocation and deallocation are done automatically. This actually leads to simplifying the code. There is no boilerplate code that just focuses on the allocation and deallocation of the memory. This also leads to fewer memory-related errors.

Let's make sure you understand what is meant by memory allocation and deallocation.

### *Memory allocation*

In code, you create variables. Sometimes, these variables are not simple values but complex objects with many data fields. When you create an object, this object needs to be stored somewhere in the memory of the device that it's running on. This is called **memory allocation**. In Java, when you create an object, device memory is automatically allocated to store the object's attributes and associated data. This is different from languages such as C and C++, where developers must manually allocate and deallocate memory. Java's automatic memory allocation streamlines the development process and reduces the chances of memory leaks or dangling pointers, which can cause unexpected behavior or crashes. It also makes the code cleaner to read, since you don't need to deal with any allocation or deallocation code.

### *Garbage collection*

When a memory block is no longer used by the application, it needs to be deallocated. The process Java uses for this is called **garbage collection**. Garbage collection is the process of identifying and reclaiming memory that is no longer in use by a program. In Java, when an object is no longer accessible or needed, the garbage collector automatically frees up the memory occupied by the object. This process ensures that the memory is efficiently utilized and prevents memory leaks and the problems that come with it.

The JVM periodically runs the garbage collector to identify and clean up unreachable objects. Java's garbage collection mechanism uses many different sophisticated algorithms to determine when an object is no longer needed.

Now that we've covered the basics, let's move on to installing Java.

## Installing Java

Before you can start writing and running Java programs, you'll need to set up the JDK on your computer. The JDK contains essential tools and libraries required for Java development, such as the Java compiler, the JRE, and other useful utilities that help development.

We will guide you through the process of installing Java on Windows, macOS, and Linux, and we'll give you some suggestions for when you don't have access to either one of those. But before proceeding with the installation of Java, it's a good idea to check whether it's already installed on your system.

### Checking whether Java is installed on your system

Java may have been pre-installed, or you may have installed it previously without realizing it. To check whether Java is installed, follow these simple steps. The first one depends on your operating system.

#### **Step 1 – Open a terminal**

For Windows, press the *Windows* key, type cmd, and press *Enter* to open the **Command Prompt**.

For macOS, press *Command + Space* to open the **Spotlight** search, type Terminal, and press *Enter* to open **Terminal**.

For Linux, open a Terminal window. The method for opening the Terminal window varies depending on your Linux distribution (for example, in Ubuntu, press *Ctrl + Alt + T*).

#### **Step 2 – Check for the Java version**

In the Command Prompt or Terminal window, type the following command and press *Enter*:

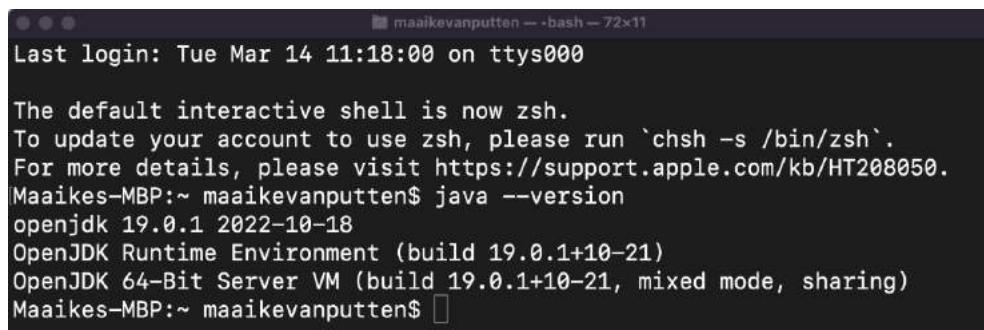
```
java -version
```

#### **Step 3 – Interpret the response**

If Java is installed, you will see the version information displayed. If not, the Command Prompt will display an error message, indicating that Java is not recognized or found.

If you find that Java is already installed on your system, make sure it's version 21 or later to ensure compatibility with modern Java features. If it's an older version or not installed, proceed with the installation process for your specific platform, as described in the following sections. If an older version is installed, you may want to uninstall this first to avoid having an unnecessarily complicated setup. You can install this the common way of uninstalling programs for your operating system.

In *Figure 1.2* and *Figure 1.6*, you'll see examples of the output you can expect when Java is installed.



The screenshot shows a terminal window titled "maaikevanputten ~ bash 72x11". It displays the following text:

```
Last login: Tue Mar 14 11:18:00 on ttys000

The default interactive shell is now zsh.
To update your account to use zsh, please run `chsh -s /bin/zsh`.
For more details, please visit https://support.apple.com/kb/HT208050.
Maaikes-MBP:~ maaikevanputten$ java --version
openjdk 19.0.1 2022-10-18
OpenJDK Runtime Environment (build 19.0.1+10-21)
OpenJDK 64-Bit Server VM (build 19.0.1+10-21, mixed mode, sharing)
Maaikes-MBP:~ maaikevanputten$
```

Figure 1.2 – The macOS terminal output where Java 19 is installed

Now, let's see how to install Java on each operating system.

## Installing Java on Windows

To install Java on a Windows operating system, follow these steps:

1. Visit the **Oracle Java SE Downloads** page at <https://www.oracle.com/java/technologies/downloads/>. This software can be used for educational purposes for free, but requires a license in production. You can consider switching to **OpenJDK** to run programs in production without a license: <https://openjdk.org/install/>.
2. Select the appropriate installer for your Windows operating system (for example, **Windows x64 Installer**).
3. Download the installer by clicking on the file link.
4. Run the downloaded installer (the .exe file) and follow the on-screen instructions to complete the installation.
5. To add Java to the system's PATH environment variable, search for **Environment Variables** in the **Start** menu and select **Edit the system environment variables**. You should see a screen similar to *Figure 1.3*.

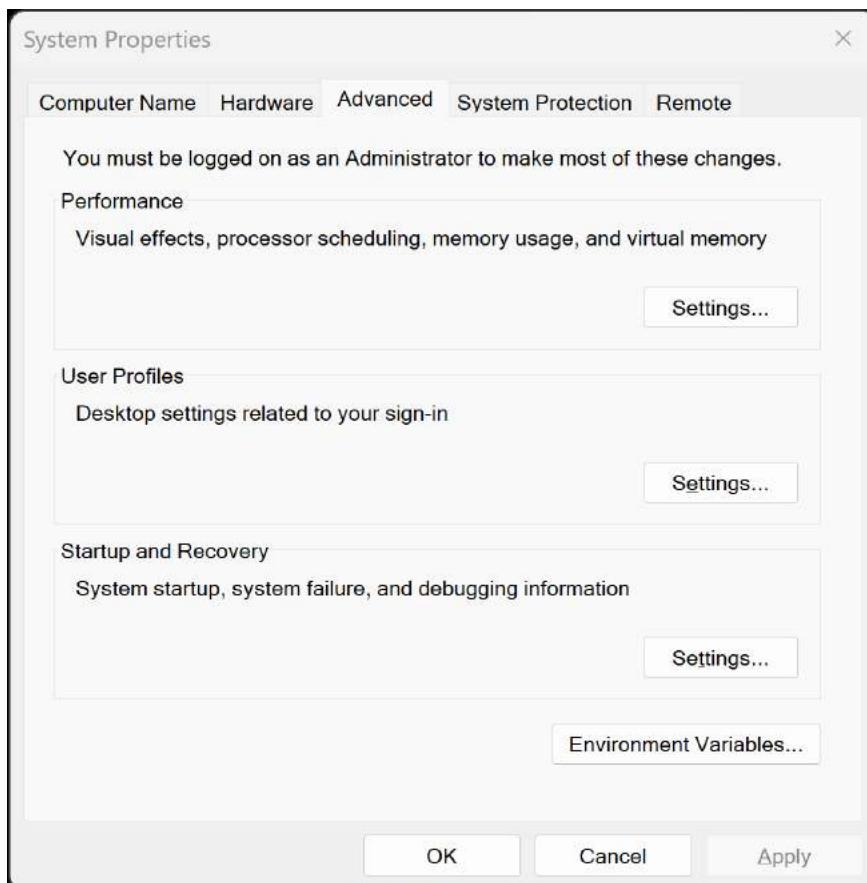


Figure 1.3 – The System Properties window

6. In the **System Properties** window, click on the **Environment Variables...** button. A screen like the one in *Figure 1.4* will pop up.
7. Under **System variables**, find the **Path** variable, select it, and click **Edit**. You can see an example of which one to select in the following *Figure 1.4*:

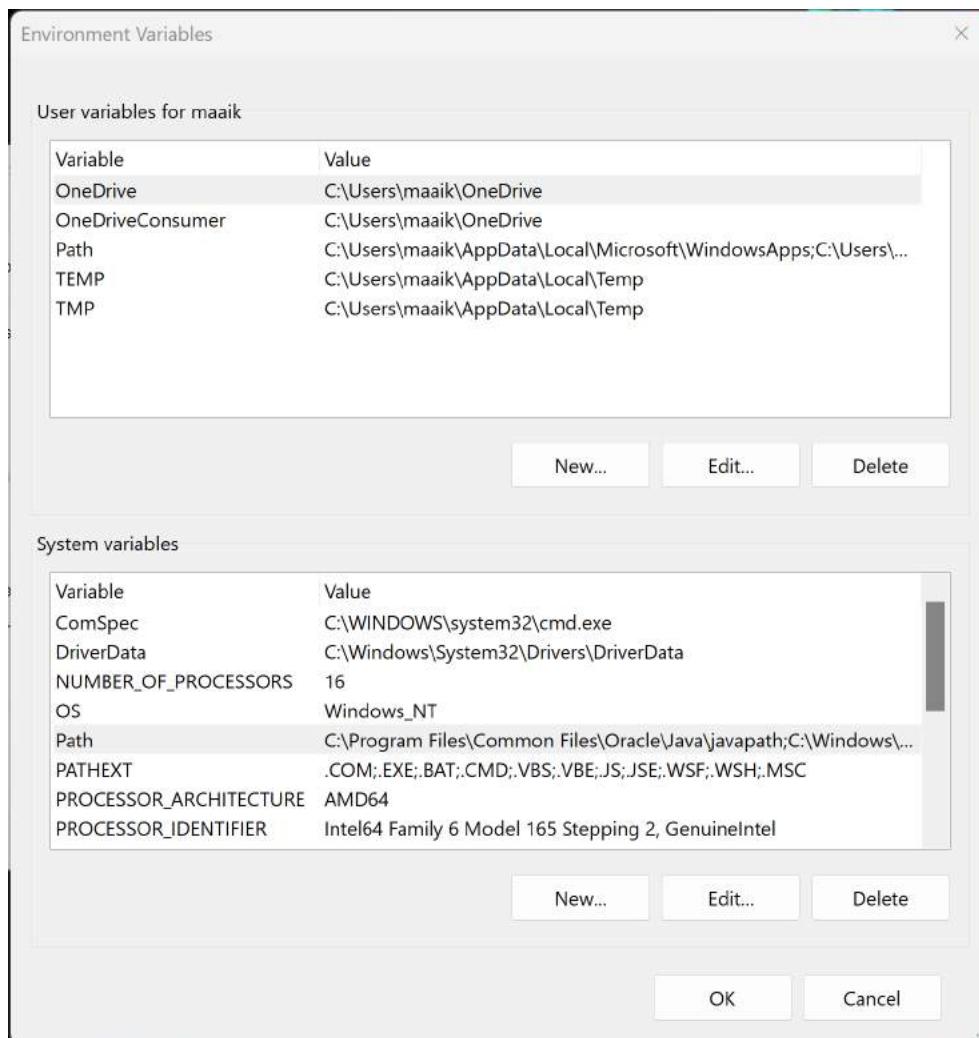


Figure 1.4 – The Environment Variables window

8. Click **New...** and add the path to the **bin** folder of your Java installation (for example, C :\ Program Files\Java\jdk-21\bin). In *Figure 1.5*, this has been done already.

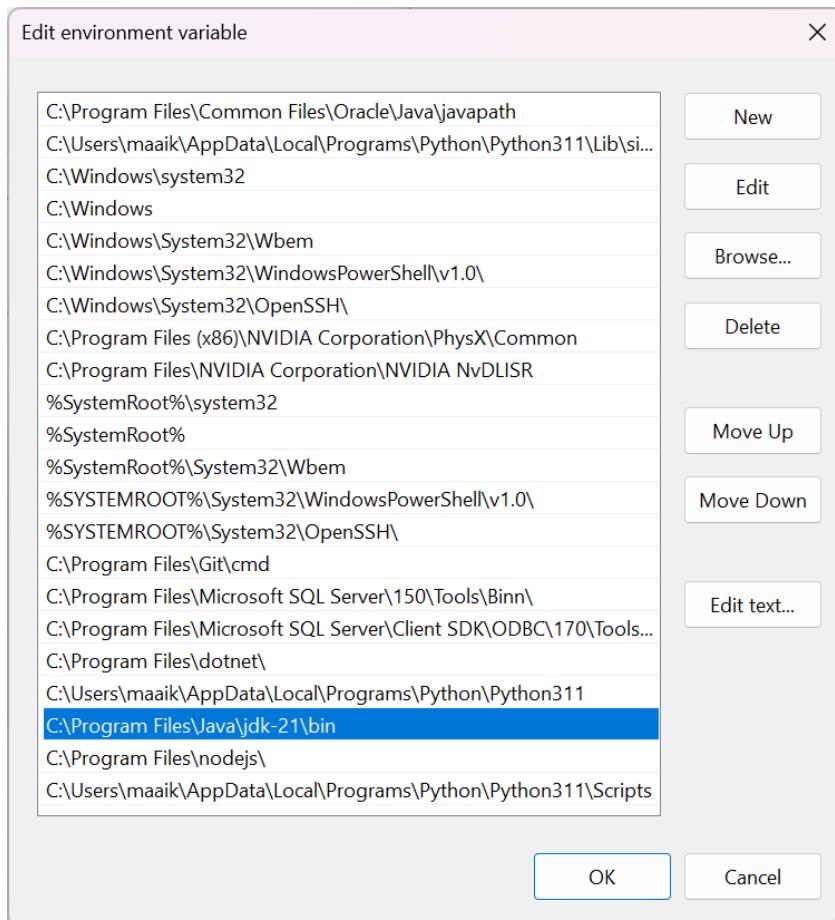


Figure 1.5 – Adding the path to Java to the Path variable

9. Click **OK** to save the changes and close the **Environment Variables** windows.
10. Verify Java is installed by opening the Command Prompt (reopen it if you have it open already) and then typing the following:

```
java -version
```



11. The output should look as shown in *Figure 1.6*. However, your version should be 21 or higher to keep up with all the snippets in this book.

A screenshot of a Windows Command Prompt window titled "Command Prompt". The window shows the following text:  
C:\Users\maaik>java -version  
java version "21" 2023-09-19 LTS  
Java(TM) SE Runtime Environment (build 21+35-LTS-2513)  
Java HotSpot(TM) 64-Bit Server VM (build 21+35-LTS-2513, mixed mode, sharing)  
C:\Users\maaik>

Figure 1.6 – Command Prompt after Java version check after installing Java

## Installing Java on macOS

To install Java on a macOS operating system, follow these steps:

1. Visit the **Oracle Java SE Downloads** page at <https://www.oracle.com/java/technologies/javase-jdk16-downloads.html>.
2. Select the macOS installer (for example, **macOS x64 Installer**).
3. Download the installer by clicking on the file link.
4. Run the downloaded installer (the .dmg file) and follow the on-screen instructions to complete the installation.
5. Java should be automatically added to your system's PATH environment variable. To verify the installation, open the Terminal and run the following command:

```
java -version
```

6. You should see the version of Java you just installed, similar to *Figure 1.2*.

## Installing Java on Linux

Installing on Linux can be a little bit tricky to explain in a few steps. Different Linux distributions require different installation steps. Here, we will see how to install Java on a Linux Ubuntu system:

1. Open the **Terminal** and update your package repository by running the following command:

```
sudo apt-get update
```

2. Install the default JDK package by running the following command:

```
sudo apt install default-jdk
```

3. To verify the installation, run the `java -version` command. You should see the version of Java you just installed.
4. If you need to set the `JAVA_HOME` environment variable (which you won't need for working your way through this book but will need for doing more complex Java projects), you first need to determine the installation path by running the following command:

```
sudo update-alternatives --config java
```

5. Take note of the path displayed (for example, `/usr/lib/jvm/java-19-openjdk-amd64/bin/java`).
6. Open the `/etc/environment` file in a text editor with root privileges:

```
sudo nano /etc/environment
```

7. Add the following line at the end of the file, replacing the path with the path you noted in *Step 4* (excluding the `/bin/java` part):

```
JAVA_HOME="/usr/lib/jvm/java-19-openjdk-amd64"
```

8. Save and close the file. Then, run the following command to apply the changes:

```
source /etc/environment
```

Now, Java should be installed and configured on your Linux operating system.

## Running Java online

If you don't have access to a computer with macOS, Linux, or Windows, there are online solutions out there. The free options are not perfect but, for example, the **w3schools** solution for trying Java in the browser is not bad at all. There are quite a few of these out there.

In order to work with multiple files there might be free tools out there, but most of them are paid. A currently free one that we would recommend is on `replit.com`. You can find it here: <https://replit.com/languages/java>.

You need to sign up, but you can work for free with multiple files and save them on your account. This is a good alternative if you would for example only have a tablet to follow along with this book.

Another option would be to use GitHub Codespaces: <https://github.com/codespaces>. They have the opportunity to enter a repository (for example the one we use for this book) and directly try the examples that are available in the repo and adjust them to try new things.

Having navigated through the installation of Java, it's time to talk about compiling and running programs.



## Writing our first program

Before diving into the process of compiling and running Java programs, let's create a simple Java program using a basic text editor. This will help you understand the structure of a Java program and how to write and save a Java source code file. For this example, we will create a **"Hello world!" program** that will be used to demonstrate the process of compilation and execution.

### Hello world

You may have heard of **"Hello world!" programs**. They are a common way to start learning a new programming language. It's a simple program that prints the message "Hello world!" to the console. Writing this program will provide you with a very basic understanding of Java syntax, and it will help you to become familiar with the process of writing, compiling, and running Java code.

### Steps to create the program

Alright, let's start coding. Here are the steps:

1. First, open a basic text editor on your computer. **Notepad** on Windows, **TextEdit** on macOS, or **Gedit** on Linux are suitable options.
2. Write the following Java code in your text editor:

```
public class HelloWorld {  
    public static void main(String[] args) {  
        System.out.println("Hello world!");  
    }  
}
```

3. Save the file as `HelloWorld.java` in a directory of your choice. Don't forget the `.java` extension when saving the file. This indicates that the file contains Java source code. The code should not have `.txt` after `.java`. This happens sometimes in Windows, so make sure to not select the text file in the filetype dropdown.

#### TextEdit – file extension issues

The later versions of macOS have some issues with **TextEdit**. You can't save it as a Java file by default. In order to enable this, you need to go to **Format | Make Plain Text** and select **UTF-8**.

After this, you can save it as a `.java` file. You may still run into encoding errors; the problem is with the encoding, and fixing it might be a lot of effort missing the goal of this exercise. It might be better to download **Notepad++**, **TextPad**, or **Sublime** for this part. Or go ahead and download the `HelloWorld.java` file from our GitHub repository.

## Understanding the program

Let's have a look at the code we just used. First of all, be aware that this is *case-sensitive*. That means that when you look at the code, most things will not work as you expect if you mix up lowercase and uppercase.

First, we created a class named `HelloWorld` with a `main` method. We'll cover classes and methods in a lot more detail, of course. But a class is the fundamental building block of Java applications, and it can contain methods. Methods can be executed to do things – *things* being executing statements.

The `main` method is a special method. It is the entry point of our Java program and contains the code that will be executed when the program is run. The line with `System.out.println("Hello world!") ;` writes the `Hello world!` message to the console. Please note, that `println` stands for print line, so it uses a lowercase *L* and not an uppercase *I*.

With the `HelloWorld.java` file saved, we are now ready to move on to the next section, where we will learn how to compile and run the Java program using the command line and an IDE.

## Compiling and running Java programs

Now that we have our first program written, let's discuss how we can compile and run it. We will cover the basics of the compilation process, the role of the JVM, and how to compile and run Java code using the command line and an IDE.

### Understanding the compilation process

The source code is written in a human-readable format using the Java programming language. Or at least, we hope that this is your opinion after this book. Before the code can be executed, it must be transformed into a format that the computer can understand. You already know that Java is a compiled language and that this process is called compilation.

During compilation, the **Java compiler (javac)** converts the source code (`.java` files) into bytecode (`.class` files). Once the bytecode is generated, it can be executed by the JVM. We have already learned that the JVM is the bytecode executor and that every platform has its own custom JVM enabling the WORA feature of Java.



## Compiling the code with javac on the command line

To compile a Java program using the command line, follow these steps:

1. Open a terminal (Command Prompt on Windows, Terminal on macOS or Linux).
2. Navigate to the directory containing your Java source code file (for example, the directory of your previously created `HelloWorld.java` file). In case you don't know how to do that, this can be done with the `cd` command, which stands for *change directory*. For example, if I'm in a directory called `documents` and I want to step into the subfolder called `java programs`, I'd run the `cd "java programs"` command. The quotes are only needed when there are spaces in the directory name. It's beyond the scope of this book to explain how to change directories for any platform. There are many excellent explanations for every platform on how to navigate the folder structure using the command line on the internet.
3. Once you're in the folder containing the Java file, enter the following command to compile the Java source code:

```
javac HelloWorld.java
```

If the compilation is successful, a new file with the same name but a `.class` extension (for example, `HelloWorld.class`) will be created in the same directory. This file contains the bytecode that can be executed by the JVM.

Let's see how we can run this compiled code.

## Running the compiled code with Java on the command line

To run the compiled Java program, follow these steps:

1. In the terminal, make sure you are still in the directory containing the `.class` file.
2. Enter the following command to execute the bytecode:

```
java HelloWorld
```

The JVM will load and run the bytecode, and you should see the output of your program. In this case, the output will be as follows:

```
Hello world!
```

It's pretty cool that we can write Java in Notepad and run it on the command line, but the life of a modern-day Java developer is a lot nicer. Let's add IDEs to the mix and see this for ourselves.



## Working with an IDE

Creating files in text editors is a little old-fashioned. Of course, you can still do it this way – it's actually an excellent way of becoming an amazing programmer, but it's also a very frustrating way. There are tools available to do quite a bit of the heavy work for us and to assist us while writing our code. These tools are called IDEs.

### What is an IDE?

An IDE is a software application that comes with everything you need to write, compile, run, and test your code. Using an IDE can make it easier to develop all sorts of programs. Not only that but also debugging and managing your code is easier. Comparatively, you can think of an IDE somewhat like Microsoft Office Word for me as I write this book. While I could have written it using Notepad, using Word provides significant advantages. It assists in checking for spelling errors and allows me to easily add and visualize layouts, among other helpful features. This analogy paints a picture of how an IDE doesn't just provide a platform to write code but also offers a suite of tools to streamline and enhance your coding experience.

### Choosing an IDE

When it comes to Java development, there are several IDEs available, each with its own set of features and capabilities. In this section, we will discuss the factors to consider when choosing an IDE and help you set up some popular Java IDEs. Throughout this book, we'll be working with **IntelliJ**. Alternatives that are also great would be **VS Code** and **Eclipse**.

#### *Factors to consider when choosing an IDE*

Most modern IDEs have features such as code completion, debugging, version control integration, and support for third-party tools and frameworks. Some of them have better versions of these than others. Compare and contrast what you prefer when choosing or switching IDEs.

Some IDEs require a heavier system to run on than others. For example, VS Code is rather lightweight and IntelliJ is rather heavy. Also, VS Code can be used for many languages, including Java. It is uncommon to do a lot of other things with IntelliJ rather than Java. Choose an IDE that provides a balance between features and performance, especially if you have limited system resources.

And of course, it's possible that the IDE you'd prefer is not available for the platform you're using. Make sure that it's available and stable for your system.

Lastly, and very importantly, think about the costs. Some IDEs are free and others require a paid license. Luckily, many of the ones that require a paid license have a free edition for non-commercial use. So, make sure to also consider your budget and the licensing you need when choosing an IDE.



In the following subsections, we'll walk you through the steps of setting up the three (currently) most common IDEs for Java development:

- IntelliJ
- Eclipse
- Visual Studio Code

#### Note

We'll be working with IntelliJ for the rest of this book.

### ***Setting up IntelliJ***

So, let's start with that one. IntelliJ IDEA is a popular Java IDE that was developed by **JetBrains**. It offers both a free **Community Edition** and a paid **Ultimate Edition**. It provides a wide range of features, including intelligent code completion, debugging tools, version control integration, and support for various Java frameworks.

Here are the steps for installing IntelliJ:

1. Visit the IntelliJ IDEA download page at <https://www.jetbrains.com/idea/download/>.
2. Choose the edition you prefer: the free **Community Edition** or the paid **Ultimate Edition**. For beginners, the Community Edition is truly great already.
3. Download the installer for your operating system (Windows, macOS, or Linux).
4. Run the installer and follow the instructions to complete the installation.
5. Launch **IntelliJ IDEA**. If you're using the Ultimate Edition, you may need to enter your JetBrains account credentials or a license key.
6. On the **Welcome** screen, you can create a new **project**, import an existing **project**, or explore the available tutorials and documentation.

### ***Setting up Eclipse***

Eclipse is a free, open source Java IDE that is widely used in the Java community. It has been around for a really long time already and quite a lot of companies work with it still. It offers a variety of features, just like IntelliJ. Eclipse can be customized to suit your needs, but its interface may be less intuitive than other IDEs.

To set up Eclipse, follow these steps:

1. Visit the Eclipse download page at <https://www.eclipse.org/downloads/>.
2. Download the Eclipse installer for your operating system (Windows, macOS, or Linux).
3. Run the installer and select **Eclipse IDE for Java Developers** from the list of available packages.
4. Choose an installation folder and follow the instructions to complete the installation.
5. Launch **Eclipse** and select a workspace directory. This is where your projects and settings will be stored.
6. On the **Welcome** screen, you can create a new Java **project**, import an existing **project**, or explore the available tutorials and documentation.

### **Setting up Visual Studio Code**

Visual Studio Code, often referred to as VS Code, is a lightweight, free, and open source code editor developed by Microsoft. It's incredibly popular for all sorts of tasks because it supports a wide range of programming languages. It is a popular choice for developers who prefer a more minimalist and fast-performing environment. All sorts of additions can be added with the use of extensions.

Here are the steps for installing VS Code and preparing it for Java development:

1. Visit the Visual Studio Code download page at <https://code.visualstudio.com/download>.
2. Download the installer for your operating system (Windows, macOS, or Linux).
3. Run the installer and follow the on-screen instructions to complete the installation.
4. Launch Visual Studio Code.
5. Open the **Extensions** view by clicking on the *Extensions* icon (four squares) on the left side of the window.
6. Search for **Java Extension Pack** in the *Extensions Marketplace* and click the **Install** button. This extension pack includes various extensions for Java development, such as **Language Support for Java (TM) by Red Hat**, **Debugger for Java**, and **Maven for Java**.
7. With the **Java Extension Pack** installed, you can now create or import Java projects. If it doesn't load directly, you may need to reopen VS Code.

Now that you've set up an IDE, let's create and run a program with it.

## **Creating and running a program with an IDE**

Working with an IDE such as IntelliJ as compared to working with a plain text editor is a breeze. We're now going to guide you through creating, running, and debugging a program with the use of IntelliJ. We'll create the same program as we did when we were using the text editor.

## Creating a program in an IDE

When you use an IDE to type code, you'll see that it helps you to complete your code constantly. This is considered very helpful by most people, and we hope you'll enjoy this feature too.

In order to get started with IntelliJ, we first need to create a project. Here are the steps for creating our `Hello World` program again:

1. Launch IntelliJ IDEA and click on **New Project** from the **Welcome** screen or go to **File | New | Project**.

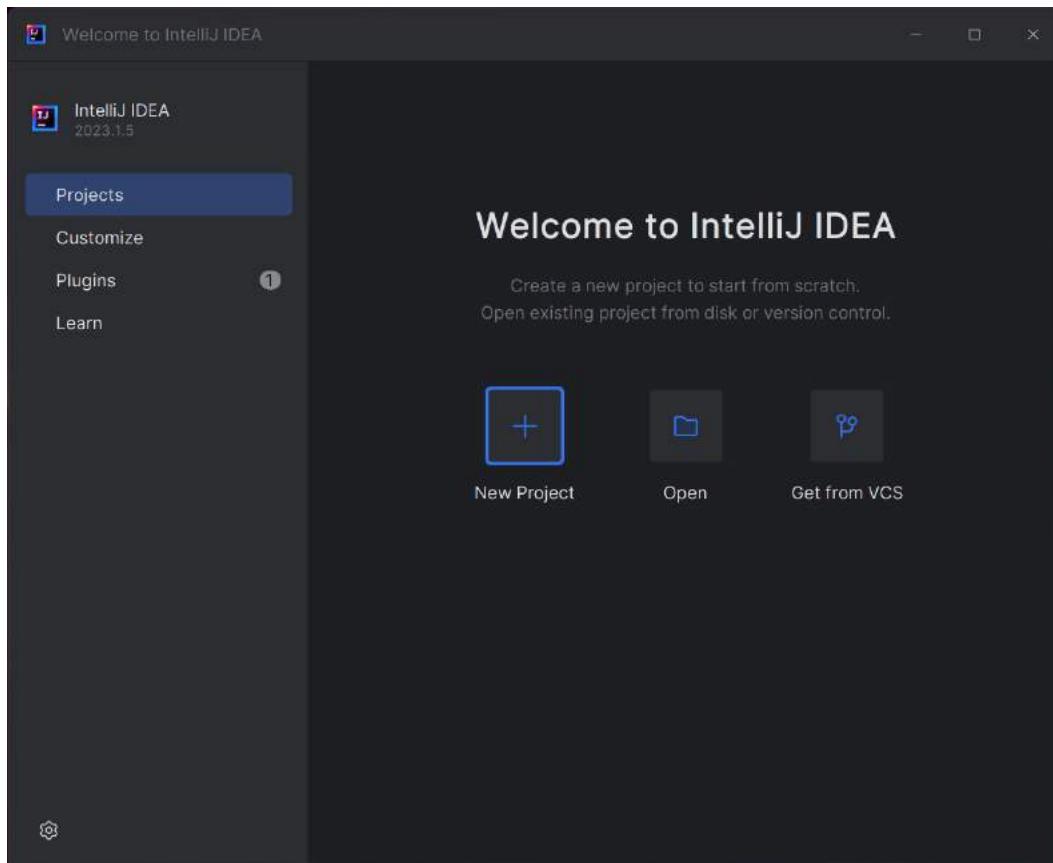


Figure 1.7 – Initial screen of IntelliJ

2. Name the project `HelloWorld`.
3. Select **Java** for the language and make sure that the correct project SDK is selected. Click **Next**.
4. Don't tick the **Create Git repository** box and don't tick the **Add sample code** box.

5. Click **Create** to create the project.

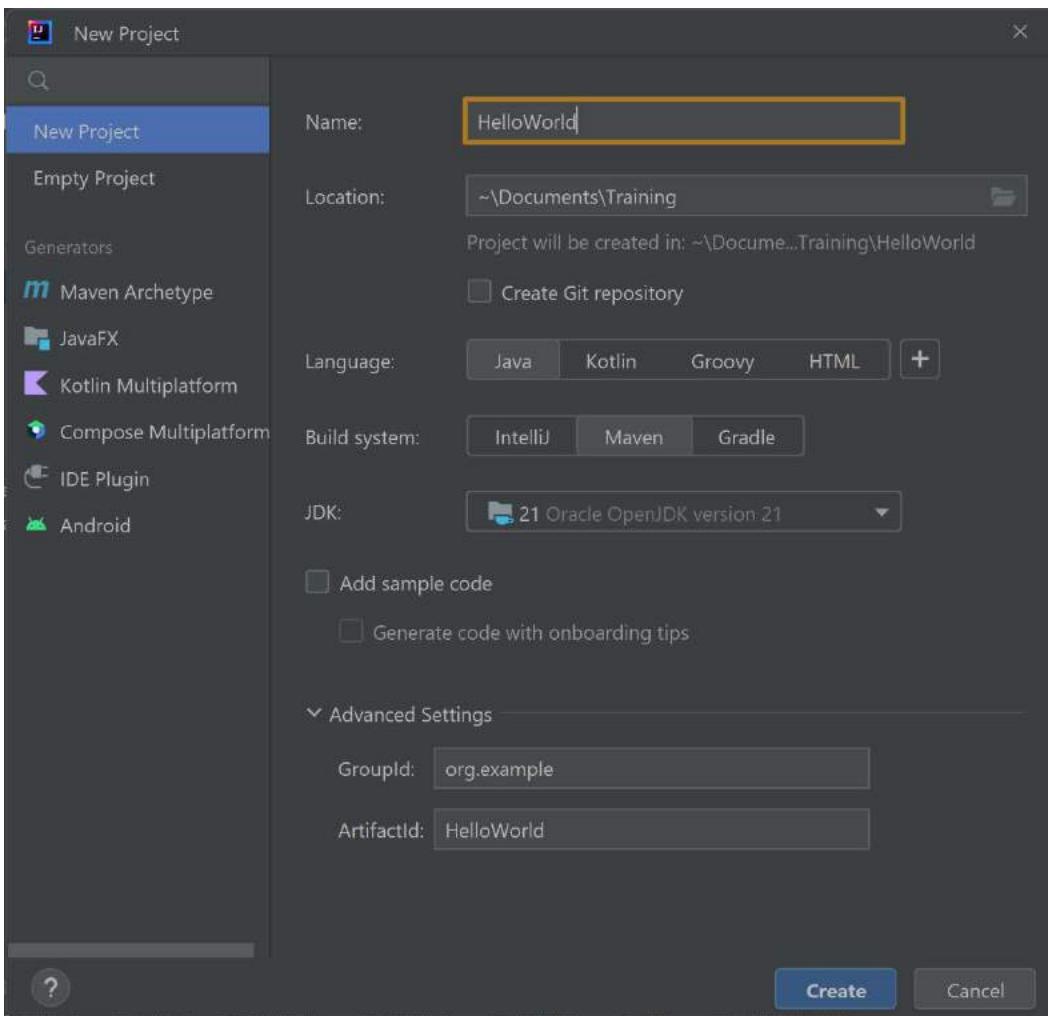


Figure 1.8 – Wizard to create a new project

6. Once the project is created, expand the `src` folder in the **Project** view on the left. If there is no other folder underneath it, right-click on the `src` folder and select **New | Java Class**. If there is another folder underneath it, there is probably a main folder with a Java folder in there. Right-click on the Java folder and select **New | Java Class**. If it's called something differently, just right-click on the blue folder.

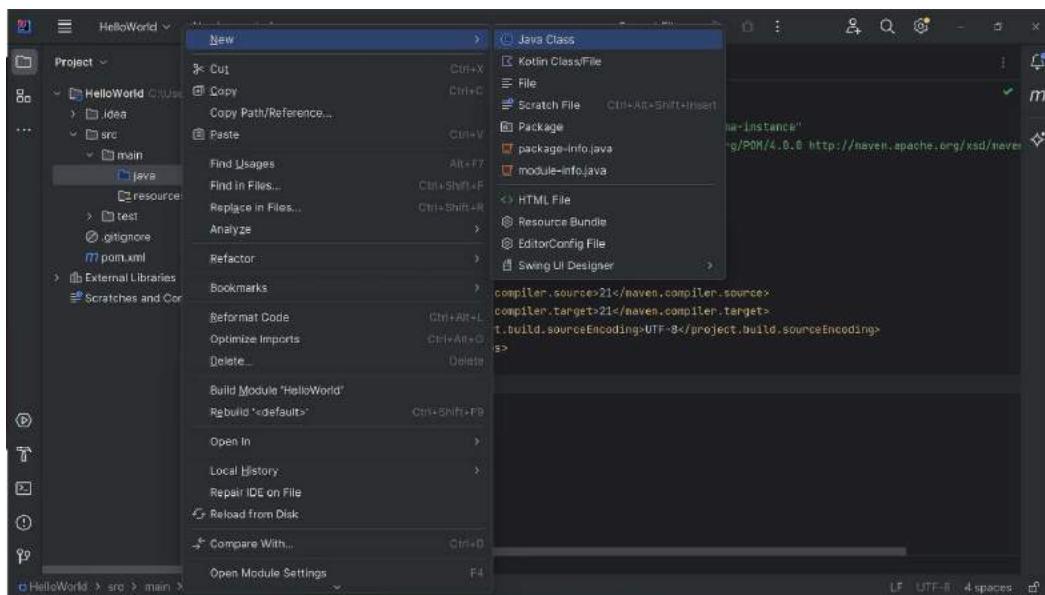


Figure 1.9 – Create a new Java Class

7. Name the new class `HelloWorld` and click **OK**. IntelliJ will create a new `.java` file with the class definition.

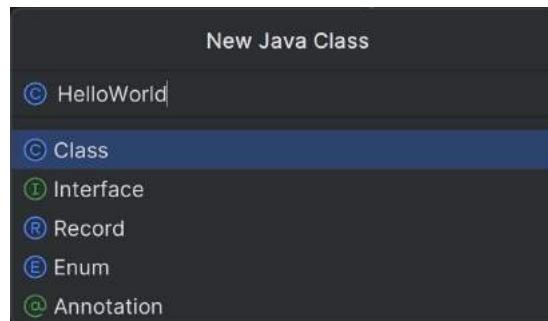


Figure 1.10 – Call the class "HelloWorld"

8. In the `HelloWorld` class, write our main method:

```
public class HelloWorld {  
    public static void main(String[] args) {  
        System.out.println("Hello world!");  
    }  
}
```



The screenshot shows the IntelliJ IDEA interface with a dark theme. A code editor window is open, titled "HelloWorld.java". The code contains a single class definition with a main method that prints "Hello world!". The code is numbered from 1 to 9 on the left. A cursor is visible at the end of the closing brace of the main method.

Figure 1.11 – Code in `HelloWorld.java`

Now that we've written our first program, make sure that it is saved. By default, IntelliJ automatically saves our files. Let's see whether we can run the program as well.

## Running your program

Admittedly, we had to take a few extra steps to create our program. We had to create a project first. The good news is, running the program is easier! Here's how to do it:

1. If you haven't done so, make sure your changes are saved by pressing *Ctrl + S* (Windows/Linux) or *Cmd + S* (macOS). By default, auto-save is enabled.
2. To run the program, right-click anywhere in the `HelloWorld` class and select Run '`HelloWorld.main()`'. Alternatively, you can click the green triangle icon next to the main method and select Run '`HelloWorld.main()`'. IntelliJ will compile and run the program.

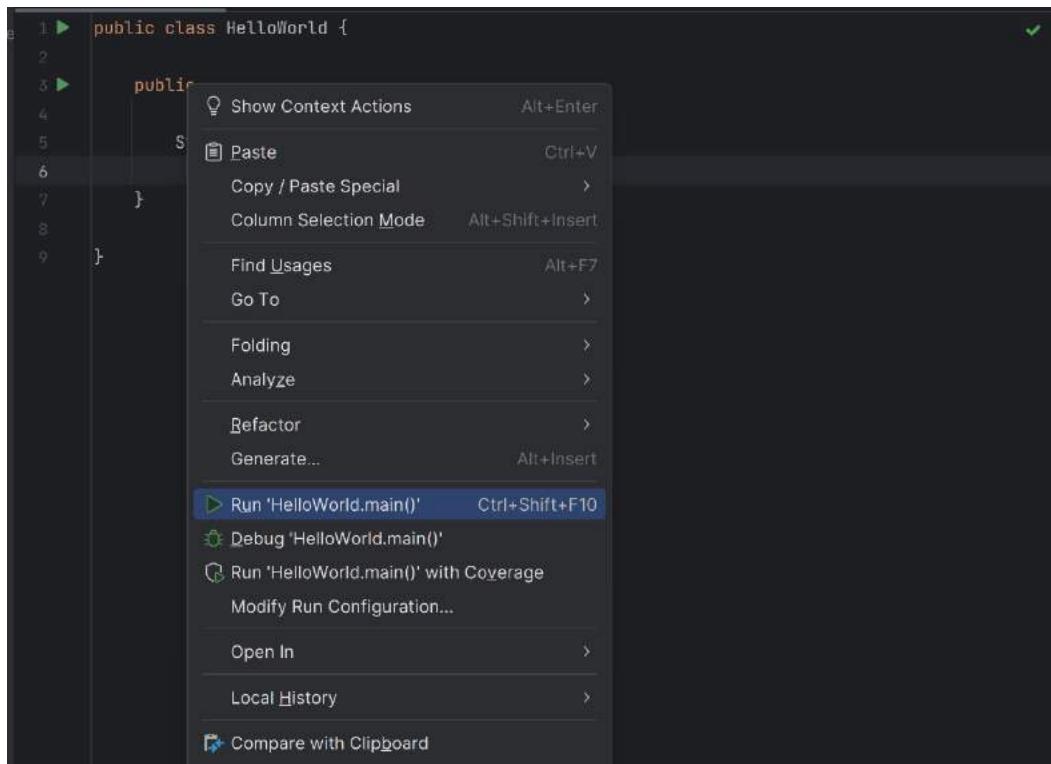


Figure 1.12 – Running the program

- Verify that the output of the program, "Hello world!", is displayed in the **Run tool** window at the bottom of the screen.

```
"C:\Program Files\Java\jdk-21\bin\java.exe" "-  
Hello world!  
  
Process finished with exit code 0
```

Figure 1.13 – Output of the program

### Saved and unsaved files

In most IDEs, you can tell whether a file is saved or not by looking at the tab of the open file. It has a dot or an asterisk next to it if it isn't saved. The dot is missing if it has been saved.

## Debugging a program

Our program is quite easy right now, but we may want to step through our program line by line. We can do that by debugging the program. Let's give our file a little extra content for debugging. This way we can see how to inspect variables, understand the execution flow, and, this way, find the flaws in our code:

1. Update the `HelloWorld.java` file with the following code:

```
public class HelloWorld {  
    public static void main(String[] args) {  
        String greeting = "Hello, World!";  
        System.out.println(greeting);  
  
        int number = 5;  
        int doubled = doubleNumber(number);  
        System.out.println("The doubled number is: " +  
            doubled);  
    }  
  
    public static int doubleNumber(int input) {  
        return input * 2;  
    }  
}
```

2. In this updated version of the program, we added a new method called `doubleNumber`, which takes an integer as input and returns its double. In the `main` method, we call this method and print the result. Don't worry if you don't fully get this – we just want to show you how you can step through your code.
  3. Save your changes by pressing `Ctrl + S` (Windows/Linux) or `Cmd + S` (macOS).
- Now, let's debug the updated program.
4. Set a breakpoint on the line you want to pause the execution at by clicking in the gutter area next to the line number in the editor. A red dot will appear, indicating a breakpoint. For example, set a breakpoint at the line `int doubled = doubleNumber(number);`. An example is shown in *Figure 1.7*.

```
1 public class HelloWorld {  
2     no usages  
3     public static void main(String[] args) {  
4         String greeting = "Hello, World!";  
5         System.out.println(greeting);  
6         int number = 5;  
7         int doubled = doubleNumber(number);  
8         System.out.println("The doubled number is: " + doubled);  
9     }  
10    1 usage  
11    |     public static int doubleNumber(int input) {  
12    |         return input * 2;  
13    |     }  
14  
15
```

Figure 1.14 – Adding a breakpoint on line 7

5. Start the debugger by right-clicking in the `HelloWorld` class and selecting `Debug 'HelloWorld.main()'` or you can click the green play icon next to the `main` method and select the **debug** option. IntelliJ will compile and start the program in debug mode.
6. When the line with the breakpoint is going to be executed, the program will pause. During the pause, you can use the **Debug** tool window, which will appear at the bottom of the screen. Here, you can view the state of the program, including the values of local variables and fields. An example is shown in *Figure 1.8*.

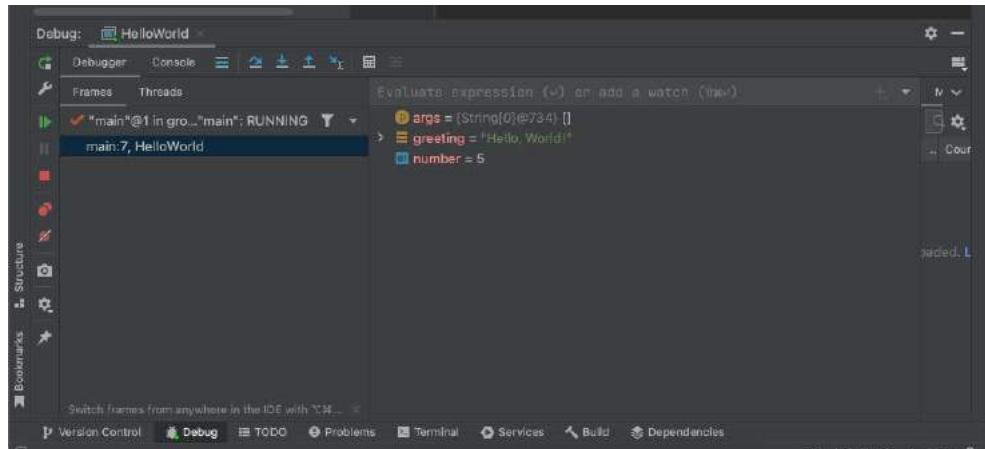


Figure 1.15 – Debug tool window in IntelliJ. The intent of this screenshot is to show the layout and text readability is not required.

7. Use the step controls in the **Debug** tool window to step through the code (blue arrow with the angle in *Figure 1.8*), step into the method that is being called (blue arrow down), or continue the execution (green arrow on the left in *Figure 1.8*).

By following these steps, you can debug Java programs using IntelliJ IDEA and step through the code to see what is happening. This is something that will come in handy to understand what is going on in your code. This process will be similar in other Java IDEs, although the specific steps and interface elements may vary.

## Exercises

And that's all theory for this chapter! So, roll up those sleeves, and let's dive into your first day at Mesozoic Eden. Welcome aboard! Mesozoic Eden is a famous zoo where dinosaurs live that have been brought to live with high end genetic manipulation techniques. Here are some exercises for you to test your knowledge so far:

1. Your first task involves welcoming our guests. Modify the following code snippet so that it outputs "Welcome to Mesozoic Eden":

```
public class Main {  
    public static void main(String[] args) {  
        System.out.println("Hello world");  
    }  
}
```

2. Complete the following programs by filling out the blanks so that it prints out your name and the position you want to have in Mesozoic Eden 5 years from now:

```
public class Main {  
    public static void main(String[] args) {  
        _____;  
        String position = "Park Manager";  
        System.out.println("My name is " + name + "  
            and I want to be a " _____ " in Mesozoic  
            Eden.");  
    }  
}
```

3. We've received some questions about opening hours. Complete the following program so that it prints the park's opening and closing hours:

```
public class Main {  
    public static void main(String[] args) {  
        String openingHours = "08:00";  
        String closingHours = "20:00";  
    }  
}
```

4. Create a Java project with a package named `dinosaur`. You can create a package by right-clicking on the `src/main/java` folder, selecting “new” and choosing “package”.
5. Modify the code from exercise 1 so that it prints out `Welcome, [YourName] to Mesozoic Eden!`, where `[YourName]` is replaced by, surprise surprise, your name. Bonus: try to create a separate String variable as shown in the second and third exercises.
6. Some guests reported feeling unsafe near the T-Rex. Let’s solve this by adding another `System.out.println` to the program of exercise 5. It should print the phrase `"Mesozoic Eden is safe and secure."` after the welcome message.

## Project

Create a program that simulates a sign at the entrance of Mesozoic Eden. The sign is simulated by printing output to the console. The sign should display a welcome message, the opening and closing hours, and a short safety message.

## Summary

You’ve made it through the first chapter! And we’ve done a lot already. We kicked off by exploring Java’s key features, such as its OOP approach, the (once unique) WORA principle, its compiled nature, and the super-helpful automatic memory management. These features make Java an incredibly versatile and powerful language – a great choice for different programming tasks, such as web development, desktop apps, mobile apps, and so much more!

Next, we walked you through the process of installing Java on various platforms: Windows, macOS, and Linux. We also discussed how to check whether Java is already installed on your system. After this part, you can be sure that you have all the essential tools to kick off your Java programming adventure.

After you had Java all setup, we demystified the compilation process and introduced you to the JVM, a vital component of the Java ecosystem that enables the portability of Java code. We then demonstrated how to compile and run Java code using the `javac` and `java` command-line tools. These tools lay the groundwork for working with Java programs at their core.

Of course, using the command line for this is great. But nowadays, we more often work with an IDE, and we can just press a button to do all this. So, we mentioned several advantages and nice features of working with an IDE, such as code completion, debugging, and project management. We discussed the factors to weigh up when choosing an IDE and provided guidance on setting up popular IDEs such as IntelliJ IDEA, Eclipse, and VS Code. In this book, we’ll be using IntelliJ throughout for the examples.

After covering the essentials of IDEs, we delved into creating and running a Java program using an IDE. We explained the structure of a typical Java program and guided you, step by step, through the process of creating, running, and debugging your very first Java program.

After this, you were ready for the first hands-on project. And now you’re here! All set and ready to take the next step on your Java journey. This next step will be working with variables and primitive data types. Good luck!

## 2

# Variables and Primitive Data Types

In *Chapter 1*, we introduced the compiler and the JVM. We learned how to use both of them from the command line when we wrote our first Java program, *HelloWorld*. We also introduced **IntelliJ**, a powerful and friendly IDE, and we ran *HelloWorld* from there as well.

All programming languages require variables and provide in-built primitive data types. They are essential for the operation of even the simplest programs. By the end of this chapter, you will be able to declare variables using Java's primitive types. In addition, you will understand the differences between the various primitive data types and which ones to use in a given situation.

In this chapter, we are going to cover the following main topics:

- Understanding and declaring variables
- Exploring Java's primitive data types

## Technical requirements

The code for this chapter can be found on GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects/tree/main/ch2>.

## Understanding and declaring variables

If you want to store a value for later use, you need a variable. Therefore, every programming language provides this feature via variables. In this section, we will learn what a variable is and how to declare one. The area in your code where you can use a particular variable is known as the variable's *scope*. This is a very important concept and will be covered in detail in *Chapter 4*.



## What is a variable?

Variables are locations in memory that have a *name* (called an identifier) and a *type*. They resemble named pigeonholes or post office boxes (see *Figure 2.1*). The variable's name is required so that we can refer to the variable and distinguish it from other variables.

A variable's *type* specifies the sort of values it can store/hold. For example, is the variable to be used for storing whole numbers such as 4 or decimal numbers such as 2.98? The answer to that question determines the variable's *type*.

## Declaring a variable

Let's suppose we want to store the number 25 in a variable. We will assume that this number represents a person's age, so we will use the `age` identifier to refer to it. Introducing a variable for the first time is known as "declaring" the variable.

A whole number (positive or negative) is an integer and Java provides an in-built primitive type especially for integers called `int`. We will discuss primitive data types in more detail in the next section. When declaring a variable in Java, we must specify the variable's type. This is because Java is known as a *strongly typed* language, which means you must specify the variable's type immediately upon declaring it.

Let's declare a variable, give it a type, and initialize it:

```
int age;  
age = 25;
```

The first line declares `age` as an `int` and the second line assigns it a value of 25.

Note that the semi-colons (`;`) at the end of the lines are delimiters that tell the compiler where a Java statement ends. The `=` sign is the assignment operator and will be covered in *Chapter 3*. For now, just realize that 25 is "assigned into" the `age` variable.

### Assignment operator

The `=` sign in Java is not the same as the equals sign, `=`, in mathematics. Java uses the `==` sign for equals, which is called equivalence.

We can write the previous two lines of code in one line:

```
int age = 25;
```

Figure 2.1 shows the in-memory representation of both code segments:

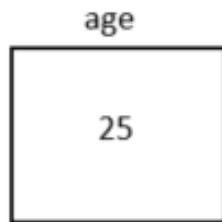


Figure 2.1 – An integer variable named age with a value of 25

As can be seen in Figure 2.1, **age** is the name of the variable and **25** is the integer value stored in the variable's *location*.

## Naming a variable

An identifier is simply a name that you give to the Java construct you are coding; for example, identifiers (names) are required for naming variables, methods, classes, and so forth.

### Identifiers

An identifier consists of letters, digits, underscores, and currency symbols. Identifiers cannot begin with a number and cannot contain whitespace (spaces, tabs, and newlines). In the following examples, the commas separate the various identifiers.

Examples of unusual but *valid* identifiers are `a€_23, _29, Z,`  
`thisIsAnExampleOfAVeryLongVariableName, €2_, and $4 ;.`

Examples of *invalid* identifiers are `9age` and `abc def ;.`

Name your variables carefully. This helps make code more readable, which results in fewer bugs and easier maintenance. *Camel case* is a very popular technique in this regard. Concerning variable names, camel case means that *all* of the first word is lowercase. In addition, the first letter in each subsequent word in the variable name starts with an uppercase letter. Here's an example:

```
int ageOfCar;  
int numberOfChildren;
```

In this code segment, we have two integer variables whose names/identifiers follow camel casing.

## Accessing a variable

To access a variable's value, just type in the variable's name. When we type in a variable's name in Java, the compiler will first ensure that a variable with that name exists. Assuming there is, the JVM will, at runtime, return the *value inside* that variable's pigeonhole. Therefore, the following code will output 25 to the screen:

```
int age = 25;  
System.out.println(age);
```

The first line declares the `age` variable and initializes it to 25. The second line accesses the variable's location and outputs its value.

### System.out.println()

`System.out.println()` displays whatever is inside the round brackets, `()`, on the screen.

## Accessing a variable that you have not declared

As stated previously, Java is known as a *strongly typed* language. This means that you have to specify the variable's type immediately upon declaring it. If the compiler comes across a variable and does not know its type, it generates an error. For example, consider the following line of code:

```
age = 25;
```

Assuming no other code declares `age`, the compiler generates an error stating `cannot resolve symbol 'age'`. This is because the compiler is looking for the *type* to associate with `age` and it cannot find it (as we did not specify the type).

Here is a slightly different example. In this example, we are attempting to output the `age` variable to the screen:

```
int length = 25;  
System.out.println(age);
```

In this example, we have declared a variable named `length` and thus, there is no declaration of `age`. When we attempt to access the `age` variable in `System.out.println()`, the compiler goes looking for `age`. The compiler cannot find `age` and generates an error stating `cannot resolve symbol 'age'`. In effect, what the compiler is saying is that we have attempted to use a variable named `age` that the compiler cannot find. This is because we did not even declare the variable, not to mention specify its type.

## Comments

Comments are very useful as they help us explain what is happening in the code.

// is a single-line comment. When the compiler sees //, it ignores the rest of the line.

/\* some text \*/ is a multi-line comment. Anything between the opening /\* and the closing \*/ is ignored. This format saves inserting // at the start of each line.

Here are some examples:

```
int age; // from here to the rest of the line is ignored
// this whole line is ignored
/* all
of
these lines
are
ignored */
```

Given that Java, as a strongly typed language, requires all variables to have a data type, we will now discuss Java's support for primitive data types.

# Understanding Java's primitive data types

Java provides eight in-built data types. In-built means that these data types come with the language. These primitive data types are the topic of this section.

## Java's primitive data types

All of the primitive data types are named using lowercase letters only; for example, int and double. When we create our own data types later on, namely classes, records, and interfaces, we will follow a different naming convention. For example, we may have a class named Person or Cat. This is simply a widely adopted coding convention and the compiler does not distinguish between naming conventions. However, it is very easy to recognize any of the primitive data types as they are always in lowercase letters only. Before we discuss the primitive data types themselves, there are a few important points to make.

### **Numeric primitive data types are signed**

In Java, all numeric primitive data types are represented as a series of bits. In addition, they are also signed. The most significant bit (leftmost bit) is used for the sign; 1 means negative and 0 means positive.



### Integer literals

A literal value is one that's typed in at the keyboard (as opposed to a computed value). An integer literal can be expressed in various numbering systems: decimal (base 10), hexadecimal (base 16), octal (base 8), and binary (base 2). However, it is no surprise that decimal is by far the most commonly used representation. For information purposes, all of the following declarations represent the decimal number 10:

```
int a = 10;      // decimal, the default
int b = 0b1010; // binary, prefixed by 0b or 0B
int c = 012;    // octal, prefixed by 0
int d = 0xa;    // hexadecimal, prefixed by 0x or 0X
```

### The sign bit affects the range

The presence of the sign bit means that `byte` has a range of  $-2^7$  to  $2^7-1$  (-128 to +127 inclusive). The -1 in the positive range is to allow for the fact that, in Java, 0 is considered a positive number. There is *not* one less positive number in any of the ranges. For example, with `byte`, you have 128 negative numbers (-1 to -128) and 128 positive numbers (0 to +127), resulting in 256 representations ( $2^8$ ). To reinforce this point with a simple example, -1 to -8 is 8 numbers and 0 to 7 (inclusive) is 8 numbers also.

With these points discussed, let's look at the various primitive types. *Table 2.1* lists the eight primitive data types, their byte sizes, and their ranges (all of which are inclusive):

primitive type	number of bytes	range
<code>byte</code>	1	-128..+127 (- $2^7$ to $2^7-1$ )
<code>short</code>	2	-32,768..+32,767 (- $2^{15}$ to $2^{15}-1$ )
<code>int</code>	4	- $2^{31}$ to $2^{31}-1$
<code>long</code>	8	- $2^{63}$ to $2^{63}-1$
<code>float</code>	4	3.4e-38 to 3.4e+38
<code>double</code>	8	1.7e-308 to 1.7e+308
<code>boolean</code>	1	<code>true</code> or <code>false</code>
<code>char</code>	2 (unsigned)	0..65,535 (0.. $2^{16}-1$ )

Table 2.1 – Java's primitive types

Here are some interesting points from the preceding table:

- `byte`, `short`, `char`, `int`, and `long` are known as *integral* types as they have integer values (whole numbers, positive or negative). For example, -8, 17, and 625 are all integer numbers.
- `char` is used for characters – for example 'a', 'b', '?' and '+'. Note that single quotes surround the character. In code, `char c = 'a';` means that the variable `c` represents the letter `a`. As computers ultimately store all characters (on the keyboard) as numbers internally (binary), we need an encoding system to map the characters to numbers and vice versa. Java uses the Unicode encoding standard, which ensures a unique number for every character, regardless of platform, language, script, and so on. This is why `char` uses 2 bytes as opposed to 1. In fact, from the computer's perspective, `char c = 'a';` is the same as `char c = 97;` where 97 is the decimal value for 'a' in Unicode. Obviously, we as humans prefer the letter representation.
- `short` and `char` both require 2 bytes but have different ranges. Note that `short` can represent negative numbers, whereas `char` cannot. In contrast, `char` can store numbers such as 65,000, whereas `short` cannot.
- `float` and `double` are for floating-point numbers – in other words, numbers that have decimal places, such as 23.78 and -98.453. These floating-point numbers can use scientific notation – for example, 130000.0 can be expressed as `double d1=1.3e+5;`, and 0.13 can be expressed as `double d2=1.3e-1;`.

### ***Representation of the various types***

Expanding from the previous callout, we can express integer literals using the following numbering systems:

- **Decimal:** Base 10; numbers 0..9. This is the default.
- **Hexadecimal:** Base 16; numbers 0..9 and letters a..f (or A..F). Prefix the literal with `0x` or `0X` to indicate that this is a hexadecimal literal.
- **binary:** Base 2; numbers 0..1. Prefix the literal with `0b` or `0B` to indicate that this is a binary literal.

Here are some sample code fragments that initialize `int` variables to 30 using the various numbering systems. Firstly, decimal is used; then hexadecimal, and finally, binary:

```
// decimal
int dec = 30;

// hexadecimal = 16 + 14
int hex = 0x1E;

// binary = 16 + 8 + 4 + 2
int bin = 0b11110;
```



Although there are several ways to initialize an `int`, using decimal is by far the most common.

A literal number, such as `22`, is considered an `int` by default. If you want to have `22` treated as `long` (instead of `int`), you must suffix either an uppercase or lowercase `L` to the literal. Here's an example:

```
int x = 10;
long n = 10L;
```

As per *Table 2.1*, using a `long` as opposed to an `int`, gives you access to much bigger and much smaller numbers. Use of uppercase `L` as opposed to lowercase `l` to signify `long` is preferred, as the lowercase `l` is similar to the number `1` (one).

Floating-point numbers behave similarly. A decimal number is, by default, `double`. To have any decimal number treated as `float` (as opposed to `double`), you must suffix the literal with either an uppercase or lowercase `F`. Assuming range is not an issue then one reason for using `float` as opposed to `double` is memory conservation (as `float` requires 4 bytes whereas `double` requires 8 bytes). Here's an example:

```
double d = 10.34;
float f = 10.34F;
```

Variables of the `char` type are initialized with single quotes around the literal. Here's an example:

```
char c = 'a';
```

Variables of the `boolean` type can store only `true` or `false`. These `boolean` literals are in lowercase only as they are reserved words in Java, and Java is case-sensitive:

```
boolean b1 = true;
boolean b2 = false;
```

That concludes this section on Java's primitive type system, where we examined the various types, their sizes/ranges, and some code segments.

Now, let's put the theory of variables and primitive types into practice! But before that, here's a bit of a cheat code to help you with the exercises.

## Screen output

As we know, `System.out.println()` outputs what is inside the `()`. To do the exercises, we want to expand on that. Firstly, here's some code:

```
String name = "James";      // line 1
int age = 23; // line 2
double salary = 50_000.0; // line 3
String out = "Details: " + name + ", " + age + ",
```



```
" + salary;//line 4
System.out.println(out); // line 5
```

Line 1 declares a string literal "James" and initializes the name variable with it. A string literal is a sequence of characters (including numbers), enclosed in double quotes. We will discuss the `String` class in detail in *Chapter 12*.

Lines 2 and 3 should be fine. We are declaring an `int` type called `age` and a `double` type called `salary` and using literal values to initialize them. The underscore used in line 3, enables us to make large numbers easier to read.

Line 4 builds the string to be output, namely `out`. We want to output the variables values, along with some helpful text to explain the output. Java builds the string from left to right.

Here, `+` is not the regular mathematical addition. We will discuss this in detail in *Chapter 3*, but for the moment, realize that when you have a string variable or literal on the left or the right of `+`, the operation becomes a `String` append (as opposed to mathematical addition).

One property this append shares with addition is that both sides of `+` must be of the same type. Since not all the variables in this example are string variables (or literals), Java has some work to do in the background (to get them all to the same type). Java copies the numeric variable values into new string locations to use them in building the string. For example, there is a location somewhere in memory that's been created for a string literal, "23" (in addition to the `int` location for `age`). This also happens for the `double` type's `salary` variable. Now, Java is ready to construct the string and assign it to `out` (line 4).

In the background, Java performs the following:

```
"Details: " + "James" => "Details: James"
"Details: James" + " " => "Details: James, "
"Details: James, " + "23" => "Details: James, 23"
"Details: James, 23" + " " => "Details: James, 23,
"Details: James, 23, " + "50000.0" => "Details: James, 23,
50000.0"
```

So, "Details: James, 23, 50000.0" is used to initialize `out`, which is what is displayed on the screen when executing line 5.

## Exercises

All is going great in our lovely dinosaur park. However, we do need to do some administrative tasks:

1. We need to keep track of the dinosaurs in the park. Declare variables to represent the breed, height, length, and weight of one dinosaur in the main method. Give the variables a value and print them.



2. Now, we want to do something similar to the program of exercise 1 and print the dinosaur's age, name, and whether it's a carnivore or not. This needs to happen in the main method. Give the variables a value and print them.
3. Our park is doing great! But it gets a bit too busy at times. The fire department advised us to introduce a maximum number of visitors that are allowed at any given time. Declare a variable to represent the maximum number of visitors allowed in the park per day. You can choose a reasonable value for the variable. Then, print it in the sentence: "There's a maximum of [x] people allowed in Mesozoic Eden."
4. Our team is an integral part of Mesozoic Eden. Let's create a profile for an employee. Declare variables to represent the name and age of a Mesozoic Eden employee. Assign values and print them.
5. We would like to know how many dinosaurs we have at any time. Declare a variable to represent the number of dinosaurs in the park. Assign it a value and print it.
6. Safety is our priority. We maintain a safety rating scale to ensure our standards. Declare a variable to represent the park's safety rating on a scale from 1 to 10. Assign a value to it and print it.
7. Now, let's bring together some dinosaur information in one statement. Create a program that uses string concatenation to print out a dinosaur's name, age, and diet (a string with a value of carnivore or herbivore).
8. Each dinosaur species has a unique name. For a quick referencing system, we use the first letter of a dinosaur species. Declare a character variable that represents the first letter of a dinosaur species, assign a value, and print it.

## Project – dinosaur profile generator

As part of your responsibilities in Mesozoic Eden, you are tasked with creating an extensive database of all the dinosaurs living in the park. For now, you only need to complete the first step: a profile generator. These profiles will not only help in keeping track of our prehistoric residents but also provide essential data for scientific study, healthcare, diet management, and visitor engagement.

In this project, we will focus on developing a program that can model an individual dinosaur's profile.

The profile should include the following characteristics:

- Name
- Age
- Species
- Diet (carnivore or herbivore)
- Weight

---

Each characteristic should be stored as a variable within the program. Here's your chance to get creative and think about the kind of dinosaur you want to describe. Is it a towering T-Rex or a friendly Stegosaurus? Maybe it's a swift, scary Velociraptor or a mighty Triceratops?

Once you have declared and assigned values to these variables, the program should print out a complete profile of the dinosaur. The output can be something like "Meet [Name], a [Age]-year-old [Species]. As a [Diet], it has a robust weight of [Weight] kilograms.".

## Summary

In this chapter, we learned that a variable is simply a memory location with a name and a value. To utilize variables, we have to know how to declare and access them.

To declare a variable, we specify the variable's name and its type – for example, `int countOfTitles=5;`. This line of code declares an `int` variable named `countOfTitles` with a value of 5. Naming them properly using camel case, is a great aid in making your code more readable and maintainable. To access the variable, we just specify the variable's name – for example, `System.out.println(countOfTitles);`.

As Java is a strongly typed language, we have to specify a variables' type when we declare it. Java provides eight in-built primitive data types for our use. They are easily recognizable due to their lowercase letters. In the preceding line of code, `int` is the primitive data type for the `countOfTitles` variable. We saw the sizes in bytes of the primitive types, which determines their range of values. All numeric types are signed, with the most significant bit being used for the sign. The `char` type is unsigned and is 2 bytes in size so that Java can support any character in any language anywhere in the world. Using code snippets, we saw variables of the different types in use.

Now that we know how to declare and use variables, let's move on to operators that enable us to combine variables.

Trial Version



Wondershare  
**PDFelement**

## 3

# Operators and Casting

In *Chapter 2*, we learned that variables are simply named pigeonholes and contain values. These values vary and Java provides eight primitive data types accordingly. These primitive types cater for whole numbers (`byte`, `char`, `short`, `int`, and `long`), decimal numbers (`float` and `double`), and the literals `true` and `false` (`boolean`).

We also learned how to declare a variable. As Java is a strongly typed language, this means you must give every variable a data type immediately upon declaration. This is where primitive data types are very useful.

Now that we know how to declare variables, let's do something interesting with them. By the end of this chapter, you will be able to combine variables using Java's various operators. In addition, you will understand Java casting, including what it is, and when and why it occurs.

In this chapter, we are going to cover the following main topics:

- Learning how Java's operators cooperate
- Understanding Java's operators
- Explaining Java casting

## Technical requirements

The code for this chapter can be found on GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects/tree/main/ch3>.

## Learning how Java's operators cooperate

Java provides numerous operators for us to work with. By way of definition, if we have an expression `3 + 4`, the `+` is the *operator*, whereas `3` and `4` are the *operands*. Since `+` has *two* operands, it is known as a *binary* operator.



Before we discuss the operators themselves, we must first discuss two important features relating to Java operators, namely **order of precedence** and **associativity**.

## Order of precedence

Order of precedence specifies how operands are grouped with operators. This becomes important when you have shared operands in a complex expression. In the following code segment, we have an expression of `2 + 3 * 4`, where `*` represents multiplication and `+` represents addition:

```
int a = 2 + 3 * 4;  
System.out.println(a);
```

In the preceding code, `3` is shared by both `2` and `4`. So, the question arises, do we group `3` with `2`, where the expression is  $(2 + 3) * 4$ , giving us `20`; or do we group `3` with `4`, where the expression is  $2 + (3 * 4)$ , giving us `14`? This is where the order of precedence applies. As `*` has higher precedence than `+`, `3` is grouped with `4` and therefore the expression evaluates to  $2 + (3 * 4)$ . Note that the evaluation order is still left to right; it is just that `3` is grouped with `4` rather than with `2`.

### Parentheses in an expression

Note that parentheses can change the default order of operator precedence. As we have seen, the default order of precedence, where `*` has higher precedence than `+`, means that  $2 + 3 * 4$  is `14`. This is the same as  $2 + (3 * 4)$ .

However,  $(2 + 3) * 4$  is `20`. In this case, the parentheses grouped `3` with `2`, so the expression evaluated to  $5 * 4 = 20$ .

This begs the question, what if you are evaluating an expression that contains operators at the same level of precedence? This is where associativity applies.

## Associativity

When an expression has two operators with the same level of precedence, operator associativity determines the groupings of operators and operands. For example, in the following code segment, we are evaluating a simple expression involving two divisions (which have the same level of precedence):

```
int x = 72 / 6 / 3;
```

As division associates left to right, `6` will be grouped with `72` and not `3`. Thus, the expression is the same as  $(72 / 6) / 3$ , which evaluates to  $12 / 3 = 4$ . Parentheses can also be used to change the default associativity order. Take, for example, the following code:

```
int x = 72 / (6 / 3);
```

In this case, `6` is now grouped with `3` and the expression evaluates to  $72 / 2 = 36$ .

*Table 3.1* outlines the order of precedence and associativity rules:

Operator Description	Operator	Associativity
Parentheses	()	left-to-right
Post-increment, post-decrement	x++, y--	left-to-right
Pre-increment, pre-decrement, negation	++x, --y, !b	right-to-left
Cast	(int)	right-to-left
Multiplication/Division/Modulus	*, /, %	left-to-right
Addition/Subtraction	+, -	left-to-right
Relational Operators	<, >, <=, >=, instanceof	left-to-right
Equality	==, !=	left-to-right
Bitwise Operators	&	left-to-right
	^	
Logical Operators	&&	left-to-right
Ternary Operator	? :	right-to-left
Assignment Operators	=, +=, -=, *=, /=	right-to-left

Table 3.1 – Order of precedence and associativity rules

Note that *Table 3.1* is simplified in that it refers to the operators that are commonly used. For example, the unsigned right shift operator, `>>>`, is omitted as it is rarely used. Also, note that the `instanceof` operator will be discussed in *Chapter 8*.

It is interesting to note that the assignment operator, namely `=`, is at the bottom of the precedence table. This means that regardless of the expression on the right-hand side of the assignment, the assignment will always be done last. This makes sense. Also, while most of the operators associate left to right, the assignment associates right to left. This is demonstrated in the following code segment:

```
boolean b1 = false;
boolean b2;
boolean b3;
b3 = b2 = b1;
System.out.println(b1);
System.out.println(b2);
System.out.println(b3);
```

The preceding code segment outputs `false` three times. The crucial line is `b3 = b2 = b1;`. Since the assignment associates right to left, the value in `b1`, which is `false`, is assigned to `b2`; then, the value in `b2`, which is now `false`, is assigned to `b3`.

Now that we understand these properties, let's examine the operators themselves.

## Understanding Java's operators

Operators can be grouped into the following categories:

- Unary operators
- Arithmetic operators
- Relational operators
- Logical operators
- Ternary operator
- Compound assignment operators

We will now discuss each category in turn.

### Unary operators

Unary operators have only one operand, hence the term *unary*. Let's examine them.

#### *Prefix and postfix unary operators*

`++` and `--` denote these operators and they increment and decrement by 1, respectively. If the operator appears before the variable, it is known as *prefix*, while if the operator appears after the variable, it is called *postfix*. For example, `++x` is prefix increment, whereas `y--` is postfix decrement.

Depending on whether `++` or `--` appears before or after the variable can, in some situations, affect the overall expression. This is best explained with a code sample, as shown in *Figure 3.1*:

The screenshot shows a code editor with two columns. The left column contains line numbers from 25 to 35. The right column contains Java code with annotations for output values. The code initializes `x=3` and `y=4`. It then performs three increments on `x` (prefix, then two postfix) and three decrements on `y` (postfix, then two prefix). The output values are indicated in the code itself.

25	<code>int x=3;</code>
26	<code>++x;</code>
27	<code>System.out.println(x); // 4</code>
28	<code>System.out.println(x++); // 4</code>
29	<code>System.out.println(x); // 5</code>
30	
31	<code>int y=4;</code>
32	<code>y--;</code>
33	<code>System.out.println(y); // 3</code>
34	<code>System.out.println(--y); // 2</code>
35	<code>System.out.println(y); // 2</code>

Figure 3.1 – Prefix and postfix increment and decrement operators

In *Figure 3.1*, on line 25, we can see that `x` is initialized to 3. On line 26, `x` is incremented by 1 to 4. Line 26 is a simple statement and because of that, whether it is prefix or postfix notation does not matter. Line 27 outputs the value of `x` to show that it is 4 at this point.

Line 28 is where things get interesting. The postfix notation on line 28 has a real effect on the screen output. As it is the postfix notation in the `System.out.println` command, the current value of `x` is output, and *afterwards*, `x` is incremented by 1. So, the output to the screen is 4, and afterwards, `x` is incremented to 5. Line 29 demonstrates that `x` is 5.

On line 31, `y` is initialized to 4. On line 32, `y` is decremented by 1 to 3. Again, as line 32 is a simple statement, prefix or postfix notation makes no difference. Line 33 outputs the value of `y` to show that it is 3 at this point.

The prefix notation on line 34 has no real effect on the screen output. As it is the prefix notation in the `System.out.println` command, the current value of `y` is decremented **before** being output. Thus, the value of `y` and the output to the screen match (both are 2). Lastly, line 35 demonstrates that the current value of `y` is 2.

### **Unary plus/minus operators**

Now that we have discussed the prefix and postfix operators, let's discuss other unary operators. The code in *Figure 3.2* will help:

```
37 int x = +6;
38 int y = -x;
39 System.out.println(x); // 6
40 System.out.println(y); // -6
41
42 int z = (int)3.45;
43 System.out.println(z); // 3
44
45 boolean b = true;
46 System.out.println(!b); // false
47 System.out.println(b); // true
```

Figure 3.2 – Other unary operators

In *Figure 3.2*, line 37 uses the unary plus sign, `+`, to initialize `x` to 6. Here, `+` is the default as numbers without a sign are assumed to be positive numbers. Line 38 uses the unary minus sign, `-`, to initialize `y` to be the negative of `x`. Lines 39 and 40 demonstrate that `x` and `y` are 6 and `-6`, respectively.

### ***Cast operator***

In *Figure 3.2*, line 42 uses the cast operator. We will discuss casting in greater detail later in this chapter. For now, `3 . 45` is a `double` literal (8 bytes) and cannot be stored in an `int` variable, `z`, as `int` variables are 4 bytes in size. The compiler spots this and generates an error. To get around this error, we can use a cast, which takes the form of (*cast type*). This cast enables us to override the compiler error. In this case, we are casting `3 . 45` to an `int` variable, which means we lose the decimal places. Thus, we store `3` in `z`, as shown by the output from line 43.

### ***Logical NOT operator***

In *Figure 3.2*, on line 45, we declare a boolean, `b`, and initialize it to `true`. On line 46, we output the inverted value of `b` by using the logical NOT operator. Note that we are not changing the value of `b` meaning, the value of `b` is still `true`. This is shown by the output from line 47.

Now, let's examine the arithmetic operators.

## **Arithmetic operators**

There are five arithmetic operators, all of which we will examine now.

### ***Addition/subtraction operators***

As in mathematics, the `+` operator represents addition and the `-` operator represents subtraction. Both are binary operators; in other words, there are two operands, one on either side of the operator. The following code example demonstrates this:

```
int res = 6 + 4 - 2;  
System.out.println(res); // 8
```

In this code segment, `res` has been assigned a value of `6 + 4 - 2`, which is `8`.

### ***Multiplication/division operators***

The `*` operator represents multiplication and the `/` operator represents division. Both are binary operators. Note that *integer division truncates*. The following code segment demonstrates this:

```
System.out.println(10/3); // 3
```

This code segment outputs `3` because integer division truncates. We are dividing one integer, `10`, by another integer, `3`. The remainder is simply discarded.

## Modulus operator

The % operator is used for calculating the modulus (remainder). The following code example demonstrates the modulus operator in action:

```
int mod1 = 10 % 3;  
System.out.println(mod1); // 1  
int mod2 = 0 % 3;  
System.out.println(mod2); // 0
```

The first line initializes mod1 to be the remainder of 10 divided by 3, which is 1. In other words, 3 goes into 10 three times and 1 is left over. Therefore, 1 is assigned to mod1.

The initialization of mod2 is interesting: 3 goes into 0 zero times and there is 0 (or nothing) left over. Hence, 0 is assigned to mod2.

## The precedence of arithmetic operators

As per *Table 3.1*, \*, /, and % have higher precedence than the + and - operators, and assignment has the lowest precedence. *Figure 3.3* shows how this affects the evaluation of expressions in code:

61	<code>int res = 3 + 2 * 4;</code>
62	<code>System.out.println(res); // 11</code>
63	<code>res = (3 + 2) * 4;</code>
64	<code>System.out.println(res); // 20</code>
65	<code>res = 6 + 4 - 2;</code>
66	<code>System.out.println(res); // 8</code>
67	<code>res = 10 / 4 * 6 % 10;</code>
68	<code>System.out.println(res); // 2</code>

Figure 3.3 – Arithmetic operators precedence

Lines 61 demonstrates that \* has higher precedence than +, in that the expression evaluates to  $3 + (2 * 4) = 3 + 8 = 11$ .

Line 63 demonstrates that parentheses change the grouping. Now, the shared value, 2, is grouped with 3 (as opposed to 4, which was the case on line 61). The expression now evaluates to  $5 * 4 = 20$ .

Line 65 demonstrates that + and - associate left to right. The expression evaluates to

$$10 - 2 = 8.$$

Lines 67 demonstrates that \*, /, and % also associate left to right. The expression evaluates to  $2 * 6 \% 10$ , which, in turn, evaluates to  $12 \% 10$ , which is 2.



### Math operations involving int variables or smaller result in an int

It is interesting to note that any math operation involving an `int` type or smaller results in `int`. This is demonstrated in the following code segment:

```
byte b1=2, b2=3;  
byte b3 = b1 + b2; // compiler error  
byte b4 = (byte)(b1 + b2); // Ok
```

The first line declares 2 bytes, namely `b1` and `b2`. Notice that, even though 2 and 3 are integer literals, the compiler is aware that these values are within the range of `byte` (-128 to +127) and, consequently, allows the declarations.

However, the next line is a problem. The compiler has a rule that all math operations involving `int` types or smaller result in an `int`. Therefore, even though the sum of the two bytes, 5, is well within the `byte` range, the compiler complains saying “possible loss of data converting from `int` to `byte`”.

The last line fixes this issue by casting the result of the addition (an `int` type) to a `byte` before the assignment. What this means is that the extra 3 bytes from the `int` (that do not fit into the `byte`) are simply discarded. Thus, the sum of `b1 + b2` is cast from `int` to `byte` and the resultant `byte` is assigned to `b4`. Casting is discussed in more detail later in the chapter.

We will finish our discussion on arithmetic operators by examining `+` in a different context.

### String append

As we have seen, Java uses `+` for mathematical addition. However, this occurs only if both operands are numbers. For example, `3 + 4` results in 7 because both operands, 3 and 4, are numbers.

However, if either operand (or both) are strings, Java performs a `String` append. A `String` literal is enclosed in double quotes – for example, `"abc"`, `"123"`, `"Sean"`, and `"Maaike"` are all `String` literals. So, just to be clear on what operation is performed and when, let's take a look at some examples:

- `3 + 4` is mathematical addition. Thus, the result is 7.
- `"3" + 4` is a string append as there is a string on the left of `+`. The result is the string “34.”
- `3 + "4"` is a string append as there is a string on the right of `+`. Again, the result is the string “34.”
- `"3" + "4"` is a string append as there is a string on both sides of `+`. The result is also the string “34.”

So, what exactly happens during a string append? *Java cannot perform any mathematical operations when the operands are of different types.* Let's examine this with an example piece of code:

```
String s = "3" + 4;  
System.out.println(s); // "34"
```

The first thing to note is that the first line of code only compiles because "3" + 4 results in a `String` literal. When Java encounters a string on the left/right/both sides of +, it performs string concatenation (append). Essentially, as + associates left to right, Java appends (adds) the string on the right of + to the end of the string on the left of +.

In this example, Java sees the `String` literal "3" and the + operator and realizes it must perform a `String` append. To do this, in memory, it creates a string version of 4 – in other words, "4". The integer 4 literal is not touched. Thus, a new variable is created under the hood – it is a `String` variable, and "4" is its value. The expression is now "3" + "4". As both operands on either side of + are now of the same type (both are strings), Java can perform the append. The new string is the result of "3" + "4", which is "34". This is what is assigned to `s`. The second line demonstrates this by outputting "34" for `s`.

In *Figure 3.4*, a more substantial example is presented:

```
78 int a=3, b=2;
79 int res = a + b;
80 System.out.println(res); // 5
81 String s = "abc";
82 String s1 = a + s; // "3abc"
83 String s2 = s + a; // "abc3"
84 System.out.println(s1 + " " + s2); // "3abc abc3"
85
86 System.out.println("Output is "+ a + b); // "Output is 32"
87 System.out.println("Output is "+ (a + b)); // "Output is 5"
```

Figure 3.4 – String append in action

On line 79, as both operands, `a` and `b`, are integers, Java initializes `res` to 5 (the sum of 3 and 2).

Line 82 is evaluated as follows: 3 + "abc" = "3" + "abc" = "3abc". In other words, Java realizes that it must do a string append due to the presence of "abc" on the right-hand side of +. Thus, somewhere in memory, a string version of the value of `a` is created. In other words, a variable with "3" is created. Note that `a` remains an `int` with 3 as its value. Now, Java can proceed since both operands are the same type (strings): "3" + "abc" results in "3abc".

Line 83 demonstrates that it does not matter which side of + the string is on. Plus, it does not matter if the string is a string literal or a string variable. The expression on line 83 is evaluated as follows: "abc" + 3 = "abc" + "3" = "abc3". This is what `s2` is initialized to. Line 84 outputs the values of both `s1` and `s2` with a space between them. Note that `System.out.println` expects a string. The string output on line 84 is constructed as follows: "3abc" + " " = "3abc " + "abc3" = "3abc abc3".

Lines 86 and 87 require special mention. The problem with line 86 is that the output string is constructed as follows: "Output is "+ 3 = "Output is " + "3" = "Output is 3" + 2 = "Output is 3" + "2" = "Output is 32". This is not what we wanted.

Line 87 rectifies this by using parentheses to ensure that a + b is grouped. Thus, the string is constructed as follows: "Output is "+ 5 = "Output is "+ "5" = "Output is 5".

That finishes the arithmetic operators. We will now examine relational operators.

## Relational operators

Java has six relational operators, all of which result in a boolean value of `true` or `false`. They are as follows:

- `==` is the equivalence operator
- `!=` is the not equivalent operator
- `>` is the greater than operator
- `>=` is the greater than or equal to operator
- `<` is the less than operator
- `<=` is the less than or equal to operator

*Figure 3.5* shows the relational operators in action in code:



```
89 int x=3, y=4;
90 System.out.println(x == y); // false
91 System.out.println(x != y); // true
92 System.out.println(x > y); // false
93 System.out.println(x >= y); // false
94 System.out.println(x < y); // true
95 System.out.println(x <= y); // true
```

Figure 3.5 – Relational operators in code

Line 89 declares two `int` variables, namely `x` and `y`, and initializes them to 3 and 4, respectively. Line 90 uses Java's equivalence operator, `==`, to check if `x` and `y` are equivalent. As they are not, line 90 outputs `false`. Line 91 checks the exact opposite. As `x` is not equivalent to `y`, line 91 outputs `true`.

Line 92 outputs whether `x` is greater than `y`. This is, of course, `false` as 3 is not greater than 4. Similarly, line 93 outputs whether `x` is greater than or equal to `y`. Again, this is `false`.

Line 94 outputs whether `x` is less than `y`. This is `true` as 3 is less than 4. Line 95 outputs whether `x` is less than or equal to `y`. Again, this is `true`.

The relational operators and their boolean return values are going to be extremely useful going forward, particularly when we look at conditional statements in *Chapter 4*.

### Implicit promotion

While Java's operators do not require the operands to be exactly the same type, the operands must be compatible. Consider the following code snippet:

```
System.out.println(3 + 4.0); // 7.0
System.out.println(4 == 4.0); // true
```

The first line tries to add an `int` variable of 3 to a `double` variable. Java realizes that the types are not the same. However, Java can figure out a safe solution without bothering us. This is where *implicit promotion* comes in. `int` requires 4 bytes of storage whereas `double` requires 8 bytes. In the background, somewhere in memory, Java declares a temporary `double` variable and promotes `int 3` to `double 3.0`, and stores `3.0` in this temporary location. Now, Java can add `3.0` to `4.0` (as both are `doubles`), resulting in the answer `7.0`.

The second line compares `int 4` with `double 4.0`. The same process happens. Java implicitly promotes `4` to `4.0` (in a new temporary location) and then compares `4.0` with `4.0`. This results in `true` being output.

Now, we will turn our attention to logical operators.

## Logical operators

Logical operators enable us to build complex `boolean` expressions by combining sub-expressions. These operators are as follows:

- `&&` is the logical AND
- `||` is the logical OR
- `&` is the bitwise AND
- `|` is the bitwise OR
- `^` is the bitwise eXclusive OR (XOR)



We will examine these in turn with code examples to help explain how they operate. But before we do that, it is worthwhile refreshing our truth tables, as shown in *Table 3.2*:

P	Q	P && Q	P    Q	P & Q	P   Q	P ^ Q
T	T	T	T	T	T	F
T	F	F	T	F	T	T
F	T	F	T	F	T	T
F	F	F	F	F	F	F

Table 3.2 – Boolean truth tables

In *Table 3.2*, the first two columns, P and Q, represent two expressions, where T means true and F means false. For example, the logical AND column (the P && Q column) represents the result of the overall expression, P && Q, depending on the values of P and Q. So, if P is true and Q is T, then P && Q is also true.

With this table in mind, let's examine the operators in turn.

### Logical AND (&&)

The logical AND states that both boolean operands must be true for the overall expression to be true. This is represented by the P && Q column in *Table 3.2*.

Note that this operator is known as a short-circuiting operator. For example, in an expression P && Q, if P evaluates to false, then && will *not* evaluate the expression Q because the overall expression will evaluate to false regardless. This is because F && F is false and F && T is also false. In effect, Java knows that once the expression P is false on the left-hand side of an && expression, the overall expression must be false. So, there is no need to evaluate the Q expression on the right-hand side, so it *short-circuits*. This is better explained with a code example:

```
boolean b1 = false, b2 = true;
boolean res = b1 && (b2=false); // F &&
System.out.println(b1 + " " + b2 + " " + res); // false true
false
```

The first line initializes two boolean variables, `b1` and `b2`, to `false` and `true`, respectively. The second line is the important one. Note that the parentheses are required around the `b2=false` sub-expression to get the code to compile (otherwise, you will get a syntax error). So, when we plug in `false` for `b1`, the expression evaluates to `F && (b2=false)`. As the evaluation order is left to right, this will lead `&&` to short-circuit, because, regardless of what remains in the expression, there is no way the overall expression can evaluate to true. This means that the `(b2=false)` sub-expression is **not** executed.

The last line outputs the values of the variables. The output is `false`, `true`, and `false` for `b1`, `b2`, and `res`, respectively. Crucially, `b2` is `true`, demonstrating that `&&` short-circuited.

### ***Logical OR (||)***

The logical OR states that either or both boolean operands can be true for the overall expression to be true. This is represented by the `P || Q` column in *Table 3.2*.

This operator is also a short-circuiting operator. For example, in an expression `P || Q`, if `P` evaluates to true, then `||` will *not* evaluate the expression `Q` because the overall expression will evaluate to true regardless. This is because `T || F` is true and `T || T` is also true. In effect, Java knows that once the expression `P` is true on the left-hand side of an `||` expression, the overall expression must be true. So, there is no need to evaluate the expression, `Q`, and hence it *short-circuits*. Again a code example will help:

```
boolean b1=false, b2=true;
boolean res = b2 || (b1=true); // T ||
System.out.println(b1 + " " + b2 + " "+res); // false true
true
```

The first line initializes two boolean variables, `b1` and `b2`, to `false` and `true`, respectively. The second line is the important one. Note again that the parentheses are required around the `b1=true` sub-expression to get the code to compile. So, when we plug in `true` for `b2`, the expression evaluates to `T || (b1=true)`. As the evaluation order is left to right, this will lead `||` to short-circuit because, regardless of what remains in the expression, there is no way the overall expression can evaluate to `false`.

The last line outputs the values of the variables. The output is `false`, `true`, and `true` for `b1`, `b2`, and `res`, respectively. Crucially, `b1` is `false`, demonstrating that `||` short-circuited.



### Order of evaluation versus precedence

This topic often causes confusion and is best explained with some sample pieces of code. Let's start with an example that can be deceptively simple:

```
int x=2, y=3, z=4;  
int res = x + y * z;      // x + (y * z)  
System.out.println(res); // 14
```

As `*` has higher precedence than `+`, the common element `y`, is grouped with `z` and not `x`. Thus, the overall expression is  $x + (y * z) = 2 + 12 = 14$ .

What is important to note here is that the evaluation order is left to right and as evaluation order trumps precedence, `x` is evaluated first before the `(y * z)` sub-expression. While this makes no difference in this example, let's look at an example where it does make a difference:

```
boolean a=false, b=false, c=false;  
// a || (b && c)  
// The next line evaluates to T ||  
boolean bool = (a = true) || (b = true) && (c = true);  
System.out.print(a + ", " + b + ", " + c); // true, false, false  
As && has higher precedence than ||, the expression evaluates to (a = true) || (b = true) && (c = true).
```

In other words, the common sub-expression `(b = true)` is grouped with `(c = true)` rather than `(a = true)`. Now comes the crucial bit: *evaluation order trumps precedence*. Therefore, `(a = true)` is evaluated first, resulting in `T || ((b = true) && (c = true))`.

As `||` is a short-circuit operator, the rest of the expression (to the right of `||`) is **not** evaluated. This is demonstrated by the output on the last line, where it outputs `true, false, false`, for `a`, `b`, and `c`, respectively. The crucial thing to note here is that `b` and `c` are still `false`!

Now that we have discussed the logical operators, we will move on to bitwise operators.

## Bitwise operators

Although some of the bitwise operators look very similar to the logical operators, they operate quite differently. The principle differences are that the bitwise operators can work with both boolean and integral (`byte`, `short`, `int`, `long`, and `char`) operands. In addition, bitwise operators do *not* short-circuit.

Let's examine the boolean bitwise operators first.

### Bitwise AND (&)

Comparing the bitwise AND (`&`) with the logical AND (`&&`), the difference is that the bitwise AND will *not* short-circuit. This is represented by the **P & Q** column in *Table 3.2*. If we take the sample code that we used for the logical AND but change it to use the bitwise AND operator, you will see the difference in the output:

```
boolean b1 = false, b2 = true;
boolean res = b1 & (b2=false); // F & F
System.out.println(b1 + " " + b2 + " " + res); // false
false false
```

In this case, the `(b2=false)` sub-expression is executed because `&` did not short-circuit. So we had `false & false`, which is `false`. Thus, the output is `false` for all the variables.

### Bitwise OR (|)

Comparing the bitwise OR (`|`) with the logical OR (`||`), the difference is that the bitwise OR will *not* short-circuit. This is represented by the **P | Q** column in *Table 3.2*. If we take the sample code that we used for the logical OR but change it to use the bitwise OR operator, you will see the difference in the output:

```
boolean b1=false, b2=true;
boolean res = b2 | (b1=true); // T | T
System.out.println(b1 + " " + b2 + " "+res); // true true
true
```

In this case, the `(b1=true)` sub-expression is executed because `|` did not short-circuit. So, we had: `true | true`, which is `true`. Thus, the output is `true` for all the variables.

### Bitwise XOR (^)

This is another non-short-circuiting operator. The bitwise XOR, represented by the `^` operator, evaluates to `true`, if and only if one of the operands is `true` but *not* both. This is represented by the **P ^ Q** column in *Table 3.2*. Let's look at some examples in terms of code:

```
boolean b1 = (5 > 1) ^ (10 < 20); // T ^ T == F
boolean b2 = (5 > 10) ^ (10 < 20); // F ^ T == T
boolean b3 = (5 > 1) ^ (10 < 2); // T ^ F == T
boolean b4 = (5 > 10) ^ (10 < 2); // F ^ F == F
// false true true false
System.out.println(b1 + " " + b2 + " " + b3 + " " + b4);
```



The boolean variable, `b1`, is initialized to `false` because both of the sub-expressions – `(5 > 1)` and `(10 < 20)` – are `true`. Similarly, `b4` is also initialized to `false` because both `(5 > 10)` and `(10 < 2)` are `false`.

However, `b2` is `true` because even though `(5 > 10)` is `false`, `(10 < 20)` is `true`, and `F ^ T` is `true`. Likewise, `b3` is `true` because `(5 > 1)` is `true`, `(10 < 2)` is `F`, and `T ^ F` is `true`.

Now that we have examined the bitwise operators when used with boolean operands, we will now briefly examine how the same operators work when the operands are integral numbers.

### ***Bitwise operators (integral operands)***

Though not commonly used, we have included them for completeness. A code example is useful here:

```
byte b1 = 6 & 8;           // both bits must be 1
byte b2 = 7 | 9;           // one or the other or both
byte b3 = 5 ^ 4;           // one or the other but not both
System.out.println(b1 + ", " + b2 + ", " + b3); // 0, 15, 1
```

When the operands are integrals (as opposed to booleans), the bit patterns become important in evaluating the result. For the `&` operator, both bits must be 1 for that bit to be 1 in the result:

```
6 & 8 (in binary) = 0110 & 1000 = 0000 = 0
```

For the `|` operator, one of the bits, or both, must be 1 for that bit to be 1 in the result:

```
7 | 9 (in binary) = 0111 | 1001 = 1111 = 15
```

For the `^` operator, one of the bits, but not both, must be 1 for that bit to be 1 in the result:

```
5 ^ 4 (in binary) = 0101 ^ 0100 = 0001 = 1
```

That completes the bitwise operators. Now, let's cover the ternary operator.

### **Ternary operator**

The ternary operator, as its name suggests, is an operator that takes three operands. The ternary operator is used to evaluate boolean expressions and assign values accordingly to a variable. In other words, as boolean expressions evaluate to `true` or `false` only, the goal of the ternary operator is to decide which of the two values to assign to the variable.

The syntax is of the following form:

```
variable = boolean expression ? value to assign if true :
                           value to assign if false
```

Let's look at an example:

```
int x = 4;
String s = x % 2 == 0 ? " is an even number" : " is an odd
number";
System.out.println(x + s); // 4 is an even number
```

In this example, the boolean expression to be evaluated is `x % 2 == 0`, which, because `x = 4`, evaluates to `true`. Thus, `is an even number` is assigned to the string, `s`, and is output. Had `x` been 5, then the boolean expression would have been `false`, and therefore, `is an odd number` would have been assigned to `s` and output.

The last group of operators we will examine are compound assignment operators.

## Compound assignment operators

These operators exist as a shorthand for more verbose expressions. For example, assuming `x` and `y` are both integers, `x = x + y` can be written as `x += y`. There are compound assignment operators for all of the mathematical operators:

- `+=` Example: `x += y` is the same as `x = x + y`
- `-=` Example: `x -= y` is the same as `x = x - y`
- `*=` Example: `x *= y` is the same as `x = x * y`
- `/=` Example: `x /= y` is the same as `x = x / y`
- `%=` Example: `x %= y` is the same as `x = x % y`

Indeed, there are compound assignment operators for the bitwise operators – for example, `x &= 3` is the same as `x = x & 3` but they are so rarely used that we will just mention that they exist.

There are one or two subtleties to be aware of. As mentioned earlier, any mathematical operation involving an `int` type or smaller results in `int`. This can result in a cast being required to get the code to compile. With the compound assignment operators, the cast is in-built, so the explicit cast is not required. Take the following code for example:

```
byte b1 = 3, b2 = 4;
// b1 = b1 + b2;           // compiler error
b1 = (byte) (b1 + b2);    // ok
b1 += b2;                 // ok, no cast required
```

The first line initializes 2 bytes, `b1` and `b2`, to 3 and 4, respectively. The second line is commented out as it generates a compiler error. The addition of `b1` and `b2` results in an `int` type that cannot be directly assigned to a `byte` variable, unless you cast it down from `int` to `byte`. This is what the third line is doing – using the cast (`byte`) to override the compiler error. We'll cover casting very soon



but for now, just realize that with the cast, you are overriding the compiler error, effectively saying “I know what I am doing, proceed.”

The last line is interesting in that, in the background, it is the same as the third line. In other words, the compiler translates `b1 += b2` into `b1 = (byte) (b1 + b2)`.

Another subtlety to be aware of is that whatever is on the right-hand side of the compound assignment operator is going to be grouped, regardless of precedence. An example will help here. Consider the following:

```
int x = 2;  
x *= 2 + 5; // x = x * (2 + 5) = 2 * 7 = 14  
System.out.println(x); // 14
```

We know that `*` has higher precedence than `+` and that the order of evaluation is left to right. That said, what is on the right-hand side of `=` is grouped by the compiler by surrounding `2 + 5` with parentheses (in the background). Thus, the expression becomes  $2 * (2 + 5) = 2 * 7 = 14$ . To further this point, had the compiler *not* inserted parentheses, the expression would have been evaluated to 9. In other words, due to operator precedence, the expression would have been  $(2 * 2) + 5 = 4 + 5 = 9$ . However, as we have seen, this is **not** the case.

Let's look at another more complicated example:

```
int k=1;  
k += (k=4) * (k+2);  
System.out.println(k); // 25
```

In this example, the right-hand side is, once again, enclosed in parentheses:

```
k += (right hand side) where the right hand side is (k=4) *  
(k+2)
```

Translating `+=` into its longer form gives us the following output:

```
k = k + (right hand side)
```

The order of evaluation is left to right, so plugging in the current value of `k`, which is 1, results in:

```
k = 1 + (right hand side)
```

Now, by plugging in the right-hand side expression, we get the following:

```
k = 1 + ( (k=4) * (k+2) )
```

As the order of evaluation is left to right, k is changed to 4 before we add 2:

```
k = 1 + ( 4 * 6 )
k = 1 + 24
k = 25
```

That concludes our treatment of Java operators. Now, let's examine Java casting, a topic we have touched on already in this chapter.

## Explaining Java casting

To discuss casting properly, we need to explain both the widening and narrowing of Java's primitive data types. With this in mind, it is helpful to remember the sizes of the primitive data types in bytes. *Table 3.3* represents this information:

Primitive Data Type	Size in Bytes
byte	1
short	2
char	2
int	4
long	8
float	4
double	8

Table 3.3 – The sizes of Java's primitive types

The preceding table presents the sizes in bytes of Java's various primitive data types. This will help us as we discuss both widening and narrowing.

## Widening

Widening is done automatically; in other words, a cast is not needed. As the promotion is done in the background, widening is also known as *implicit promotion*. With *Table 3.3* in mind, the widening rules are as follows:

byte → short/char → int → long → float → double



Given the sizes from *Table 3.3*, most of these rules should make sense. For example, a `byte` can automatically fit into a `short` because 1 byte fits into 2 bytes automatically. The only interesting one is `long` → `float`, which is *widening* from 8 bytes to 4 bytes. This is possible because even though a `long` requires 8 bytes and a `float` requires only 4 bytes, their ranges differ – that is, a `float` type can accommodate any `long` value but not vice versa. This is shown in the following code snippet:

```
System.out.println("Float: " + Float.MAX_VALUE); // Float:  
3.4028235E38  
System.out.println("Float: " + Float.MIN_VALUE); // Float:  
1.4E-45  
System.out.println(Long.MAX_VALUE); // 9223372036854775807  
System.out.println(Long.MIN_VALUE); // -9223372036854775808
```

Note the scientific notation E used for floating point. `float` takes up less space, but due to its representation, it can hold larger and smaller numbers than `long`.

### Scientific notation

Scientific notation is a shorthand way to represent decimal numbers and can be useful for representing very large and/or very small numbers. Here are some examples:

```
double d1 = .00000000123;  
double d2 = 1.23e-9;  
System.out.println(d1==d2); // true  
double d3 = 120_000_000;  
double d4 = 1.2e+8;  
System.out.println(d3==d4); // true
```

As the comparisons both return `true`, this means that `d1` is the internal representation of `d2`. Similarly, both `d3` and `d4` are equivalent.

Let's examine widening in code. *Figure 3.6* demonstrates this:

14	<code>char c = 'a'; // normal</code>
15	<code>int i = c; // widening, char to int</code>
16	<code>float f = i; // widening, int to float</code>
17	<code>double d = f; // widening, float to double</code>
18	<code>float f2 = 1L; // widening, long to float</code>

Figure 3.6 – Implicit widening examples

Line 14 is a regular `char` assignment – in other words, no widening is involved. Note that characters (represented by `char`) are simply small numbers (0..65,535). To represent a character, we enclose the character in single quotes. In contrast, a `String`, which is a sequence of characters, is represented in double quotes. Therefore, "`a`" is a `String`, whereas '`a`' is a character.

Line 15 is a widening from `char` (2 bytes) to `int` (4 bytes). Line 16 is a widening from `int` to `float`. Although both `int` and `float` require 4 bytes, as discussed earlier with `long`, `float` has a greater range, so there is no issue here. Line 17 is a widening from `float` to `double`. Lastly, line 18 is a widening from `long` to `float`. Note that there are no compiler errors anywhere and that the cast operator is not needed in any of the assignments.

Now, let's discuss narrowing, where the cast operator *is* required.

## Narrowing

The cast operator is a type enclosed in parentheses – for example, `(int)` and `(byte)` are both cast operators that cast to `int` and `byte`, respectively. With *Table 3.3* in mind, the following figure, *Figure 3.7*, presents assignments that require casting:

```
22 // Narrowing
23 int i = (int)3.3;           // narrowing, double to int
24 byte b = (byte) 233;       // narrowing, int to byte
25 float f = (float) 3.5;     // narrowing, double to float
26 System.out.println(i + " " + b + " " +f); // 3 -23 3.5
```

Figure 3.7 – Casting examples

In the preceding figure, line 23 is attempting to assign `3 . 3`, a `double` type (8 bytes), to an `int` type (4 bytes). Without the cast, this would be a compiler error. With the cast, you are overriding the compiler error. So, on line 23, we are casting `3 . 3` to `int` and assigning this `int` to the `i` variable. Therefore, after the assignment completes, `i` has a value of 3.

Line 24 is casting the `int` type, `233`, into the `byte` variable, `b`. This literal value is outside the range of `byte` (-128 to +127), so a cast is required. Line 25 is casting the `double` type, `3 . 5`, to `float`. Remember that, by default, a decimal number is `double`; to have it considered as a `float` as opposed to a `double`, you must suffix `f` or `F`. For example, `3 . 3f` is `float`.

The output on line 26 is 3, -23, and 3.5 for `i`, `b`, and `f`, respectively. Note that in the output, the `float` variable appears without `f`.

How we arrived at -23 is explained in the following callout.

### Overflowing the byte

Remember that the range of `byte` is -128 (10000000) to +127 (01111111). The leftmost bit is the sign bit, with 1 representing negative and 0 representing positive.

In the preceding example, we did the following:

```
byte b = (byte) 233;
```

The literal value of 233 (an integer) is too big for `byte` but how was `b` assigned the value of -23? Mapping 233 as an `int` type gives us the following bit pattern:

$$11101001 = 1 + 8 + 32 + 64 + 128 = 233 \text{ (int)}$$

Note that as an `int` is 4 bytes, 233 is 00000000000000000000000011101001. Mapping that bit pattern as a `byte` (the high order 3 bytes are truncated) gives us the following output:

$$11101001 = 1 + 8 + 32 + 64 + (-128) = -23 \text{ (byte)}$$

Remember that the leftmost bit is the sign bit. That is why  $-128$  is in the calculation. It is  $-(2^7) = -128$ .

We will conclude this section by looking at some unusual examples where casting is/is not required.

### To cast or not to cast, that is the question

There are certain situations where, because the compiler applies rules in the background, a cast is *not* required. Let's examine some of these situations with code examples. *Figure 3.8* presents the code:

32	<code>char c = 12;</code>
33	<code>char c2 = 90_000;</code>
34	<code>short s = 12;</code>
35	<code>s = c;</code>
36	<code>s = (short) c;</code>
37	<code>c = s;</code>
38	<code>c = (char) s;</code>
39	
40	
41	
42	

```
char c = 12;
char c2 = 90_000;
short s = 12;
s = c;
s = (short) c;
c = s;
c = (char) s;

final char c1 = 12;
short s1 = 12;
s1 = c1;
```

Figure 3.8 – Situations where casting is not always necessary

Line 32 declares and initializes a `char` variable `c`, to an `int` value of 12. Remember that `char` variables are essentially small positive numbers. Although we are assigning an `int` value (4 bytes) to a `char` variable (2 bytes), because the literal value is within the range of `char` (0 to 65,535), the compiler allows it. Had the literal value been out of the range of `char`, the compiler would have generated an error – this is what is happening on line 33.

---

Line 34 declares and initializes a `short` variable, `s`, to an `int` value of 12. Again, although `short` can hold only 2 bytes, the compiler realizes it can store the literal value, 12, and allows it.

Note that, from the compiler's perspective, assigning literal values into variables is different to assigning *variables* to variables. For example, lines 32 and 37 are quite different. This will become apparent as we discuss the next few lines in the figure.

Lines 35 to 38 demonstrate that while both `char` and `short` require 2 bytes, they have different ranges: `char` (0 to 65,535) and `short` (-32,768 to +32,767). This means that a `short` variable can hold a negative value such as -15, whereas a `char` variable cannot. Conversely, a `char` variable can hold a value such as 65,000 but a `short` variable cannot. Therefore, as lines 35 and 37 demonstrate, you cannot directly assign a `char` variable to a `short` variable and vice versa. You need a cast in both scenarios. Lines 36 and 38 demonstrate this.

### Compile-time constants

However, lines 40 to 42 show a way around the requirement for the cast we just outlined. If you declare your variable as a *compile-time constant* (and assuming the value is in range), the compiler will allow the variable-to-variable assignment. Line 40 uses the `final` keyword to declare a compile-time constant. We will discuss `final` in detail in *Chapter 9*, but in this context, it means that `c1` will always have a value of 12. The value is fixed (or *constant*) for `c1` and this is done at *compile time*. If you try to change the value of `c1`, you will get a compiler error. Now that the compiler knows that `c1` will always have 12 as its value, the compiler can apply the same rules that it applies to literal values; in other words, is the value in range? This is why line 42 does *not* generate a compiler error.

This concludes our discussion on operators. Now, let's apply them!

## Exercises

Mesozoic Eden is doing great. The dinosaurs are healthy and the guests are happy. Now that you have some new skills, let's go ahead and perform slightly more complicated tasks!

1. The caretakers want to be able to keep track of dinosaur weights. It's your task to write a program that calculates the average weight of two dinosaurs. This will help our team of nutritionists in planning the correct food portions.
2. Proper nutrition is essential for the health of our dinosaurs. The caretakers want to have a rough guideline of how much to feed a dinosaur. Write a program that determines the amount of food required for a dinosaur based on its weight. You can come up with the amount of food needed per weight unit of the dinosaurs.
3. For our park, we need to have a leap year checker. In our commitment to scientific accuracy, use the modulus operator to determine if the current year is a leap year. We want to make sure our calendar-themed exhibits are always up to date.



4. Create a program that checks whether the park's maximum capacity has been reached. The program only needs to print true or false after the words "Max capacity reached:". This is crucial in maintaining safety standards and ensuring a positive visitor experience.
5. Sometimes visitors want to compare dinosaurs' ages. And we get it – this could be interesting for educational purposes. Write a program that calculates the age difference between two dinosaurs.
6. In Mesozoic Eden, we have a very strong safety-first policy. Write a program that checks whether the park's safety rating is above a certain threshold. Maintaining a good safety rating is our utmost priority.

## Project – Dino meal planner

As a zookeeper in Mesozoic Eden, the crucial tasks include planning the meals for our beloved dinosaurs. While we're not using conditionals and loops yet, we can still calculate some basic requirements!

Develop a simple program to help the zookeepers plan the meal portions for different dinosaurs. The program should use the dinosaur's weight to calculate how much food it needs to eat per meal.

If you need a bit more guidance, here's how you can do it:

- Declare variables for the dinosaur's weight and the proportion of its weight it needs to eat per day. For instance, if a dinosaur needs to eat 5% of its body weight daily, and it weighs 2,000 kg, it would need to eat 100 kg of food.
- Now, let's say you feed the dinosaur twice a day. Declare a variable for the number of feedings and calculate how much food you need to serve per feeding. In this example, it would be 50 kg per feeding.
- Print out the result in a meaningful way – for example, "Our 2,000 kg dinosaur needs to eat 100 kg daily, which means we need to serve 50 kg per feeding."

## Summary

In this chapter, we learned about how Java's operators work and how they cooperate. In addition, we learned how to cast in Java.

Initially, we discussed two important properties relating to operators: precedence and associativity. We saw that precedence dictates how common terms are grouped. Associativity comes into play when the operators have the same order of precedence.

We then examined the operators themselves. We started by looking at unary operators, which have one operand such as the prefix/postfix increment/decrement operators, `++` and `--`.

---

We then moved on to the arithmetic operators: `+`, `-`, `*`, `/`, and `%`. We noted that integer division truncates. In addition, we discussed that any math operations involving `int` types or smaller results in `int`. Lastly, we discussed in detail how the `+` operator works when one or both operands are strings. In these cases, a string append is performed.

Next, we discussed relational operators. The results of these operators are always boolean values and will be used when we construct conditional statements in *Chapter 4*.

As Java cannot perform operations where the types are different, where possible, Java performs implicit promotion. This is where Java promotes the smaller type to the larger type somewhere in memory. This is Java's way of invisibly continuing with the operation.

We then discussed the logical operators: `&&`, `||`, `&`, `|`, and `^`. Truth tables were presented to aid in understanding. Both the logical `&&` and logical `||` operators are short-circuiting operators. Understanding this is important because the order of evaluation trumps precedence.

The bitwise operators, bitwise AND (`&`) and bitwise OR (`|`), are similar except that in contrast to `&&` and `||`, both `&` and `|` never short-circuit and can also work with integral operands.

The ternary operator takes three operands. It evaluates a boolean expression and assigns one of two values to a variable, depending on whether the boolean expression was `true` or `false`.

Regarding operators, the last group we covered were the compound assignment operators, of which there is one for each mathematical operator.

In our discussion on casting, we covered both widening and narrowing. Widening is done in the background and is often called *implicit promotion*. There is no risk here as the type being promoted fits easily into the target type.

Narrowing is where the cast is required. This is because, given that you are going from a type that requires more storage space to a type that requires less, there is a potential loss of data.

Now that we know how to use operators, in the next chapter, we will move on to conditional statements, where operators are commonly used.

Trial Version



Wondershare  
**PDFelement**

## 4

# Conditional Statements

In *Chapter 3*, we learned about Java operators. We discussed two important properties of operators, namely, precedence and associativity. Precedence helps group shared operands. When precedence levels match, associativity is then used for grouping.

We discussed the unary operators – prefix and postfix increment/decrement, cast, and logical NOT. We also covered the binary operators – arithmetic, relational, logical, bitwise, and compound assignment. We learned about the behavior of the + symbol when one (or both) operands is a string. We discussed the logical AND (`&&`) and logical OR (`||`) and their short-circuiting property. Finally, the ternary operator, with its three operands, was covered.

We also learned about Java casting. This can be done implicitly, known as **implicit promotion** or **widening**. The other alternative is explicit casting, known as **narrowing**. When narrowing, we must cast to the target type in order to remove the compiler error. Lastly, we discussed compile-time constants, which, because their values never change, enable the compiler to apply different rules.

Now that we know about operators, let us do something interesting with them. By the end of this chapter, you will be able to use Java's operators to create conditional statements. Conditional statements enable us to make decisions. In addition, you will understand a fundamental concept in Java, namely, scope.

In this chapter, we are going to cover the following main topics:

- Understanding scope
- Exploring `if` statements
- Mastering `switch` statements and expressions

## Technical requirements

The code for this chapter can be found on GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects/tree/main/ch4>.

## Understanding scope

In programming, scope defines where a variable is/is not usable within a program. This is often referred to as the visibility of the variable. In other words, where in the code is the variable “visible”. Java uses **block scope**. In order to explain Java’s scope, we must first understand what a block is.

### What is a block?

Curly braces delimit a block of code. In other words, a block starts with the opening curly brace, {, and ends with the closing curly brace, }. Note that the braces face each other, as in { }. A variable is visible and available for use, from where it is declared in the block, to the closing } of that block. *Figure 4.1* presents a code example to help explain:

```
3 ► public class Scope {
4 ►   public static void main(String[] args) {
5     int x = 1;
6     x++;
7
8     { // start of block
9       int y = 2;
10      y++;
11      x++;
12    } // end of block
13    x++;
14    y++; // out of scope
15  }
16}
```

Figure 4.1 – Block scope in Java

In the preceding figure, we declare an `int` variable, `x`, on line 5 and initialize it to 1. The current block of code is the group of Java statements surrounded by `{ }`. Therefore, the `x` variable’s block of code starts on line 4, where the opening curly brace is, and ends on line 15, where the closing curly brace is. Thus, the scope of `x` is from line 5, where it is declared, to line 15 (the closing curly brace of the current scope). When we refer to `x` on line 6, there is no issue, as `x` is in scope.

On line 8, we start a new block/scope with `{`. Though somewhat unusual, as there is no code preceding `{` on line 8, lines 8 to 12 define a valid code block. Note that variables from the outer scope are visible within this inner (nested) scope. This is shown on line 11 where, in the inner scope, we refer to a variable declared in the outer scope, namely `x`, without any issue.

The inverse is not true, however; a variable defined in an inner scope is not visible in the outer scope. The `y` variable is in scope from where it is declared (line 9 of the inner scope) to line 12 (the closing curly brace of that scope). Thus, we get a compiler error on line 14 where, in the outer scope, we refer to the `y` variable.

### Indentation

Indentation really helps with the identification of code blocks and, consequently, scopes. The style we use is to start the code block at the end of the line. For example, in *Figure 4.1*, note the opening curly brace `{` on line 3. The closing curly brace `}` for that code block (and thus the scope) is on line 16. From an indentation point of view, the closing curly brace `}` is directly under the `public` keyword from line 3. More specifically, the closing curly brace `}` is directly under the '`p`' in `public`. While not necessary for compilation, it does make your code easier to read and maintain – the scope starts on line 3 and to find where the scope ends, one just scans down the program to find the matching curly brace `}` which lines up under `public` (from line 3).

Line 4 also defines a block and thus a scope. Line 15 contains the matching curly brace `}` for that scope – note that the closing curly brace is lined up under the keyword `public` (from line 4). Thankfully, the editors are a great help in keeping your code properly indented.

In summary, blocks are defined with `{ }` . As Java uses block scope, the code blocks define the scope of where a variable can be used. Variables are visible in nested scopes, but not vice versa.

Now that we understand scope in Java, let us examine conditional logic in Java. We will start with `if` statements.

## Exploring if statements

As their name suggests, conditional statements are based on the evaluation of a condition. The result of this condition is either `true` or `false` – in other words, a boolean. *Figure 4.2* introduces the syntax of the overall `if` statement:

```
if (booleanExpression) statement/block
[ else if (booleanExpression) statement/block ] ...
[ else statement/block ]
```

Figure 4.2 – The if statement syntax

The square brackets `[ ]` in the preceding figure denote something as optional. For example, both the `else if` statements and the `else` statement are optional. The `if` statement itself is mandatory. The three ellipses, `... ,` indicate that you can have as many `else if` statements as you like (or none at all).

Now that we have the overall syntax, let us break it down into smaller pieces.

## The if statement itself

As stated earlier, an `if` statement evaluates a boolean expression. This boolean expression is enclosed in parentheses. Curly braces that delimit a block of code are optional if there is only one statement after the `if` clause. However, it is considered good practice to always explicitly declare a block of code. *Figure 4.3* demonstrates both styles:



```
8 int x=5, y=4;
9 if(x > y)
10     System.out.println(x + " > " + y);
11 if(x < y)
12     System.out.println(x + " < " + y);
13 if(x == y){
14     String s = x + " == " + y;
15     System.out.println(s);
16 }
```

Figure 4.3 – Simple if statements

First, let us explain the code. Note the indentation, which is automatically facilitated by the code editors. This makes it easier to see the statements that are governed by the various `if` statements. Both lines 9 and 11 demonstrate simple `if` statements that control just one statement. Line 9 controls line 10. This means that if line 9 is true, line 10 is executed. If line 9 is false then line 10 is skipped. Similarly, if line 11 is true: line 12 is executed; if line 11 is false, line 12 is skipped.

However, if you wish to execute two or more statements when an `if` statement is true, a code block is required. This is demonstrated by lines 13 to 16. If the boolean expression on line 13 evaluates to true, then both of the statements on lines 14 and 15 will be executed. This is because they are in a code block.

As regards the running program, with `x` initialized to 5 and `y` to 4, when line 9 executes, it is true (as `5 > 4`). Therefore, line 10 executes and therefore the output from *Figure 4.3* is `5 > 4`. Lines 11 and 13 are both executed but as they both evaluate to false, nothing else is output to the screen.

What if, on line 8, we initialized the variables as follows:

```
int x=4, y=5;
```

Now, `if (x > y)` is false and line 10 is not executed; `if (x < y)` is true and line 12 outputs `4 < 5` to the screen; `if (x == y)` is also false, so lines 14 and 15 are not executed.

Lastly, let us make the variables equal by changing line 8 as follows:

```
int x=4, y=4;
```

Now, `if (x > y)` is false so line 10 is not executed; `if (x < y)` is also false so line 12 is not executed; however, `if (x == y)` is true, so line 14 builds up the string `s` to be “`4 == 4`” and line 15 outputs it.

Note the indentation, which is automatically facilitated by the code editors. This makes it easier to see the statements that are governed by the various `if` statements.

## else if statements

In *Figure 4.3*, there is no code to cater to the situations where an `if` expression evaluates to false. This is where the `else if` statement comes in. *Figure 4.4* shows an `else if` in code:



```
int x=4, y=5;
if(x > y) {
    System.out.println(x + " > " + y);
} else if(x < y) {
    System.out.println(x + " < " + y);
} else if(x == y){
    System.out.println(x + " == " + y);
}
System.out.println("Here");
```

Figure 4.4 – else if statements

As `x` is 4 and `y` is 5 (line 19), the `if` expression on line 20 evaluates to false and thus control jumps to line 22 where the first `else if` is evaluated. As this evaluates to true, line 23 is executed. Now, *no other branch will be evaluated*. In other words, the next line of code executed after line 23 is line 27. Note that, as per good coding practice, each branch is coded as a block, even though there is only one statement in each block.

Had `x` been initialized to 5, then lines 20 and 22 would both have evaluated to false. Line 24 would be true, and thus, line 25 would be executed.

For situations where the `if` and `else if` statements do not match, we can use the `else` statement. Let us discuss that now.

## else statements

The code in *Figure 4.4* evaluates all possible scenarios when comparing `x` and `y`. Either `x` is greater than, less than, or equal to `y`. This logic lends itself nicely to introducing the `else` statement. The `else` statement is a *catch-all*. As per [] in *Figure 4.2*, the `else` clause is optional. If present, it must be coded at the end after any `if` and/or `else if` clauses. *Figure 4.5* is *Figure 4.4* refactored using an `else` clause, except that the values of `x` and `y` in *Figure 4.5* are now the same.

```

30
31   if(x > y) {
32     System.out.println(x + " > " + y);
33 } else if(x < y) {
34   System.out.println(x + " < " + y);
35 } else {
36   System.out.println(x + " == " + y);
37 }
38 System.out.println("Here");

```

Figure 4.5 – else statement

In *Figure 4.5*, as both `x` and `y` are 4, lines 31 and 33 evaluate to `false`. However, there is no condition on line 35 as it is simply an `else` statement (as opposed to `else if`). This means that the code block beginning on line 35 is executed automatically and line 36 is executed.

With regard to `if else` statements, it is important to understand a subtle problem that can arise known as the “dangling `else`” problem.

### Dangling else

Consider the following unindented code, which uses no code blocks:

```

boolean flag=false;           // line 1
if (flag)                   // line 2
if (flag)                   // line 3
System.out.println("True True"); // line 4
else                         // line 5
System.out.println("True False"); // line 6

```

This code has two `if` statements but only one `else` statement. Which `if` statement does `else` match with? The rule is: when an `else` statement is looking to match with `if`, it will match with the nearest *unmatched if* as it progresses back up through the code.

Following this rule, the `else` statement matches with the second `if` (line 3) as that `if` statement has not yet been matched. This means that the `if` statement on line 2 remains unmatched. The code, when written using proper indentation, is much clearer:

```

if (flag)                   // line 2
  if (flag)                 // line 3
    System.out.println("True True"); // line 4
  else                      // line 5
    System.out.println("True False"); // line 6

```

This is confirmed by the output. If `flag` is true, the output is "True True"; if `flag` is false, however, nothing is output to the screen. Interestingly, there is no way line 6 can now be reached (as boolean variables have only two values: `true` and `false`).

Using code blocks makes the code even easier to understand, as can be seen as follows:

```
if (flag) {  
    if (flag) {  
        System.out.println("True True");  
    }  
    else {  
        System.out.println("True False");  
    }  
}
```

This is why using code blocks, even for one statement, is very helpful. Throughout the book, we will use proper indentation and code blocks to aid clarity and ease of understanding.

Now, let us look at a more involved example. This example (and others to follow) use the pre-defined `Scanner` class from the Java **Application Programming Interface (API)**. The API is a suite of pre-defined types (for example classes) that are available for our use. We will cover these topics as we progress through the book but suffice to say that the API is extremely useful as it provides pre-defined and well-tested code for our use.

The `Scanner` class resides in the `java.util` package. Therefore, we need to briefly discuss both packages and the `Scanner` class.

## Packages

A package is a group of related types, such as classes, that we can use. Conveniently, many are already available for us to use in the API.

To gain access to these types, we need to “import” them into our code. For this purpose, Java provides the `import` keyword. Any `import` statements go at the top of your file. We can import a whole package using the `*` wildcard; for example: `import java.util.*;`. We can also import a particular type explicitly by naming it in the `import` statement; for example: `import java.util.Scanner;`.

When you precede the type with its package name, such as `java.util.Scanner`, this is known as the “fully qualified name”. We could omit the `import` statement and simply refer each time to `Scanner` using its fully qualified name; in other words, everywhere `Scanner` is mentioned, replace it with `java.util.Scanner`. Generally speaking, however, importing the type and using its non-qualified name is preferred.

There is one package that is automatically available (imported) for us and that is the `java.lang` package. For example, the `String` class resides in `java.lang` and that is why we never have to import anything to get access to the `String` class.

### Scanner class

It is helpful to know that while `Scanner` is a versatile class, for our purposes, we will simply use `Scanner` to enable us to retrieve keyboard input from the user. With this in mind, note that `System.in` used in the examples refers to the standard input stream which is already open and ready to supply input data. Typically, this corresponds to the keyboard. `System.in` is therefore perfect for getting input data from the user via the keyboard. The `Scanner` class provides various methods for parsing/interpreting the keyboard input. For example, when the user types in a number at the keyboard, the method `nextInt()` provides that number to us as an `int` primitive. We will avail of these methods in our examples.

### Nested if statements

Now that we have discussed packages and `Scanner`, in *Figure 4.6* and *Figure 4.7*, we discuss a more involved `if-else` example that deals with input from the user (via the keyboard). Both figures relate to the one example. *Figure 4.6* focuses on the declaration of constants to make the code more readable. In addition, *Figure 4.6* also focuses on declaring and using `Scanner`. On the other hand, *Figure 4.7* focuses on the subsequent `if-else` structure.

```
11     final int JAN = 1; final int FEB = 2; final int MAR = 3; // constants
12     final int APR = 4; final int MAY = 5; final int JUN = 6;
13     final int JUL = 7; final int AUG = 8; final int SEP = 9;
14     final int OCT = 10; final int NOV = 11; final int DEC = 12;
15
16     Scanner sc = new Scanner(System.in); // import java.util.Scanner;
17     System.out.print("Enter month --> ");
18     int month = sc.nextInt();
```

Figure 4.6 – Using `Scanner` to get input from the keyboard

Lines 11 to 14 define constants using the keyword `final`. This means that their values cannot change. It is good practice to use uppercase identifiers for constants (where words are separated by underscores). This will make the code in *Figure 4.7* more readable; in other words, instead of comparing `month` with 1 (which, in this context, means January), we will compare `month` with `JAN`, which reads better. For brevity's sake, we have declared three constants per line but you can easily declare one per line also.

Line 16 in *Figure 4.6* creates our `Scanner` object reference, `sc`. Essentially, we are creating a reference, namely `sc`, which refers to the `Scanner` object created using the `new` keyword. As stated previously, `System.in` means that `sc` is looking at the keyboard. This reference is what we will use to interact with `Scanner`, much like a remote control is used to interact with a television.

Line 17 prompts the user to enter a month (1 .. 12) on the keyboard. This is actually very important because, without the prompt, the cursor will just blink and the user will wonder what they should type in. Line 18 is where `Scanner` really comes into its own. Here we use the `nextInt()` method to get in a number. For the moment, just know that when we call `sc.nextInt()`, Java does not

return to our code until the user has typed in something and hit the return key. For the moment, we will make the (convenient) assumption that it is an integer. We will store the `int` primitive returned in our own `int` primitive, `month`. Now, in our code, we can use what the user typed in. *Figure 4.7* shows this in action.

```

20     int numDays=0;
21     if(month == JAN || month == MAR || month == MAY || month == JUL
22         || month == AUG || month == OCT || month == DEC) {
23         numDays=31;
24     } else if (month == APR || month == JUN || month == SEP || month == NOV) {
25         numDays=30;
26     } else if (month == FEB) {
27         System.out.print("Enter year --> ");
28         int year = sc.nextInt();
29
30         if( (year % 400 == 0) || (year % 4 == 0 && !(year % 100 == 0)) ){
31             numDays = 29; // leap year e.g. 2000, 2012, 2016
32         }else{
33             numDays = 28; // 1900 (divisible by 100)
34         }
35     } else {
36         System.out.println("Invalid month: "+month);
37     }
38     if(numDays > 0){
39         System.out.println("Number of days is: "+numDays);
40     }

```

Figure 4.7 – A complex if statement

Note that the code presented in the preceding image is a *continuation* of *Figure 4.6*. On line 20, we declare an `int` variable, namely `numDays`, and initialize it to 0. Lines 21 to 22 are the start of the `if` statement. Using the boolean logical OR operator, the `if` statement checks to see whether the `month` value matches any of the constants defined in *Figure 4.6*. In the background, the constant values are used, so in reality, the `if` statement is as follows:

```
if(month == 1 || month == 3 || month == 5 || month == 7 ||
month == 8 || month == 10 || month == 12)
```

Note that the `month` variable must be specified each time. In other words, `if (month == JAN || MAR || MAY || JUL || AUG || OCT || DEC)` will *not* compile.

So, assuming that the user typed in 1 (representing January), `month` becomes 1 and, as a result, lines 21 to 22 evaluate to true and `numDays` is set to 31 on line 23. The logic is the same if the user types in 3, 5, 7, 8, 10, and 12, representing March, May, July, August, October, and December, respectively.

If the user types in 4 (representing April), lines 21 to 22 evaluate to false and the `else if` statement on line 24 is evaluated. Line 24 evaluates to true and `numdays` is set to 30 on line 25. The logic is the same if the user types in 6, 9, and 11, representing June, September, and November, respectively.

Now, let us deal with the user typing in 2, representing February. The `if` condition on lines 21 to 22 and the `else if` condition on line 24 both evaluate to false. Line 26 evaluates to true. Now, we need to review the leap year logic. Of course, February has 28 days every year (as do all the months!) but has one extra day when the year is a leap year. The logic for determining whether a year is a leap year is as follows:

- **Scenario A:** `year` is a multiple of 400 => leap year
- **Scenario B:** `year` is a multiple of 4 AND `year` is *not* a multiple of 100 => leap year

The following are all leap years: 2000 (satisfies scenario A), 2012, and 2016 (both satisfy scenario B).

As the leap year algorithm depends on the year, we first need the year from the user. Lines 27 to 28 accomplish this. We then encounter a nested `if` statement from lines 30 to 34, which determines, according to the logic just outlined, whether the year entered by the user is a leap year. Line 30 implements the logic for both scenarios A and B. Whether `year` is a multiple of 400 is achieved with `(year % 400 == 0)`. The condition of `year` being a multiple of 4 and not a multiple of 100 is achieved with

`(year % 4 == 0 && !(year % 100 == 0))`. The fact that either condition satisfies the leap year calculation is achieved by using the logical OR operator between them. Assuming a year of 2000, line 30 would be true and `numDays` is set to 29. Assuming a year of 1900, the `if` statement on line 30 is false and, as there is no condition on line 32 (it is just an `else` statement), line 33 is executed, setting `numDays` to 28.

An invalid `month` value of, for example, 25 or -3 would result in the `else` branch on line 35 being executed. An error message would be output to the screen on line 36 as a result.

Lines 38 to 40 output the number of days provided that `numDays` was changed from its initial value of 0. The `if` statement on line 38 prevents the message "Number of days is: 0" from appearing on the screen if the user typed in an invalid `month` value.

That concludes our treatment of the `if` statement. Now, let us examine both `switch` statements and expressions, which can, in certain situations, be a more elegant option.

## Mastering switch statements and expressions

Complicated `if` statements, with many `else if` branches and an `else` branch can be verbose. The `switch` structure can, in many situations, be more concise and elegant. Let us start with `switch` statements.

## switch statements

Firstly, let us examine the syntax of the `switch` statement. *Figure 4.8* introduces the syntax.

```
switch (expression) {  
    [case label:  
        statement(s);  
        break;] ...  
    [default:  
        statement(s);  
        break;]  
}
```

Figure 4.8 – The switch statement syntax

A `switch` statement evaluates an expression. As of Java 21, the expression can be an integral primitive (excluding `long`) or any reference type. This means that we can `switch` on primitive variables of type `byte`, `char`, `short`, or `int` and also `switch` on class types, enum types, record types and array types. The `case` labels can now include a `null` label. Java 21 also brought in *pattern matching for switch*. We will present another `switch` example demonstrating this feature when we have those topics covered (*Chapter 9*). Until then, we will focus on the more traditional `switch`.

### Wrapper types

For each of the primitive types, there is a corresponding class, known as a “wrapper type”: `byte` (wrapped by `Byte`), `short` (`Short`), `int` (`Integer`), `long` (`Long`), `float` (`Float`), `double` (`Double`), `boolean` (`Boolean`), and `char` (`Character`). They are so called because they represent objects that encapsulate the primitive. As they are class types, useful methods are available. For example, `int val = Integer.parseInt("22");` converts the `String "22"` into the number 22, stored in `val`, where we can perform arithmetic.

The expression just evaluated is compared against the `case` labels. The `case` labels are compile-time constants of the same type as the `switch` expression. If there is a match with a `case` label, the associated block of code is executed (note: no need for curly braces in the `case` or `default` blocks). To exit the `case` block, ensure you insert a `break` statement. The `break` statement exits the `switch` block. However, the `break` statement is optional. If you omit the `break` statement, the code *falls through* to the next `case` label (or `default`), even though there is no match.

The `default` keyword is used to specify a code block to execute if none of the `case` labels match. Typically, it is coded at the end of the `switch` block but this is not mandatory. In effect, `default` can appear anywhere in the `switch` block with similar semantics (this is a poor programming practice, however).

Figure 4.9 presents an example.

```
18 Scanner sc = new Scanner(System.in);
19 System.out.print("Enter a sport --> ");
20 String sport = sc.next();
21 switch(sport){
22     case "Soccer":
23         System.out.println("I play soccer");
24         break;
25     case "Rugby":
26         System.out.println("I play Rugby");
27         break;
28     default:// can be moved around
29         System.out.println("Unknown sport");
30         break;
31 }
```

Figure 4.9 – switch on a String example

Using the `Scanner` class, lines 18 to 20 ask and retrieve from the user a sport. Note that the `Scanner` method used on this occasion is `next()`, which returns a `String` type (as opposed to a primitive type).

Lines 21 to 31 present the `switch` block. Note that the `case` labels on lines 22 and 25 are both `String` compile-time constants. If the user types in "Soccer", the `case` label on line 22 matches, and both lines 23 and 24 will execute. Interestingly, even though there are two statements to be executed, there is no need for curly braces here. This is a feature of `switch` blocks. As line 22 matched, line 23 will execute, and "I play soccer" will be echoed to the screen. The `break` statement on line 24 ensures that the `switch` block is exited and that the "Rugby" section is not executed.

If, on the other hand, the user types in "Rugby", line 25 matches, and "I play Rugby" is echoed to the screen. Again, the `break` statement, this time on line 27, ensures that the `switch` block is exited and that the `default` section is not executed.

If the user typed in "Tennis", then neither of the `case` labels on lines 22 and 25 will match. This is when the `default` section comes into play. When there are no matches with any `case` label, the `default` section is executed. Typically, the `default` section is coded at the end of the `switch` block and the `break` statement is traditionally inserted for completeness.

#### Case labels are case-sensitive

Note that the `case` labels are case-sensitive. In other words, in *Figure 4.9*, if the user types in "soccer", the `case` label on line 22 will *not* match. The "Rugby" `case` label will not match either, naturally. Thus, the `default` section will execute, and "Unknown sport" will echo (print) to the screen.

Let us look at another example. *Figure 4.10* is a `switch` statement based on integers.

```
34 Scanner sc = new Scanner(System.in);
35 System.out.print("Enter a number (1..10) --> ");
36 int number = sc.nextInt();
37 final int two = 2; // compile-time constant
38 switch(number){
39     case 1:
40     case 3:
41     case 5:
42     case 7:
43     case 9:
44         System.out.println(number + " is odd.");
45         break;
46     case two:
47     case 4: case 6: case 8: case 10:
48         System.out.println(number + " is even.");
49         break;
50     default:
51         System.out.println( number + " is outside range (1..10).");
52     break;
53 }
```

Figure 4.10 – switch on an integer example

In the preceding figure, line 37 declares a compile-time constant, `two`, and initializes it to the integer literal, 2. Line 38 starts the `switch` block by switching on an `int` variable named `number`, which was declared and initialized based on user input, on line 36. All of the `case` labels in the `switch` block must now be integers – be they literal values as on lines 39 to 43 and line 47, or compile-time constants as on line 46. Note that if the `two` variable is not `final`, a compile-time error is generated (as it is no longer a constant).

In this example, we have multiple labels together, such as, lines 39-43. This section of code can be read as *if number is 1 or 3 or 5 or 7 or 9, then do the following*. So, for example, if the user types in 1, the `case` label on line 39 matches. This is known as the *entry point* of the `switch` statement. As there is no `break` statement on line 39, the code *falls through* to line 40, even though line 40 has a `case` label for 3. Again, line 40 has no `break` statement and the code *falls through* to line 41. In fact, the code keeps executing from the entry point until it reaches a `break` statement (or the end of the `switch` statement itself). This *fall-through* behavior is what enables the *if it's 1 or 3 or 5 or 7 or 9 type logic* to work. If this *fall-through* behavior were not present, we would have to duplicate lines 44 to 45 for each of the `case` labels! Line 44 uses the `String` append to output that the number entered is odd, by appending "is odd" to the number – for example, "7 is odd". Line 45 is the `break` statement that ensures we exit the `switch` block.

Line 46 is just to demonstrate that compile-time constants work for `case` labels. Line 47 shows that the `case` labels can be organized in a horizontal fashion if so desired. Remember, the use of indentation and spacing is just for human readability – the compiler just sees one long sequence of characters. So, lines 46 to 47 match for the numbers 2, 4, 6, 8, and 10. Again, the *fall-through* logic is used to keep the code concise. Lines 48 to 49 output that that number is even and `break` out of the `switch` block.

Line 50 is the `default` section, which caters to any numbers outside of the `1..10` range. Line 51 outputs that the number entered is out of range, and line 52 is the `break` statement. While the `break` statement is not strictly needed here (as `default` is at the bottom of the `switch` statement), it is considered good practice to include it.

Let us rewrite the code in *Figure 4.7* using a `switch` statement instead of a complicated `if-else` statement. Note that *Figure 4.6* is still relevant in declaring `Scanner` and the constants used; this is why we separated *Figure 4.6* (so we could use it with both the `if` and `switch` code). *Figure 4.11* represents *Figure 4.7* refactored using a `switch` statement.

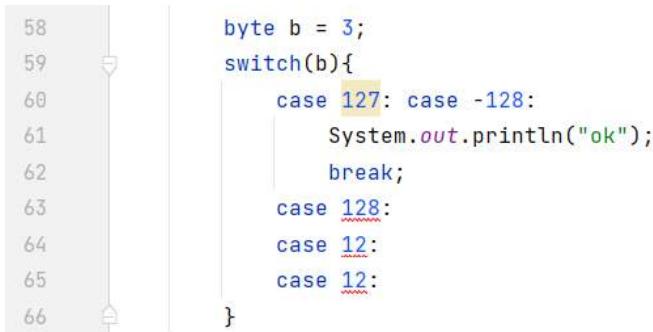
```
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
```

```
int numDays=0;
switch(month){
    case JAN:case MAR:case MAY:case JUL: case AUG:case OCT:case DEC:
        numDays=31;
        break;
    case APR:case JUN:case SEP:case NOV:
        numDays=30;
        break;
    case FEB:
        System.out.print("Enter year --> ");
        int year = sc.nextInt();
        if( (A)           || (B           && C) ) {
            if( (year % 400 == 0) || (year % 4 == 0 && !(year % 100 == 0)) ){
                numDays = 29; // leap year e.g. 2000, 2012, 2016
            }else{
                numDays = 28; // 1900 (divisible by 100)
            }
        }
        break;
    default:
        System.out.println("Invalid month: "+month);
        break;
}
if(numDays > 0){
    System.out.println("Number of days is: "+numDays);
}
```

Figure 4.11 – Refactoring an `if` statement with a `switch` statement

In the preceding figure, we can see the `switch` statement starts on line 23. The first group of `case` labels is on line 24 and the second group is on line 27. February, the odd one out, has a `case` label to itself (line 30). Finally, the `default` label is on line 40. Personally speaking, I find the use of `switch` in this example preferable due to the absence of the multiple logical OR expressions required in *Figure 4.7* (lines 21 to 22 and 24).

Now that we have covered valid `switch` statements, let us examine, in *Figure 4.12*, a few scenarios where compiler errors can arise.



```
58
59
60
61
62
63
64
65
66
```

```
byte b = 3;
switch(b){
    case 127: case -128:
        System.out.println("ok");
        break;
    case 128:
    case 12:
    case 12:
```

Figure 4.12 – Some switch compiler errors

In the preceding figure, we are switching on a `byte` variable, `b`. Recall that the valid `byte` range is `-128` to `+127`. Line 60 demonstrates that the minimum and maximum values are fine. Line 63 shows that, as `128` is out of range for our `byte` type `b`, the `case` labels that are out of range of the `switch` variable cause a compiler error. Line 64 was fine until line 65 used the same `case` label – `case` label duplicates are not allowed.

We will finish our discussion on `switch` by discussing `switch` expressions.

## switch expressions

Like all expressions, `switch` expressions evaluate to a single value and, therefore, enable us to return values. All the `switch` examples so far have been `switch statements`, which return nothing. Note that in a `switch expression`, `break` statements are not allowed.

On the other hand, `switch expressions` return something – either implicitly or explicitly (using `yield`). We will explain `yield` shortly, but note that `yield` cannot be used in a `switch statement` (as they do not return anything). In addition, `switch statements` can *fall through*, whereas `switch expressions` do not. These differences are encapsulated in *Table 4.1*.

	<b>break</b>	<b>Returns a value</b>	<b>yield</b>	<b>Fall through</b>	<b>New case label</b>	<b>Regular case label</b>
<b>switch statement</b>	Yes	No	No	Yes	Yes	Yes
<b>switch expression</b>	No	Yes	Yes	No	Yes	Yes

Table 4.1 – Comparison of switch statements versus switch expressions

Let us look at some example code to demonstrate the differences. *Figure 4.13* is the traditional `switch` statement.

```

16 int nLetters=0;
17 String name="Jane";
18 switch(name){
19     case "Jane":
20     case "Sean":
21     case "Alan":
22     case "Paul":
23         nLetters = 4;
24         break;
25     case "Janet":
26     case "Susan":
27         nLetters = 5;
28         break;
29     case "Maaike":
30     case "Alison":
31     case "Miriam":
32         nLetters = 6;
33         break;
34     default:
35         System.out.println("Unrecognized name: "+name);
36         nLetters = -1;
37         break;
38 }
39 System.out.println(nLetters);

```

Figure 4.13 – A traditional switch statement

In the preceding figure, we are switching on the `String` variable, `name`. As it is initialized to "Jane", line 19 is true and line 23 sets `nLetters` to 4 (the number of letters in "Jane"). The `break` statement on line 24 ensures that there is no fall-through to line 27. Line 39 outputs 4 to the screen.

Notice that the code is quite verbose and requires the correct use of the break statement to prevent fall through. Plus, these break statements are tedious to write and easy to forget. *Figure 4.14* represents *Figure 4.13* written using a switch expression.

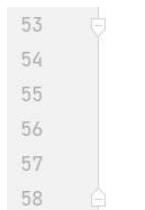


```
41 nLetters = switch(name){  
42     case "Jane", "Sean", "Alan", "Paul" -> 4;  
43     case "Janet", "Susan" -> 5;  
44     case "Maaike", "Alison", "Miriam" -> 6;  
45     default -> {  
46         System.out.println("Unrecognized name: "+name);  
47         yield -1; // 'nLetters' initialized to -1  
48     }  
49 };  
50 System.out.println(nLetters);
```

Figure 4.14 – A switch expression

The preceding figure shows the new case label where the labels are comma-delimited and an arrow token separates the labels from the expression (or code block). There is no break statement required anywhere as there is no fall-through behavior to worry about. As name is still "Jane" (from *Figure 4.13*, line 17), line 42 is executed, which initializes/returns 4 into nLetters. Thus, line 50 outputs 4. Note that the list of case labels in a switch expression must be exhaustive. In almost all cases, this means that a default clause is required. The default clause in this example executes a code block (lines 45-48), where an error is output to the screen and using the yield keyword, nLetters is initialized to (an error value of) -1.

We can omit the need for the nLetters variable by returning the expression value straight into the System.out.println() statement. *Figure 4.15* demonstrates this.



```
53 System.out.println(switch(name){  
54     case "Jane", "Sean", "Alan", "Paul" -> 4;  
55     case "Janet", "Susan" -> 5;  
56     case "Maaike", "Alison", "Miriam" -> 6;  
57     default -> "Unrecognized name: "+name;  
58 });
```

Figure 4.15 – A switch expression returning straight to System.out.println()

In the preceding figure, there is no variable used to store the result of the switch expression. This makes the code even more concise. The result of the switch expression is returned straight into the System.out.println() statement. Again, 4 is output to the screen.

### The `yield` keyword

*Figure 4.14* and *Figure 4.15* were simple `switch` expressions, where (for the most part), to the right of the arrow token was the value to be returned. However, instead of simply returning a value, what if you wished to execute a code block? This is where `yield` is used. *Figure 4.16* shows the use of `yield` in a `switch` expression.

```

60 nLetters = switch(name){
61     case "Jane", "Sean", "Alan", "Paul" -> {
62         System.out.println("There are 4 letters in: " + name);
63         yield 4;
64     }
65     case "Janet", "Susan" -> {
66         System.out.println("There are 5 letters in: "+name);
67         yield 5;
68     }
69     case "Maaike", "Alison" , "Miriam" -> {
70         System.out.println("There are 6 letters in: "+name);
71         yield 6;
72     }
73     default -> {
74         System.out.println("Unrecognized name: "+name);
75         yield -1;
76     }
77 };
78 System.out.println(nLetters);

```

Figure 4.16 – A switch expression using `yield`

The preceding figure highlights the fact that, if you need to execute more than one statement in a `switch` expression, you must provide a code block. This is shown with the curly braces for lines 61-64, 65-68, 69-72 and 73-76. To return an expression result from a code block we use `yield`. This is what is done on lines 63, 67, 71 and 75, where 4, 5, 6 and -1 are returned, respectively. As `name` is still "Jane", line 63 returns 4 from the `switch` expression, initializing `nLetters` to 4. Therefore, line 78 outputs 4.

To aid flexibility, you can, up to a point, mix the syntaxes. In other words, regular `case` labels can be used in `switch` expressions, and the new `case` labels syntax can be used in `switch` statements. However, as stated earlier, `break` only appears in `switch statements` and `yield` only appears (if required) in `switch expressions`. *Figure 4.17* is a refactor of the verbose `switch` statement in *Figure 4.13*, where the new `case` labels are used in a `switch` statement.

```

82     switch(name){
83         case "Jane", "Sean", "Alan", "Paul" -> nLetters = 4;
84         case "Janet", "Susan" -> nLetters = 5;
85         case "Maaike", "Alison" , "Miriam" -> nLetters = 6;
86         default -> {
87             System.out.println("Unrecognized name: "+name);
88             nLetters = -1;
89         }
90     }
91     System.out.println(nLetters);

```

Figure 4.17 – A switch statement using new case labels and an arrow token

In the preceding figure, the labels are comma-delimited and the arrow token is present. As before, on the right of the arrow token, we have the initialization of `nLetters` to the number of letters in the name. However, in contrast to *Figure 4.13*, as we are using the arrow token, no `break` statements are required. Note however that curly braces are required for a code block as per the `default` clause (lines 86-89).

We can also use the regular `case` labels with `switch` expressions. This is shown in *Figure 4.18*, which is a refactored version of *Figure 4.16*.

```

94     nLetters = switch(name){
95         case "Jane":
96         case "Sean":
97         case "Alan":
98         case "Paul":
99             System.out.println("There are 4 letters in: " + name);
100            yield 4;
101        case "Janet":
102        case "Susan":
103            System.out.println("There are 5 letters in: "+name);
104            yield 5;
105        case "Maaike":
106        case "Alison":
107        case "Miriam":
108            System.out.println("There are 6 letters in: "+name);
109            yield 6;
110        default:
111            System.out.println("Unrecognized name: "+name);
112            yield -1;
113    };
114    System.out.println(nLetters);

```

Figure 4.18 – A switch expression using old-style case labels



In the preceding figure, the old-style case labels are used. This means that the keyword `case` must precede each label. The curly braces for the code blocks can, however, be omitted. As it is a `switch` expression, where there is more than one statement to be executed when a match is found, we need to use `yield` to return the expression result. As the name variable has never been changed from "Jane" throughout all the examples, a match is made on line 95, resulting in line 99 outputting "There are 4 letters in: Jane" to the screen. The `yield` on line 100 returns 4, and thus, `nLetters` is initialized to 4. Finally, line 114 outputs 4 to the screen.

This completes our treatment of `switch` expressions and, indeed, `switch` statements in general. We will now put what we have learned into practice.

## Exercises

We finally have the coding capability to make decisions. Mesozoic Eden is going to be benefiting from this so much. Let's show off our newly acquired skills, shall we?

1. We need to determine whether a dinosaur is a carnivore or herbivore. Write an `if` statement that prints whether a dinosaur is a carnivore or herbivore based on a `boolean` variable. This information is critical for feeding and care guidelines.
2. Different species require different care strategies and exhibit unique behavior traits. Write a `switch` statement that prints a description of a dinosaur based on its species. This will help educate both the staff and park visitors.
3. Some dinosaurs are tougher to handle than others. Write an `if` statement that checks whether a number of years of experience is enough experience to work with a certain type of dinosaur. This ensures the safety of both our dinosaurs and employees.
4. We are working with beautiful but dangerous creatures. So, safety first. Write a program that prints a warning message if the park's safety rating falls below a certain threshold. We must always be alert to potential issues that could harm our staff, visitors, or dinosaurs.
5. Proper housing is essential for the dinosaurs' well-being. Write a `switch` statement that assigns a dinosaur to a specific enclosure based on its size (XS, S, M, L, or XL).
6. Proper nutrition is crucial for maintaining our dinosaurs' health. Write an `if` statement that determines the number of feeds a dinosaur requires per day based on its weight.
7. It is important to delegate tasks properly to keep operations running smoothly. Create a program that assigns different duties to employees based on their job titles using a `switch` statement.
8. The park is not open to day visitors 24/7. Write an `if` statement that checks whether the park is open for them based on the time. They are open for day visitors from 10 A.M. to 7 P.M. This helps in managing visitor expectations and staff schedules.

## Project – Task allocation system

The manager of Mesozoic Eden needs a systematic way of managing the team and ensuring all tasks are efficiently accomplished.

Design a simple program that assigns tasks to the Mesozoic Eden employees based on their roles (for example, feeding, cleaning, security, and tour guiding). The program should decide tasks based on time, the employee's role, and other factors, such as the park's safety rating.

This program would not only help streamline operations but also ensure the safety and satisfaction of our staff, visitors, and, most importantly, our dinosaurs!

## Summary

In this chapter, we started by explaining that Java uses block scope. A block is delimited by { }. A variable is visible from the point of declaration to the closing } of that block. As blocks (and therefore, scopes) can be nested, this means that a variable defined in a block is visible to any inner/nested blocks. The inverse is not true, however. A variable declared in an inner block is not visible in an outer block

Conditional statements enable us to make decisions and are based on the evaluation of a condition resulting in true or false. The if statement allows several branches to be evaluated. Once one branch evaluates to true and is executed, no other branch is evaluated. An if statement can be coded on its own without any else if or else clause. The else if and else clauses are optional. However, if an else clause is present, it must be the last clause. We saw how a complex if example can lead to code verbosity.

We briefly discussed packages and the Scanner class. The Scanner class resides in the java.util package and is very useful for retrieving keyboard input from the user.

We also discussed both switch statements and switch expressions. An expression can return a value but a statement cannot. We saw how switch statements can make complicated if statements more concise and elegant. The expression you switch on (typically, a variable) can be a primitive byte, char, short, or int type; or a reference type. The case labels must be compile-time constants and must be in range for the switch variable. switch statements have a fall-through feature, which enables multiple case labels to use the same section of code without repetition. However, this fall-through behavior requires a break statement to exit the switch statement. This requires care, as break statements are easy to forget.

switch expressions can return a value. They have no fall-through logic so break statements are not required. This makes the code more concise and less error-prone. If you wish to execute a code block in a switch expression, use yield to return the value.



`switch` statements do not support `yield` (as they do not return anything), and `switch` expressions do not support `break` (as they must return something). However, both of the `case` labels, namely the old-style `case X:` and the newer style `case A, B, C ->`, can be used with either `switch` statements or `switch` expressions.

Now that we know how to make decisions, we will move on to iteration in the next chapter, where we will examine the Java structures that enable us to repeat statements.

## 5

# Understanding Iteration

In *Chapter 4*, we learned about scope and conditional statements in Java. Scope determines the visibility of identifiers – in other words, where you can use them. Java uses block scope, which is defined by curly braces, { }. Scopes can be nested but not vice versa.

We discussed variations of the `if` statement. Each of these statements evaluates a boolean condition, resulting in true or false. If true, then that branch is executed and no other branch is evaluated. If false, then the next branch is evaluated. Unless an `else` clause is present, it is possible that no branch at all will be executed.

For complex `if` statements, Java supports the more elegant `switch` structure. We examined `switch` statements, with their *fall-through* behavior, and the use of the `break` statement. In addition, we discussed `switch` expressions, where a value can be returned, and their use of `yield`.

Now that we understand conditional logic, let us examine iteration (looping). Looping constructs enable us to repeat statements and/or blocks of code a finite number of times while a boolean condition is true or while there are more entries in the array/collection. By the end of this chapter, you will be able to use Java's looping constructs.

In this chapter, we are going to cover the following main topics:

- `while` loops
- `do-while` loops
- `for` loops
- Enhanced `for` (`for-each`) loops
- `break` and `continue` statements

## Technical requirements

The code for this chapter can be found on GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects/tree/main/ch5>.

## while loops

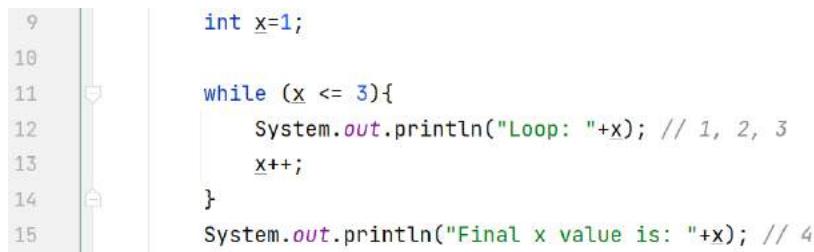
An important feature of any programming language is the ability to perform an action repeatedly. This is known as “looping”. We may want to repeat a piece of code a finite number of times or until some condition is met; for example, the user typing in a value that signifies that the loop should terminate. In most cases, a boolean expression can be used to determine whether the loop continues or not.

A `while` loop is one such looping construct. It repeatedly executes a statement or a block of code as long as a boolean expression is true. As soon as the boolean expression is false, the loop exits, and the next statement after the `while` loop executes.

```
while (booleanExprIsTrue) {  
    // do something  
}
```

Figure 5.1 – The while loop syntax

In the preceding figure, we are assuming a block of code, hence the curly braces `{ }`. You could, of course, omit the curly braces `{ }` and the loop will just repeatedly execute one statement (which ends with a semi-colon). Interestingly, as the boolean expression could be false to begin with, the `while` loop may not execute at all. More formally, a `while` loop executes *zero* or more times. Let us look at some examples. *Figure 5.2* presents a simple `while` loop:



```
int x=1;  
  
while (x <= 3){  
    System.out.println("Loop: "+x); // 1, 2, 3  
    x++;  
}  
System.out.println("Final x value is: "+x); // 4
```

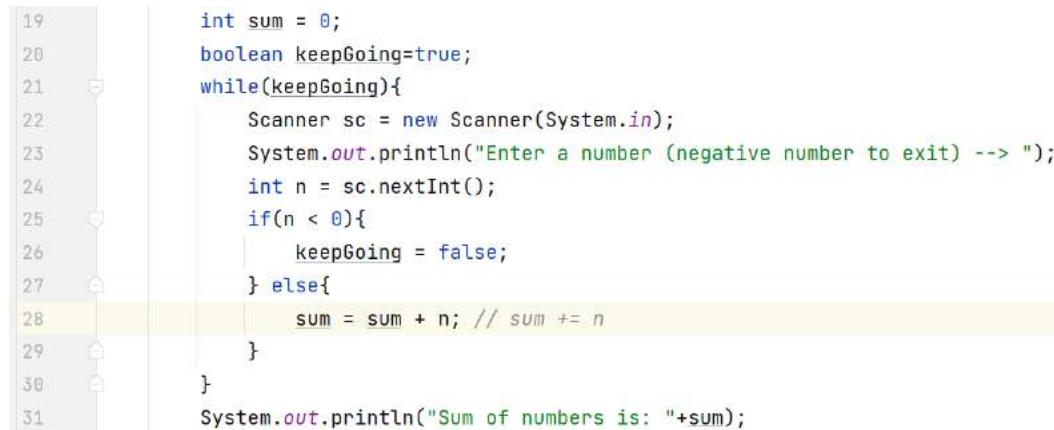
Figure 5.2 – A simple while loop

On line 9 in the preceding figure, a local variable, `x`, is initialized to 1. Line 11 evaluates the boolean expression, `x <= 3`. As `x` is 1, the boolean expression is true and the loop executes. Line 12 outputs "Loop: 1" and line 13 increments `x` to 2. The `}` symbol on line 14 is reached and the loop condition (on line 11) is automatically rechecked to see whether it still holds. As `x` is 2 and `2 <= 3`, the condition is true and we re-enter the loop. Line 12 outputs "Loop: 2" and line 13 increments `x` to 3. The end of the block is reached on line 14, and again, the loop continuation expression is re-evaluated. As the expression `3 <= 3` is true, the loop is executed again. Line 12 outputs "Loop: 3" and line 13 increments `x` to 4. Once again, the end of the code block is reached and the loop continuation

expression is re-evaluated. As `x` is now 4 and as `4 <= 3` is false, the loop exits. This is shown by line 15 outputting "Final x value is: 4".

Note that if, on line 9, `x` had been initialized to 11 (as opposed to 1), then the initial boolean expression on line 11 would have been evaluated to false and the loop would never have been executed at all.

A `while` loop can be very useful when you do *not* know how many times the loop is going to iterate. For example, the loop continuation expression may be predicated on user input. *Figure 5.3* is one such loop.



```
19 int sum = 0;
20 boolean keepGoing=true;
21 while(keepGoing){
22     Scanner sc = new Scanner(System.in);
23     System.out.println("Enter a number (negative number to exit) -->");
24     int n = sc.nextInt();
25     if(n < 0){
26         keepGoing = false;
27     } else{
28         sum = sum + n; // sum += n
29     }
30 }
31 System.out.println("Sum of numbers is: "+sum);
```

The screenshot shows a code editor with the above Java code. The code uses a while loop to read integers from the user via a Scanner. It adds each integer to a running total 'sum'. The loop continues as long as the user enters a non-negative number. A negative number exits the loop. The code is color-coded for syntax: int, boolean, while, Scanner, System.out, println, Scanner.nextInt(), if, else, and sum are in blue; keepGoing is in green; and the condition in the while loop is in purple. Line numbers 19 through 31 are on the left. A vertical toolbar on the left has icons for file operations like Open, Save, and Print.

Figure 5.3 – A while loop that ends based on user input

In the preceding figure, an algorithm for summing up a sequence of positive, user-inputted numbers is presented. The loop will keep going, totaling up the numbers entered by the user, until a negative number is entered. This negative number is naturally, not part of the total. Let us discuss it in more detail.

Line 19 declares a local `int` variable, `sum`, and initializes it to 0. Line 20 declares a local boolean variable, `keepGoing`, and sets it to true. The boolean expression on line 21 evaluates to true (due to line 20) and, as a result, the loop block executes. Line 22 declares our `Scanner` reference, `sc`, pointing at the keyboard. Line 23 prompts the user to enter a number while informing the user that any negative number terminates the loop. Line 24 uses the `Scanner` method, `nextInt()`, to get an integer (whole number) from the user. This number is stored in the local variable, `n`. Line 25 checks to see whether a negative number has been entered. If so, line 26 sets the `keepGoing` flag to false so that the loop will not execute again. If a non-negative integer was entered by the user, then the number entered, `n`, is added to the running total, `sum`.

Let us walk through an example. We will add the following numbers: 1, 2, and 3, totaling 6. This is what the screen output looks like:

```
Enter a number (negative number to exit) -->
1
```

```
Enter a number (negative number to exit) -->
2
Enter a number (negative number to exit) -->
3
Enter a number (negative number to exit) -->
-1
Sum of numbers is: 6
```

Let us examine what is happening in the code. The loop (line 21) is entered because the `keepGoing` boolean was set to true on line 20. We are then prompted for our first number (line 23). We type in 1, resulting in `n` being initialized to 1 on line 24. As `n` is 1, the `if` statement on line 25 is false and the `else` block (lines 27-29) is executed; setting `sum` to 1 ( $0 + 1$ ).

The loop block end is reached (line 30) and the loop continuation expression is automatically re-evaluated (line 21). As `keepGoing` is still true, the loop continues. We are prompted for our second number; we enter 2 and `sum` is changed to 3 ( $1 + 2$ ).

The loop block end is reached again and, as `keepGoing` is still true, the loop continues. We are prompted for our next number; we enter 3 and `sum` is changed to 6 ( $3 + 3$ ).

Again, the loop block end is reached and, as `keepGoing` is still true, the loop continues. We are prompted for our next number. This time we enter -1. As `n` is now negative, the `if` statement on line 25 is true and `keepGoing` is set to false (line 26). Now, when the loop continuation expression is next evaluated, as `keepGoing` is false, the loop exits.

Lastly, line 31 outputs "Sum of numbers is: 6".

Now that we have covered the `while` loop, let us examine its close relative, the `do-while` loop.

## do-while loops

As we have seen with the `while` loop, the boolean loop continuation expression is at the start of the loop. Though similar to the `while` loop, the `do-while` loop is different in one critical aspect: in the `do-while` loop, the loop continuation expression is at the *end* of the loop. Thus, the `do-while` loop is executed at least *once*. More formally, a `do-while` loop executes *one or more times*.

*Figure 5.4* presents the syntax of the `do-while` loop.

```
do {
    // do something
} while (booleanExprIsTrue);
```

Figure 5.4 – The do-while loop syntax

As can be seen in the preceding figure, the loop continuation expression is at the end of the loop, after one loop iteration. Also note the semi-colon, ; after ).

*Figure 5.5* presents a do-while version of the while loop in *Figure 5.2*.



```
19 int sum = 0;
20 boolean keepGoing=true;
21 do {
22     Scanner sc = new Scanner(System.in);
23     System.out.println("Enter a number (negative number to exit) --> ");
24     int n = sc.nextInt();
25     if(n < 0){
26         keepGoing = false;
27     } else{
28         sum = sum + n; // sum += n
29     }
30 } while(keepGoing);
31 System.out.println("Sum of numbers is: "+sum);
```

Figure 5.5 – A do-while loop that ends based on user input

In the preceding figure, the only differences with *Figure 5.2* are lines 21 and 30. On line 21, we simply enter the loop as, unlike in the while loop, there is no condition preventing us from doing so. Line 30 checks to see whether it is okay to re-enter the loop. The rest of the code is the same and the execution is the same.

While (pardon the pun) the two examples given have no material difference in the outcome, let us examine a situation where using a while loop as opposed to a do-while loop is preferable.

## while versus do-while

As already stated, a do-while loop executes at least once, whereas a while loop may not execute at all. This can be very useful in certain situations. Let us look at one such example. *Figure 5.6* presents a while loop that checks to see whether a person is of the legal age to purchase alcohol (which is 18 in Ireland).

```
49 Scanner sc = new Scanner(System.in);
50 System.out.println("Please enter your age -- > ");
51 int age = sc.nextInt();
52 while(age >= 18){
53     // purchase alcohol...
54     System.out.println("As you are "+age+" years of age, " +
55         "you can purchase alcohol.");
56
57     System.out.println("Please enter your age -- > ");
58     age = sc.nextInt();
59 }
```

Figure 5.6 – A while loop to prevent underage purchasing of alcohol

In the preceding figure, line 49 declares the `Scanner` and points it at the keyboard so we can get user input. Line 50 prompts the user to enter their age. Line 51 takes in the user input and stores it in a local variable, namely `age`. Line 52 is important. The condition prevents the loop from being executed with an invalid age. The loop itself is trivial and simply outputs a message that includes the age so we can validate that the loop is executing properly.

Lines 57 to 58 are very important in that they enable us to prompt and get a new age from the user. The code deliberately overwrites the `age` variable. If we did not, then `age` would remain as the first value entered by the user and we would have an infinite loop. So, the first age is entered before the `while` loop is entered and every other age is entered at the end of the loop. This is a common pattern in `while` loops. The condition on line 52 prevents any `age` input that is  $< 18$ , from entering the loop.

Here is the first run of the code in *Figure 5.6*:

```
Please enter your age -- >
21
As you are 21 years of age, you can purchase alcohol.
Please enter your age -- >
12
```

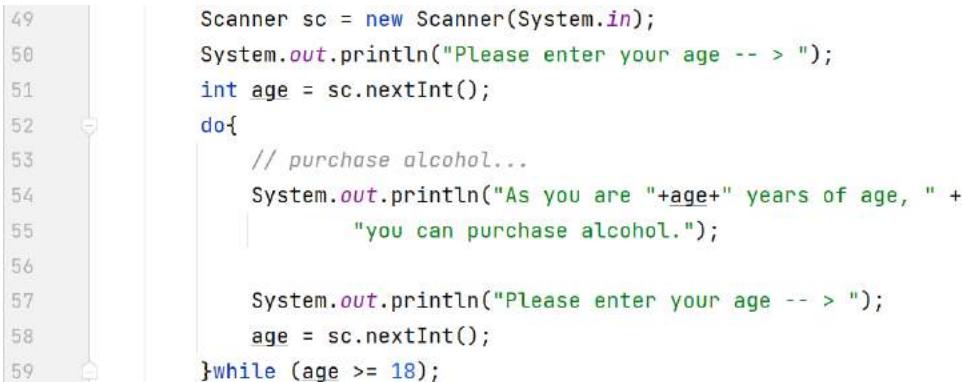
The first two lines: the prompt and user input, are before the `while` loop and, as  $21 \geq 18$ , we enter the loop. The message `As you are 21 years of age, you can purchase alcohol.` is perfectly correct. The last two lines: repeating the prompt and user input, are from the bottom of the loop. We have entered 12, which causes the `while` loop to terminate.

The following is the output if, when prompted for the first age, we enter 12:

```
Please enter your age -- >
12
Process finished with exit code 0
```

Importantly, the message about purchasing alcohol does *not* appear.

Now, let us look at the do-while version. *Figure 5.7* presents the do-while version of the while loop in *Figure 5.6*.



```
49     Scanner sc = new Scanner(System.in);
50     System.out.println("Please enter your age -- >");
51     int age = sc.nextInt();
52     do{
53         // purchase alcohol...
54         System.out.println("As you are "+age+" years of age, " +
55             "you can purchase alcohol.");
56
57         System.out.println("Please enter your age -- >");
58         age = sc.nextInt();
59     }while (age >= 18);
```

Figure 5.7 – A do-while loop to prevent underage purchasing of alcohol

In the interests of having as much of the code as similar as possible, lines 49 to 51 are untouched. Lines 52 and 59 are all that have changed. The condition is now at the end of the loop, after one iteration of the loop. This has implications when we start with an age of 12, as can be seen in the output:

```
Please enter your age -- >
12
As you are 12 years of age, you can purchase alcohol.
Please enter your age -- >
12
Process finished with exit code 0
```

The third line in the output is the issue. Obviously, 12 is too young to buy alcohol but the do-while loop would require an if statement to protect its code, whereas the while loop provides that protection automatically. Therefore, in this example, there is a material advantage in using the while loop over the do-while loop.

Now that we have covered while and do-while loops, let us now discuss for loops.

## for loops

The `for` loop comes in two styles: the *traditional* `for` loop and the *enhanced* `for` loop. The enhanced `for` loop is also known as the `for-each` loop and is specifically designed to work with arrays and collections. We will start by examining the traditional `for` loop.

### Traditional for loop

This type of `for` loop is extremely useful when you know how many iterations you wish to perform beforehand. Its syntax is detailed in *Figure 5.8*.

```
for (initialization; booleanExpression; incr/decr) {  
    // do something  
}
```

Figure 5.8 – The traditional for loop

The code block in the preceding figure is optional. We could simply control one statement, such as `System.out.println("Looping");`, and omit `{ }`. The `for` header is the section inside `()`. It consists of three parts, delimited by semi-colons:

- **Initialization section:** This is where you initialize your loop control variables. The variables declared here have the scope of the loop block *only*. Traditionally, the variables declared here are named `i`, `j`, `k`, and so forth.
- **Boolean expression:** This determines whether the loop should be executed and is checked before every iteration, including the first one. Sound familiar? Yes, you are correct, a `while` loop and a traditional `for` loop are interchangeable.
- **Increment/decrement section:** This is where you increment/decrement your loop control variables (declared in the initialization section) so that the loop terminates.

We must understand the order of execution of the loop. In other words, which part is executed and when. *Figure 5.9*, which presents a simple `for` loop, will help in this regard.

```
// for(init; booleanExpr; incr/decr)  
for(int i=1; i<=3; i++){  
    System.out.println(i); // 1,2,3  
}
```

Figure 5.9 – A simple traditional for loop

In this figure, the order of execution of the code is represented in numerical order as follows:

1. **Initialization section:** The loop control variable, *i*, is declared and initialized to 1.
2. **Boolean expression:** Evaluate the boolean expression to see whether it is okay to execute the loop. As  $1 \leq 3$ , it is okay to enter the loop.
3. **Execute the loop block:** This outputs 1 to the screen.
4. **Increment/Decrement section:** *i* is incremented (by 1) from 1 to 2 and then execution pops over to the boolean expression.
5. **Evaluate the boolean expression:** As  $2 \leq 3$ , the loop is executed.
6. **Execute the loop block:** This outputs 2 to the screen.
7. **Increment *i* from 2 to 3:** and then pop over to the boolean expression.
8. **Evaluate the boolean expression:** As  $3 \leq 3$ , the loop is executed.
9. **Execute the loop block:** This outputs 3 to the screen.
10. **Increment *i* from 3 to 4:** and then pop over to the boolean expression.
11. **Evaluate the boolean expression:** As 4 is not  $\leq 3$ , the loop exits.

In summary, the initialization section is executed only once, at the start of the loop. The boolean expression is evaluated and, assuming it is true, the loop body is executed, followed by the increment/decrement section. The boolean expression is again evaluated and, again, assuming it is true, the loop body is executed, followed by the increment/decrement section. This repetition of the execution of the loop body followed by the increment/decrement section continues until the boolean expression fails and the loop exits.

Figure 5.10 presents a `for` loop that goes from 3 down to 1 in decrements of 1:

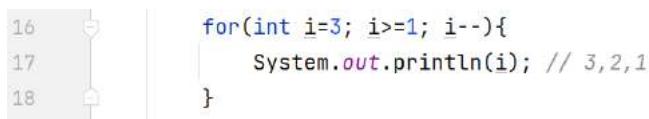
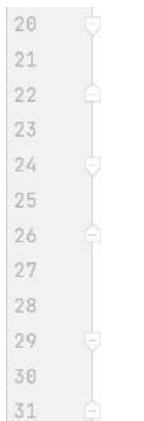


Figure 5.10 – A simple for loop that operates in descending order

In the preceding figure, we initialize *i* to 3 and check the boolean expression. As  $3 \geq 1$ , we enter the loop and output 3. We then decrement *i* by 1 to 2 and check the boolean expression again. As  $2 \geq 1$ , we output 2 and then decrement *i* to 1. As the boolean expression is still true; we output 1 and *i* is decremented to 0. At this point, as *i* is 0, the boolean expression is false and the loop terminates.

Figure 5.11 presents some code samples enabling us to discuss this looping construct further.



```
20
21     for(int i=1; i<=3; i++){
22         System.out.println("Looping"); // only appears once!
23     }
24
25     for(int i=10; i<=50; i+=10){
26         System.out.println(i); // 10, 20, 30, 40, 50
27     }
28     System.out.println(i); // i is out of scope
29
30     for(int i=0, j=0; i<1 && j<1; i++, j++){
31         System.out.println(i + " " + j); // 0 0
}
```

Figure 5.11 – Additional traditional for loops

In the first loop of the preceding figure (lines 20-22), the important thing to notice is the ; symbol, which is just after the ) symbol of the `for` header. This loop controls an empty statement! Even though the indentation may suggest otherwise, the block of code that follows has nothing to do with the loop at all, and as a result, "Looping" appears only once in the output. In effect, the loop iterates three times, doing nothing each time. The block of code surrounding line 21 is not predicated on any condition and just executes once (as normal).

In the second loop (lines 24-26), the loop control variable, `i`, starts out at 10 and goes up in increments of 10 until it reaches 60, at which point the loop terminates. Each valid value of `i` is output to the screen – in other words, 10, 20, 30, 40, and 50. Note that line 27 does *not* compile, as each of the `i` variables declared in the preceding loops only have the scope of their individual loop. For example, the `i` variable declared on line 20 is only available until line 22; similarly, the `i` variable declared on line 24 is only available until line 26. Note: obviously, line 27 must be commented out for the code to compile and run.

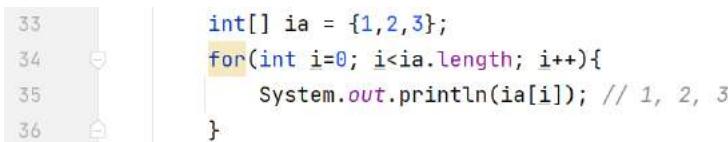
The last loop (lines 29-31) shows that we can declare multiple loop control variables and use them throughout the loop. In this loop, we declare `i` and `j` and initialize them both to 0. The boolean expression is true as both `i < 1` and `j < 1` are true (true && true == true). Thus, we execute the loop and output 0 and 0. Both `i` and `j` are then incremented to 1. The loop condition fails and the loop terminates.

While arrays will be discussed in detail in *Chapter 6*, `for` loops are such a natural fit for arrays that we have inserted some examples here as well. Let us first examine how a traditional `for` loop can be used to process an array.

## Processing an array

Any `for` loop is useful for iterating over an array. An array is simply an area of memory set aside and given an identifier name for ease of reference. An array consists of elements which are organized in consecutive memory locations – in other words, the array elements are right beside each other in memory. This makes it easy to process arrays using loops.

Each element in an array is accessed by an index. Crucially, array indices start at 0 and go up in steps of 1. Therefore, the last valid index is the size of the array minus one. For example, an array of size 5 has valid indices of 0, 1, 2, 3, and 4. *Figure 5.12* is a loop processing an array.



```
33 int[] ia = {1,2,3};  
34 for(int i=0; i<ia.length; i++){  
35     System.out.println(ia[i]); // 1, 2, 3  
36 }
```

Figure 5.12 – Processing an array using a traditional for loop

In this figure, line 33 declares an `int` array containing the values 1, 2, and 3 in indices 0, 1, and 2, respectively. The length of the array, accessible using the `length` property, is 3. The `for` loop (lines 34-35) processes the array, outputting each location one by one. Thus, when `i` is 0, `ia[0]` outputs 1 to the screen; when `i` is 1, `ia[1]` outputs 2, and when `i` is 2, `ia[2]` outputs 3.

Now that we have covered the traditional `for` loop, let us examine the enhanced `for` loop.

## Enhanced for loop

As stated earlier, the enhanced `for` loop, also known as the `for-each` loop, is ideal for processing arrays and/or collections. We will discuss collections in detail in *Chapter 13*. For the moment, just imagine a collection as a *list* of items. The enhanced `for` loop enables you to iterate over the list one element at a time. The syntax of the enhanced `for` loop is outlined in *Figure 5.13*.

```
for (dataType variableName : array or collection){  
    // do something  
}
```

Figure 5.13 – Enhanced for loop syntax

In the preceding figure, we can see that a variable is declared. The variables type matches the type of array/collection. For example, if the array is an array of `String`, then `String` is the data type of the variable. The variable name is of course, up to us. Again, the code block is optional.

Let us look at an example to help explain further. *Figure 5.14* is an enhanced `for` loop version of the traditional `for` loop presented in *Figure 5.12*.

```

38 int[] ia = {1,2,3};
39   for(int n:ia){
40     System.out.println(n); // 1, 2, 3
41   }

```

Figure 5.14 – Processing an array using an enhanced for loop

In this figure, line 38 reads as follows: *for each int n in (the array) ia*. Thus, on the first iteration, n is 1; on the second iteration, n is 2, and on the last iteration, n is 3. In the enhanced `for` loop, we do not have to keep track of a loop control variable ourselves. While this is useful, be aware that you are limited to starting at the beginning of the array/collection and progressing one element at a time, until you reach the end. With the traditional `for` loop, none of these restrictions apply.

However, with the traditional `for` loop, if you code the increment/decrement section incorrectly, you could end up in an infinite loop. This is not possible in the enhanced `for` version.

## Nested loops

Loops can, of course, be nested. In other words, loops can be coded within other loops. *Figure 5.15* presents one such example.

```

12 int[] data = {9, 3, 5, 7};
13
14   System.out.println("[\t[n]\tHistogram");
15   for (int i=0; i<data.length; i++){
16     System.out.print(i + "\t" + data[i] + "\t");
17     for(int j=1; j<=data[i]; j++) { // write out data[i] stars
18       System.out.print("*");      // print() not println()
19     }
20     System.out.println();          // go onto next line
21   }

```

Figure 5.15 – Nested for loops

The output from this program is presented in *Figure 5.16*. In the preceding figure, we are representing an array of `int` values, namely `data`, as a histogram (represented as a row of stars). The array is declared on line 12. Line 14 outputs a line of text so the output from the program is easier to interpret. The output has three columns: the current array index, the value in the `data` array at that index, and the histogram. Note that the output is tab-delimited. This is achieved by the use of the `\t` escape sequence.

## Escape sequences

An escape sequence is a character preceded by a backslash. For example, `\t` is a valid escape sequence. When the compiler sees `\`, it peeks ahead at the next character and checks to see whether the two characters together form a valid escape sequence. Popular escape sequences are as follows:

- `\t`: Insert a tab at this point in the text
- `\b`: Insert a backspace at this point in the text
- `\n`: Insert a newline at this point in the text
- `\ "`: Insert a double quote at this point in the text
- `\ \`: Insert a backslash at this point in the text

They can be very useful in certain situations. For example, if we wanted to output the text *My name is "Alan"* (including the double quotes) to the screen, we would say:

```
System.out.println("My name is \"Alan\"");
```

If we did not escape the double quote before the A in Alan (in other words, if we tried `System.out.println("My name is "Alan")`), then the double quote before the A would have been matched with the first `"` at the start of the string. This would have resulted in a compiler error with the A in Alan.

By escaping the double quote before the A in Alan, the compiler no longer treats that double quote as an end-of-string double quote and instead inserts `"` into the string to be output. The same happens to the double quote after the n in Alan—it is also escaped and therefore ignored as an end-of-string double quote and inserted into the string to be output. The double quote just before the `)` is not escaped however, and is used to match the opening double quote for the string, namely the one just after the `(`.

The outer loop (lines 15-21) loops through the `data` array. As the array has 4 elements, the valid indices are 0, 1, 2, and 3. These are the values that the `i` loop control variable, declared on line 15, will represent. Line 16 outputs two of the columns: the current array index and the value at that index in the `data` array. For example, when `i` is 0, `data[0]` is 9, so `"0\t9\t"` is output; when `i` is 1, `data[1]` is 3, so `"1\t3\t"` is output, and so forth.

The inner loop (lines 17-19) outputs the actual histogram as a horizontal row of stars. The inner loop control variable, `j`, goes from 1 to the value of `data[i]`. So, for example, if `i` is 0, `data[i]` is 9; therefore, `j` goes from 1 to 9, outputting a star each time. Note that the `print()` method is used as opposed to `println()` – this is because `println()` automatically brings you on to the next line, whereas `print()` does not. As we want the stars to output horizontally, `print()` is exactly what we need. When we have our row of stars output, we execute `System.out.println()` (line 20), which brings us on to the next line.

Figure 5.16 represents the output from the code in Figure 5.15.

```
[] [n] Histogram
0 9 *****
1 3 ***
2 5 ****
3 7 *****
```

Figure 5.16 – Output from the code in Figure 5.15

In this figure, you can see that the first column is the array index. The second column is the value in the `data` array at that index, and the third column is the histogram of stars based on the second column. So, for example, when `i` is 2, `data[2]` is 5, and we output a histogram of 5 stars.

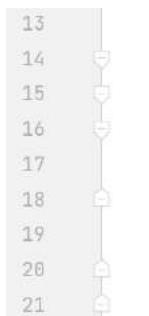
Now that we understand loops, we will move on to two keywords that are particularly relevant to loops, namely `break` and `continue`.

## break and continue statements

Both the `break` and `continue` statements can be used in loops but with very different semantics. In the code examples presented, nested loops will be used to contrast the labeled versions with the non-labeled versions. We will start with the `break` statement.

### break statements

We have already encountered `break` in `switch` statements. When used in a loop, the loop exits immediately. Figure 5.17 presents nested for loops with a `break` in the inner loop.



```
13
14
15
16
17
18 System.out.println("i, j");
19 for (int i = 1; i <= 3; i++) {
20     for (int j = 1; j <= 5; j++) {
21         if (j == 3) {
22             break; // breaks out of inner loop
23         }
24         System.out.println(i + ", " + j);
25     }
26 }
```

Figure 5.17 – Showing break inside a loop

---

In this figure, the outer loop, controlled by `i`, loops from 1 to 3 in steps of 1. The inner loop, controlled by `j`, loops from 1 to 5 in steps of 1.

The `if` statement on line 16 becomes true when `j` is 3. At this point, the `break` statement on line 17 is executed. A `break` without a label exits the nearest enclosing loop. In other words, the `break` on line 17 refers to the loop on line 15 (controlled by `j`). As there is no code between the closing `}` of both loops (lines 20 and 21), when `break` is executed in this program, the next line of code executed is the `}` for the outer loop (line 21). Automatically, the next iteration of the outer loop, `i++` (line 14), starts. In effect, there is never any `j` value of 3 or higher in the output because, when `j` is 3, we break out of the inner loop and start with the next value of `i`. The output reflects this:

```
i, j
1, 1
1, 2
2, 1
2, 2
3, 1
3, 2
```

Without any `break` statement, in other words, if we had commented out lines 16 to 18, the output would be as follows (note the values of `j` go from 1 to 5):

```
i, j
1, 1
1, 2
1, 3
1, 4
1, 5
2, 1
2, 2
2, 3
2, 4
2, 5
3, 1
3, 2
3, 3
3, 4
3, 5
```

Before we discuss the labeled `break`, we will quickly discuss the label itself.

## Label

A label is a case-sensitive identifier followed by a colon that immediately precedes the loop being identified. For example, the following code defines a valid label, OUTER, for the loop controlled by `i`:

```
OUTER:
for (int i = 1; i <= 3; i++) {
    for (int j = 1; j <= 5; j++) {
```

Now let us look at the labeled `break` itself.

## Labeled break

A `break` that uses a label exits the loop identified by that label. The labeled `break` statement must be in the scope of the loop identified. In other words, you cannot `break` to a loop somewhere else in the code, completely unrelated to the current scope. *Figure 5.18* is closely related to the code in *Figure 5.17*, except this time, a label and a labeled `break` are used.

```
25 System.out.println("i, j"); // placed BEFORE label!!
26 OUTERLOOP:
27 for (int i = 1; i <= 3; i++) {
28     for (int j = 1; j <= 5; j++) {
29         if (j == 3) {
30             break OUTERLOOP; // case sensitive
31         }
32         System.out.println(i + ", " + j);
33     }
34 }
35 System.out.println("here");
```

Figure 5.18 – Labeled break

In the preceding figure, we have labeled, on line 26, the outer loop as `OUTERLOOP`. Yes, it took a while to come up with that identifier! Note that it is a compiler error to have any code between the label and the loop. That is why line 25 precedes the label.

The loop control variables, `i` and `j`, behave as before; `i` goes from 1 to 3 in steps of 1, and within each step of `i`, `j` goes from 1 to 5 in steps of 1. This time, however, when `j` is 3 in the inner loop, rather than breaking out of the inner loop, we are breaking out of the outer loop. After the labeled `break` (line 30) is executed, there are no more iterations of `i` and the next line executed is `System.out.println("here")` on line 35. As a result, the output is as follows:

```
i, j
1, 1
```

```
1, 2
here
```

As can be seen, once *j* reaches 3, the outer loop exits, and *here* is output.

Now, let us look at *continue* statements.

## continue statements

A *continue* statement can only occur inside a loop. When executed, *continue* says “skip to the *next iteration*” of the loop. Any other statements remaining in the current iteration are bypassed. There is a labeled version also. We will examine the unlabeled version first. *Figure 5.19* presents an example of *continue*.

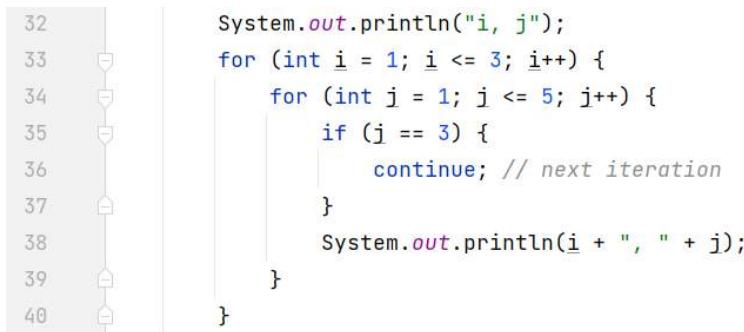


Figure 5.19 – A *continue* example

In the preceding figure, the nested loops are the same as before – the outer loop iterates from 1 to 3; within that, the inner loop iterates from 1 to 5. On this occasion, when *j* is 3, we execute *continue*. What that means is that we jump to the end of the loop and the next statement executed is *j++*. This means that as line 38 is skipped, *j* with a value of 3 will never be output. The output demonstrates this:

```
i, j
1, 1
1, 2
1, 4
1, 5
2, 1
2, 2
2, 4
2, 5
3, 1
3, 2
3, 4
3, 5
```

As can be seen, *j* with a value of 3 is never output. Now, let us examine the labeled *continue*.

### ***labeled continue***

A *continue* that uses a label continues the next iteration of the loop identified by that label. All other statements are bypassed. As with the labeled *break*, the labeled *continue* must be in the scope of the loop identified. *Figure 5.20* is closely related to the code in *Figure 5.19*, except this time, a label and a labeled *continue* are used.

```
28 System.out.println("i, j\n===="); // placed BEFORE label!!
29
30    OUTERLOOP:
31        for (int i = 1; i <= 3; i++) {
32            for (int j = 1; j <= 5; j++) {
33                if (j == 3) {
34                    continue OUTERLOOP; // continues with OUTERLOOP
35                }
36                System.out.println(i + ", " + j);
37            }
}
```

Figure 5.20 – A labeled *continue* example

In this figure, line 29 gives the *OUTERLOOP* label to the outer loop starting on line 30. Now, when *j* is 3 and *continue OUTERLOOP* executes, the next line to code to execute is *i++*. Thus, every time *j* reaches 3, we start with the next value of *i*. So, there are no values of *j* greater than 2 output, as can be seen in the output:

```
i, j
1, 1
1, 2
2, 1
2, 2
3, 1
3, 2
```

That completes our explanations on the various looping constructs and the *break* and *continue* statements used with them. Let us now put that knowledge into practice to reinforce the concepts.

## Exercises

Now that we can iterate, it's time to do some similar tasks to the chapters before but iterate for multiple values!

Be creative on how to implement these ones and add context where you need it. As always, there's not one right answer:

1. All of our dinosaurs are unique. Okay, we cloned their DNA, but still. Let's say they have unique personalities. That's why the IDs of all our dinosaurs are unique too: they are called `dino1`, `dino2`, `dino3`, and so on. Write a `for` loop that prints out the IDs of the first 100 dinosaurs in the park.
2. Some of our dinosaurs have large appetites! Write a `do-while` loop that continues to feed a dinosaur until it is no longer hungry.
3. We all love the thrill of waiting for the park to open. Use a `while` loop to print out a countdown to the park's opening time.
4. For planning purposes, it's essential to know the total weight of all dinosaurs in a specific enclosure. Write a `for` loop that calculates this.
5. Ticket selling can get hectic during the peak season. Write a `while` loop that simulates the park's ticket-selling process until tickets are sold out.
6. Security is our topmost priority. Use a `do-while` loop to simulate a security check process that continues until all security measures are met.

## Project – Dino meal planner

Dinosaurs are not easy animals to keep. This is very advanced pet ownership. The right nutrition is difficult to manage, but it's vital to their health and well-being. Therefore, you are asked to create a system that can manage the feeding schedule of our various dinosaur residents.

The project's primary goal is to create a program that calculates the meal portions and feeding times for each dinosaur. Since we haven't covered arrays yet, we'll focus on a single dinosaur for now.

Here's how we can do it:

1. Start by declaring a variable to hold the current time; let's say it's an integer and it goes from 0 (midnight) to 23 (last hour of the day).
2. Define variables for each dinosaur species with different feeding times. For example, T-Rex could eat at 8 (morning), 14 (afternoon), and 20 (evening), while the Brachiosaurus could eat at 7 (morning), 11 (mid-morning), 15 (afternoon), and 19 (evening).

3. Next, establish a conditional statement (such as an `if-else` block) to check whether it's feeding time for each species, comparing the current time with their feeding times.
4. Now, let's define the feeding portions for our dinosaurs. We can assume that each species requires a different amount of food, depending on their sizes. For instance, the T-Rex requires 100 kg of food per meal, while the Brachiosaurus requires 250 kg of food per meal.
5. Similarly, using an `if-else` block, check which species you are dealing with and assign the food portions accordingly.
6. Finally, print the result. For instance, "It's 8:00 - Feeding time for T-Rex with 100kg of food".
7. Wrap all of the preceding information inside a loop that runs from 0 to 23, simulating the 24 hours in a day.

On behalf of the hungry dinosaurs in the park: thank you so much for putting your Java skills to use!

## Summary

In this chapter, we discussed how Java implements iteration (looping). We started with the `while` loop, which, because the condition is at the start of the loop, will execute zero or more times. In contrast, the `do-while` loop, where the condition is at the end of the loop, will execute one or more times. The `while` and `do-while` loops are very useful when you do not know how many times a loop will iterate.

In contrast, the traditional `for` loop is extremely useful when you do know how often you want a loop executed. The traditional `for` loop's header consists of three parts: the initialization section, the boolean expression, and the increment/decrement section. Thus, we can iterate a discrete number of times. This makes the traditional `for` loop ideal for processing arrays.

The enhanced `for` (`for-each`) loop is even more suitable for processing arrays (and collections), provided you are not interested in the current loop iteration index. Being concise, succinct, and easy to write, it is a more elegant `for` loop.

In effect, if you need to loop a specific number of times, use the traditional `for` loop. If you need to process an array/collection from the beginning all the way through to the end, with no concern for the loop index, use the enhanced `for` version.

All loops can, of course, be nested, and we looked at one such example. We defined a label as a case-sensitive identifier followed by a colon that immediately precedes a loop.

---

Nested loops and labels prepared us for our discussion regarding the `break` and `continue` keywords. Where `break` can also be used in a `switch` statement, `continue` can only be used inside loops. There are labeled and non-labeled versions of both. Regarding `break`, the unlabeled version exits the current loop, whereas the labeled version exits the identified loop. With regard to `continue`, the unlabeled version continues with the next iteration of the current loop, whereas the labeled version continues with the next iteration of the identified loop.

That completes our discussion on iteration. In this chapter, we touched upon arrays. Moving on to our next chapter, *Chapter 6*, we will cover arrays in detail.

Trial Version



Wondershare  
**PDFelement**

# 6

# Working with Arrays

Arrays are an essential data structure that you can use to store multiple values in one variable. Mastering arrays will not only make your code more organized and efficient but also open the door to more advanced programming techniques. Once you add arrays to the mix, you can level up the data structures of your applications.

In this chapter, we'll explore arrays and equip you with the skills needed to effectively work with this fundamental data structure. You'll learn how to create, manipulate, and iterate over arrays to solve a wide range of programming challenges.

Here's an overview of what we'll cover in this chapter:

- What arrays are and how to use them
- Declaring and initializing arrays
- Accessing array elements
- Getting the length of an array and understanding the bounds
- Different ways to loop through arrays and process their elements
- Working with multidimensional arrays
- Performing common operations with arrays using the `Arrays` class

By the end of this chapter, you'll have a solid foundation in working with arrays, enabling you to tackle more complex programming tasks with confidence. So, let's dive in!

## Technical requirements

The code for this chapter can be found on GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects/tree/main/ch6>.



## Arrays – what, when, and why?

So far, we've only seen single values, such as `int`, `double`, and `String`. Imagine we want to calculate an average result. That would look something like this:

```
double result1 = 7.0;
double result2 = 8.6;
double result3 = 9.0;

double total = result1 + result2 + result3;
double average = total / 3;
System.out.println(average);
```

This code isn't very scalable. If we were to add a fourth result, we would need to do three things in order to make this work:

- Declare and initialize a fourth variable
- Add this fourth variable to the total
- Divide by 4 instead of 3

This is a hassle, and it is error-prone. If we knew arrays, we could alter this by only changing one element of our code. Let's see what arrays are. Then, we will rewrite this example once we get to iterate over arrays.

### Java can't do basic math!?

If you were to run the previous code snippet, you'd see something interesting. If I asked you to calculate the average, you'd say 8.2, and you would be right. If we ask Java to do it, it says 8.200000000000001.

You may wonder whether there is any use in learning Java at all if it can't do basic calculations. This is not just a Java problem; this is a general computer problem. It has to translate decimal numbers into binary numbers – much like you can't express  $\frac{1}{3}$  in decimal numbers exactly (0.33333).

## Arrays explained

Alright, so **arrays** can be a solution to structure our code better in specific situations. But what are they? An array is a data structure that can store a fixed-size, ordered collection of elements of the same data type. The elements in an array are stored in contiguous memory locations, making it easier for the computer to access and manipulate the data.

So far, we haven't seen a lot of situations yet where we would need them. We are really going to level up the complexity of our logic now as we learn how to work with arrays.

## When to use arrays

So, let's talk about when to use arrays. In our example earlier, where we calculated the average, an array would mean we wouldn't need three separate variables to store our three results. We would store them in one variable of the double array type instead. This makes it easier to handle the data.

Arrays (as well as other types of ways to store multiple values in one variable, which we'll see later) are used for various reasons:

- **Organizing data:** Arrays can help organize and manage large amounts of data in a structured way
- **Simplifying code:** Using arrays can simplify your code by reducing the number of variables needed to store and manipulate data
- **Improving performance:** Accessing and modifying elements in an array is faster than using other data structures because elements are stored in contiguous memory locations

Being able to work with arrays is going to be a great tool in your Java toolbox! Let's see how we can declare and initialize them.

## Declaring and initializing arrays

There are different ways to declare and initialize arrays in Java. What you'll need will depend a lot on the specific situation. So, let's just start with the basics of declaring arrays.

### Declaring arrays

To declare an array in Java, you need to specify the data type of the elements, followed by square brackets ( [ ] ) and the array's name. Take the following example:

```
int [] ages;
```

Here, `int []` is the data type of the array, and `ages` is the name of the array. Right now, we can't add any values to the array, because it hasn't been initialized yet. This is different from initializing variables, which we have seen so far. Let's see how to initialize arrays next.

### Initializing arrays

After declaring an array, it needs to be initialized. We do this by specifying its size and allocating memory for the elements. We can use the `new` keyword to do this, followed by the data type, and then specify the size of the array inside the square brackets. Take the following example:

```
ages = new int [5];
```

This code initializes the `ages` variable to hold an array of integers with a size of 5.

We can also declare and initialize an array in a single line of code:

```
int [] ages = new int [5];
```

Here, we first declare the array on the left-hand side and initialize it on the right-hand side. We can also assign its values directly with a special short syntax, which we will explore next.

## Short syntax for array initialization

We can use Java's shortcut syntax for declaring and initializing arrays with specific values. Instead of declaring and initializing the array separately, we can use curly braces ({} ) to specify the elements directly. Take a look at the following example:

```
int [] ages = {31, 7, 5, 1, 0};
```

This code creates an array of integers and initializes it with the specified values. The size of the array is determined by the number of elements inside the curly braces.

Actually, our previous arrays had values already as well, because when you create an array using the new keyword, Java automatically initializes the elements with default values based on their data type. The default values are as follows:

- Numeric types (byte, short, int, long, float, double): 0 or 0.0
- char: '\u0000' (the Unicode null character)
- boolean: false
- Reference types (objects and arrays): null

For example, say you create an array of integers with a size of 3:

```
int [] results = new int [3];
```

Java initializes the elements with the default value of 0, because int is numeric. So far, we have seen how to declare and initialize arrays. It's now time to learn how to access the elements in an array and update the values.

## Accessing elements in an array

In order to access elements in an array, we need to use their **index**. The index represents the position in the array. This allows us to retrieve the value at a certain position and assign it a new value. Let's first talk about indexing.

## Understanding indexing

In Java, arrays use zero-based indexing, which means the first element has an index of 0, the second element has an index of 1, and so on. Take a look at our example of the `ages` array:

```
int[] ages = {31, 7, 5, 1, 0};
```

This means that the first element (31) has an index of 0 and the last element has an index of 4.

Index:	0	1	2	3	4
	31	7	5	1	0

Figure 6.1 – Indexing explained with the `ages` array

We count the length of an array like we normally do, starting with 1. So, the length of this array would be 5. The last element in the array has an index equal to the array's length minus 1. For an array with a length of N, the valid indexes are in the range of 0 to N-1.

It is important to know how to use the index because that way we can access the elements in the array.

## Accessing array elements

To access an element in an array, you can use the array's name, followed by the index of the desired element inside square brackets. For example, to access the first element of our array named `ages`, you can use the following code:

```
int maaikesAge = ages[0];
```

This will store the value 31 in the `age` variable. In order to access the second element, you'd have to do the following:

```
int gaiasAge = ages[1];
```

We can also access the element and store another value in the element using the index.

### Printing arrays

If we print the variable holding the array, we can get something like this: [I@28a418fc

This is not going to be very helpful. So, mind that you're printing what the `toString()` method is returning. This is not customized for the array and is not very useful. What we most likely want to see is the elements inside the array. There is a way to print the content of arrays. We'll see this when we cover the built-in methods for dealing with arrays.



## Modifying array elements

Modifying the elements is also done with the index. It looks a lot like assigning a variable as we did before. For example, to change the value of the last element in our array, named `ages`, we can use the following code:

```
ages[4] = 37;
```

We can only access elements that are there. If we try to get an element that is not there, we get an exception (error) message.

## Working with length and bounds

To avoid getting exceptions, we need to stay within the bounds of the array. Indexes always start at 0, and they end at the length of the array minus 1. If you try to access an element outside this range, you'll get `ArrayIndexOutOfBoundsException`. The key to avoiding this is working with the length of the array.

## Determining the length of an array

We can determine the length of an array using the `length` property. The `length` property returns the number of elements in the array. For example, to get the length of our `ages` array, we can use the following code:

```
int arrLength = ages.length;
```

The length of the array starts counting at 1. Therefore, the length of our `ages` array is 5. The maximum index is 4.

## Dealing with the bounds of an array

If you try to access or modify an array element using an invalid index (an index that is less than 0 or greater than or equal to the array's length), Java throws `ArrayIndexOutOfBoundsException`. This exception is a runtime error, which means it occurs when the program is running, not when we compile it. We'll learn more about exceptions in *Chapter 11*.

To prevent `ArrayIndexOutOfBoundsExceptions`, we should always validate array indexes before using them to access or modify array elements. We can do this by checking whether the index is within the valid range (0 to array length - 1). Here's an example that demonstrates how to validate an array index:

```
String[] names = {"Maria", "Fatiha", "Pradeepa", "Sarah"};
int index = 5;
if (index >= 0 && index < names.length) {
```

```
        System.out.println("Element at index " + index + ": " +
                           names[index]);
    } else {
        System.out.println("Invalid index: " + index);
    }
```

The output will be as follows:

```
Invalid index: 5
```

This code snippet checks whether the index is within the valid range before accessing the array element. If the index is invalid, the program prints an error message instead of throwing an exception.

We can also use the loops we learned about in the previous chapter to iterate over the elements in an array and access or modify their values.

## Iterating over arrays

There are different methods to iterate over arrays. We will have a look at the use of the traditional `for` loop and the enhanced `for` loop (also known as the `for-each` loop).

### Using the `for` loop

We can use the traditional `for` loop to iterate over an array by using an index variable. The loop starts at index 0 and continues until the index reaches the length of the array. Here's an example that demonstrates how to use a `for` loop to iterate over an array and print its elements:

```
int[] results = {10, 20, 30, 40, 50};
for (int i = 0; i < results.length; i++) {
    System.out.println("Element at " + i + ": " +
                       results[i]);
}
```

The output will be as follows:

```
Element at 0: 10
Element at 1: 20
Element at 2: 30
Element at 3: 40
Element at 4: 50
```

At this point, we know enough to revisit the example that we saw at the beginning of the chapter, calculating the average of several results. Instead of having separate primitives, we're now going to have an array. Here is what it will look like:

```
double[] results = {7.0, 8.6, 9.0};  
double total = 0;  
for(int i = 0; i < results.length; i++) {  
    total += results[i];  
}  
double average = total / results.length;  
System.out.println(average);
```

If we now want to add a result, we only need to alter it in one place. We just add the result to the `results` array. Since we loop over all the elements, we don't need to add an extra variable to calculate the total result. Also, since we use the length, we don't need to change 3 to 4.

We can also use loops to modify the values of the array. Here's an example that demonstrates how to double the value of each element in an array using a `for` loop:

```
int[] results = {10, 20, 30, 40, 50};  
// Double the value of each element  
for (int i = 0; i < results.length; i++) {  
    results[i] = results[i] * 2;  
}  
// Print the updated array elements  
for (int i = 0; i < results.length; i++) {  
    System.out.println("Element at " + i + ": " +  
        results[i]);  
}
```

The output will be as follows:

```
Element at 0: 20  
Element at 1: 40  
Element at 2: 60  
Element at 3: 80  
Element at 4: 100
```

As you can see, the elements in the array are doubled in the first `for` loop. In the second `for` loop, they are printed. As you can tell by the output, the values did double!

Let's have a look at the enhanced `for` loop and how we can use that to iterate over arrays.

## Using the for each loop

We can also use the `for-each` loop, also known as the enhanced for loop, to iterate over arrays. This special `for` loop simplifies the process of iterating over arrays (and other iterable objects). The `for-each` loop automatically iterates over the elements in the array and does not require an index variable. Here's an example that demonstrates how to use the `for-each` loop to iterate over an array and print its elements:

```
int[] results = {10, 20, 30, 40, 50};
for (int result : results) {
    System.out.println("Element: " + result);
}
```

The output will be as follows:

```
Element: 10
Element: 20
Element: 30
Element: 40
Element: 50
```

The `for-each` loop requires a temporary variable that is used to store the current element during each iteration. In our example, this is `int result`. It is logical to call it `result`, since it is one element in the `results` array. But this is not necessary for the functionality; I could have also called it `x`, as follows:

```
int[] results = {10, 20, 30, 40, 50};
for (int x : results) {
    System.out.println("Element: " + x);
}
```

The output would have been exactly the same. I like to read the line of code that says `for (int x : results)` in my head like this: for every element `x` in `results`, do whatever is in the code block.

So there are two ways to loop over arrays, let's talk about which one to choose when.

## Choosing between the regular loop and the enhanced for loop

We can use the regular `for` loop and the (enhanced) `for-each` loop to iterate over an array. These two approaches have some differences and there's a reason for choosing one or the other.

When you need to have the index available, you should use the traditional `for` loop since this uses an index variable to access the elements in the array, while the `for-each` loop directly accesses the elements without using an index variable.



The `for-each` loop does not allow you to modify the array elements during iteration, as it does not provide access to the index variable. If you need to modify the array elements during iteration, you should use the traditional `for` loop.

If you only want to read the variables and you don't need the index, you typically want to go for the `for-each` loop because the syntax is easier.

Alright, so now we know how to iterate over arrays. Let's make the data structure slightly more complicated and learn about multidimensional arrays.

## Handling multidimensional arrays

A **multidimensional array** is an array of arrays. In Java, you can create arrays with two or more dimensions. The most common type of multidimensional array is the two-dimensional array, also known as a matrix or a table, where the elements are arranged in rows and columns.

Let's see how to create multidimensional arrays.

### Declaring and initializing multidimensional arrays

To declare a two-dimensional array, you need to specify the data type of the elements, followed by two sets of square brackets (`[] []`) and the name of the array. Take the following example:

```
int [] [] matrix;
```

Just like the one-dimensional array, we initialize a two-dimensional array with the use of the `new` keyword, followed by the data type and the size of each dimension inside the square brackets, like this:

```
matrix = new int [3] [4];
```

This code initializes a matrix of 3 rows and 4 columns. The type is `int`, so we know that the values of the matrix are integers.

We can also declare and initialize a multidimensional array in a single line:

```
int [] [] matrix = new int [3] [4];
```

We can use the short syntax as well. To initialize a multidimensional array with specific values, we use the nested curly braces (`{ }` ):

```
int [] [] matrix = {  
    {1, 2, 3, 4},  
    {5, 6, 7, 8},  
    {9, 10, 11, 12}  
};
```

Just like the one-dimensional arrays, Java determines the length by the provided values. This matrix has three inner arrays (three rows) each with four elements (four columns). Accessing and modifying the elements in a multidimensional array is similar, but we now need to provide two indices.

## Accessing and modifying elements of multidimensional arrays

To access or modify the elements of a multidimensional array, you need to specify the indexes of each dimension inside square brackets. For example, to access the element in the first row and second column of a two-dimensional array named `matrix`, you can use the following code:

```
int element = matrix[0][1];
```

To modify the same element, you can use the following code:

```
matrix[0][1] = 42;
```

*Figure 6.2 shows how the indexing works for our two-dimensional array, `matrix`.*

Index:	0	1	2	3
0	1	2	3	4
1	5	6	7	8
2	9	10	11	12

Figure 6.2 – The index of the rows and columns for array `matrix`

So, if we want to get to the element with value 12 and store it in a `last` variable, our code will be as follows:

```
int last = matrix[2][3];
```

We can also iterate over all the variables in a multidimensional array. Let's see how that is done.

## Iterating over multidimensional arrays

Since a multidimensional array is just an array in an array, we can use nested loops to iterate over multidimensional arrays. Here's an example that demonstrates how we can use a nested `for` loop to iterate over a two-dimensional array:

```
int[][] matrix = {
    {1, 2, 3, 4},
    {5, 6, 7, 8},
    {9, 10, 11, 12}
};

for (int i = 0; i < matrix.length; i++) {
    for (int j = 0; j < matrix[i].length; j++) {
        System.out.print(matrix[i][j] + " ");
    }
    System.out.println();
}
```

The output will be as follows:

```
1 2 3 4
5 6 7 8
9 10 11 12
```

All we do at this point is just print the element. This is something we can also do with the enhanced `for` loop to iterate over multidimensional arrays. Here's an example that demonstrates how to do that:

```
int[][] matrix = {
    {1, 2, 3, 4},
    {5, 6, 7, 8},
    {9, 10, 11, 12}
};

for (int[] row : matrix) {
    for (int element : row) {
        System.out.print(element + " ");
    }
    System.out.println();
}
```

The output will be the same as in the previous example:

```
1 2 3 4
5 6 7 8
9 10 11 12
```

As you can see, the outer `for-each` loop iterates over the rows of the two-dimensional array. The row is an array itself as well, which is why the type is `int []`. The inner `for-each` loop iterates over the elements within each row. These are integers.

Both traditional nested `for` loops and nested `for-each` loops can be used to iterate over multidimensional arrays. It's a matter of preference and whether you need to access the index of the elements.

Arrays can go very many levels deep, but that doesn't really change the basic principles. For example, for a four-dimensional array, you'll have `[] [] [] []` behind the type and you need a nested loop of four levels to iterate over all the elements.

Java helps us deal with arrays in different ways. Let's look at some built-in methods for arrays that we can use.

## Using Java's built-in methods for arrays

Working with arrays is very common. Usually, for very common things, Java has built-in functionality. We can do many common things we'd like to do with arrays with the use of the methods on the built-in `Arrays` class.

### Built-in Arrays class for working with arrays

The built-in `Arrays` class is a helper class in the `java.util` package. It offers many utility methods to help us efficiently work with arrays. We'll explore some common array manipulation tasks using the `Arrays` class.

#### *The `toString()` method*

A highly useful operation you may want to perform on an array is to convert it into a `String`, which can be invaluable for debugging and logging purposes. To achieve this, the `Arrays` class offers a dedicated method called `toString()`. It's important to note that this method is static, allowing us to call it directly on the `Arrays` class.

```
import java.util.Arrays;

public class ArrayHelperMethods {
    public static void main(String[] args) {
```



```
        int[] results = {30, 10, 50, 20, 40};

        // Convert the array to a string representation
        String arrayAsString = Arrays.toString(results);
        System.out.println("Array: " + arrayAsString);
    }
}
```

The output will be as follows:

```
Array: [30, 10, 50, 20, 40]
```

As you can see, the `results` array is converted to a string that represents the array's elements, enclosed by square brackets and separated by commas. There are many such methods on the `Arrays` class! Let's explore the `sort` method next.

### ***The `sort()` method***

A common operation you want to do on an array is to sort the array. Here's an example that shows how to sort the values of an array with the `sort` method from the `Arrays` class:

```
import java.util.Arrays;

public class ArrayHelperMethods {
    public static void main(String[] args) {
        int[] results = {30, 10, 50, 20, 40};

        // Sort the array
        Arrays.sort(results);
        System.out.println(Arrays.toString(results));
    }
}
```

The output will be as follows:

```
[10, 20, 30, 40, 50]
```

As you can see, the `results` array is unsorted at first. We can call the methods on the `Arrays` class directly on the `Arrays` class because they're static. For integer values, it sorts them from low to high by default. We can alter this behavior, but we don't have the knowledge we need to do that just yet.

We print the array with another built-in method, namely `toString`. This translates the array into something that we can understand.

When the array is sorted, we can use the `binarySearch` method to find a value.

### The `binarySearch()` method

We can also search for a value in an array. We are going to use the built-in `binarySearch` method to do this. Very importantly, this can only be done with sorted arrays because of how the search algorithm works. Here's an example of how to do this:

```
import java.util.Arrays;

public class ArrayHelperMethods {
    public static void main(String[] args) {
        int[] results = {10, 20, 30, 40, 50};

        int target = 30;
        int index = Arrays.binarySearch(results, target);
        System.out.println("Index of " + target + ": " +
                           index);
    }
}
```

The output will be as follows:

```
Index of 30: 2
```

The `binarySearch` method requires the input array to be sorted beforehand. The `binarySearch` algorithm is meant for finding a target value within a sorted array. Instead of searching the array element by element, it divides the array into halves repeatedly until it finds the target or the remaining portion to search becomes empty. When the value at the half is bigger, it knows it needs to move towards the left side of the array, when it's smaller it knows it needs to move towards the right. That's why it's a must that the array is sorted. The `binarySearch` method returns the index of the target value if found. If the target wasn't found, it returns a negative value, which represents the insertion point. So, say we updated our code to this:

```
int[] results = {10, 20, 30, 40, 50};

int target = 31;
int index = Arrays.binarySearch(results, target);
System.out.println("Index of " + target + ": " +
                           index);
```

This would result in the following:

```
Index of 31: -4
```

This is because it would have been at the fourth position in the array (not the fourth index!).

Let's see how we can give all the elements in the array a specific value with the `fill` method.



### The `fill()` method

Sometimes, you want to create an array of the same values programmatically. Here's an example of how this can be done. We use the `fill` method from the `Arrays` class. Here's how to do it:

```
import java.util.Arrays;

public class ArrayHelperMethods {
    public static void main(String[] args) {
        int[] results = new int[5];

        Arrays.fill(results, 42);
        System.out.println(Arrays.toString(results));
    }
}
```

The output will be as follows:

```
[42, 42, 42, 42, 42]
```

The `fill` method sets all elements in the array to the specified value. Sometimes we need to create a copy of our array or resize it. In that case, we can use the `copyOf` method.

### The `copyOf()` method

Sometimes you need to create a copy of an array, for example, when you want to end it to another location in the application, but you don't want this to affect your original array.

This is an example of how we can create a copy of an array:

```
import java.util.Arrays;

public class ArrayHelperMethods {
    public static void main(String[] args) {
        int[] results = {10, 20, 30, 40, 50};

        int[] copiedResults = Arrays.copyOf(results,
            results.length);
        System.out.println(Arrays.toString(copiedResults));
    }
}
```

The output will be as follows:

```
[10, 20, 30, 40, 50]
```

We can prove we copied the array with the following code:

```
copiedResults[0] = 1000;  
System.out.println(Arrays.toString(copiedResults));  
System.out.println(Arrays.toString(results));
```

If we didn't create a copy but just stored it in another variable instead, it would alter both arrays. The preceding code snippet will give the following output:

```
[1000, 20, 30, 40, 50]  
[10, 20, 30, 40, 50]
```

But say we have this code:

```
int[] copiedResults = results;  
  
copiedResults[0] = 1000;  
System.out.println(Arrays.toString(copiedResults));  
System.out.println(Arrays.toString(results));
```

It will give us this output:

```
[1000, 20, 30, 40, 50]  
[1000, 20, 30, 40, 50]
```

As you can see, this alters both the variables holding the array. This is because both variables, `copiedResults` and `results`, have the same array object that they're pointing to. So, if you change it in one place, it changes for both. That's why you sometimes need to create copies of arrays.

So, the `copyOf` method creates a new array with the same elements as the original array, whereas this second method just creates a new variable that points to the same array object. We can also use it to resize the array by passing in a second argument.

### ***Resizing arrays with `copyOf()`***

Arrays have a fixed size, but sometimes you need to alter the size nonetheless. The `Arrays.copyOf()` method that we just saw is also useful for resizing arrays. To resize an array, you can create a new array with the desired size and copy the elements from the original array to the new array. All you need to do is give it a second argument.

Here's an example that demonstrates how to resize:

```
import java.util.Arrays;  
int[] originalArray = {10, 20, 30, 40, 50};  
int newLength = 7;
```



```
int[] resizedArray = Arrays.copyOf(originalArray, newLength);

System.out.println("Original array: " + Arrays.
toString(originalArray));
System.out.println("Resized array: " + Arrays.toString(resizedArray));
```

The output will be as follows:

```
Original array: [10, 20, 30, 40, 50]
Resized array: [10, 20, 30, 40, 50, 0, 0]
```

In this example, we resized `originalArray`, which had a length of 5, to a new length of 7. The new array contains the elements of the original array, followed by default values (0 for `int`) to fill the remaining positions.

This is not something you should be doing constantly. It can be inefficient in terms of performance. If you would need to resize your array a lot, it's worth having a look at *Chapter 13* where we learn about **collections**.

### ***The equals() method***

The last built-in method we're going to discuss is the `equals()` method. This method can determine whether two arrays have the same values. With this built-in method, you can compare two arrays for equality. Here's how to go about it:

```
import java.util.Arrays;

public class ArrayHelperMethods {
    public static void main(String[] args) {
        int[] results1 = {10, 20, 30, 40, 50};
        int[] results2 = {10, 20, 30, 40, 50};

        boolean arraysEqual = Arrays.equals(results1,
            results2);
        System.out.println("Are the arrays equal? " +
            arraysEqual);
    }
}
```

The output will be as follows:

```
Are the arrays equal? true
```

The `equals()` method compares two arrays element by element to check whether they have the same values in the same order. It returns `true` if the arrays are equal; otherwise, it returns `false`.

**Well done!**

You've done a great job learning arrays! At this point, you're ready to understand this programming joke:

Why did the Java developer quit their job?

Because they couldn't get "arrays!"

## Exercises

Arrays are incredibly useful for storing and managing similar types of data, such as a list of dinosaur names, dinosaur weights, and visitors' favorite snacks. Arrays are helpful and they enable us to manage more complex data in Mesozoic Eden. Try out the following:

1. The unique appeal of our park lies in the diversity of our dinosaur species. (And also in that we have dinosaurs at all.) Create an array that holds the names of all the dinosaur species in the park. This list will help us in inventory management.
2. Every visitor has their favorite dinosaur, and for many, it's the heaviest one. Write a program that finds this star's weight in an array of dinosaur weights. This information can then be highlighted in our park tours and educational programs.
3. Dinosaurs come in all sizes, and the smallest ones hold a special place in the hearts of children. Write a program that finds this smallest dinosaur in an array of dinosaur weights.
4. Running a dinosaur park is not a one-man show and requires a dedicated team of employees. Create an array of park employee names and print out the names using an enhanced `for` loop. This will help us to appreciate and manage our staff more effectively.
5. To ensure the well-being of our dinosaur inhabitants, it's essential to monitor their average age. This data can help inform our care and feeding programs to better suit the age profile of our dinosaurs. Write a program that calculates this using an array of dinosaur ages.
6. Our park is meticulously divided into various sections to facilitate visitor navigation and dinosaur housing. Create a two-dimensional array representing the park map, with each cell containing an array of Strings indicating an enclosure or facility for a certain section.
7. The enjoyment of a park tour depends significantly on comfortable seating arrangements. Use nested loops to print out a seating chart for a park tour bus from a two-dimensional array. This will help us ensure that every guest has a pleasant journey throughout the park.

## Project – Dino tracker

Safety always comes first. That's why keeping track of all our dinosaur residents is of utmost importance. The park managers need to have an easy-to-use system for managing information about their slightly exotic pets.

For this project, you'll be creating a Dino tracker. This is a simple tracking system that maintains records of each dinosaur's name, age, species, and enclosure number. This will be done using fixed arrays – four arrays in total, one for each attribute.

Assume you have room for 10 dinosaurs in your park for now, so each array should have a length of 10. Each dinosaur will correspond to an index in the array. For example, if the dinosaur "Rex" is in the first position of the name array, his age, species, and enclosure number will also be in the first position of their respective arrays.

You're going to print information about all the dinosaurs and print their average age and weight after that.

I realize this might be a lot. If you need some extra guidance, here are some steps to guide you through the process:

1. **Initialization:** Start by creating four arrays: `dinoNames`, `dinoAges`, `dinoSpecies`, and `dinoEnclosures`. Each should have a size of 10.
2. **Data entry:** Manually enter the details for ten dinosaurs into the arrays. This is to populate your arrays with some initial data. If you're feeling lazy, like the Brachiosaurus, you could have `Dinosaur1`, `Dinosaur2`, and so on as names.
3. **Displaying details:** Write a loop that goes through the arrays and prints out the details of each dinosaur in a readable format.
4. **Average calculations:** Add the end, print the average age and weight of the dinosaurs. For the ages, you will need to sum up all the ages in the `dinoAges` array and divide by the number of dinosaurs. And, of course, this process is similar for weight, but using the weight array.

## Summary

In this chapter, we have explored arrays in Java. Arrays are data structures that allow us to store multiple values of the same data type in a contiguous block of memory. They provide an efficient way to organize lists of data.

We began by discussing the declaration and initialization of arrays. We learned about different ways to declare and initialize arrays, including using the shortcut syntax for array initialization. We also covered how to initialize arrays with default values.

After that, we discussed how to access and modify array elements using indexes. We learned about the importance of the array length and that we can find out the length by using the `length` property. We also talked about avoiding `ArrayIndexOutOfBoundsExceptions` by validating array indexes.

We then looked at iterating over arrays using both the traditional `for` loop and the enhanced `for` loop (the `for-each` loop).

After this, we explored multidimensional arrays, which are arrays of arrays, and learned how to declare, initialize, and access their elements. We also discussed how to iterate over multidimensional arrays.



---

Finally, we covered common array operations with the use of the `Arrays` class and its built-in methods. We saw how to sort arrays, search for elements in a sorted array, fill an array with a specific value, copy and resize an array, and compare arrays.

By mastering these concepts, you now have a solid foundation for working with arrays in Java. This understanding will help you store and manipulate data more efficiently in your Java programs. We ended by looking at some built-in methods. In the next chapter, you're going to learn how to write your own methods.

Trial Version



Wondershare  
**PDFelement**



## 7

# Methods

In *Chapter 6*, we learned about arrays in Java. We learned that arrays are data structures that are fixed in size. They are stored in contiguous memory locations where each location is of the same type. We also saw how to declare, initialize, and process arrays. Both the traditional and enhanced `for` loops are ideal for processing arrays.

In addition, we discussed multi-dimensional arrays, including how they are organized and how to process them. Lastly, as arrays are very common, we discussed the `Arrays` class, which has several useful methods for processing arrays.

In this chapter, we will cover methods. Methods enable us to create a named block of code that can be executed from elsewhere in the code. Firstly, we will explain why methods are so commonplace. You will learn the difference between the method definition and the method invocation. We will explore what a method signature is and how method overloading enables methods to have the same name, without conflict. We will also explain variable arguments (`varargs`), which enable a method to be executed with 0 or more arguments. Lastly, Java's principle of call-by-value for passing arguments (and returning values) will be outlined. By the end of this chapter, you will be well able to code and execute methods. In addition, you will understand method overloading, `varargs`, and Java's call-by-value mechanism.

This chapter covers the following main topics:

- Explaining why methods are important
- Understanding the difference between method definition and method execution
- Exploring method overloading
- Explaining `varargs`
- Mastering call by value

## Technical requirements

The code for this chapter can be found on GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects/tree/main/ch7>.

## Explaining why methods are important

Methods are code blocks that are given a name for ease of reference. They can accept inputs and return an output. Both the inputs and output are optional. A method should do one task and do it well. It is considered good practice to keep your methods short (less than 20 lines). The longer the method, the more likely it is that the method is doing too much. The maxim of “keep it simple” applies here.

## Flow of control

Simply put, when a method is called (executed), the normal flow of control of execution is changed. Let us discuss a simple example that will help demonstrate this. This is an important point to appreciate, especially for inexperienced developers. *Figure 7.1* presents the code:

```
3 > public class Methods {
4   >   public static void main(String[] args) {
5     System.out.println("main: before call to simpleExample()");
6     simpleExample(); // method call
7     System.out.println("main: after call to simpleExample()");
8   }
9   >   public static void simpleExample(){ // method definition
10    System.out.println("\tExecuting simpleExample() method...");
11  }
12 }
```

Figure 7.1 – A very simple method

In this example, we have two methods: the `main()` method (lines 4 to 8) and the `simpleExample()` method (lines 9 to 11). Both exist inside the `Methods` class (lines 3 to 12).

In Java, every program starts with the `main()` method. The JVM calls it on our behalf; we do not have to call (or execute) it ourselves. Therefore, in this example, the first line in `main()`, line 5, is the first line to execute.

Line 6 is important – it is what we refer to as a method call. There is a direct correlation between the `simpleExample()` method definition on line 9 and the method call on line 6. We will discuss this relationship shortly. For the moment, just understand that the method call changes the order of execution of the program. Normally, Java executes lines of code from top to bottom and this is true. However, *method calls alter that*. In this example, when line 6 executes, the next line to execute is line 10 (inside the `simpleExample()` method).

So, the `main()` method has now handed over control to the `simpleExample()` method, and control will not return to `main()` until `simpleExample()` exits. This can occur when execution hits the closing } at the end of the `simpleExample()` method (line 11). This is what happens in this example. Alternatively, a method can exit by using the `return` keyword.

So, line 6 calls the `simpleExample()` method, causing its code to execute. Line 10 outputs some text to the screen. The closing } on line 11 causes `simpleExample()` to exit and control now returns to `main()`, where execution resumes at line 7.

In summary, the order of execution in this program is illustrated by the output:

```
main: before call to simpleExample()
    Executing simpleExample() method...
main: after call to simpleExample()
```

Here, you can see that `println()`, inside the `simpleExample()` method, is sandwiched between the two `println()` statements from `main()`. This demonstrates that the flow of control was altered by the method call on line 6.

### The stack

So, how can a caller method, such as `main()`, simply *resume* where it left off after the `simpleExample()` method returns? What about the local variables of `main()`?

The ability of a method to resume exactly where it left off, after the method it called returns, requires the use of a memory structure called the *stack*. We will discuss the stack later in this chapter.

Returning to our *why methods are important*' discussion, two major advantages of methods are that they provide abstraction and avoid code duplication. Let's examine these in turn.

## Abstraction

Abstraction is a principle in software engineering where clients of a service, are abstracted from the service implementation. This decouples clients, who use the service, from knowing how the service is implemented. Thus, if the service implementation is changed, the clients are not impacted.

Take, for example, a McDonald's drive-thru where you drive up and place your order. In this situation, you are the client of the McDonald's service. You do not care how McDonald's process your order; you simply want to place an order and receive the food/drinks. If McDonald's changes its internal implementation, you are shielded (abstracted) from those changes. This is known as abstraction.

For our purposes, the method itself is the McDonald's service. The method call is the McDonald's customer. The method call is abstracted from internal changes to the method code.

## Code duplication

Methods can help us avoid code replication. This has the added benefit of easing debugging. Let's look at a simple example of this. *Figure 7.2* demonstrates duplicated code:



```
6 Scanner sc = new Scanner(System.in);
7
8 System.out.print("Enter a number (1..10) --> ");
9 int number = sc.nextInt();
10 if(number < 1 || number > 10){
11     System.out.println("Invalid number! "+number);
12 }
13
14 System.out.print("Enter a number (1..10) --> ");
15 number = sc.nextInt();
16 if(number < 1 || number > 10){
17     System.out.println("Invalid number! "+number);
18 }
19
20 System.out.print("Enter a number (1..10) --> ");
21 number = sc.nextInt();
22 if(number < 1 || number > 10){
23     System.out.println("Invalid number! "+number);
24 }
```

Figure 7.2 – Duplicated code

In the preceding figure, lines 8 to 12 are repeated on lines 14 to 18 and lines 20 to 24. Each of these sections prompts the user for a number, stores the user input in a variable named `number`, and checks to see if the number is in range. If the number is out of range, then an error is flagged. While a loop would be an obvious improvement, bear in mind that these lines of code could well be in separate parts of the program. In addition, for this simple example, we are only interested in highlighting code duplication. We simply prompt for a number, accept the user's input, and validate it. The result is five lines of code repeated three times.

Now, let's assume that we want to adjust the upper valid range from 10 to 100. We have to change the prompts on lines 8, 14, and 20. In addition, the `if` statements on lines 10, 16, and 22 need to change. Thus, a simple range adjustment has resulted in quite a few code changes and we could easily forget to make one or more of the changes required. Let's refactor the code into a method. *Figure 7.3* shows the refactored code:

```
3 ► public class Methods {  
4 ►     public static void main(String[] args) {  
5         int number = getNumber();  
6         number = getNumber();  
7         number = getNumber();  
8     }  
9     public static int getNumber(){  
10        Scanner sc = new Scanner(System.in);  
11        System.out.print("Enter a number (1..10) --> ");  
12        int number = sc.nextInt();  
13        if(number < 1 || number > 10){  
14            System.out.println("Invalid number! "+number);  
15        }  
16        return number;  
17    }  
18}
```

Figure 7.3 – The code from Figure 7.2 refactored to use a method

In the preceding figure, the method itself is coded from lines 9 to 17 and will be explained in detail in the next section. The five lines of repeated code from *Figure 7.2* are coded only once, on lines 11 to 15. The execution calls of the method are on lines 5, 6, and 7; one execution call per line. If we want to change the upper valid range from 10 to 100, we just need to change the method – that is, lines 11 and 13. These two changes are *automatically* reflected throughout the code. In effect, the three method calls on lines 5, 6, and 7 automatically reflect the changes made in the method.

As you can imagine, this situation scales very well. For example, if, in *Figure 7.2*, we had duplicated the code 10 times, we would have to make changes in 10 areas of the code. However, with the method implementation, there is still *only one* location to make the change and that is in the method itself.

Now that we have justified why methods exist, let's examine the difference between the method itself and the method call.

## Understanding the difference between method definition and method execution

For those new to programming, it may surprise you to know that there are two parts to having a method *do something*. Firstly, we must code the method (the method definition). This is similar to a bank machine on the street – it just sits there, doing nothing, waiting to be used. Secondly, we must execute the method (the method execution). This is similar to a customer “using” the bank machine.

Remember that the `main` method is the only method that is automatically executed by the JVM. Any other method calls have to be explicitly coded.

Now, let's examine the method definition and method execution in turn.

## Method definition

The method definition (declaration) is the method code itself - this is the block of code that is executed when the method is called. *Figure 7.4* presents the syntax:

```
[access-modifier] [static] return-type methodName([parameters]) [throws someException] {
    // method code
}
```

Figure 7.4 – The syntax of the method definition

In the preceding figure, as in other figures, square brackets signify optional elements. The `access-modifier` and `static` elements will be discussed in *Chapter 8*. The `throws someException` element will be covered in *Chapter 11*. In this chapter, we will focus on the elements in bold; namely, `return-type` (mandatory), `methodName` (mandatory), and `parameters` (optional).

The return type of the method can be a primitive type, a reference type, or `void`. The `void` keyword means that the method is not returning anything. If that is the case, you *cannot* simply leave out the return type; you must specify `void`. In addition, when you're not returning anything from a method, you can specify `return ;` or simply leave out the `return` keyword altogether (which is what we have done for all the `main()` methods).

Let's examine a method that accepts input and returns a result. *Figure 7.5* presents such an example:

```
14      // method definition/declaration
15  @
16      public static int performCalc(int x, int y, String operation){ // "parameters"
17          int result = switch(operation){
18              case "+" -> x + y;
19              case "-" -> x - y;
20              case "*" -> x * y;
21              case "/" -> x / y;
22              case "%" -> x % y;
23              default -> {
24                  System.out.println("Unrecognized operation: "+operation);
25                  yield -1; // error
26              }
27          };
28          return result;
}
```

Figure 7.5 – Sample method definition

In the preceding figure, we have a method that takes in two integers and a mathematical operation to be performed using the two integers as operands. For example, if "+" is passed in, the two numbers are added and the result is returned. Let's review how the method does this.

Line 15 is very important. For the moment, as stated earlier, *Chapter 8* will explain both `public` (access-modifier) and `static`. The return-type is an `int` – meaning, this method returns whole numbers. The name of the method is `performCalc`. Method names often begin with verbs and follow camel casing style.

Note that round brackets follow the method name. The round brackets are delimiters for the optional input parameters to the method. For each parameter, you must specify the data type of the parameter (as Java is a strongly typed language) and the parameter's identifier name. If you have two or more parameters, comma-separate them. These parameters are how the method accepts input. In *Figure 7.5*, we have two integers namely `x` and `y`, followed by a `String` called `operation`. The scope of any method parameters, in this case, `x`, `y`, and `operation`, is the whole of the method.

Lines 16-26 encapsulate a `switch` expression. In effect, depending on the mathematical `operation` passed in, that operation is performed on the two inputs, `x` and `y`. The local `int` variable, `result`, is initialized accordingly. The `result` variable is returned on line 27. As the return type declared on line 15 is an `int`, returning `result`, which is also an `int`, is fine.

A method definition in and of itself does not do anything. It just defines a block of code. As stated previously, this is similar to a bank machine on the street – it just sits there, doing nothing, waiting to be used. For the bank machine to be useful, you must “use” it. Similarly, we must “use” the method – this is what we call executing the method.

## Method execution

Executing the method is also known as calling or invoking the method. The method that calls the method is known as the “calling” (or caller) method. So, you have the calling method and the called method. When you call a method, you pass down the required arguments, if there are any. The called method will execute at this point. When the called method finishes, control returns to the caller method. The called method’s result, if there is one, is also returned. This enables the called method to return data to the caller method, where it can be output to the screen, stored in a variable, or simply ignored.

### Method parameters versus method arguments

The method definition defines parameters, whereas the method call passes down arguments. These terms are often used interchangeably.

*Figure 7.6 presents a code example to help explain this further:*

```

4  public static void main(String[] args) {
5      int result = performCalc(x: 10, y: 2, operation: "+"); // method call; passing down "arguments"
6      System.out.println(result); // 12
7      System.out.println(performCalc(x: 10, y: 2, operation: "-")); // 8
8      System.out.println(performCalc(x: 10, y: 2, operation: "*")); // 20
9      System.out.println(performCalc(x: 10, y: 2, operation: "/")); // 5
10     performCalc(x: 10, y: 2, operation: "%");// return value ignored
11     System.out.println(performCalc(x: 10, y: 2, operation: "&")); // Unrecognized operation: &, -1
12 }
13 @
14 public static int performCalc(int x, int y, String operation){ // "parameters"
15     int result = switch(operation){
16         case "+" -> x + y;
17         case "-" -> x - y;
18         case "*" -> x * y;
19         case "/" -> x / y;
20         case "%" -> x % y;
21         default -> {
22             System.out.println("Unrecognized operation: "+operation);
23             yield -1; // error
24         }
25     };
26     return result;
}

```

Figure 7.6 – Sample code demonstrating method calls

### IntelliJ IDEA inlay hints

Note that the IntelliJ editor inserts inlay hints when you are coding. In the previous figure, the `performCalc` method signature (line 13) specifies that the parameters are namely `x`, `y`, and `operation`. That is why on each method call, the inlay hint uses these parameter names. For example, on line 5, we typed in 10 as the first argument; however, IntelliJ, upon inspecting the method signature, realized that 10 was mapping to 'x' and that is why you see "x:" before the 10. We did not type in "x:" at all! It is not part of the Java language to do that (IntelliJ is just trying to help us). In actual fact, for line 5, we typed in `performCalc(10, 2, "+")` and IntelliJ converted that to `performCalc(x: 10, y: 2, operation: "+")`.

In *Figure 7.6*, the `performCalc` method (lines 13-26) is unchanged from *Figure 7.5*. However, we can now see the various method calls (lines 5 and 7-11).

Let's start with line 5. On the right-hand side of the assignment, we have the `performCalc(10, 2, "+")` method call. This method call has higher precedence than the assignment, so it is executed first. The IntelliJ IDE does a very nice job of highlighting that 10 will be passed into the method as `x`, 2 will be passed into the method as `y`, and "`+`" will be passed in as `operation`. It is very important

to realize that once we get to the method call on line 5, the next line of code that's executed is line 14 – so, from line 5, we jump into the `performCalc` method and start executing the `switch` expression on line 14.

Since `operation` is "+" for this method invocation, line 15 assigns  $10 + 2$  (12) to `result`. Line 25 returns `result` back to the calling method (line 5), where the value 12 is assigned into `result`. Line 6 outputs the return value from the `performCalc` invocation on line 5, which is 12.

### Different scopes

Note that the two `result` variables (lines 5 and 14) are completely different as they are in two separate scopes – one is in the `main()` method and the other is in the `performCalc` method. As a result, there is no conflict or ambiguity whatsoever.

Line 7 executes `System.out.println()` with a method call inside the () of `println`. In this scenario, Java will execute the method call inside the () of `println`, and whatever the method returns will then be output to the screen. So, for line 7, the arguments passed to `performCalc` are 10, 2, and "-". Therefore, in `performCalc`, `x` is 10, `y` is 2, and `operation` is "-". The `switch` expression now executes line 16, resulting in `result` becoming 8 ( $10 - 2$ ). This `result` is returned (line 25) back to the calling method (line 7), where 8 is output to the screen.

Lines 8 and 9 operate similarly to line 7 except that the lines of code executed in the `switch` expression are different. The method call on line 8 executes line 17 in the `switch` expression, resulting in `result` being initialized to 20. This value is returned to the calling method (line 8), where 20 is output to the screen. The method call on line 9 executes line 18 in the `switch` expression, resulting in `result` being initialized to 5, and thus 5 is output to the screen.

Line 10 causes line 19 in the `switch` expression to be executed, initializing `result` to 0 ( $10 \% 2$ ). This `result` is returned back to the calling method, where, because it is not stored in a variable, it is simply lost/ignored.

The `performCalc` call on line 11 passes in "&", which executes the `default` branch of the `switch` expression. This results in the error message "Unrecognized operation: &" being displayed on the screen and -1 being returned. The -1 is then output on the screen.

Now that we know how to define and execute methods, we will move on to discussing method overloading, where distinct methods can have the same identifier name.

## Exploring method overloading

Consider a scenario where you have an algorithm, implemented by a method, that operates similarly on various input types – for example, `String` and `int`. It would be a shame to have two separately contrived method names, one for each input type, such as `doStuffForString(String)` and `doStuffForInt(int)`. It would be much better if both methods had the same name – that is,

doStuff – differentiated by their input types, which are doStuff (String) and doStuff (int). Thus, there will be no contrived method names. This is what method overloading provides. To discuss method overloading properly, we must first define the method signature.

## Method signature

The method signature consists of the method's name and the optional parameters. It does *not* consist of the return type. Let's look at an example to explain this further:

A code snippet showing a Java method definition. The method is named 'performCalc' and takes three parameters: 'int x', 'int y', and 'String operation'. A red dashed rectangle highlights the entire parameter list ('int x, int y, String operation'). A red arrow points from the text 'method signature' to the start of this highlighted area.

```
public static int performCalc(int x, int y, String operation){}
```

Figure 7.7 – Method signature

In the preceding figure, the method signature is highlighted in a dashed rectangle. It consists of the name of the method, followed by both the type and the order of the parameters. What this means is that the signature for the method in *Figure 7.7* is `performCalc`, which takes in two integers and a `String`, *in that order*. Note that the parameter names do not matter. So, in effect, from the perspective of the compiler, the method signature is `performCalc(int, int, String)`.

## Overloading a method

A method is overloaded when two or more methods share the same name but the parameters are different in type and/or order. This makes sense if you consider this from the viewpoint of the compiler. If you call a method that has two or more definitions, how will the compiler know which one you are referring to? To locate the correct method definition, the compiler compares and matches the method call with the overloaded method signatures. *Figure 7.8* presents an overloaded method with various signatures:

A code snippet showing an overloaded method 'someMethod'. The method is defined 13 times with different parameter lists. The first 12 definitions have different parameter types (int, double, String), while the 13th definition returns an integer value. Red dashed rectangles highlight each of the 13 definitions.

```
6 public static void someMethod(){}
7 public static void someMethod(int x){}
8 public static void someMethod(double x){}
9 public static void someMethod(String x){}
10 public static void someMethod(double x, int y){}
11 public static void someMethod(int x, double y){}
12 public static void someMethod(int a, double b){}
13 public static int someMethod(int x, double b){ return 0;}
```

Figure 7.8 – The method signature's impact on overloading

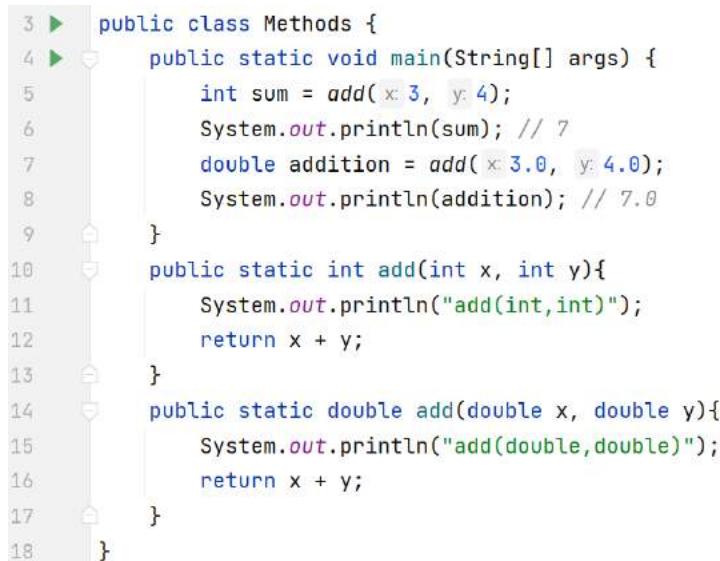
In this figure, the `someMethod` method is overloaded several times. The method signatures on lines 6 to 10 are `someMethod()`, `someMethod(int)`, `someMethod(double)`, `someMethod(String)`, and `someMethod(double, int)`, respectively.

The interesting cases are the compiler errors on lines 11-13. The error on line 11 is a misleading error from the compiler. In other words, if we comment out lines 12 and 13, the compiler error on line 11 disappears. There is nothing wrong with line 11 as this is the first time the compiler has seen this particular method signature – that is, `someMethod(int, double)`. The problem is that lines 12 to 13 have the same signatures and the compiler is flagging all lines with that signature.

Line 12 reinforces the point that the parameter names do not matter as they are not part of the method signature. Therefore, the fact that they are named `x` and `y` on line 11 and `a` and `b` on line 12 makes no difference whatsoever.

Similarly, line 13 demonstrates that the return type is not part of the method signature. Line 13 is a compiler error because its signature, `someMethod(int, double)`, is the same as on lines 11 and 12, even though the two methods have different return types (`int` and `void`, respectively).

In summary, the return type and parameter names are *not* part of the method signature. Now that we understand what is (and what is not) part of the method signature, let's look at a simple example of method overloading. *Figure 7.9* presents the code:



```
3: public class Methods {  
4:     public static void main(String[] args) {  
5:         int sum = add(x: 3, y: 4);  
6:         System.out.println(sum); // 7  
7:         double addition = add(x: 3.0, y: 4.0);  
8:         System.out.println(addition); // 7.0  
9:     }  
10:    public static int add(int x, int y){  
11:        System.out.println("add(int,int)");  
12:        return x + y;  
13:    }  
14:    public static double add(double x, double y){  
15:        System.out.println("add(double,double)");  
16:        return x + y;  
17:    }  
18: }
```

Figure 7.9 – Method overloading example

In this figure, we have an overloaded `add` method. The first version (lines 10 to 13) takes in two `int` parameters; the second version (lines 14-17) takes in two `double` parameters. Their respective

signatures are captured on lines 10 and 14, respectively. Thus, when we call add on line 5 and pass down two integers, the compiler matches the call with the version of add on line 10 because that version of add takes two integers. Similarly, the call to add on line 7 matches add on line 14 because both the call and method signature match (two double types in both).

Now that we understand how method parameter types and their order affect method overloading, let's examine how Java enables us to execute methods where the number of arguments is variable.

## Explaining varargs

Consider the following situation: you want to call a method, m1, but the number of arguments may vary. Do you overload the method with each version of the method taking in one extra parameter? For example, assuming the argument types are of the `String` type, do you overload m1 when each new version takes in an extra `String` parameter? In this case, you would have to code `m1 (String)`, `m1 (String, String)`, `m1 (String, String, String)`, and so forth. This is not scalable.

This is where `varargs` comes in. `varargs` is a very flexible language feature in Java, specifically provided for this use case. The syntax is that the type name is followed by an ellipsis (three dots). *Figure 7.10* shows `varargs` in action:

```

4 ►      public static void main(String[] args) {
5       m1();                      // 0
6       m1( ...args: 1);          // 1
7       m1( ...args: 1, 2);        // 3
8       m1( ...args: 1, 2, 3);    // 6
9     }
10 @      public static void m1(int... args){ // varargs
11     int sum = 0;
12     for(int i:args){
13         sum += i;
14     }
15     System.out.println(sum);
16   }

```

Figure 7.10 – varargs example

In this figure, on line 10, `m1 (int... )` defines a method signature for the `m1` method, defining 0 or more `int` parameters. This is quite different from `String []` defined on line 4 for `main`. In effect, you don't have to pass in any argument to `m1` at all; or you can pass in 1, 2, 3, or more integers. This is shown by the method calls (lines 5-8). Internally, in the `m1` method, `varargs` is treated as an array. The `for` loop (lines 12-14) demonstrate that.

The output from *Figure 7.10* is as follows:

```
0  
1  
3  
6
```

Line 5 generates no output at all. Line 6 generates 1; line 7 generates 3; and line 8 generates 6.

Let's examine some edge cases with varargs. *Figure 7.11* will help:

```
3 ► public class Methods {  
4 ►     public static void main(String[] args) {  
5         m1();  
6         m1("A");  
7         m1("A", "B");  
8         m1( n: "A", ...args: "B", "C");  
9     }  
10    public static void m1(int n, String... args){}  
11    public static void m1(String... args, int n){}  
12    public static void m1(String[] args){} // this is not varargs  
13  
14 //    public static void m1(String... args){ // varargs  
15 //        for(int i=0; i<args.length; i++){  
16 //            System.out.println(args[i]);  
17 //        }  
18 //        for(String s:args{  
19 //            System.out.println(s);  
20 //        }  
21 //    }  
22 }
```

Figure 7.11 – varargs compiler errors

In the preceding figure, we can see that varargs must be the last parameter in the method definition. Line 10 is fine as it defines the varargs parameter as the last parameter. However, line 11 is a compiler error because it attempts to define a parameter *after* the varargs parameter. This makes sense as all other parameters are mandatory; so, if varargs can define 0 or more arguments, it must be the last parameter.

Given that `varargs` is treated as an array, this begs the question, can we use an array instead of `varargs`? The answer is no. The compiler errors (lines 5-8) all relate to the fact that, despite the presence of `m1 (int [] )` on line 12, the compiler cannot find the method definition that matches any of these method calls.

The last major topic for methods is an important one: call by value. We will discuss that now.

## Mastering call by value

Java uses call by value when passing arguments to methods and returning results from methods. Concisely, this means that Java *makes a copy of something*. Effectively, when you are passing an argument to a method, a copy is made of that argument and when you are returning a result from a method, a copy is made of that result. Why do we care? Well, depending on what you are copying – a primitive or a reference has **major** implications. An example of a primitive type is `int` and an example of a reference type is an array.

In a method, there is a clear difference between the effect of changes when the parameter is a primitive type versus when the parameter is a reference type. We will demonstrate this shortly with a code example but first, to appreciate the differences, we need to understand what is happening in memory.

## Primitives versus references in memory

An array is an object, whereas a primitive is not. We will discuss objects in detail in *Chapter 8*, but for now, let's examine the code in *Figure 7.12*:



```
3 ► public class Methods {
4 ►     public static void main(String[] args) {
5         int x = 19;           // primitive
6         int[] arr = {1, 2}; // array
7     }
8 }
```

Figure 7.12 – Sample code containing a primitive and an array

To understand what the code in the preceding figure looks like in memory, we need to discuss the stack, the heap, and references.

### Stack

The stack is a special area of memory used by methods. Each time a new method A, is called, a new frame is *pushed* (created) onto the stack. The frame contains, among other things, A's local variables and their values. Each frame is stacked one on top of the other, like plates. If A calls another method, B, the existing frame for A is saved and a new frame for B is pushed onto the stack, creating a new context. When B finishes, its stack frame is *popped* (removed) from

the stack, and the frame for A is restored (with all its local variables and their values as they were, prior to the call to B). This is why a stack is called a **Last-In, First Out (LIFO)** structure. For further detail on the stack and Java Memory Management in general, please see our previous book: [https://www.amazon.com/Java-Memory-Management-comprehensive-collection/dp/1801812853/ref=sr\\_1\\_1?crid=3QUEBKJP46CN7&keywords=java+memory+management+maaike&qid=1699112145&sprefix=java+memory+management+maaike%2Caps%2C148&sr=8-1](https://www.amazon.com/Java-Memory-Management-comprehensive-collection/dp/1801812853/ref=sr_1_1?crid=3QUEBKJP46CN7&keywords=java+memory+management+maaike&qid=1699112145&sprefix=java+memory+management+maaike%2Caps%2C148&sr=8-1)

For our discussion here, what we need to be aware of is that local variables (primitives and/or references) are stored on the stack. Objects are *not* stored on the stack; objects are stored on the heap.

## Heap

The heap is an area of memory reserved for objects and arrays are objects. This means that arrays are stored on the heap. To access an object, we use a reference.

## References

The named identifier used to access an object is known as a reference. A reference is similar to a pointer. Consider a TV that has no buttons on it to change the channel but does have a remote control. The reference is the remote control and the TV is the object.

With these definitions in mind, let's review the code in *Figure 7.12*. Line 5 declares a primitive `int` type called `x` and initializes it to 19. Line 6 declares an `int` array, namely `arr`, and initializes `arr[0]` to 1 and `arr[1]` to 2. The array reference is `arr`. *Figure 7.13* shows the in-memory representation of *Figure 7.12* as we reach line 7:

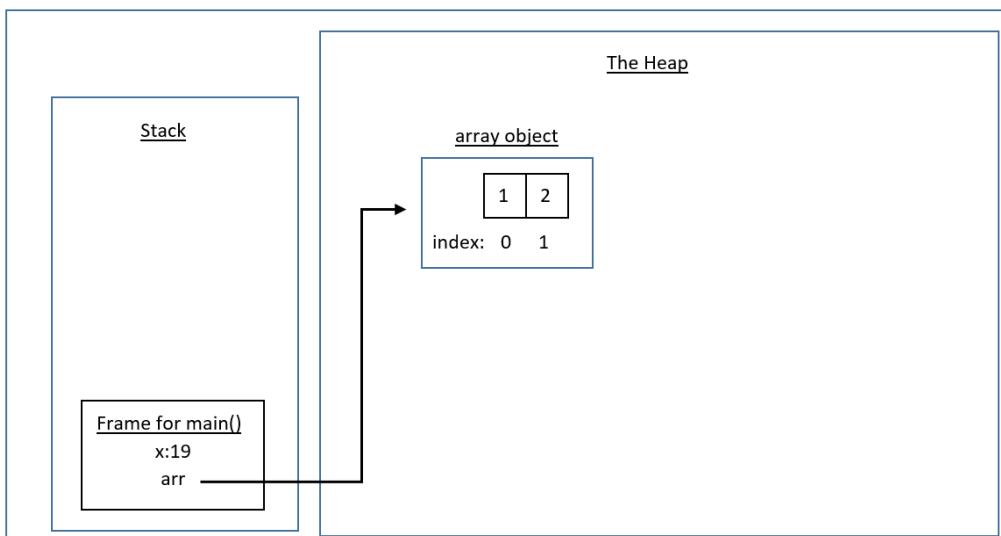
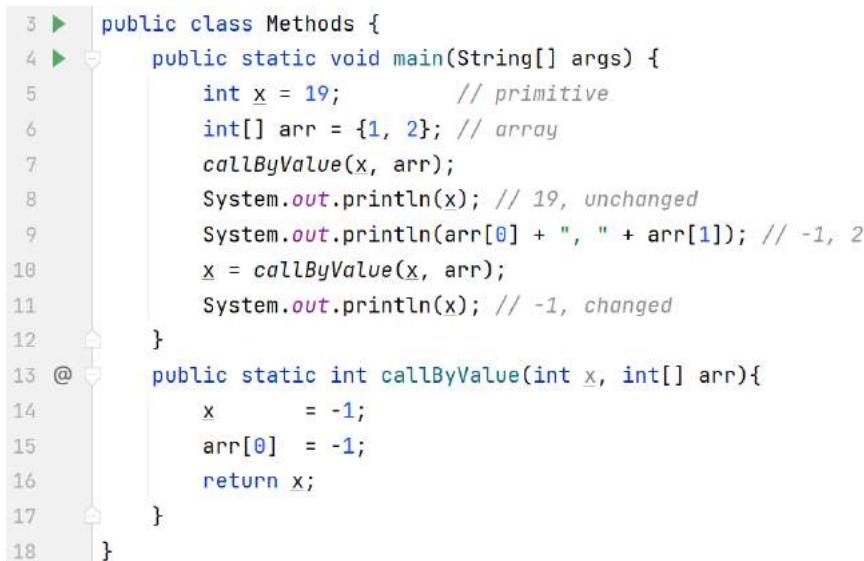


Figure 7.13 – In-memory representation of the code from Figure 7.12

In the preceding figure, we can see that there is a stack frame for the `main` method containing the local variables, `x` and `arr`. Note that, for simplicity, the `String[] args` parameter in `main` is omitted. Immediately, you can see the difference between the way primitives, namely `x`, and references, namely `arr`, are stored – `x` and its value are stored on the stack; whereas the value of `arr` refers to the object on the heap.

With this in mind, we are now in a position to examine a proper code example demonstrating the real impact of call by value when passing primitives versus passing references. *Figure 7.14* represents the code example we will be using:



```
3 ► public class Methods {  
4 ►   ► public static void main(String[] args) {  
5     int x = 19;           // primitive  
6     int[] arr = {1, 2}; // array  
7     callByValue(x, arr);  
8     System.out.println(x); // 19, unchanged  
9     System.out.println(arr[0] + ", " + arr[1]); // -1, 2  
10    x = callByValue(x, arr);  
11    System.out.println(x); // -1, changed  
12  }  
13 @  ► public static int callByValue(int x, int[] arr){  
14     x      = -1;  
15     arr[0] = -1;  
16     return x;  
17 }  
18 }
```

The code shows a Java class named `Methods` with a `main` method. Inside `main`, `x` is set to 19 and `arr` is an array [1, 2]. It then calls `callByValue(x, arr)`. After the call, `x` remains 19 (unchanged) and `arr` is printed as [-1, 2]. Next, `x` is set to the result of `callByValue`, which is -1. Finally, `x` is printed again as -1 (changed). The `callByValue` method itself takes `x` and `arr` as parameters and changes their values to -1. It then returns `x`.

Figure 7.14 – Call by value passing primitives and references

In this figure, the `callByValue` method is defined on lines 13-17: the method accepts an `int` type and an `int` array in that order and returns an `int`. Line 14 changes the value of the `int` parameter to -1 and line 15 changes index 0 of the array to -1. Lastly, the method returns the value of `x` on line 16.

Let us examine the first call to `callByValue`, passing down the `x` and `arr` arguments. It is important to note that the `x` and `arr` variables declared in `main` are completely separate variables from the `x` and `arr` parameters declared in the method `callByValue`. This is because they are in two separate scopes (methods). As Java uses call by value, copies of the primitive, `x`, and the reference, `arr`, are made and it is the *copies* that are passed into the method `callByValue`.

Making a copy of a primitive is like photocopying a blank sheet of paper - if you pass the photocopy to someone and they write on it, your original blank sheet is still blank. Making a copy of a reference is like copying a remote control – if you give the second remote control (the copied one) to someone

else, they can change the channel on the TV. Crucially, there is only one TV in all of this – the copy is made of the remote control, *not* the TV.

#### Note

This saves memory as copying a reference has a much smaller memory footprint than copying a potentially large object.

*Figure 7.15* represents the in-memory representation of the code as we are about to return from the *first* invocation of `callByValue`:

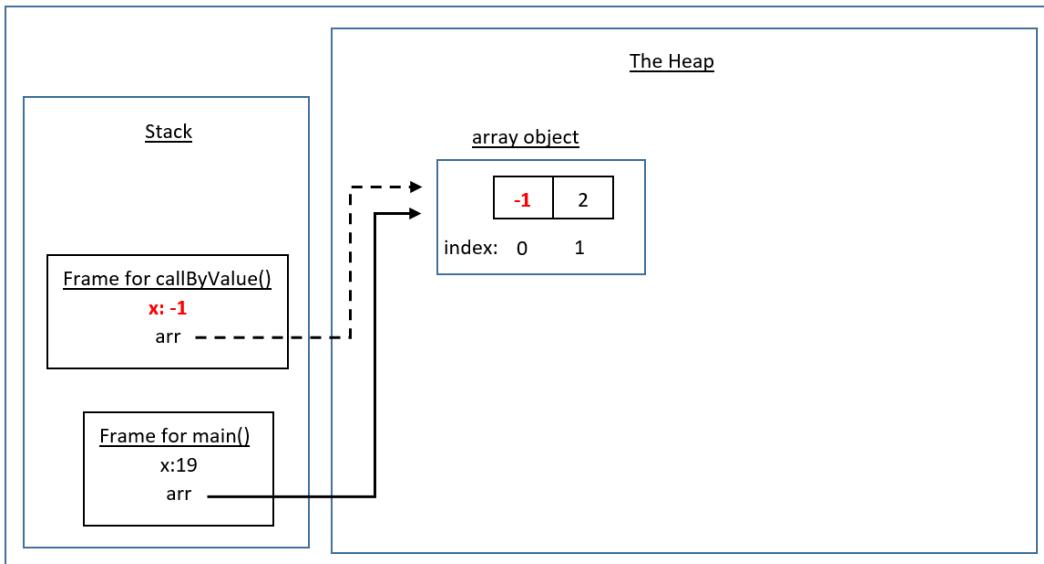


Figure 7.15 - In-memory view of Figure 7.14 (line 16) based on first call to `callByValue` (line 7)

As can be seen from the stack in the preceding figure, when we call `callByValue(x, arr)` from `main`, the existing frame for `main` is saved and a frame for `callByValue` is pushed onto the stack (on top of the frame for `main`). Then the code for `callByValue` is executed:

```
x = -1;
arr[0] = -1;
return x;
```

Firstly, `x` is changed in the `callByValue` frame. This is the copy of `x` from `main`. Note that the value of `x` in (the frame for) `main` remains untouched (still 19). Consequently, 19 is output for `x` in `main` (line 8). Thus, a called method cannot (directly) change the caller method's primitive values. We will revisit this point shortly.



However, the line `arr[0] = 1;` in `callByValue` does have a material impact on `main`. When `callByValue` uses its `arr` reference, which is a copy of the `arr` reference from `main`, it changes the one object that both methods share. In effect, the array object that `main` is looking at is changed. This is demonstrated in `main` after the `callByValue` method returns:

```
System.out.println(arr[0] + ", " + arr[1]); // -1, 2
```

Crucially, `-1` is output for the value at `arr[0]`. Therefore, be aware that when passing a reference to a method, the method can change the object that you are looking at.

Let's revisit the primitive situation. What if we wanted the called method to change the primitive value that's passed down? This is why `callByValue` is returning `x`. The first method call to `callByValue` completely ignores the return value:

```
callByValue(x, arr);
```

However, the second call does not:

```
x = callByValue(x, arr);
```

The `-1` that's returned from `callByValue` is used to overwrite the value of `x` in `main`. As a result, `-1` is output for `x` in `main` (line 11).

That completes our discussion on methods. Now, let's put that knowledge into practice to reinforce these concepts.

## Exercises

Maintaining a dinosaur park takes a lot more than just raw passion. It involves regular health check-ups for our dinosaurs, ensuring our guests are comfortable, and that the park is well-staffed. All these tasks involve methodical processes. Luckily, we now know about methods!

You can add these methods to the same class:

1. The stage of life a dinosaur is in can significantly impact its behavior and needs. Write a method that takes a dinosaur's age and returns whether it's a hatchling, juvenile, or adult.
2. It's important to remember that our dinosaurs aren't actually pets – they're large, often hefty creatures with large appetites. Write a method that accepts a dinosaur's weight and calculates how much food it needs daily.
3. Being on top of our dinosaurs' average age helps us plan for the future. Design a method that accepts an array of dinosaur ages and calculates the average age.
4. The park isn't open 24/7 to day visitors. We need some time to clean up the popcorn and repair any minor damages to the enclosures. Write a method that checks if the park is open or closed based on the current time. (Hint: this method doesn't require any input.)

5. Personalization is key to making our guests feel special. Create a method that uses a dinosaur's name and a visitor's name to craft a personal greeting message.
6. As you're well aware, safety is our top priority. We need a method to return whether we can let in another group of guests (a certain number of people) to the park based on the current number of visitors and the maximum number of visitors allowed.

## Project – Mesozoic Eden assistant

This is going to be our biggest project so far. So, buckle up!

Let's start with a high-level description. The Mesozoic Eden assistant is an interactive console application that assists in managing a dinosaur park. The assistant should have features to do the following:

- Add or remove dinosaurs
- Check the park's opening hours
- Greet guests and provide park information
- Track visitor counts to ensure the park isn't overcrowded
- Manage park staff details

Since we won't let you drown, if you need it, here is a step-by-step guide. A starting project will follow these steps:

1. **Create a data structure:** Create the necessary classes for Dinosaur, Guest, and Employee. Include appropriate properties and methods.
2. **Initialize the data:** You can choose to hardcode initial data or provide a mechanism to input data using the Scanner class.
3. **Implement interaction:** Implement a simple console-based interaction with the user using the Scanner class.
4. **Create a menu:** Create a menu that allows the user to interact with the park management system.
5. **Handle actions:** Each menu item should trigger a certain action, such as adding a dinosaur, checking park hours, or greeting a guest.
6. **Exit the program:** Provide an option for the user to exit the program.

Here is a code snippet to get you started:

```
import java.util.Scanner;

public class Main {
```



```
// Use Scanner for reading input from the user
Scanner scanner = new Scanner(System.in);

public static void main(String[] args) {
    Main main = new Main();
    main.start();
}

public void start() {
    // This is the main loop of the application. It
    // will keep running until the user decides to exit.
    while (true) {
        displayMenu();
        int choice = scanner.nextInt();
        handleMenuChoice(choice);
    }
}

public void displayMenu() {
    System.out.println("Welcome to Mesozoic Eden
                        Assistant!");
    System.out.println("1. Add Dinosaur");
    System.out.println("2. Check Park Hours");
    System.out.println("3. Greet Guest");
    System.out.println("4. Check Visitors Count");
    System.out.println("5. Manage Staff");
    System.out.println("6. Exit");
    System.out.print("Enter your choice: ");
}

public void handleMenuChoice(int choice) {
    switch (choice) {
        case 1:
            // addDinosaur();
            break;
        case 2:
            // checkParkHours();
            break;
        case 3:
            // greetGuest();
            break;
    }
}
```

```
        case 4:  
            // checkVisitorsCount();  
            break;  
        case 5:  
            // manageStaff();  
            break;  
        case 6:  
            System.out.println("Exiting...");  
            System.exit(0);  
    }  
}  
}
```

So, this is a great starting point! But it's not done. In the preceding code snippet, `addDinosaur()`, `checkParkHours()`, `greetGuest()`, `checkVisitorsCount()`, and `manageStaff()` are placeholders for methods you need to implement according to your data structures and functionality. The `Scanner` class is used to read the user's menu choice from the console.

You can make the project as sophisticated as you like by adding additional features and enhancements.

## Summary

In this chapter, we started our discussion on methods by stating that methods are simply code blocks that are given a name for ease of reference. Methods are important because they enable us to abstract away the implementation, while at the same time helping us to avoid unnecessary code duplication.

There are two parts to a method: the method definition (or declaration) and the method call (or invocation). The method definition declares (among other things) the method name, the input parameters, and the return type. The method name and the parameter types (including their order) constitute the method signature. The method call passes down the arguments (if any) to be used as inputs in the method. The return value from a method (if there is one) can be captured by assigning the method call to a variable.

Method overloading is where the same method name is used across several different methods. What distinguishes the various methods is that they have different signatures – the parameter types and/or their order will be different. The parameter names and return types do not matter.

A `varargs` (variable arguments) parameter is specified in a method declaration using an ellipsis (three dots). This means that when calling this method, the arguments corresponding to that parameter are variable – you can pass down 0 or more arguments for that parameter. Internally, in the method, the `varargs` parameter is treated as an array.

When passing arguments to a method, Java uses call by value. This means that a copy of the argument is made. Depending on whether you are passing down a primitive or a reference has major implications regarding the effect of the changes that are made by the called method on the calling method. If it's



a primitive, the called method cannot change the primitive that the caller method has (unless the caller method deliberately overwrites the variable with the return value). If it's a reference, the called method can change the object that the caller method is looking at.

Now that we have finished looking at methods, let's move on to our first strictly **object-oriented programming (OOP)** chapter, where we will look at classes, objects, and enums.

# Part 2:

# Object-Oriented Programming

In this part, we will take a deep dive into **Object-Oriented Programming (OOP)**. We will start by looking at classes and their relationship with objects. We will discuss the first core pillar of OOP, namely Encapsulation. Enums and records, both of which are closely related to classes, will also be covered. We will then move on to the remaining two major pillars in OOP, namely Inheritance and Polymorphism. Next up are `abstract` classes and the hugely important `interface` construct. Following that, we will examine Java's exception framework. Lastly, we will explore selected classes from the Java Core API, such as `String` and `StringBuilder`.

This section has the following chapters:

- *Chapter 8, Classes, Objects, and Enums*
- *Chapter 9, Inheritance and Polymorphism*
- *Chapter 10, Interfaces and Abstract Classes*
- *Chapter 11, Dealing with Exceptions*
- *Chapter 12, Java Core API*

Trial Version



Wondershare  
**PDFelement**



## 8

# Classes, Objects, and Enums

In *Chapter 7*, we learned about methods in Java. After understanding why methods are useful, we learned that there are two parts to methods – the method definition and the method call. We saw that the method definition is the code that's executed when the method is invoked via the method call. We discussed how method signatures enable method overloading. We also learned how `varargs` helps us call a method with zero or more arguments. Finally, we discussed Java's call by value mechanism, where arguments that are passed to a method are copied in memory. Depending on the type of argument passed, primitive or reference, will have implications as to the effect of the changes made in the called method to those arguments passed from the caller method.

*Chapter 7* concluded the Java fundamentals section of this book. The topics in that section are common across many programming languages, including non-**object-oriented programming (OOP)** languages such as C. *Chapter 8* starts the OOP section of this book.

In this chapter, we will cover classes, objects, records and enums. Classes and objects are unique to OOP languages (such as Java); in other words, non-OOP languages (such as C) do not support them. Though closely related, understanding the difference between a class and an object is important. We will discuss the relationship between the class and objects of the class. To access an object, we must use a reference. Separating the reference from the object will prove very useful going forward. Instance versus class members will be discussed, as well as when to use either/both. This chapter will also explain the '`this`' reference and how it relates to the object responsible for the instance method currently executing.

We will also explain the access modifiers in Java. These access modifiers enable one of the key cornerstones in OOP, namely encapsulation. Though basic encapsulation can be easily achieved, properly encapsulating your class requires extra care. This will be covered in the Advanced encapsulation section.

Understanding the object life cycle, with regard to what is happening in memory as your program executes, is crucial to avoiding many subtle errors. This topic will be explained with the aid of diagrams.

Toward the end of the chapter, given our understanding (and separation!) of references from the objects they refer to, we will discuss the `instanceof` keyword. Lastly, we will cover a variation of classes, namely enums, whereby the number of object instances is restricted.



This chapter covers the following main topics:

- Understanding the differences between classes and objects
- Contrasting instance with class members
- Exploring the 'this' reference
- Applying access modifiers
- Achieving encapsulation
- Mastering advanced encapsulation
- Delving into the object life cycle
- Explaining the instanceof keyword
- Understanding enums
- Appreciating records

## Technical requirements

The code for this chapter can be found on GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects/tree/main/ch8>.

## Understanding the differences between classes and objects

As classes and objects are integral to OOP, it is vital to understand their differences. We will discuss the relationship between a class and its objects in this section. As creating objects requires the use of the new keyword, this will also be covered. Understanding constructors and what they do will also be examined. All of these topics are linked: objects are the in-memory representation of the class (template); to create an object, a constructor is used and to call the constructor, we use the new keyword. Let's examine these in turn.

### Classes

A class is so integral in Java that you cannot write any program without defining one! A class is a blueprint or template for your object. It is similar to a plan of a house – using a house plan, you can discuss the house all you want; however, you cannot go into the kitchen and make a cup of tea/coffee. The house plan is abstract in that regard and so is the class. The class defines fields (properties) and methods which operate on those fields. The fields are your data and the methods enable manipulation of that data.

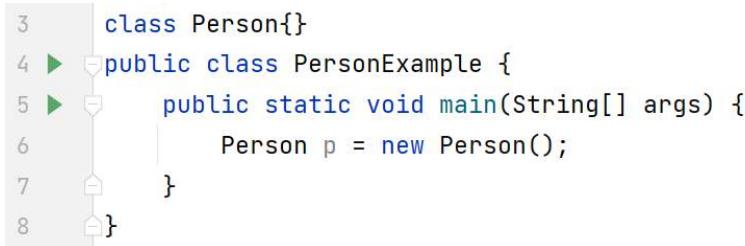
## Objects

An object is your in-memory representation of your class. If the class is your house plan, then the object is your built house. Now, you can go into the kitchen and make that cup of tea/coffee. As with houses and house plans, you can create many objects based on the class. These objects are known as object *instances*, emphasizing that each object is its own unique instance.

In summary, the class is the template and the object is the in-memory representation of the class. You need an object (instance) if you want to execute its (instance) methods. So, how do we create an object? We use the new keyword.

## Getting familiar with the new keyword

The new keyword in Java enables us to create objects. The object is created on the heap, a special area of memory reserved for objects. A reference (similar to a pointer) to the object is returned. This reference enables us to manipulate the object; for example, to execute the instance methods. Let's examine the code example shown in *Figure 8.1*:



```
3  class Person{}
4  public class PersonExample {
5      public static void main(String[] args) {
6          Person p = new Person();
7      }
8  }
```

Figure 8.1 – Creating an object

In the preceding figure, line 3 defines a Person class. It contains nothing at the moment; we will expand it as we progress. Line 6 is important – we are creating a Person object using the new keyword. Apart from the new keyword, line 6 is very similar to any method call. The p reference (on the stack) is initialized to refer to an object of type Person (on the heap). **It is very important to separate the reference from the object**. For example, as we shall see when we discuss Inheritance (*Chapter 9*), a reference of type X does not have to refer to an object of type X; the reference can refer to an object of type Y, once Y is related to X. On line 6, the Person reference named p is referring to a Person object; however, going forward, that will rarely be the case. When “constructing” objects using the new keyword, the method that's invoked is a special method known called a *constructor*.

### Constructors

A constructor is a special method that's invoked by the `new` keyword. It has two distinct properties that differentiate it from other methods: it has the same name as the class and defines no return type, not even `void`. (Java returns the reference to the object in the background).

Every class contains a constructor, even if you do not code one yourself. If you do not code a constructor for your class, Java will synthesize (or define) a “default constructor” for you. The default constructor will have the same properties as regular constructors; namely, the same name as the class and no return type. However, the default constructor will not define any parameters; it will have the same access modifier as the class and will contain only one line of code, which is `super();`. We will discuss access modifiers later in this chapter and `super()` in *Chapter 9*.

Note that if you insert even one constructor, the default constructor is not synthesized. It's as if the compiler says, “Okay, you have a constructor(s), you know what you are doing, so I won't get involved.”

Now that we know when default constructors are synthesized by the compiler, we can see that default constructors are required for both `Person` and `PersonExample` in *Figure 8.1*. *Figure 8.2* represents the code *after* the compiler has inserted the default constructors:

```
3  class Person{  
4      Person(){  
5          super();  
6      }  
7  }  
8  public class PersonExample {  
9      PersonExample(){  
10         super();  
11     }  
12     public static void main(String[] args) {  
13         Person p = new Person();  
14     }  
15 }
```

Figure 8.2 – Default constructors inserted

The red rectangles in the preceding figure represent the default constructors inserted by the compiler. This happened to both classes because neither class defined any constructor at all and every class requires a constructor. The default constructors, in addition to having the same name as the class and not returning anything (not even `void`), define no parameters (lines 4 and 9) and simply call `super();`. As stated in the previous callout, `super()` will be discussed when we discuss Inheritance in *Chapter 9*.

We will discuss access modifiers in detail later but note that the access for the default constructors match the access for their respective classes. For example, `PersonExample` is a `public` class and so is its constructor (lines 8 and 9 respectively). The `Person` class mentions no *explicit* access modifier at all and neither does its constructor (lines 3 and 4 respectively).

Now, you can see why `new Person();` on line 13 does not generate a compiler error. To be clear, there is no compiler error on line 13 because the compiler inserted the default constructor for the `Person` class (lines 4 to 6) and thus `new Person()` was able to locate the constructor and therefore compile.

The default constructor for `PersonExample` (lines 9 to 11) has no material effect in this program. The JVM starts every program in the `main` method.

We will now move on to discuss instance members versus class members. Note that local variables (in a method) are neither.

## Contrasting instance with class members

An object can be more correctly termed an object *instance*. This is where *instance* members (methods/data) get their names: every object gets a copy of an instance member. Class members, however, are different in that there is only one copy per class, regardless of the number of object instances created. We'll discuss both of these topics now.

### Instance members (methods/data)

This is more easily explained by presenting a code example first. *Figure 8.3* presents a class with instance members:

```

3   class Person{
4       private String name;      // instance variable
5       private int count;        // instance variable
6
7       Person(String aName) { // constructor
8           name = aName;
9           count++;
10      }
11      public String getName() { // instance method
12          return name;
13      }
14      public void setName(String aName) { // instance method
15          name = aName;
16      }
17      public int getCount() { // instance method
18          return count;
19      }
20  }
21  public class PersonExample {
22      public static void main(String[] args) {
23          Person p1 = new Person( aName: "Maaike");
24          Person p2 = new Person( aName: "Sean");
25          System.out.println(p1.getName()); // Maaike
26          System.out.println(p2.getName()); // Sean
27          p1.setName("Maaike van Putten");
28          p2.setName("Sean Kennedy");
29          System.out.println(p1.getName()); // Maaike van Putten
30          System.out.println(p2.getName()); // Sean Kennedy
31      }
32  }

```

Figure 8.3 – A class with instance members

When you create an object using `new`, you are creating an object *instance*. Each instance gets a copy of the instance members (variables and methods). Regarding instance variables, we need to define where instance variables are declared and their resultant scope. An instance variable is defined within the class but outside every method coded in the class. Thus, the scope of an instance variable is the class itself; meaning, every instance method in the class can access the instance variables.

Now let us discuss the code example. In the preceding figure, the `Person` class defines both instance variables and instance methods. As the instance variables are declared outside every method, they have the scope of the class. The fact that the instance variables are marked `private` and the instance methods are marked `public` will be explained later in this chapter. The constructor is as follows:

```

Person(String aName) { // constructor
    name = aName;
    count++;
}

```

This constructor enables us to pass in a `String` and initialize the instance variable based on that `String`. For example, when we instantiate an object as follows:

```
Person p1 = new Person("Maaike");
```

we are passing "Maaike" into the constructor, so the `name` instance variable in the object referred to by `p1` refers to "Maaike". The constructor is also keeping a count of the number of objects that are created by incrementing `count` each time the constructor is invoked. Note that no default `Person` constructor was inserted by the compiler in this example as a constructor was already coded in the class.

We also invoke the `getName()` instance method using the `p1` and `p2` references as follows:

```
System.out.println(p1.getName()); // Maaike
System.out.println(p2.getName()); // Sean
```

This syntax of `refName.instanceMethod()` is known as *dot notation*. As per the comments in the code, "Maaike" and "Sean" are output to the screen (in that order). *Figure 8.4* shows the in-memory representation of the code after we have created both objects, referenced by `p1` and `p2` respectively:

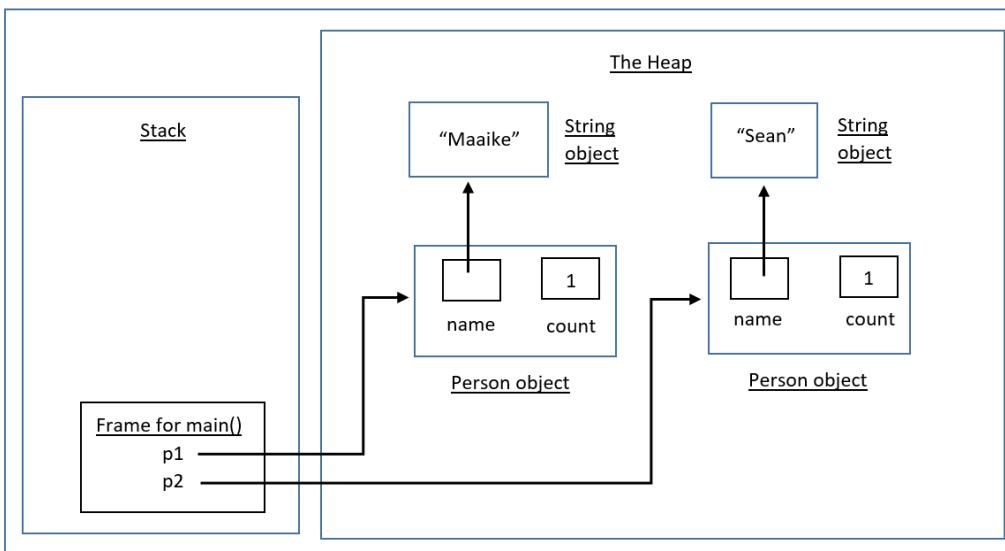


Figure 8.4 – In-memory representation of Figure 8.3 (start of line 27)

As the preceding figure shows, we have two references on the stack, namely p1 and p2. p1 refers to the first Person object on the heap – that is, the object that was created on line 23. The instance variable values of p1 (its “state”) are "Maaike" and 1 for name and count, respectively. As strings are objects, name is a reference to another object, a String object, which has a value of "Maaike". Similarly, the p2 reference refers to the object that was created on line 24. As can be seen from the diagram, the instance variable values of p2 are "Sean" and 1 for name and count, respectively.

Note that each Person object *instance* on the heap has a copy of the *instance* variables. That is why they are called instance variables.

Lines 27 and 28 change the values of the name instance variables to "Maaike van Putten" and "Sean Kennedy" for p1 and p2, respectively. *Figure 8.5* shows these changes:

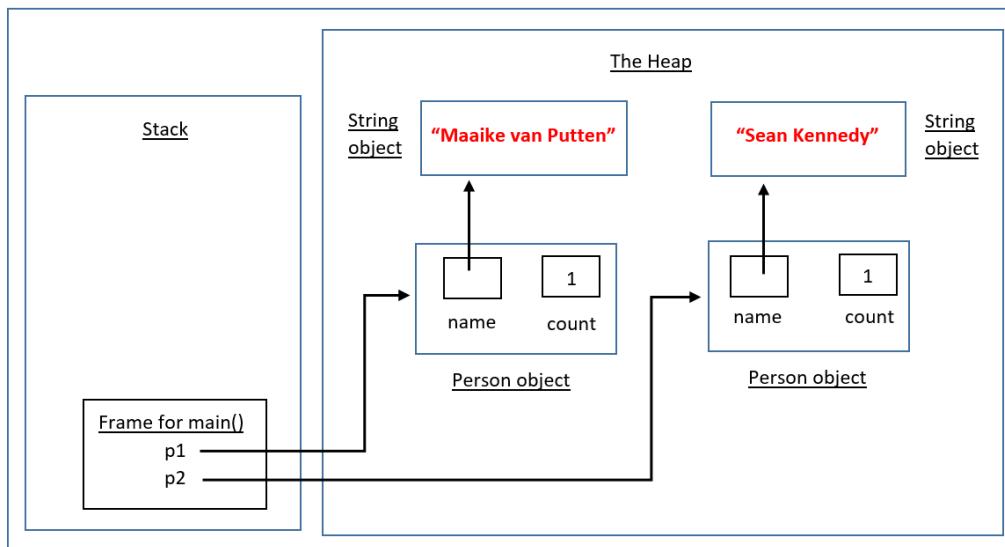


Figure 8.5 – In-memory representation of Figure 8.3 (start of line 29)

This figure shows that the two String objects have been changed: p1’s instance variable name refers to "Maaike van Putten" and p2’s instance variable name refers to "Sean Kennedy". Consequently, lines 29 to 30 output "Maaike van Putten" and "Sean Kennedy", respectively.

### String immutability

Strings are immutable objects. This means that `String` objects, once created, cannot be changed. Ever. It may look like they have changed, as in the effect is created of a change, but a completely new object has been created and the original is left untouched. We will revisit `String` immutability in greater detail in *Chapter 12*.

So, the original `String` objects, "Sean" and "Maaike", are still on the heap taking up space. They are of no use because, as we have no references to them, we have no way to get to them. Remember, the name instance variables for both `p1` and `p2` refer to the newly created `String` objects containing "Maaike van Putten" and "Sean Kennedy", respectively.

So, what happens to these no-longer-used objects? They are "garbage collected." We will discuss this soon but for now, just know that the JVM runs a process called a garbage collector in the background to tidy up (reclaim) all the objects that can no longer be reached. We have no control over when this process runs but the fact that there is a garbage collector saves us from having to tidy up after ourselves (whereas in other OOP languages such as C++, you have to!).

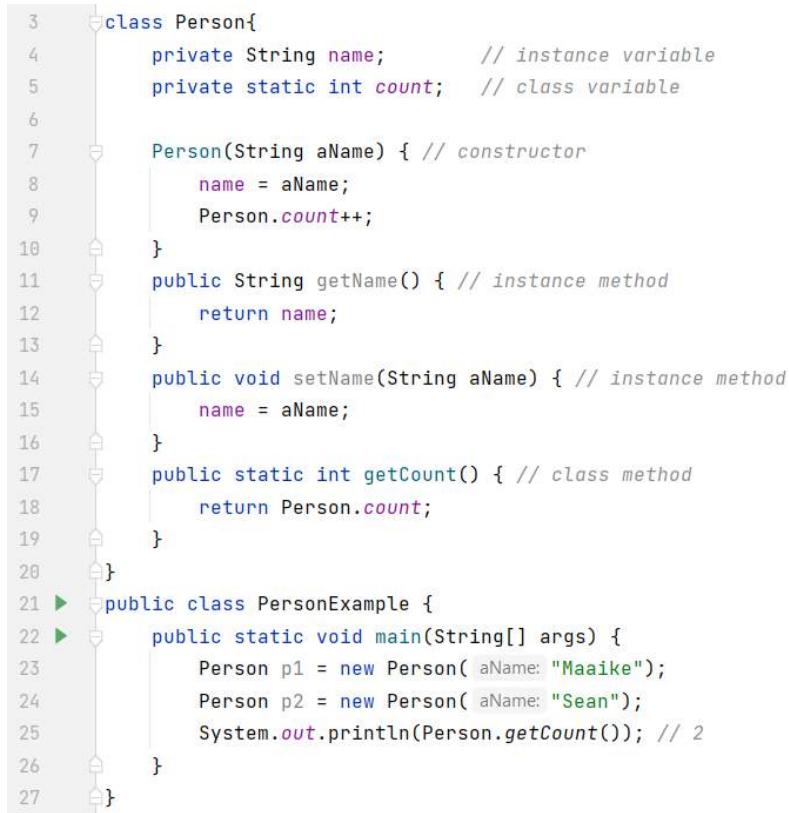
The code in *Figure 8.3* has an issue – `count` is 1 and it should be 2. Instance variables that are integers are initialized to 0 by default. In each of the constructor calls, we increment `count` from 0 to 1. We would like the first constructor call to increment `count` from 0 to 1 and the second constructor call to increment `count` from 1 to 2. This is where class members come in.

## Class members (methods/data)

To mark a field and/or method as a class member, as opposed to an instance member, you can insert the `static` keyword into the declaration of the member. Class members are shared by all instances of the class. This means that you do not have to create an object instance to access the `static` members of the class.

The syntax for accessing a `static` member is different from accessing an instance member. Rather than use the reference, the class name is used, as in `className.staticMember`. This emphasizes the class nature of the member being accessed. For example, the JVM starts the program in *Figure 8.3* with `PersonExample.main()`. This is how the JVM starts every program as it saves on constructing an object and its resulting memory footprint.

Let's get back to our problem with `count` (which is 1 instead of 2). *Figure 8.6* represents the changes that must be made to fix this issue:



```
3  class Person{
4      private String name;          // instance variable
5      private static int count;    // class variable
6
7      Person(String aName) { // constructor
8          name = aName;
9          Person.count++;
10     }
11     public String getName() { // instance method
12         return name;
13     }
14     public void setName(String aName) { // instance method
15         name = aName;
16     }
17     public static int getCount() { // class method
18         return Person.count;
19     }
20 }
21 public class PersonExample {
22     public static void main(String[] args) {
23         Person p1 = new Person( aName: "Maaike");
24         Person p2 = new Person( aName: "Sean");
25         System.out.println(Person.getCount()); // 2
26     }
27 }
```

Figure 8.6 – Making “count” static

Contrasting the code in *Figure 8.6* with the code in *Figure 8.3*, we can see that `count` is declared `static` (line 5). Thus, there is only one copy of `count`, which is shared across all instances of `Person`. Thus, `p1` and `p2` are looking at the same `count`.

In the constructor (line 9), while not necessary, we use the correct syntax to emphasize the `static` nature of `count`. Similarly, as `getCount` (line 17) is simply returning a `static` member, we marked it as `static`. In addition, we used the `Person.count` static syntax (line 18). Lastly, line 25 accesses the `static` method using the correct syntax, `Person.getCount`, to retrieve the `private` class variable, `count`. We can see that it outputs 2, which is correct. Comparing the other differences in code, some of the extra code in `main` (*Figure 8.3*) has been removed to help us focus on what we are discussing here.

### Instance to static but not vice versa

If you are in an instance method, you can access a `static` member but not vice versa. We will discuss the reason why when we explain the `this` reference. This means that, in *Figure 8.6*, you could use the `p1` reference to access the `getCount` method (line 25). As such, `p1.getCount()` is valid but this is a *poor* programming practice as it conveys the impression that `getCount` is an instance method when it is a `static` method - use `Person.getCount()` as per the code.

*Figure 8.7* shows the in-memory representation of the code in *Figure 8.6*:

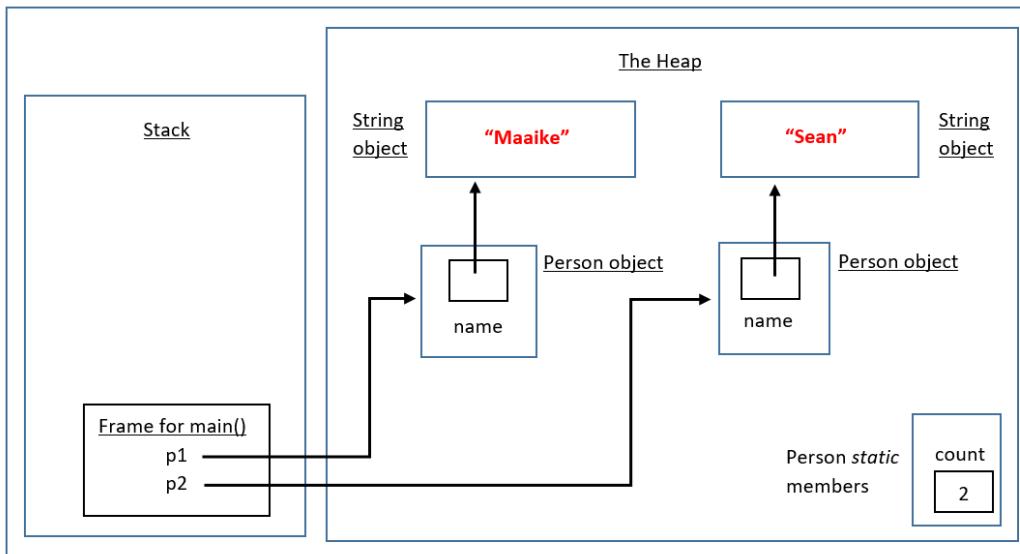


Figure 8.7 – In-memory representation of the code in Figure 8.6

As can be seen in the bottom-right corner of the preceding figure, the `static/class` members of the `Person` class are stored separately from the instances themselves. There is now only one copy of `count` and it is shared between `p1` and `p2`. Thus, the `count` value of 2 is correct.

### Default values for class and instance variables

Instance variables are initialized to default values every time a class is new'ed.

Class variables are initialized to default values the very first time a class is loaded. This could occur when using `new` or when referring to a class member (using the class syntax).



The default values for class and instance variables are as follows:

Type	Default value
byte, short, and int	0
long	0L
float	0.0f
double	0.0d
char	'\u0000' (Unicode zero)
String (or any reference to an object)	null
boolean	false

Table 8.1 – Default values for class and instance variables

In a previous callout, we highlighted that you can access class members from an instance method but not vice versa. Let's delve into that now.

## Exploring the “this” reference

When you call an instance method, the compiler secretly passes into the method a copy of the object reference that invoked the method. This reference is available to the instance method as the `this` reference.

Class methods do not get a `this` reference. This is why, if you are in a `static` method (context) and you try to access an instance member directly (without an object reference), you will get a compiler error. In effect, every instance member requires an object reference when accessing it. This makes sense because instance members are instance-specific and therefore, you need an instance (reference) to say, “*I want to access this particular instance/object as opposed to that particular one.*”

Let's refactor the code in *Figure 8.3* so that the `Person` class uses the `this` reference explicitly. In addition, all references to the incorrectly working `count` instance variable have been removed so that we can focus on the `this` reference. *Figure 8.8* includes the refactored `Person` class (the `PersonExample` class remains untouched):

```

3  class Person{
4      private String name; // instance variable
5
6      Person(String aName) { // constructor
7          // name = aName;
8          this.name = aName;
9      }
10     public String getName() { // instance method
11         return name;
12         return this.name;
13     }
14     public void setName(String aName) { // instance method
15         // name = aName;
16         this.name = aName;
17     }
18 }
19 ► public class PersonExample {
20 ►     public static void main(String[] args) {
21         Person p1 = new Person( aName: "Maaike");
22         Person p2 = new Person( aName: "Sean");
23         System.out.println(p1.getName()); // Maaike
24         System.out.println(p2.getName()); // Sean
25         p1.setName("Maaike van Putten");
26         p2.setName("Sean Kennedy");
27         System.out.println(p1.getName()); // Maaike van Putten
28         System.out.println(p2.getName()); // Sean Kennedy
29     }
30 }

```

Figure 8.8 – Using the “this” reference

In the preceding figure, lines 7, 11, and 15 are commented out and replaced by lines 8, 12, and 16, respectively. Let’s contrast both the commented-out line 7 and the new line 8 more closely:

```
// name = aName; // line 7
this.name = aName; // line 8
```

Firstly, assume line 7 is uncommented. How does line 7 reconcile its variables? Initially, the compiler checks the current scope (the constructor block of code) and reconciles `aName` as a parameter to the constructor. However, the compiler still has not reconciled `name`, so it checks the next outer scope, the class scope, where the instance/class variables are defined. Here, it finds an instance variable called `name`, and therefore line 7 compiles.

Line 8 operates somewhat differently. Yes, it reconciles `aName` similarly but now, it comes across `this.name` (as opposed to `name`). Upon seeing `this`, the compiler immediately checks the instance variables that have been declared. It finds an instance variable called `name`, and therefore line 8 compiles. Lines 7 and 8 are, in effect, the same.

Line 16 is the same as line 8 as we used the same parameter identifier, `aName`. Line 12 is simply returning the `name` instance variable.

So, that covers how to use `this` in a class, but how do we associate the instance with `this`?

### Associating an instance with the “this” reference

Thankfully, the compiler does this automatically. As stated previously, the `this` reference is only passed (in secret) to instance methods and it refers to the instance that is invoking the method at that time. For example, when executing `p1.getName()`, the `this` reference in `getName` refers to `p1`, whereas when executing `p2.getName()`, the `this` reference in `getName` refers to `p2`. Thus, the `this` reference varies, depending on the instance that invokes the method. *Figure 8.9* represents the dynamic nature of the `this` reference in action:

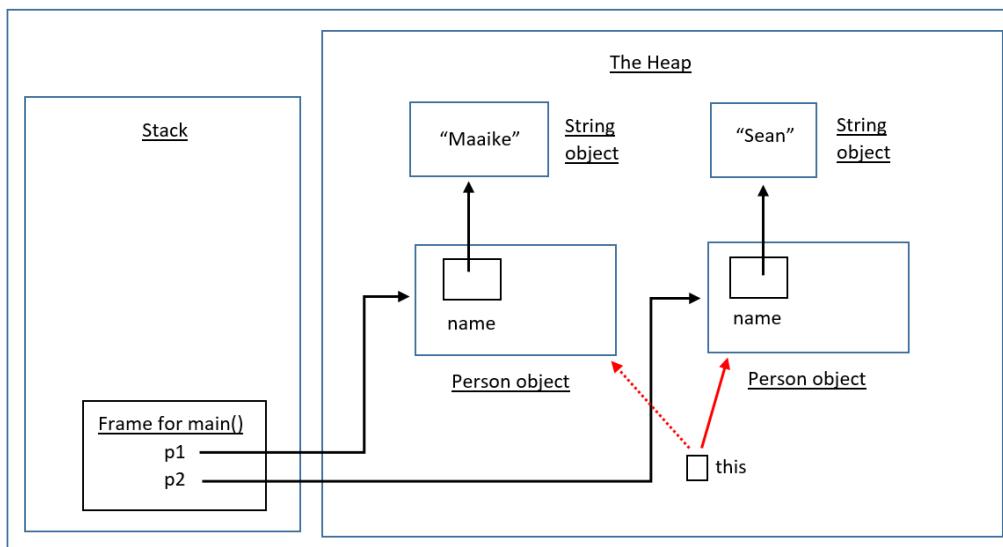


Figure 8.9 – The dynamic nature of the “this” reference

This figure represents the in-memory representation of the code in *Figure 8.8* as we execute line 12 from the method call on line 24. As `getName` on line 24 is invoked on `p2` – in other words, `p2.getName()` – the `this` reference inside `getName` refers to the same object that `p2` is

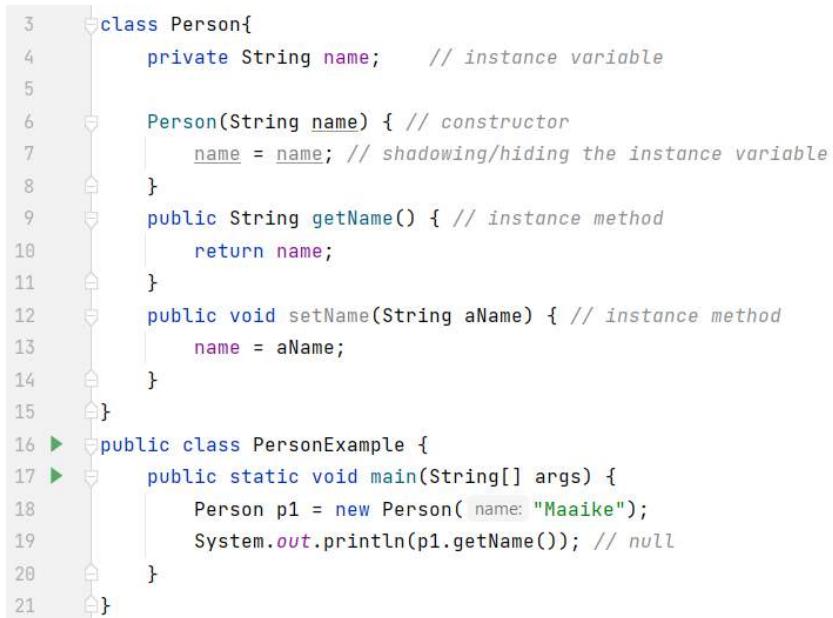
referring to. This is represented by the solid line from the `this` reference referring to the same object that `p2` is referring to.

The dashed line represents what the `this` reference was referring to on line 12, from the method call on line 23, namely `p1`. Thus, the `this` reference is dynamically referring to the instances referred to by `p1` or `p2`, depending on which invoked the instance method.

As we saw in the code in *Figure 8.8*, the `this` reference was not needed. Let's examine a situation where the `this` reference is needed.

## Shadowing or hiding an instance variable

Shadowing an instance variable occurs when a variable has the same identifier as the instance variable. *Figure 8.10* presents code where this occurs so that we can observe the issue it creates:



```
3  class Person{
4      private String name; // instance variable
5
6      Person(String name) { // constructor
7          name = name; // shadowing/hiding the instance variable
8      }
9      public String getName() { // instance method
10         return name;
11     }
12     public void setName(String aName) { // instance method
13         name = aName;
14     }
15 }
16 public class PersonExample {
17     public static void main(String[] args) {
18         Person p1 = new Person( name: "Maaike");
19         System.out.println(p1.getName()); // null
20     }
21 }
```

The code shows a `Person` class with a private instance variable `name`. It has a constructor that takes a parameter `name` and assigns it to the local variable `name` (line 7), thus shadowing the instance variable. The `main` method creates a `Person` object with `name: "Maaike"` and prints its name, which is expected to be `"Maaike"` but instead outputs `null`.

Figure 8.10 – Shadowing an instance variable

In the preceding figure, the constructor has a logical issue; in other words, the code compiles but the code is not working as expected. Line 7 is the issue. Remember that, if a variable is not qualified with `this`, the current scope is checked to see if the variable is declared there. On line 6, we have declared a constructor parameter that uses the `name` identifier, which is the same identifier as the instance variable on line 4. Thus, line 7 is essentially assigning the local variable to itself and the instance variable remains untouched. As the instance variable is a `String` type, its default value is `null`. As a result, `null` is output on line 19 instead of `"Maaike"`.

To fix this issue, we have two options. The first option is to use a different identifier for the constructor parameter and use that new identifier. This is what `setName` does (lines 12-13): a method parameter called `aName` is used that does not shadow the `name` instance identifier. The second option is to use the `this` reference to specify that the variable being initialized is an instance variable. *Figure 8.11* shows this:

```
3  class Person{
4      private String name;      // instance variable
5
6      Person(String name) { // constructor
7          this.name = name;
8      }
9      public String getName() { // instance method
10         return name;
11     }
12     public void setName(String aName) { // instance method
13         name = aName;
14     }
15 }
16 public class PersonExample {
17     public static void main(String[] args) {
18         Person p1 = new Person( name: "Maaike");
19         System.out.println(p1.getName()); // Maaike
20     }
21 }
```

Figure 8.11 – Using “this” to fix the shadowing issue

In this figure, line 7 is important: `this.name` refers to the `name` instance variable, while `name`, on its own, refers to the method parameter. Thus, shadowing has been removed and line 19 now outputs "Maaike" as expected.

We know that only non-static (instance) methods receive the `this` reference. Let's examine how this issue can affect us and how to resolve it. *Figure 8.12* presents code where we are in a `static` context (method) and are trying to directly access an instance variable:

```
3 ► public class PersonExample {  
4     int x;           // instance variable  
5     public void m(){} // instance method  
6 ► public static void main(String[] args) {  
7         // Non-static field 'x' cannot be referenced  
8         // from a static context  
9         x = 9; // same as 'this.x = 9;'  
10        this.x = 99;  
11        m(); // same as 'this.m();'  
12        this.m();  
13  
14        // this works  
15        PersonExample pe = new PersonExample();  
16        pe.x=999; // ok  
17        pe.m(); // ok  
18        System.out.println(pe.x); // 999  
19    }  
20 }
```

Figure 8.12 – Accessing instance variables from a “static” context

In the preceding figure, we have an instance variable called `x` (line 4), an instance method called `m` (line 5), and a `static` method called `main` (lines 6–19). As we know, `static` methods such as `main`, do not get the `this` reference automatically (as they are class methods as opposed to instance methods).

There are compiler errors on lines 9, 10, 11, and 12. When you access an instance member directly, as on lines 9 and 11, the compiler inserts `this` before the member. In other words, by the time the compiler is finished with lines 9 and 11, internally, they look the same as lines 10 and 12. Consequently, as `main` does not have a `this` reference, the compiler complains about lines 9, 10, 11 and 12.

Lines 15–18 encapsulate how to resolve this issue. When you’re in a `static` context and you want to access an instance member (variable or method), you need to create an object instance to refer to the instance member. Therefore, on line 15, we create an (object) instance of the class containing the instance member, namely `PersonExample`, and store the reference in an identifier, `pe`. Now that we have an instance, we can access the instance members, which we do on lines 16, 17 and 18. Line 16 successfully changes `x` from (its default value of) 0 to 999. This is what is output on line 18. Line 17 shows that access to `m` is not an issue either. Note that you must comment out lines 9–12 before the code will compile and run.

Throughout these examples, we have used the `private` and `public` access modifiers. Let’s discuss these in more detail.



## Applying access modifiers

One of the cornerstones of OOP is the principle of *encapsulation* (data abstraction). Encapsulation can be achieved using access modifiers. Before we discuss encapsulation, we must understand the access modifiers themselves.

Access modifiers determine where a class, field, or method is visible and therefore available for use. The level you are annotating at, determines the available access modifiers:

- **Top level:** Classes, enums, records and interfaces – `public` or package-private (no keyword)
- **Member level:** The access modifiers are, in order from most restrictive to least restrictive, `private`, package-private, `protected`, and `public`

Let's discuss these in turn.

### **private**

A member marked as `private` is accessible within its own class only. In other words, the block scope of the class defines the boundary. When in a class (scope), you cannot access the `private` members of another class, even if you have an object reference to the class containing the `private` member.

### **Package-private**

There is no special keyword for `package-private`. If a type (class, interface, record, or enum) has no access modifier then `package-private` is applied. Types that are `package-private` are only visible within the same package. Recall that a package is simply a named group of related types.

The `Person` class (line 3) in *Figure 8.11* is a `package-private` class, meaning `Person` cannot be imported into another package. In addition, the `Person` constructor (line 6) is `package-private`, meaning that you cannot create an object of the `Person` type from within a different package.

At a **member** level, there are a few exceptions to the preceding text that you need to be aware of when you omit the access modifier:

- Class/record members are, as above, `package-private` by default.
- Interface members are `public` by default.
- Enum constants (members) are `public static` and `final` by default. Enum constructors are `private` by default. We will discuss enums later in the chapter.

### The default package

The default package is also known as the package with no name or the unnamed package. Types that have no explicit package statement at the top of the file are put into this package. This is the package where the `Person` and `PersonExample` classes from *Figure 8.11* are placed.

The implications of this are that, given the package has no name, if we are in a different (named) package, we have no way of importing `Person` and `PersonExample`. The fact that `PersonExample` is `public` (line 16) makes no difference. Therefore, only other types in the same (default) package can access them.

## **protected**

A member marked as `protected` means that it's visible within its own package (as with package-private) but also visible to subclasses outside of the package. We will discuss subclasses and `protected` in more detail when we cover inheritance in *Chapter 9*.

## **public**

A type or member marked as `public` is visible everywhere. Thus, no boundaries apply.

*Table 8.2* summarizes the access modifiers and their visibility:

Access Modifier	Class	Package	Subclass (anywhere)	Everywhere
<code>private</code>	Y	N	N	N
package-private	Y	Y	N	N
<code>protected</code>	Y	Y	Y	N
<code>public</code>	Y	Y	Y	Y

Table 8.2 – Access modifiers and their visibility

Let's examine *Table 8.2* horizontally. Only the class has access to members marked as `private`. If a class or member has no access modifier (package-private), then that class or member is only visible within the class and the package. If the member is marked as `protected`, then the member is visible to the class, package, and subclasses of that class, regardless of the package. Finally, if a class or member is marked as `public`, then the class or member is visible everywhere.

To further help explain *Table 8.2*, let's diagram an example suite of classes and their associated packages and draw up another visibility table specifically for it. *Figure 8.13* shows this:

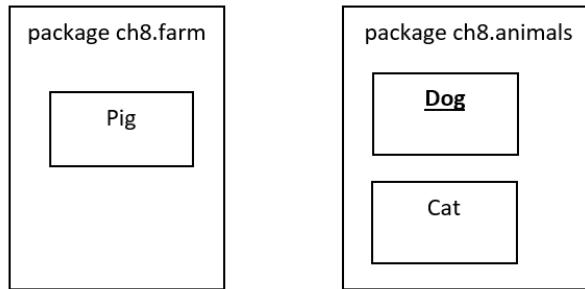


Figure 8.13 – Sample example access modifiers diagram

In this figure, the Dog class is in bold and underlined because the following table, *Table 8.3*, represents the visibility of *its* members. For example, when reading the `private` row, assume that we have marked a member in Dog as `private` and are determining its visibility in the other classes. Let's examine *Table 8.3*:

Access Modifier	<u>Dog</u>	Cat	Pig
<code>private</code>	Y	N	N
<code>package-private</code>	Y	Y	N
<code>protected (*)</code>	Y	Y	N
<code>public</code>	Y	Y	Y

Table 8.3 – Visibility when modifiers are applied to a Dog member

Thus, if a Dog member is `private`, only Dog can see it. If the Dog member is `package-private`, only Dog and Cat can see it. If the Dog member is `protected`, Dog and Cat can see it. Lastly, if the Dog member is `public`, every class can see it.

(\*) We will complete this table when revisiting `protected` in the inheritance chapter.

#### How do access levels affect you?

Access levels will affect you in two ways. Firstly, you could be using an external class (from the Java API, for example) and want to know if you can use that class and/or its members in your code. Secondly, when writing a class, you will want to decide the access level each class and member will have. A good rule of thumb is to keep your members as `private` as possible to avoid misuse. Additionally, avoid `public` fields unless they are constants. We will discuss this further when we discuss encapsulation.

Let's look at these access modifiers in code. In particular, we will focus on the package and learn how to create one and how access is affected by its boundary.

## packages

Recall that the fully qualified type name includes the package name. A *package defines a namespace*. For example, in *Figure 8.13*, the Dog class in the ch8.animals package is fully qualified as ch8.animals.Dog. Therefore, a Dog class in a package named kennel would have a qualified name of kennel.Dog; which is completely different to ch8.animals.Dog. Thus, Java can distinguish between the two Dog types and no name collisions occur. As we shall see, the package structure is also used as a directory structure for your java files. Oracle gives very good guidelines (see <https://docs.oracle.com/javase/tutorial/java/package/namingpkgs.html>) on how to name your packages so that your types do not conflict with someone else's. Package names are written in all lowercase letters to differentiate them from type names. Following that, companies should use reverse internet domain names to begin their package names. For example, if you work at a company called somecompany.com and you are creating a package called somepackage, then the full package name should be com.somecompany.somepackage. Within a company, naming can then follow company conventions, such as including the region: com.somecompany.region.somepackage.

Let's examine the packages from *Figure 8.13*. We will start with ch8.animals:

```
1  package ch8.animals;
2
3  public class Dog {
4      private String dogName;
5      protected int age;
6      public Dog(String dogName) { this.dogName = dogName; }
7      public String getDogName() { return dogName; }
8      void pkgPrivate(){}
9  }
10 class Cat{ // package private
11     Cat(){}
12     public void testDogAccess(){
13         Dog d = new Dog(dogName: "Rex");
14         d.dogName = "Abc"; // dogName is private to Dog
15         d.age = 2;
16         d.pkgPrivate(); //ok
17     }
18 }
19 }
```

Figure 8.14 – The “ch8.animals” package from Figure 8.13

In this figure, note that, for simplicity, we have grouped the two classes in the package into one Java file. This file is called `Dog.java` (as the `public` class is `Dog`). The first line is important: `package ch8.animals` states that the types (classes and so forth), that are defined here, go into this package. In addition, the file `Dog.java` will be put into a folder on the hard disk named `ch8\animals`.

In this figure, line 4 defines a `private` instance variable called `dogName`. This is accessible within the class only (as per lines 6 and 9) but not outside the class (as per line 18).

Line 5 defines a `protected` instance variable called `age` which we can access from another class within the package (line 19). Line 12 defines a package-private method called `pkgPrivate()` and line 20 shows that we can access it from another class in the same package. Note also that the `Cat` class and its constructor are both package-private (lines 14 and 15, respectively).

*Figure 8.15* shows the other package, `ch8.farm`:

```
1  package ch8.farm;
2
3  import ch8.animals.Dog; // class is public, ok
4  //import ch8.animals.Cat; // class is pkg-private, error
5
6  public class Pig{
7      void testDog(){
8          Dog d = new Dog(dogName: "Shep"); // constructor is public
9          // d.pkgPrivate(); // package-private method, error
10     }
11 }
```

Figure 8.15 – The “ch8.farm” package from Figure 8.13

Again, note that line 1 states the package name – this is the `ch8.farm` package. The filename is `Pig.java` (as the `public` class is `Pig`) and the file will be put into a folder on the hard disk named `ch\farm`.

Note the use of the fully qualified names when importing (lines 3 and 4). As we want access to the `Dog` class which resides in a separate package, we must import it. There is no issue importing `Dog` since it is `public`. However, we are unable to import `Cat` as `Cat` is package-private (and we are in a different package).

Line 8 demonstrates that `Pig` can create a `Dog` object. Note that there are two access points here: the `Dog` class is `public` (so we can import it); and the `Dog` constructor is also `public` (so we can create an instance of `Dog` from code in a different package). This is why the access modifiers for the class and constructors should match. Line 9 shows that, when we are in a different package, we do not have access to package-private members from another package.

Now that we understand access modifiers, we are in a position to discuss encapsulation.

## Encapsulation

As previously stated, encapsulation is a key concept in OOP. The principle here is that you protect the data in your class and ensure that the data can only be manipulated (retrieved and/or changed) via your code. In other words, you have control over how external classes interact with your internal state (data). So, how do we do this?

### Achieving encapsulation

Basic encapsulation is very easy to achieve. You simply mark your data as `private` and manipulate the data via `public` methods. Thus, external classes cannot access the data directly (as it is `private`); these external classes must go through your `public` methods to retrieve or change the data.

These `public` methods make up the class's "interface"; in other words, how you interact with the class. This "interface" (group of `public` methods) is very different from and not to be confused with the `interface` language construct (*Chapter 10*). *Figure 8.16* presents a code example to help us further develop this topic:

```
1  package ch8;
2
3  class Adult{
4      private String name;
5      private int age;
6
7      Adult(String name, int age) {
8          this.age = age;
9          this.name = name;
10     }
11
12     public String getName() { return name; }
13     public void setName(String name) { this.name = name; }
14     public int getAge() { return age; }
15     public void setAge(int age) { this.age = age; }
16
17 }
18
19 public class BasicEncapsulation {
20     public static void main(String[] args) {
21         Adult john = new Adult(name: "John", age: 20);
22         System.out.println(john.getName() + " "
23                             + john.getAge()); // John 20
24         //john.age = -99; // 'age' is private
25         john.setAge(-99); // uh-oh!
26         System.out.println(john.getName() + " "
27                             + john.getAge()); // John -99
28     }
29 }
30
31 }
```

Figure 8.16 – Basic encapsulation in action



In the preceding figure, the `Adult` class has two `private` instance variables, namely `name` and `age` (lines 4 and 5, respectively). Thus, these instance variables only have access within the `Adult` block of code. Note that even having an `Adult` object reference cannot bypass this access rule – the compiler error on (the commented out) line 29 demonstrates this.

#### public class name and filename relationship

In Java, the name of the `public` class must match the filename. In *Figure 8.16*, the `public` class is `BasicEncapsulation`. This means that the filename must be named `BasicEncapsulation.java`, which it is. This rule implies that you cannot have two `public` classes in the same file – that is why the `Adult` class is not `public`.

What if we were in a different package and we wanted to create an `Adult` object, as defined in *Figure 8.16*? This is an issue because `Adult` is package-private (line 3). To fix this issue, we need to make the `Adult` class `public` so that we can `import` it when in a different package. This means that we need to move the `Adult` class into a separate file, named `Adult.java`. In addition, both `Adult` and its constructor would need to be `public`. Why? Well, when we're in a different package, the class being `public` enables us to `import` the class and the constructor being `public` enables us to create objects of the `Adult` type.

The `Adult` constructor (line 7) has no access modifier and is therefore package-private. Thus, only classes within the same package can invoke this constructor. In other words, only classes in the `ch8` package (line 1) can create `Adult` objects. As `BasicEncapsulation` is also in `ch8`, the object creation on line 26 is fine.

The rest of the `Adult` class (lines 11-20) provides the getter/setter method pairs for manipulating the object state (the instance variables). These getter/setter methods are also known as accessor/mutator methods, respectively. There is usually a pair for each instance variable and they follow this format (note that this is just an example):

```
public int getAge() {
    return age;
}
public void setAge(int age) {
    this.age = age;
}
```

After creating the `Adult` object on line 26 in *Figure 8.16*, we output the object state using the `public` accessor methods, `getName` and `getAge` (lines 27-28). As these accessor methods are `public`, these methods are available to any class in any package. Given that 'John' and 20 are output, we know our object was created correctly.

Let's assume that we are the developers of the `Adult` class and require an adult to be 18 years or older. In addition, we will assume that the developer of the `BasicEncapsulation` class is unknown to us. Line 29 demonstrates that as our `Adult` data is `private`, it is protected from direct external corruption. This is exactly what encapsulation provides; it is its *raison d'être*!

Line 30 demonstrates that the object's state can still be corrupted. However, the corruption that's done via the `set`/mutator method on line 30 is very different from the direct corruption on line 29. As the author of the `Adult` class, we can control and therefore fix the corruption error in our `set` methods. The issue with our `set` (mutator) method is replicated in the constructor. *Figure 8.17* addresses this (internal) corruption issue in both the constructor and mutator methods:

```
1  package ch8;
2
3  class Adult{
4      private String name;
5      private int age;
6
7      Adult(String name, int age) {
8          setAge(age);
9          this.name = name;
10     }
11
12     public String getName() { return name; }
13     public void setName(String name) { this.name = name; }
14     public int getAge() { return age; }
15     public void setAge(int age) {
16         if(isAgeOk(age)){
17             this.age = age;
18         }else{
19             this.age = -1; // error state
20         }
21     }
22
23     private boolean isAgeOk(int age){ // private
24         return age >= 18 ? true : false ; // ternary operator
25     }
26 }
```

Figure 8.17 – Ensuring “age” is at least 18

As `BasicEncapsulation` remains unchanged, it is not included in the preceding figure. Note that a new `isAgeOk` method has been introduced (lines 27-29). This method takes in an `int` parameter `age` and checks to see if it is  $\geq 18$ . If so, the method returns `true`; otherwise, it returns `false`.

The `isAgeOk` method is invoked from the `setAge` mutator method (line 21). As the constructor calls `setAge` (line 8), it also avails of the age check logic. If an invalid age is passed into the constructor or `setAge`, an error value of `-1` is set. Note that there are better ways to do this, but for now this is fine. When we run the program now, since the `age` value that is being passed into `setAge` is `-99` (`john.setAge(-99)`), the instance variable `age` is set to the error value of `-1`.

That covers basic encapsulation. We will now discuss a particular issue with basic encapsulation and how advanced encapsulation resolves it.

## Mastering advanced encapsulation

The simple maxim of “private data, public methods” (where the `public` methods manipulate the data) goes a long way to ensuring proper encapsulation of your data. However, you are not completely safe just yet. In this section, we will review Java’s call by value principle, which is used when passing arguments to and returning values from methods. We will examine how this can present a subtle issue. Lastly, we will examine how to protect your code from encountering this issue in the first place.

### Call By value revisited

In *Chapter 7*, we discussed how, when passing arguments to methods, Java’s *call by value* mechanism creates *copies* of those arguments. We saw the need to be aware that when the argument is a reference, such as to an array, the called method can now manipulate the array object that the caller method is looking at.

Similarly, when a method is *returning* something, call by value applies again. In other words, a copy is made of what you are returning. Depending on what the copy is of, this can result in encapsulation being broken or not. If you are returning `private` primitive data, then there is no issue – a copy of the primitive is returned and the client can do whatever it likes to the copy; your `private` primitive data is safe. As you may recall from *Chapter 7*, copying primitives is like photocopying a sheet of paper. The photocopied sheet can be written on without it affecting the original copy.

### The issue

The issue arises if your `private` data is a reference (to an object). If the client receives a copy of the reference, then the client can manipulate your `private` object! From *Chapter 7*, you may recall that copying a reference is like copying a remote control to a TV. The new remote can change the channels on the same TV. *Figure 8.18* presents code that breaks encapsulation:

```
5  class Seniors {
6      private int[] ages = new int[2];
7      private int num;
8
9      Seniors(){
10         num = 2;
11         ages[0] = 30;
12         ages[1] = 40;
13     }
14     public int getNum() { return num;}
15     public int[] getAges() { // breaks encapsulation
16         return ages;
17     }
18 }
19 ➤ public class AdvancedEncapsulation {
20 ➤     public static void main(String[] args) {
21         Seniors seniors = new Seniors();
22         // 1. Returning primitives is okay.
23         int num = seniors.getNum();
24         System.out.println(num); // 2
25         num = -100;
26         num = seniors.getNum();
27         System.out.println(num); // 2, ok, primitives are encapsulated once 'private'
28
29         // 2. Returning references requires care.
30         int[] copyAges = seniors.getAges(); // 'copyAges' and 'ages' refer to the same array object!
31         System.out.println(copyAges[0] + ", " + copyAges[1]); // 30, 40
32         // As we have a copy of the internal array reference, we can, from HERE
33         // change the "private" internal Seniors array! This breaks encapsulation.
34         copyAges[0] = -9;
35         copyAges[1] = -19;
36         int[] copyAges2 = seniors.getAges();
37         System.out.println(copyAges2[0] + ", " + copyAges2[1]); // -9, -19
38     }
39 }
```

Figure 8.18 – Code that breaks encapsulation

In the preceding figure, the `Seniors` class has two `private` instance variables (lines 6-7), namely `ages` and `num`. The constructor (lines 9-13) initializes the instance variables. We have a public `getNum` accessor method, which returns the `private` instance variable, `num` (line 14). Note that we have put this method on one line in the interest of space.

We have another accessor method called `getAges` (lines 15-17) that returns a `private` array called `ages`. *Line 16 is the problem* as it breaks encapsulation. We will explain why when we discuss the code in `main`.

In `main`, the first thing we do is create an instance of `Seniors` (line 21). This is so we can access the instance methods defined in `Seniors`. The rest of `main` is divided into two sections: one section (lines 23-27) demonstrates that returning `private` primitive data is fine; the other section (lines 30-37) demonstrates that simply returning `private references` breaks encapsulation.

Let's examine the first section. Line 23 initializes the local variable, `num`, based on the return value from `seniors.getNum()`. As the `private` `Seniors` instance variable, `num`, was initialized to 2 in the `Seniors` constructor (line 10), the (completely separate) local variable, `num`, is initialized to 2. We output this fact on line 24. We then change the local `num` variable's value to -100 (line 25). The question now is, when we changed the local variable `num`, was the `private` `Seniors` instance variable, `num`, changed also? To find out, we can simply retrieve `num` again using the `public` accessor method, `getNum` (line 26). Line 27 outputs 2, proving that the `private` primitive, `num`, was safe from changes made in `main`.

The second section is where things get interesting. Line 30 initializes a local variable called `copyAges` based on the return value from the `public` accessor method, `seniors.getAges()`. As `getAges` simply returns (a copy of) the `private` `ages` reference, we now have two references referring to the one array object. These references are the `private` instance variable, `ages`, and the local variable, `copyAges`. Line 31 outputs the values of `copyAges`, which are 30 and 40 for the indices 0 and 1, respectively. These are the same values that the `private` `ages` array was initialized to in the `Seniors` constructor (lines 11-12).

Now, on lines 34-35, we change the values of the `copyAges` array: index 0 is set to -9 and index 1 is set to -19. As with the first section, we are now wondering, did changing the local array have any effect on the `private` instance array in `Seniors`? The answer is yes! To prove this, we can retrieve the `private` array again using `getAges` (line 36) and output its values (line 37). The output values of -9 and -19 demonstrate that the client, `AdvancedEncapsulation`, was able to manipulate (change) the so-called `private` data of `Seniors`. Therefore, `Seniors` is not encapsulated after all.

*Figure 8.19* shows the situation in memory, shedding light on why `Seniors` is not encapsulated:

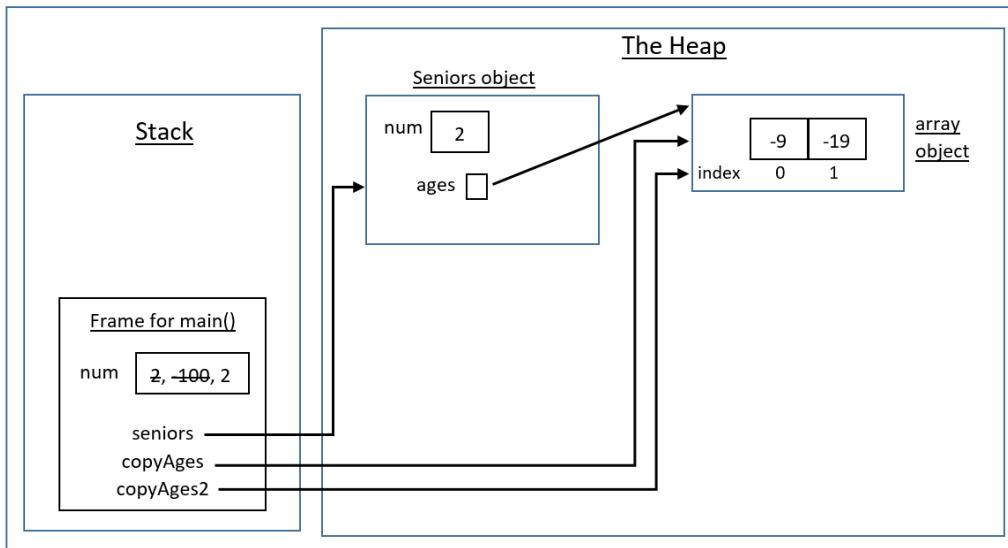


Figure 8.19 – In-memory representation of Figure 8.18 (at line 37)

In the preceding figure, the local variable, num, is on the stack. It is a copy of the private Seniors num instance variable, and its different values as we progress through main are reflected in strikethrough font. Line 25 (*Figure 8.18*) changes the local variable to -100. As can be seen, this change does not affect the private instance variable, num, in Seniors, which remains 2.

The issue is with the reference to the private array object, ages. Because getAges (line 15) simply returns the reference, a copy of that reference is stored in the local variable, copyAges (line 30). As the local reference, copyAges, and the private reference, ages, now refer to the same object, the copy reference can change the private array object. That is why the array object has values of -9 and -19 for indices 0 and 1, respectively. The copyAges2 reference is just there to prove that point.

## The solution

Now that we know the issue, fixing it is quite straightforward. The key is to, when returning a reference, ensure that you simply do not return the private reference (as call by value will return a copy of that reference). The solution is to *make a copy of the object you wish to return and return a reference to the new object*. Thus, the external class (client) can manipulate this new object without affecting your private internal object. *Figure 8.20* is the properly encapsulated, refactored version of *Figure 8.18*:

```
3 import java.util.Arrays;
4
5 class Seniors {
6     private int[] ages = new int[2];
7     private int num;
8
9     Seniors(){
10         num = 2;
11         ages[0] = 30;
12         ages[1] = 40;
13     }
14     public int getNum() { return num;}
15     public int[] getAges() { // properly encapsulated
16         int newArr[] = Arrays.copyOf(ages, newLength: 2);
17         return newArr;
18     }
19 }
20 public class AdvancedEncapsulation {
21     public static void main(String[] args) {
22         Seniors seniors = new Seniors();
23
24         int[] copyAges = seniors.getAges(); // 'copyAges' and 'ages' refer to 2 different arrays
25         System.out.println(copyAges[0] + ", " + copyAges[1]); // 30, 40
26         copyAges[0] = -9;
27         copyAges[1] = -19;
28         int[] copyAges2 = seniors.getAges();
29         System.out.println(copyAges2[0] + ", " + copyAges2[1]); // 30, 40
30     }
31 }
```

Figure 8.20 – Properly encapsulated code

In the preceding figure, we have replaced the accessor method, `getAges`, with a new version (lines 15-18). This new version is properly encapsulated. On line 16, instead of simply returning the (reference to the) array instance variable, we are copying the array, `ages`, into a new array, namely `newArr`. We achieve this using the `Arrays . copyOf` method. We return a (copy of the) reference to the new array object.

Now, on line 24, when we initialize `copyAges`, it is referring to the copy array that was created on line 16. That reference, `newArr`, has gone out of scope (since we returned from `getAges`) but the new array object is still on the heap, with `copyAges` referring to it. The important point here is that on line 25, we have two distinct array references: the `ages` instance and the local `copyAges`. These references now refer to two *different* objects.

Line 25 outputs the details of the copy array; 30 for index 0 and 40 for index 1. This is as expected. Lines 26 and 27 change the contents of the copy array indices, 0 and 1, to -9 and -19, respectively. Now, we need to check something: when we changed the contents of the `copyAges` array, were the contents of the private internal `Seniors` array's ages changed? To check, on line 28, we can initialize a `copyAges2` array with the (copy of the) contents of the private array, `ages`. When we output the details of `copyAges2` on line 29, we get 30 and 40, thereby proving that the private internal array, `ages`, was *not* changed when we changed the local `copyAges` array (lines 26-27). Now, `Seniors` is properly encapsulated.

*Figure 8.21* show this situation in memory as we execute line 29:

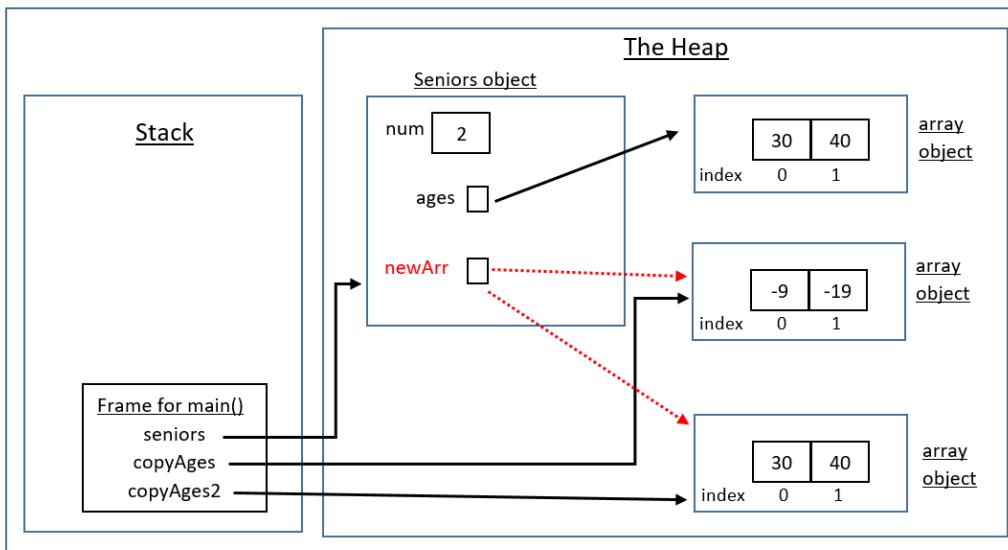


Figure 8.21 – In-memory representation of Figure 8.20

In the preceding figure, just after the `Seniors` object is constructed (line 22), we have a `seniors` reference on the stack referring to a `Seniors` object on the heap. The `Seniors` object contains a `num` primitive set to 2 (line 10) and an `ages` array reference referring to the array object (lines 11-12).

When we call `getAges` (line 24), the copy array, `newArr`, is created (line 16) and although not shown here, the new array initially contains the values of 30 and 40 (indices 0 and 1, respectively), as per line 25. When `newArr` is returned from `getAges` (line 17), the (copy of the) reference is assigned to `copyAges` (line 24). As shown in the preceding diagram, the `copyAges` local variable and the `ages` instance variable refer to two different array objects. This is what we want. *Any changes made using copyAges will not affect the private array ages.*

This is what the changes on lines 26-27 demonstrate. The changes that were made using the `copyAges` reference are reflected in the diagram. To prove that the changes on lines 26-27 did not affect the `private` array, `ages`, we call `getAges` again. A new array, representing a copy of the `private` array, is again created (line 16) and the (copy of) the new array reference is returned and assigned to the local reference, `copyAges2`. When we output the new array's contents on line 29, we get 30 and 40, demonstrating that the `private` array is unaffected by changes to the local array (lines 26-27).

Now that we understand call by value and advanced encapsulation, we are in an excellent position to discuss the object life cycle.

## Delving into the object life cycle

To understand Java, it is extremely helpful to have an appreciation of what is happening in the background, in memory. This section will help cement what is happening on the stack and the heap when we call methods, declare local/instance variables, and so forth.

Local variables are kept on the stack (for fast access), whereas instance variables and objects live on the heap (a large area of memory). As we know, we use the `new` keyword to create a Java object. The `new` keyword allocates space on the heap for the object and returns the reference to the object. What happens if the object is no longer accessible? For example, the reference may have gone out of scope. How do we reclaim that memory? This is where garbage collection comes into play.

## Garbage collection

As mentioned previously, garbage collection reclaims memory taken up by objects that are no longer being used; as in the objects have no references pointing to them. This garbage collection process is a JVM process that runs in the background. The JVM may decide during an idle time to run garbage collection and then it may not. Simply put, we have no control over when garbage collection runs. Even if we invoke `System.gc()`, this is but a suggestion to the JVM to run garbage collection – the JVM is free to ignore this suggestion. The major advantage of garbage collection is that we do not have to do the tidy-up ourselves; whereas in languages such as C++, we do.

For further detail on Java Memory Management please see our previous book: [https://www.amazon.com/Java-Memory-Management-comprehensive-collection/dp/1801812853/ref=sr\\_1\\_1?crid=3QUEBKJP46CN7&keywords=java+memory+management+maaike&qid=1699112145&sprefix=java+memory+management+maaike%2Caps%2C148&sr=8-1](https://www.amazon.com/Java-Memory-Management-comprehensive-collection/dp/1801812853/ref=sr_1_1?crid=3QUEBKJP46CN7&keywords=java+memory+management+maaike&qid=1699112145&sprefix=java+memory+management+maaike%2Caps%2C148&sr=8-1).

## Object life cycle example

A sample program will help at this point. *Figure 8.22* presents a program to suit our purposes:

```
3  class Tag{}
4 ►  public class Cow {
5      Tag tag;
6      String country;
7
8 ►   public static void main(String[] args) {
9      Cow cow1 = new Cow();
10     Cow cow2 = cow1;// reassignment
11     cow2.tagAnimal(cow1);
12 }
13 @ ► void tagAnimal(Cow cow){
14     tag = new Tag();
15     cow.setCountry("France");
16 }
17 ► void setCountry(String country){
18     this.country = country;
19 }
20 }
```

Figure 8.22 – Sample program to explain an object's life cycle

As this (simple and very contrived) program executes, three methods are pushed onto the stack, namely `main`, `tagAnimal`, and `setCountry`. *Figure 8.23* represents the in-memory representation when we are just about to exit the `setCountry` method (line 19):

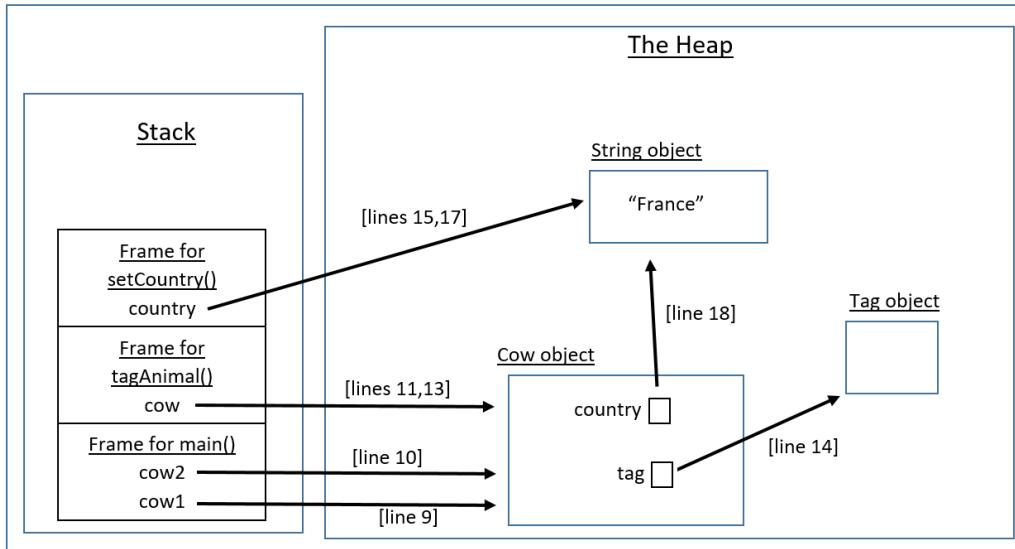


Figure 8.23 – In-memory representation of code in Figure 8.22

Let's look at this in more detail.

### ***The main method***

As can be seen from the previous two figures, line 9 creates the Cow object on the heap, and the local reference, `cow1`, on the stack in the frame for `main`, refers to it. The instance variables in the Cow object on the heap, namely `tag` and `country`, will be null at this point.

Line 10 assigns the value in `cow1` to another local reference in `main`, namely `cow2`. Now, at line 11, we have a frame for `main` on the stack with two local reference variables, namely `cow1` and `cow2`, both referring to the one Cow object on the heap.

Line 11 uses the `cow2` reference to execute the instance method, `tagAnimal`. Thus, when inside the `tagAnimal` method (during this invocation), the `this` reference will be referring to whatever `cow2` is referring to (which is the Cow object on the heap). In addition, the `cow1` reference is passed as an argument to the `tagAnimal` method. This is not necessary as `tagAnimal` already has a reference to the Cow object (using `this`) but this program is just for example purposes.

### ***The tagAnimal method***

As with any method invocation, a stack frame for `tagAnimal` is pushed on the stack. As per call by value rules, `tagAnimal` (line 13) aliases the method parameter `cow` for `cow1` from line 11 (the method call). Thus, the `cow` reference in `tagAnimal` and the `cow1` reference in `main` are pointing at the same `Cow` object, which was created on line 9.

As we know, the `this` reference refers to the object instance responsible for the method call – in this case, `cow2` (line 11). Therefore, the reference to `tag` on line 14 (which is `this.tag` in effect) is referring to the `tag` instance variable that can be accessed via `cow2`. As a result, line 14 creates a new `Tag` object on the heap and stores its reference in the `tag` instance variable of the `Cow` object, overwriting its previous default value of `null`. Note that at this point, given the contrived nature of this example, the `Cow` object is referred to by three different references: `cow1` and `cow2` in `main`; and `cow` in `tagAnimal`.

Line 15 specifies a `String` literal of "France". As `String` literals are objects, a `String` object is created on the heap. Using the `cow` reference, the `setCountry` method is called, passing down the `String` literal, "France".

### ***The setCountry method***

A stack frame for `setCountry` is pushed onto the stack. The `setCountry` declaration aliases the method parameter `country` to refer to the `String` literal, "France", which is passed down in the method call (line 15). Line 18 initializes the `country` instance variable to the argument passed down, namely "France". Line 18 explicitly uses the `this` reference because the parameter name and instance variable have the same identifier, `country`. The `this` reference refers to whatever `cow` is referring to, which is the `Cow` object on the heap. This is because the `setCountry` method call (line 15) was executed on the reference `cow`.

Now that we know how methods are pushed onto the stack, let's examine the memory as we return from these method calls – in other words, as we pop the stack. *Figure 8.24* represents memory after we have exited the `setCountry` method but before we exit the `tagAnimal` method:

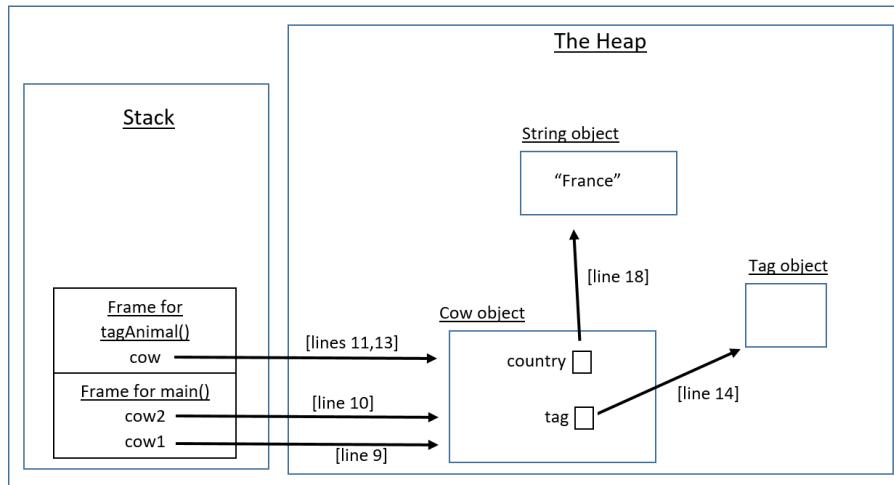


Figure 8.24 – In-memory representation after the “setCountry” method completes

As can be seen from the preceding figure, the `setCountry` frame has been popped from the stack. However, the `String` object, `"France"`, remains on the heap because the `country` instance variable from the `Cow` instance object still refers to it. Only objects that have no references pointing to them are eligible for garbage collection.

*Figure 8.25 represents the in-memory representation just after `tagAnimal` finishes but before `main` completes:*

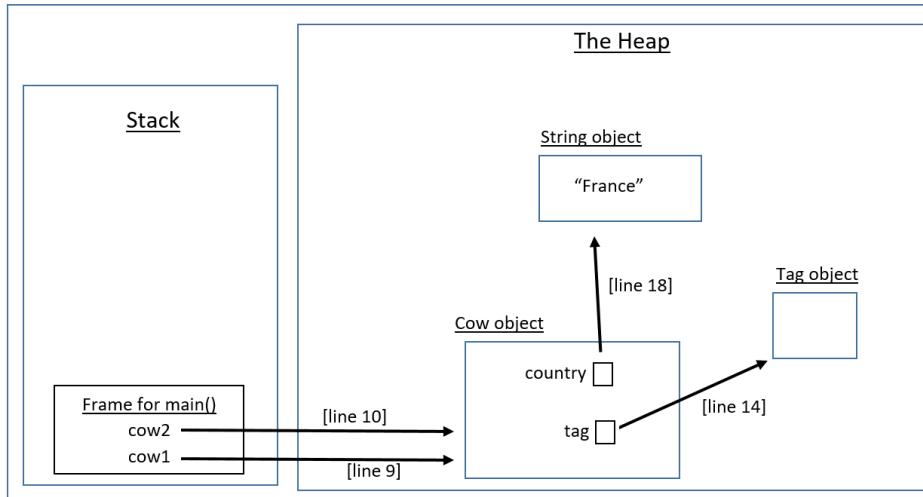


Figure 8.25 – In-memory representation after the ‘tagAnimal’ method completes

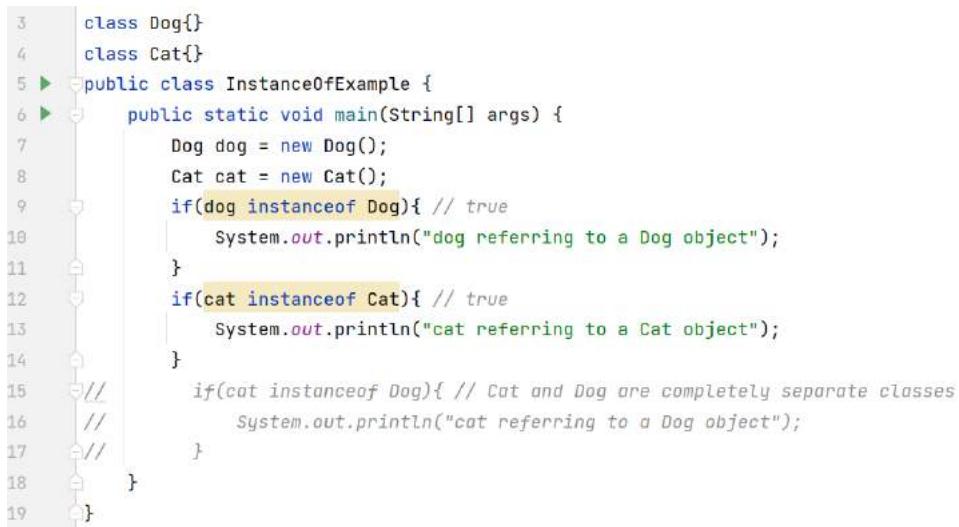
There is very little change in this figure from the previous figure, except that the stack frame for `tagAnimal` has been popped. The `Cow` object on the heap cannot be garbage collected because both the references, `cow1` and `cow2`, in `main` refer to it. In addition, because the `Cow` object cannot be removed, neither can the `Tag` or `String` objects. This is because the `Cow` instance variables, `tag` and `country`, refer to them. This figure represents the situation in memory until `main` exits, at which point everything can be reclaimed.

That concludes our discussion on an object's life cycle. We will now move on and discuss the `instanceof` keyword.

## Explaining the `instanceof` keyword

The `instanceof` keyword enables us to determine the object type that a reference is referring to. That is why it is so critical to separate the reference from the object. The reference's type and the object's type are often very different. In fact, in most cases, they are different. We will discuss `instanceof` in greater detail when we cover inheritance (*Chapter 9*) but also when we discuss interfaces (*Chapter 10*).

So, for the moment, we will keep it simple – where the reference type and object type are the same. *Figure 8.26* presents one such code example:



```
3  class Dog{}
4  class Cat{}
5  public class InstanceOfExample {
6      public static void main(String[] args) {
7          Dog dog = new Dog();
8          Cat cat = new Cat();
9          if(dog instanceof Dog){ // true
10              System.out.println("dog referring to a Dog object");
11          }
12          if(cat instanceof Cat){ // true
13              System.out.println("cat referring to a Cat object");
14          }
15          if(cat instanceof Dog){ // Cat and Dog are completely separate classes
16              System.out.println("cat referring to a Dog object");
17          }
18      }
19  }
```

Figure 8.26 – Basic “`instanceof`” example

In this figure, line 7 creates a Dog object referred to by a Dog reference named dog. Line 8 creates a Cat object referred to by a Cat reference named cat. Line 9 checks if the object at the end of the dog reference is “an instance of” Dog. It is, so line 10 executes. Similarly, line 12 checks to see if the object referred to by cat is of the Cat type. It is, so line 13 executes.

Line 15 is commented out as it generates a compiler error. As Cat and Dog are completely unrelated classes (lines 3-4), the compiler knows that there is no way a Cat reference, namely cat, can refer to a Dog object. Conversely, a Dog reference, such as dog, cannot refer to a Cat object.

We will come back to `instanceof` later in this chapter. For now, let us move on to our next topic, which is closely related to classes: namely, enumerations.

## Understanding enums

**Enumerations**, or **enums** for short, are a special type of class. Whereas with a class, you can have as many instances (of the class) as you wish; with enums, the instances are predefined and therefore restricted. Enums are very useful for situations where a finite set of values apply – for example, days of the week, seasons of the year, and directions.

This ensures *type-safety* because, with the help of the compiler, only the instances defined are allowed. It is always better to find an issue at compile time than runtime. For example, if you had a method that defined a `String` parameter, namely `direction`, then someone could invoke the method with "WESTT" (note the incorrect spelling). The compiler would not catch this error as it is a valid `String`, so the error would manifest at runtime. If, however, the method parameter were an enum instead, the compiler would catch it. We will see this shortly.

There are two types of enums: simple and complex. We will discuss them now.

### Simple enums

A simple enum is named as such because it is, well, simple. This is in the sense that when you look at the enum, there is very little code present. *Figure 8.27* presents code using a simple enum:

```

3 enum Water {
4     STILL, SPARKLING;
5 }
6 public class SimpleEnums {
7     public static void main(String[] args) {
8         // Water stillWater = new Water(); // compiler error
9         Water stillWater = Water.EXTRA_SPARKLING; // type safety
10
11         Water stillWater = Water.STILL;
12         System.out.println(stillWater == Water.STILL);           // true
13         System.out.println(stillWater.equals(Water.STILL));    // true
14         switch(stillWater){
15             case STILL:
16                 System.out.println("Still water");
17                 break;
18             case Water.STILL: // unqualified enum value required
19             case 0: // cannot use an int
20         }
21         if(Water.STILL == 0){} // Water == int
22         Water sparklingWater = Water.valueOf( name: "SPARKLING");
23         System.out.println(sparklingWater); // SPARKLING
24
25         for(Water water: Water.values()){
26             // Ordinal value of: 0 is STILL
27             // Ordinal value of: 1 is SPARKLING
28             System.out.println("Ordinal value of: " + water.ordinal() + " is " + water.name());
29         }
30
31     }
32 }
```

Figure 8.27 – A simple enum

In the preceding figure, the `Water` enum is defined (lines 3-5). The values of an enum are expressed in capital letters (similar to constants). It is not mandatory to do this but it is common practice. What this enum is saying is that we have an enum named `Water` and there are only two instances allowed, namely `STILL` and `SPARKLING`. In effect, `STILL` and `SPARKLING` are references to the only object instances allowed. The semicolon at the end of line 4 is optional for simple enums. The corresponding semicolon for complex enums is mandatory. The enum values are given ordinal values starting at 0. So, for `Water`, `STILL` has an ordinal value of 0 and `SPARKLING` has an ordinal value of 1.

As stated previously, enums are a special type of class. However, there are some differences. One is that enum constructors are `private` by default. This includes the default constructor generated by the compiler (as in *Figure 8.27* for `Water`). Contrast this with the default constructor of a class, which has the same access as the class itself. Thus, you cannot instantiate an enum as you would a normal object. This is why line 8 will not compile – the default enum constructor generated by the compiler is `private` and therefore inaccessible to external types.

So, if we cannot new an enum, how do we create an enum instance? In other words, where are the constructor calls? *The declaration of the enum values, STILL and SPARKLING, (line 4) are the constructor calls!* As they are within the class, they have access to the private constructor. These enum values are initialized only once – that is, when the enum is first used.

So, to create an enum (object), use the relevant enum value. This is done on line 11, where we now have a reference, `stillWater`, referring to the `STILL` instance. Contrast line 11 with line 9 (which does not compile). Attempting to use any other value such as `EXTRA_SPARKLING` will not compile. This is the type safety we discussed previously. Only two instances of `Water` are allowed, `STILL` and `SPARKLING`, and the compiler enforces this rule.

Lines 12 and 13 demonstrate that only one instance of `Water . STILL` is created. As the equivalence operator and the `equals` method both return `true`, there can be only one instance.

### Inherited methods

Although inheritance will be discussed in detail in *Chapter 9*, we need to dip into the topic to understand enums. Every class in Java implicitly inherits from a class called `Object`. This means that there are methods in `Object` that you get by default. This is how Java ensures every class has certain important methods. You can accept the version from `Object` or replace it (known as *overriding* the method).

One of these methods that's inherited from `Object` is `equals`. The version in `Object` compares the references to see if they are equal and returns `true` or `false` depending on that comparison. Essentially, this is the same as using `==` to compare the references.

Enums implicitly inherit from the `Enum` class (and `Enum` inherits from `Object`, so there is no escaping `Object`!). Thus, enums have access to methods such as `valueOf`, `values`, `ordinal`, and `name`.

The `switch` statement (lines 14-20) switches on the `Water` reference, namely `stillWater` (line 14). The `case` label is the unqualified enum value (`STILL`, line 15). Line 18 shows that the qualified enum value is incorrect. Line 19 (and line 21) demonstrate that even though enum values have ordinal values, enums are types and not integers.

Several interesting methods in the `Enum` type are available to us due to inheritance. Let's start with `valueOf(String)`.

### ***The valueOf(String) method***

This is an implicitly declared method, which, when given one of the enum constant names, returns that enum instance (line 22). Thus, this method provides a quick and easy way to create an enum instance, once you know the constant name.

Let's examine how we can iterate over all the enum instances using the `values()` method.

### ***The values() method***

This is another implicit method. On line 25, we use an enhanced `for` loop to iterate over the enums in the order they are declared on line 4, namely `STILL` followed by `SPARKLING`. Once we have an enum instantiated, we can use other methods to get details of that particular enum.

Let's see how the `ordinal()` method provides the ordinal number for the enum.

### ***The ordinal() method***

The `ordinal()` method (line 28) returns the ordinal value of this enum. The initial enum constant is given an ordinal value of 0; therefore, `ordinal()` for `STILL` returns 0, and `ordinal()` for `SPARKLING` returns 1.

To determine an enum's name, we can use the `name()` method.

### ***The name() method***

The `name()` method (line 28) returns the name of this enum, exactly as declared in the enum (line 4). For example, `name()` for `STILL` returns "STILL" and `name()` for `SPARKLING` returns "SPARKLING". Note that rather than use the `name` method, the better option would be to override the `toString()` method as you can customize the `String` that's displayed (to the user) to be more user-friendly. We will do a lot of this in inheritance (*Chapter 9*).

Now that we have examined simple enums, let's move on and discuss complex enums.

## **Complex enums**

As stated earlier, enums are a special type of class where the instances are finite. As simple enums are so straightforward, it can be a little harder to see the class/enum relationship. With complex enums, identifying the relationship between an enum and a class is much easier.

Complex enums have instance variables, constructors, and methods, so they are quite similar to classes. *Figure 8.28* presents a complex enum for discussion:

```
3  enum WorkDay {
4      // values must be first
5      MONDAY( hoursOfWork: "9-5"), // constructor calls
6      TUESDAY( hoursOfWork: "9-5"),
7      WEDNESDAY( hoursOfWork: "9-5"),
8      THURSDAY( hoursOfWork: "9-5"),
9      FRIDAY( hoursOfWork: "9-5"),
10     SATURDAY( hoursOfWork: "10-1"){
11         // constant specific class body
12         public String getWorkLocation(){ return "Home";}
13     }; // ; required at end
14
15     private String hoursOfWork;
16     WorkDay(String hoursOfWork) { // constructor is 'private'
17         this.hoursOfWork = hoursOfWork;
18     }
19     public String getHoursOfWork() {
20         return hoursOfWork;
21     }
22     public String getWorkLocation() {
23         return "Office";
24     }
25 }
26 public class ComplexEnums {
27     public static void main(String[] args) {
28         WorkDay monday = WorkDay.MONDAY;
29         System.out.println(monday.getHoursOfWork() + ", " // 9-5,
30                             +monday.getWorkLocation()); // Office
31         System.out.println(WorkDay.SATURDAY.getHoursOfWork() + ", " // 10-1
32                             +WorkDay.SATURDAY.getWorkLocation()); // Home
33     }
34 }
```

Figure 8.28 – A complex enum

In this figure, we declare the `WorkDay` enum (lines 3-25). This enum encapsulates that we work 9 to 5, Monday to Friday at the office and 10 to 1 on Saturday from home. Presumably, we try to rest on Sunday!

The enum constants are declared from lines 5-13. There is a `private` instance variable called `hoursOfWork` (line 15), which is initialized by the constructor (lines 16-18). Note that the constructor

---

is private by default. The accessor method, `getHoursOfWork` (lines 19-21), is how external classes gain access to the private instance variable, `hoursOfWork`. The other accessor method, `getWorkLocation` (lines 22-24), assumes that we work from the office every day (a pre-pandemic assumption for sure!). The SATURDAY constant (lines 10-13) merits discussion and we will come to that shortly.

Let's examine line 5 closely: this is a *constructor call* to the constructor that's declared (lines 16-18). In other words, the `hoursOfWork` instance variable is set to "9-5" for MONDAY. The other constants – TUESDAY, WEDNESDAY, THURSDAY, and FRIDAY (lines 6-9) – are initialized similarly.

What about SATURDAY? Since we haven't covered inheritance yet, this may be a little tricky. What we are saying is that for Saturday, we only work from home. To do this, we have to replace ("override") the default `getWorkLocation` method (lines 22-24). The default `getWorkLocation` method returns "Office" but our custom `getWorkLocation` (line 12) returns "Home" for SATURDAY. The SATURDAY constant defines a "constant specific class body," which starts with the curly brace on line 10 and ends with the curly brace on line 13.

Note that the semicolon on line 13 is required at the end of the complex enum constants, regardless of whether they declare a constant specific class body or not. That particular semicolon (line 13) tells the compiler, "We have now finished defining the enum constants, so you can expect instance variables or constructors or methods from here on."

Now that we have defined our enum, let's use it. Line 28 instantiates MONDAY, resulting in the enum constant (line 5) executing the constructor (lines 16-18), thereby initializing `hoursOfWork` for the MONDAY instance to "9-5". Line 29 proves this fact by outputting "9-5". Line 30 calls the (default) version of `getWorkLocation` (lines 22-24), thereby outputting "Office" to the screen.

Line 31 instantiates SATURDAY and outputs "10-1" for `hoursOfWork` as that is what is passed into the constructor from line 10. Line 32 invokes the constant-specific version of `getWorkLocation` for SATURDAY, which outputs "Home" to the screen.

That completes our discussion on enumerations. Let us now discuss a very useful feature, namely records.

## Appreciating records

Records are a special type of class, and are considered "data carriers". They help us avoid typing in copious amounts of boilerplate code. Records are specified using a record declaration where you list the *components* of the record. Implicitly generated in the background are a canonical constructor; `toString`, `equals`, and `hashCode` methods and `public` accessor methods for each of the components specified. The accessor methods take on the same names as the components themselves (as opposed to the more traditional `get` methods). Records are best explained by contrasting them to regular classes. *Figure 8.29* presents a normal class with a lot of boilerplate code:

```
5  public final class Person {  
6      private final String name;  
7      private final Integer age;  
8  
9      public Person(String name, Integer age) {  
10         this.name = name;  
11         this.age = age;  
12     }  
13     public String name() {  
14         return name;  
15     }  
16     public Integer age() {  
17         return age;  
18     }  
19     @Override  
20     public boolean equals(Object obj) {  
21         if (obj == this) return true;  
22         if (obj == null || obj.getClass() != this.getClass()) return false;  
23         var that = (Person) obj;  
24         return Objects.equals(this.name, that.name) &&  
25             Objects.equals(this.age, that.age);  
26     }  
27     @Override  
28     public int hashCode() {  
29         return Objects.hash(name, age);  
30     }  
31     @Override  
32     public String toString() {  
33         return "Person[" +  
34             "name=" + name + ", " +  
35             "age=" + age + ']';  
36     }  
37 }
```

Figure 8.29 - A class with a lot of boilerplate code

The Person class in the preceding figure is customized somewhat to map to a record more easily. For example, the class itself is `final` (line 5) and the instance variables, namely `name` and `age` (lines 6-7), are also `final`. The fact that the instance variables are *blank final's* (declared as `final` but not initialized at declaration time) means that the instance variables must be initialized in the constructor. This is what the constructor does (lines 10-11).

There are two accessor methods for retrieving the instance variables, namely name (lines 13-15) and age (lines 16-18). Note that the method names are deliberately not preceded by `get`, in other words, `getName` and `getAge`. This is because, records use the components identifiers for both naming the instance variables *and* the accessor methods.

In addition, this class also has custom versions of `equals`, `hashCode`, and `toString`, lines 20-26, 28-30 and 32-36 respectively. Each of these methods is overriding an inherited version by providing a specific, custom version. This topic of overriding is discussed in detail in Inheritance (*Chapter 9*). The job of `toString` is to return a string containing the instance variables values (the component values). The `equals` method ensures that two records are considered equal if they are of the same type and contain equal component values. The `hashCode` method ensures that equal objects return the same hashcode value (more on this in *Chapter 13*).

Now let us examine the equivalent record in *Figure 8.30*:

```
5 public record Person(String name, Integer age) { }
```

Figure 8.30 - Equivalent record of class from Figure 8.29

Yes – just one line of code! As you can see, this saves us from a lot of boilerplate code. In fact, *Figures 8.29 and 8.30* are equivalent (by the time the compiler is finished). The two parameters are called components and the preceding one liner leads to the following code being generated in the background:

- A final class named after the record (`Person` in this example).
- private final instance variables, one for each component, named after the components.
- A canonical constructor for initializing the components (instance variables).
- Accessor methods, one for each component, named after the components.
- Custom `toString`, `equals` and `hashCode` methods.

Records are customizable. In other words, we can override (replace) all the default versions if we so wish. *Figure 8.31* presents such a situation.

```

5  public record Person(String name, Integer age) {
6      // compact canonical constructor
7      //    public Person(String name, Integer age) {
8      //        if(age < 18){
9      //            this.name = "Error"; this.age = -1;
10     //        }
11     //        this.name = name;
12     //        this.age = age;
13    //    }
14    // compact constructor
15    public Person {
16        if(age < 18){
17            name = "Error"; age = -1;
18        }
19    }
20 }
21 class PersonTest{
22     public static void main(String[] args) {
23         Person p1 = new Person( name: "Joe Bloggs", age: 20);
24         System.out.println(p1);          // Person[name=Joe Bloggs, age=20]
25         System.out.println(p1.name());   // Joe Bloggs
26         System.out.println(p1.age());    // 20
27     }
28 }

```

Figure 8.31 - canonical and compact constructors

In this figure, we are customizing the canonical constructor (lines 7-13) as we want to validate the age component of the person – if they are younger than 18, that is an error and we generate custom error values. Note again that there are better ways to handle error values but for now, this is fine. Otherwise, the components are initialized to the values passed in.

However, this canonical constructor can be written in an even more concise fashion. The compact constructor (lines 15-19) is replacing the canonical constructor. Compact constructors are a variation of the canonical constructor and are specific to records. Note that there is not even a pair of round brackets on line 15 – the components can be inferred from the component list (line 5). Also, there is no need to initialize the components as per lines 11-12; again, the compiler can do this for us.

Lines 23-26 demonstrate how to use the record Person that we have declared. Line 23 declares a Person instance referred to by p1. Line 24 calls the implicit `toString` provided by the Record class (which every record inherits from). Lines 25-26 invokes the two accessor methods; note their names are `name()` and `age()` respectively. The output is in comments on the right of each line (lines 24-26).

As records are so closely related to classes, it is no surprise that records can be used with the `instanceof` keyword. This is what we will examine in record patterns.

## Record patterns

Over the years, the `instanceof` keyword has evolved past the simple `instanceof-and-cast` idiom to support both type patterns and record patterns. Let us first discuss what a “type pattern” is and “pattern matching”.

### Type patterns and pattern matching

In Java 16, `instanceof` was extended to take a type pattern and perform pattern matching. Prior to Java 16 the following code was commonplace:

```
if(obj instanceof String){ // 'obj' is of type Object
    String s = (String)obj;
    System.out.println(s.toUpperCase());
}
```

This code is checking to see if the `Object` reference `obj` is referring to a `String` object and if so, to (safely) cast the reference to a `String` so we can access the `String` methods. Remember, the methods you can access are based on the reference type. However, if the object at the end of the reference is a `String` object then we can safely cast the reference to a `String` and thus access the `String` methods using the new `String` reference. We will discuss this in more detail in Inheritance (*Chapter 9*).

As of Java 16, we can write the previous code segment more concisely and safely:

```
if(obj instanceof String s){ // "String s" - type pattern
//     String s = (String)obj; // no longer needed
    System.out.println(s.toUpperCase());
}
```

There are two changes to note. The first one is the use of a type pattern `String s` as part of the `instanceof`. Pattern matching occurs at runtime whereby `instanceof` checks the type against the provided type pattern and if there is a match, performs the cast for us as well. The second change is that, as `instanceof` performs the cast on our behalf, we no longer need to do the cast ourselves. This leads to a more declarative style (where you state what you want rather than how to get what you want).

This leads on nicely to record patterns which were introduced in Java 21. Prior to record patterns, the following code was required (assuming the `Person` record from *Figure 8.30*):

```
if(obj instanceof Person p){ // type pattern
    String name = p.name(); // accessor
```

```
int age      = p.age(); // accessor
System.out.println(name + "," + age);
}
```

Using record patterns, the previous code can be expressed more concisely:

```
if(obj instanceof Person(String sName, Integer nAge))
    System.out.println(sName + "," + nAge);
}
```

In this code, `Person(String sName, Integer nAge)` is a record pattern. A record pattern consists of a type, a component pattern list (which may be empty) and an optional identifier. A record pattern does two things for us: firstly, it checks to see if the object passes the `instanceof` test and secondly, disaggregates the record instance into its components. So, in our example, assuming `obj` is referring to a `Person` object, then the local variable `sName` will be initialized to the return value of the `name()` accessor method and the local variable `nAge` will be initialized to the return value from the `age()` accessor method. We deliberately used different identifiers for our local variables to highlight the fact that they do not have to match the component identifiers used in *Figure 8.30*. Note however that the order of the types must match; in other words, the record pattern must specify a `String` variable followed by an `Integer` variable, as that is the order of the component list in *Figure 8.30*.

That completes our discussion on records and indeed concludes *Chapter 8*. Now, let's put that knowledge into practice to reinforce the concepts we've learned.

## Exercises

Classes, objects, and enums are great for enhancing our Mesozoic Eden software. In these exercises, you will be creating classes to represent different entities in our park and using enums to define fixed sets of constants:

1. We have many types of dinosaurs in our park, each with unique characteristics. Define a class called `Dinosaur` with properties such as name, age, and species.
2. Our park's heart and soul lie in its employees. Create a class called `Employee` that encapsulates properties such as name, job title, and years of experience.
3. With these classes in place, create some instances of `Dinosaur` and `Employee` and practice manipulating these objects. It's hard for me to provide more details for this exercise, but for example, you could create a new class called `App`. Then, in this class, you could create a few instances of `Dinosaur` and `Employee`. If you want to go wild, you can add a method that takes `Dinosaur` as an argument and then prints the information (such as its name, age, and so on) of this dinosaur. Of course, you could do the same thing for `Employee`.

- 
4. The “park” itself can be thought of as an object with its own properties and behavior. Design a Park class that contains methods for opening and closing the park, adding or removing dinosaurs, and so on. You can also consider giving it an array of employees and an array of dinosaurs.
  5. The food we serve to our dinosaurs varies greatly. Define a class for Food with properties such as name, nutritional value, and cost.
  6. As you know, safety is our main priority. For obvious safety reasons, our dinosaurs are housed in different enclosures. Create an Enclosure class that contains an array of Dinosaur objects.
  7. To add more clarity, let’s define an enumeration for dinosaur types, such as herbivore, carnivore, and omnivore.
  8. A park visit isn’t complete without a ticket. Create a Ticket class with properties such as price, visitor’s name, and visit date.

## Project – Mesozoic Eden park manager

In this project, you’ll be creating a fully interactive console application known as Mesozoic Eden park manager. This application allows the park manager to oversee and manage the various aspects of the dinosaur park. The park manager can use this application to efficiently manage multiple dinosaurs, park employees, and park tickets. Some of the key features of this system should be as follows:

1. The ability to create, edit, or remove dinosaur profiles, park employee profiles, and park tickets.
2. A real-time tracking system that monitors the location and status of the dinosaurs within the park.
3. A fundamental roster system to organize and manage park employee schedules.
4. A robust ticketing system to manage guest admissions and ensure the park maintains optimal capacity.
5. The system should also handle special scenarios such as emergencies or VIP guest visits.

This might sound like a lot. So, here’s a step-by-step guide to achieve this:

1. **Expand the data structures:** Start working from the Dinosaur and Employee classes. Also, add a class called Guest. Each class should include more properties and methods.
2. **Enhance initialization:** Create the necessary data initialization to support multiple dinosaurs, employees, and ticket types. This could involve creating arrays or lists of Dinosaur, Guest, and Employee objects.
3. **Implement interaction:** Implement an interactive console-based interface using the Scanner class. This interface should provide the park manager with a variety of options to manage the park.
4. **Enhance menu creation:** The menu should now include options to manage multiple dinosaurs, employees, and tickets. Each option should correspond to a particular function in the program.



5. **Handle actions:** Each menu item should trigger a function. For example, selecting the **Manage Dinosaurs** option could trigger a function to add, remove, or edit dinosaur profiles.
6. **Exit the program:** Provide an option for the user to exit the program.

Here is a starting code snippet:

```
import java.util.Scanner;

public class Main {
    // Use Scanner for reading input from the user
    Scanner scanner = new Scanner(System.in);

    public static void main(String[] args) {
        Main main = new Main();
        main.start();
    }

    public void start() {
        // This is the main loop of the application. It
        // will keep running until the user decides to exit.
        while (true) {
            displayMenu();
            int choice = scanner.nextInt();
            handleMenuChoice(choice);
        }
    }

    public void displayMenu() {
        System.out.println("Welcome to Mesozoic Eden Park
                           Manager!");
        System.out.println("1. Manage Dinosaurs");
        System.out.println("2. Manage Park Employees");
        System.out.println("3. Manage Tickets");
        System.out.println("4. Check Park Status");
        System.out.println("5. Handle Special Events");
        System.out.println("6. Exit");
        System.out.print("Enter your choice: ");
    }

    public void handleMenuChoice(int choice) {
```

```
switch (choice) {  
    case 1:  
        // manageDinosaurs();  
        break;  
    case 2:  
        // manageEmployees();  
        break;  
    case 3:  
        // manageTickets();  
        break;  
    case 4:  
        // checkParkStatus();  
        break;  
    case 5:  
        // handleSpecialEvents();  
        break;  
    case 6:  
        System.out.println("Exiting...");  
        System.exit(0);  
    }  
}  
}
```

The commented-out method calls are placeholders for methods you need to implement according to your data structures and functionality.

## Summary

In this chapter, we started our discussion by differentiating objects and classes. Classes are similar to a plan of a house, whereas an object is the (built) house itself. We create an object using the new keyword and manipulate the object using its reference. Differentiating the reference from the object is very important going forward. A useful analogy is that the reference is like a remote control and the object is the TV.

Constructors are special methods that are used when constructing an object. The constructor is a method that has the same name as the class but with no return type. There is always a constructor present – if you don't provide one, the compiler intervenes and inserts the default constructor. The constructor is typically used to initialize the instance variables.

Every object gets a copy of the instance members (variables and methods). Class members are marked as static, and are shared by all instances. When accessing an instance member, we use the reference but when accessing a class member, we use the class name. Dot notation applies to both syntaxes.

The `this` reference is a special reference available to us in instance methods. It refers to the object instance responsible for the method call. Consequently, it is dynamic since its value depends on the reference used to invoke the method. It is not available to class (`static`) methods.

Access modifiers apply at both the top (class/interface/record) level, and the member level. At the top level, `public` or package-private access applies. Package-private is achieved by not specifying any keyword at all and ensures that the top-level construct is visible within the same package only. If the top-level construct is `public`, then it is available everywhere; there are no restrictions.

Members (variables/methods) can, in addition to `public` and package-private (with the same semantics), be `private` and `protected`. `private` means that the member is visible within the class only. `protected` is similar to package-private except that subclasses, regardless of package, can access the member.

Encapsulation is one of the cornerstones of OOP. It means that a class can hide its data from external misuse; this is often called “data hiding.” In Java, it is achieved by marking data as `private` and providing `public` accessor/mutator (get/set) methods to manipulate the data. The important concept here is that external code has to access `private` data via your `public` methods. Thus, by using conditional logic in your `public` methods, you can prevent your data from being corrupted.

However, the principle of “private data, public methods” only goes so far. When returning a reference to a `private` object, Java’s call by value mechanism returns a copy of that reference. Thus, the `private` object is now *directly* accessible via external code. Advanced encapsulation combats this by copying the `private` object and returning the reference to the copy object. Thus, your `private` object is still private and safe from external interference.

Understanding an object’s life cycle is extremely beneficial. Local variables live on the stack, whereas objects and instance variables reside on the heap. When an object no longer has any references referring to it, it is eligible for garbage collection. Garbage collection is an automatic process run by the JVM, at a time of the JVM’s choosing. When the garbage collector runs, objects eligible for garbage collection are removed and the heap space is reclaimed.

The `instanceof` keyword enables us to determine the object type that a reference is referring to. This will be very useful going forward.

Enumerations (enums) are closely related to classes in that enums are simply classes, where the number of instances are finite and specified. They are very useful for ensuring type safety, whereby the compiler flags an error as opposed to discovering the error at runtime.

Enums are categorized into two separate types: simple and complex. Simple enums just specify the constant values; the compiler synthesizes the default constructor. All enum constructors are `private` by default. Thus, external classes cannot new them – the constants that are defined are, in fact, the constructor calls. Complex enums look very similar to classes as they have instance variables, (explicit) constructors, and methods.

---

Records are useful when you have classes with a lot of boilerplate code. The components of the record are specified in the record declaration. The compiler, in the background, generates the instance variables, canonical constructor, accessor methods, `toString`, `equals`, and `hashCode` methods. Records are `final`, as are the instance variables (components). A compact constructor is a more concise variation of the canonical constructor.

That completes our discussion of classes, objects, and enums. We will now move onto another important OOP chapter: inheritance.

Trial Version



Wondershare  
**PDFelement**

## 9

# Inheritance and Polymorphism

In *Chapter 8*, we learned about classes, objects, and enums. Initially, we explored the relationship between classes and objects and the need to separate the reference type from the object type. We contrasted instance versus class members and saw that using the `static` keyword applies class scope to a member. We discussed the `this` reference and demonstrated that inside an instance method, the `this` reference refers to the object instance responsible for the method call. We also covered various access modifiers: `private`, package-private (no keyword), `protected`, and `public`. These modifiers enable us to apply one of the cornerstones of OOP, namely encapsulation. While encapsulation is commonly referred to as “private data, public methods,” we demonstrated that this does not go far enough due to Java’s call by value mechanism when passing references into and out of methods. We showed how a technique called “defensive copying” can be used to apply proper (advanced) encapsulation. To improve our understanding of what is happening in the background, we detailed the object life cycle and gently touched on garbage collection. We also covered the `instanceof` keyword, which is used to determine the object type a reference is referring to. We covered a variation of a class, namely **enumerations (enums)**. Enums enable us to limit the number of instances created, thereby facilitating type safety. We covered both simple and complex enums. Lastly, we covered another class variation, namely records, which saves us from typing a lot of boilerplate code.

In this chapter, we will explore inheritance, another core principle of OOP. Initially, we will outline the benefits of inheritance and the Java keywords to use. This leads to polymorphism, another core pillar of OOP. We will explain polymorphism and, with the aid of examples, how polymorphism is achieved. As polymorphism requires “method overriding,” we will explain how to use `instanceof`, to ensure type safety when downcasting.

We will also contrast method overriding with method overloading. We will explain the `super` keyword and how it is used. As promised in *Chapter 8*, we will revisit `protected`, the most misunderstood of Java's access modifiers.

After that, we will discuss both the `abstract` and `final` keywords and their place in inheritance. We will also show how `sealed` classes enable us to scope inheritance. In addition, we will cover both `static` and instance blocks in an inheritance hierarchy. Lastly, we will discuss upcasting and downcasting the inheritance tree, and how a simple rule-of-thumb helps prevent `ClassCastException` errors.

This chapter covers the following main topics:

- Understanding inheritance
- Applying inheritance
- Exploring polymorphism
- Contrasting method overriding and method overloading
- Exploring the `super` keyword
- Revisiting the `protected` access modifier
- Explaining the `abstract` and `final` keywords
- Applying `sealed` classes
- Understanding instance and `static` blocks
- Mastering upcasting and downcasting

## Technical requirements

The code for this chapter can be found on GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects/tree/main/ch9>.

## Understanding inheritance

There are three core pillars in Java: polymorphism, inheritance, and encapsulation (data hiding). It is easy to remember them using the acronym "PIE" (*Polymorphism, Inheritance, and Encapsulation*). Let us now examine inheritance.

---

Inheritance is a code reusability mechanism where common properties between related types are exploited by forming relationships between those types. Inheritance relationships in Java are created by extending from a class or by implementing an interface. We will cover interfaces in *Chapter 10*, so for the moment, we will assume classes throughout. To understand why inheritance in OOP is important, we will examine its advantages (and disadvantages). As we have not covered the terminology used yet, this discussion will be somewhat abstract.

## Advantages of inheritance

One principle advantage of inheritance is code reuse. A new class can be written based on an existing class rather than writing the new class from scratch. In other words, the new class can inherit code that has been already written (and tested). This is called *code reuse* and reduces redundancy.

Inheritance naturally promotes polymorphism, which we discuss later. This feature gives your code flexibility. For example, you could have a method that deals with an `Animal` reference but at runtime, the code executed is in the `Dog` type (or `Cat` or any other type of `Animal` in the hierarchy). In effect, one method works with all `Animal` types.

Inheritance organizes code into a hierarchy. This can improve productivity and simplify the maintenance of code as changes made to inherited code are immediately reflected throughout the hierarchy.

## Disadvantages of inheritance

Despite its advantages, inheritance does have its disadvantages. Tight coupling between the base (source) type and the derived (target) type is one such drawback. Any changes made to the base type affect all the derived types.

Code bloat is another disadvantage. Changes may be made to the base type that many derived types do not need and this can result in an unnecessarily large code base.

Now that we have an appreciation of inheritance and why it is used, let's discuss the nomenclature (terms) used when discussing inheritance.

## Base class

The “base” class is also known as the “super” or “parent” class. This is where the inherited members are defined. As a class is a type, the term *type* is often used interchangeably for class. Note that in Java, the `Object` class is at the top of every hierarchy.

## Subclass

The subclass is also known as the “child” or “derived” class. So, the subclass inherits functionality (and/or data) from the base class. Again, as a class is a type, the term *subtype* is often used interchangeably for subclass. A class can be both a base class and a subclass. Java ensures that *Object* is at the top of every (inheritance) hierarchy. Thus, every class we write is implicitly a subtype already (even if you do not say so).

### The “is-a” relationship

Inheritance generates what is called an “*is-a*” relationship. *Figure 9.1* will help us explain this. We will expand on this diagram as this chapter progresses:

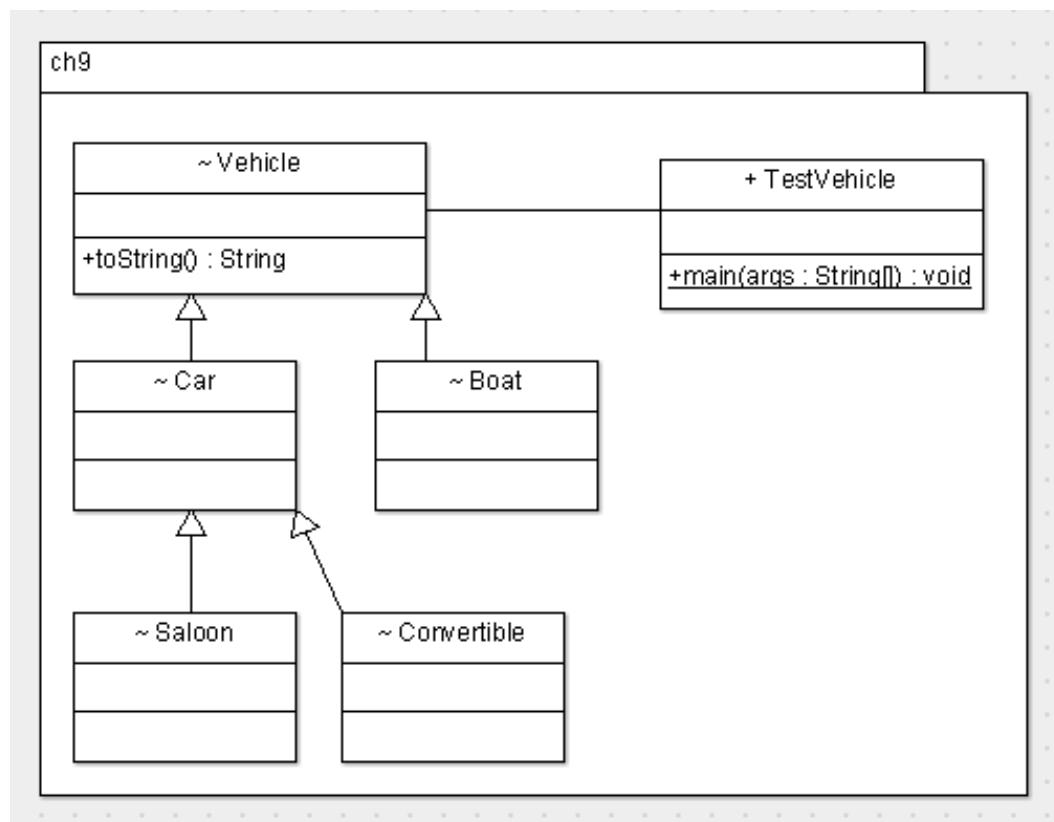


Figure 9.1 – UML diagram for the Vehicle hierarchy

### Unified Modeling Language (UML)

UML is a modeling language used in software design availing of the maxim that “a picture speaks a thousand words.” UML makes understanding topics such as inheritance very straightforward, so we will present a very brief overview of UML here. Further detail is available here: [https://en.wikipedia.org/wiki/Unified\\_Modeling\\_Language](https://en.wikipedia.org/wiki/Unified_Modeling_Language).

With *Figure 9.1* in mind, here is an overview of the symbols used:

- *Package name*: The package name is at the top left (ch9)
- *Classes*: Classes are in boxes with three sections – the top box is the class name; the middle box is for the instance/class variables; the bottom box is for the methods
- *Access modifiers*: public (+), private (-), package-private (~), and protected (#)
- *Static*: The underline is used to signify that a member is static
- *Method return type*: The last part of the method signature in UML
- *Class inheritance*: An arrow with a solid line; for example, Car inherits from Vehicle
- *Interfaces*: These are shown in boxes with dashed lines (*Chapter 10*)
- *Interface inheritance*: An arrow with a dashed line (*Chapter 10*)
- *Association*: A solid line; for example, TestVehicle is associated with Vehicle for the simple reason that we will be creating objects based on the Vehicle hierarchy in main()

As shown in *Figure 9.1*, we have a package, namely ch9. There are five classes in the Vehicle hierarchy: Vehicle, Car, Saloon, Convertible, and Boat. In this hierarchy, from a base class viewpoint, Vehicle is the base class for Car and Boat; and Car is the base class for Saloon and Convertible. Interpreting the diagram from the sub-class perspective, Car and Boat are sub-classes of Vehicle, whereas Saloon and Convertible are sub-classes of Car. Regardless of which perspective you use, every Car “is-a” Vehicle, and every Boat “is-a” Vehicle too. In addition, every Saloon “is-a” Car, and every Convertible “is-a” Car.

It also follows that because Saloon “is-a” Car and Car “is-a” Vehicle, Saloon “is-a” Vehicle as well. The same applies to Convertible; in other words, Convertible “is-a” Car, Car “is-a” Vehicle; therefore, Convertible “is-a” Vehicle also.

However, the “is-a” relationship works in one direction only (reading the diagram from the bottom up). For example, while *every* Car “is-a” Vehicle, *not* every Vehicle “is-a” Car; some are Boats. There is a very good reason for this, which we will explore further when we discuss upcasting and downcasting.

There is one method in Vehicle, namely `toString()`, which, because it is `public`, is inherited by all the subtypes; namely, Car, Saloon, Convertible, and Boat. Thus, the version of `toString()` in Vehicle is available throughout the whole hierarchy. Lastly, the other class, TestVehicle contains the `main()` method so that we can test the hierarchy.

Now that we understand the concept of inheritance, let’s apply it in code.

## Applying inheritance

As we learned in the previous section, inheritance creates an “is-a” relationship hierarchy. This enables base class functionality to be inherited and therefore available to subclasses, without any extra coding. Java uses two keywords in applying inheritance: `extends` and `implements`. Let’s discuss them now.

### **extends**

This is the principle keyword that’s used and relates to both classes and interfaces. Regarding classes, we state that `class Sub extends Base { }`. In this case, all of the non-private members from the Base class will be inherited into the Sub class. Note that private members and constructors are not inherited – this makes sense as both private members and constructors are class-specific. In addition, Java prohibits multiple class inheritance. This means that you cannot extend from more than one class at a time. Regarding interfaces, we state that `interface ChildInt extends ParentInt { }`.

## implements

While we will discuss interfaces in detail in *Chapter 10*, a brief overview here is appropriate. An interface is a construct that enables Java to ensure that if a class implements an interface, the class is, in effect, signing a contract. The contract states, generally speaking, that the class will have code for the `abstract` methods in the interface. An `abstract` method, which we will discuss in more detail later, is a method that has no implementation code; in other words, no curly braces.

Concerning inheritance, unlike classes, Java allows interfaces to extend from more than one interface at a time. So, for example, `interface C extends A, B {}`, where A, B, and C are all interfaces, is fine. Note that, as of Java 8, both the `default` and `static` methods in interfaces have implementation code.

A class implements an interface using the `class Dog implements Walkable` syntax. With this, the `static` and `default` methods in `Walkable` are available to `Dog`.

Now, let's look at inheritance in action. *Figure 9.2* shows the Java code for the UML in *Figure 9.1*:

```
1 package ch9;
2
3 // class Vehicle extends Object
4 @class Vehicle{
5     public String toString(){
6         return "Vehicle::toString()";
7     }
8 }
9 @class Car extends Vehicle{}
10 class Boat extends Vehicle{}
11 class Saloon extends Car {}
12 class Convertible extends Car {}
13
14 public class TestVehicle {
15     public static void main(String[] args) {
16         Vehicle vehicle = new Vehicle();
17         System.out.println(vehicle.toString()); // Vehicle::toString()
18         Car car = new Car();
19         // next line invokes car.toString()
20         System.out.println(car); // Vehicle::toString()
21         Saloon saloon = new Saloon();
22         System.out.println(saloon); // Vehicle::toString()
23
24         System.out.println(new TestVehicle().toString()); // ch9.TestVehicle@378bf509
25     }
26 }
```

Figure 9.2 – Inheritance in action

In this figure, lines 3 and 4 are equivalent. `Vehicle` is at the top of this particular hierarchy and to ensure that `Object` is inherited by every class, the compiler simply inserts `extends Object` after `class Vehicle`, as per line 3. Lines 5-7 are a custom implementation of the `toString()` method inherited from `Object`. This is known as *overriding*, a topic we will discuss in detail shortly. Lines 9-12 represent the rest of the inheritance hierarchy: a `Car` “is-a” `Vehicle`; a `Boat` “is-a” `Vehicle`; a `Saloon` “is-a” `Car`; and a `Convertible` “is-a” `Car`.

On line 16, we create a `Vehicle` object and use a `Vehicle` reference called `vehicle` to refer to it. On line 17, we call the `toString()` method, defined on lines 5-7, outputting `Vehicle::toString()`.

On line 18, we create a `Car` object and use a `Car` reference called `car` to refer to it. On line 20, we simply insert the `car` reference inside `System.out.println()`. When Java encounters a reference like this inside `System.out.println()`, it looks up the object type (`Car`, in this instance) and calls its `toString()`. As every class inherits from `Object`, and `Object` defines a basic (unfriendly) `toString()`, a version of `toString()` will exist. However, in this hierarchy, `Vehicle` has replaced (overridden) `toString()` inherited from `Object` with its own custom one (lines 5-7). This custom one from `Vehicle` is inherited by `Car`. What happens is that Java checks if there is a custom `toString()` defined in `Car`; as there isn’t one, Java then checks its parent, namely `Vehicle`. If `Vehicle` has no `toString()`, the version from `Object` would be used. Since `toString()` is defined in `Vehicle`, this is the version inherited by `Car` and used on line 20. Thus, the output is, again, `Vehicle::toString()`.

On line 21, we create a `Saloon` object and use a `Saloon` reference called `saloon` to refer to it. Again, on line 22, we simply insert the `saloon` reference inside `System.out.println()`. As `Saloon` has no custom `toString()` defined, and its parent, `Car`, has no custom version either, the one inherited from `Vehicle` is used. This results in `Vehicle::toString()` being output to the screen.

Line 24 is used to demonstrate the output when the `toString()` method from `Object` is used. On line 24, we are creating an instance of `TestVehicle` and calling its `toString()` method. As `TestVehicle` is not explicitly inheriting from any class (using `extends`), it implicitly inherits from `Object`. In addition, as `TestVehicle` is not overriding `toString()` with its own custom version, the one inherited from `Object` is used. This is demonstrated by the output from line 24: `ch9.TestVehicle@378bf509`. The output from the `toString()` method in `Object` is formatted as `package_name.class_name@hash_code`. The package name in this case is `ch9` (line 1), the class name is `TestVehicle` (line 24), and the hash code is a hexadecimal number that’s used in hashing collections (*Chapter 13*).

Now that we have seen basic inheritance in action, let’s move on to another cornerstone of OOP, namely polymorphism.

## Exploring polymorphism

Polymorphism has its origins in the Greek terms poly (many) morphe (forms). Any object that passes more than one “is-a” test can be considered polymorphic. Therefore, only objects of the `Object` type are not polymorphic as any type passes the “is-a” test for both `Object` and itself.

In this section, we will discuss why separating the reference type from the object type is so important. In addition, we will examine method overriding and its critical role in enabling polymorphism.

### Separating the reference type from the object type

Now that we have inheritance hierarchies, we will regularly differentiate the reference type from the object type. The reference type can be a class, record, enum or interface. In other words, we have flexibility with regard to the reference type. The object type is more restrictive: the object type is based on non-abstract classes, records, and enums only. In other words, we cannot create objects based on abstract classes or interfaces.

For example, given the hierarchy in *Figure 9.2*, it is perfectly legal to say the following:

```
Vehicle v = new Car();
```

This is because every `Car` “is-a” `Vehicle` (reading it right to left, as assignment associates right to left). In this instance, the reference, `v`, is of the `Vehicle` type and it is referring to an object of the `Car` type. This is known as *upcasting*, as we are going *up* the inheritance tree (again, reading it from right to left, from `Car` *up* to `Vehicle`). We are upcasting the `Car` reference, created by `new Car()`, and casting it to a `Vehicle` reference, `v`.

Why does this work? This works because, due to inheritance, every inheritable method available to `Vehicle` will exist in `Car`. That is a guarantee. Whether `Car` has overridden (replaced) any/all `Vehicle` methods with its own custom ones is immaterial. Given that the compiler looks at the reference type (and not the object type), the methods we can call using the `Vehicle` reference, `v`, are defined in `Vehicle` (and `Object`) and will be present in `Car` (the object type).

So, that is the first point to keep in mind - the compiler is always looking at the reference type. As we will see shortly, the object type comes into play at runtime. So, a simple but effective rule of thumb is that *a reference can refer to objects of its own type or objects of subclasses*. In effect, a reference can point “across and down” the (UML) hierarchy.

If a reference is ever pointing “up” the hierarchy, that is when you get `ClassCastException` errors. Why is this? Well, a subclass inherits from its parent. In addition to replacing inherited functionality (overriding), the subclass can also add extra methods. So, if you have a reference of the subclass type, you can invoke these *extra* added methods. But if your object is of the parent type, these methods will not exist! This is a serious issue for the JVM and it throws an exception (*Chapter 11*) immediately.

So, the reference type determines the methods that can be called on the object. In addition, while the reference type cannot change, the object type it refers to can.

Now, let’s address how to avail of polymorphism.

## Applying polymorphism

Polymorphism applies only to instance (non-static) methods, as only instance methods can be overridden. At compile time, the compiler decides which method signature to bind to; however, the object that will provide the actual method to execute is decided at runtime! That is what polymorphism is. This is why polymorphism is also known as “runtime binding” or “late binding.”

### What if you are accessing a static member?

A `static` member (method or data) is associated with the class and therefore not involved in polymorphism. The following applies: if you are accessing any type of data (`static` or non-`static`) or `static` methods, the JVM uses the reference type. Only if it’s an instance method is the object type used (polymorphism).

Thus, for polymorphism to work, we need instance methods in the base and subclass where the subclass overrides the base version. For this to happen, the subclass must code a method that has the same signature as the parent.

Okay, that’s enough theory – let’s look at an example that reinforces everything we’ve learned thus far.

### ***Polymorphism in code – example 1***

*Figure 9.3 shows the UML for the code to follow:*

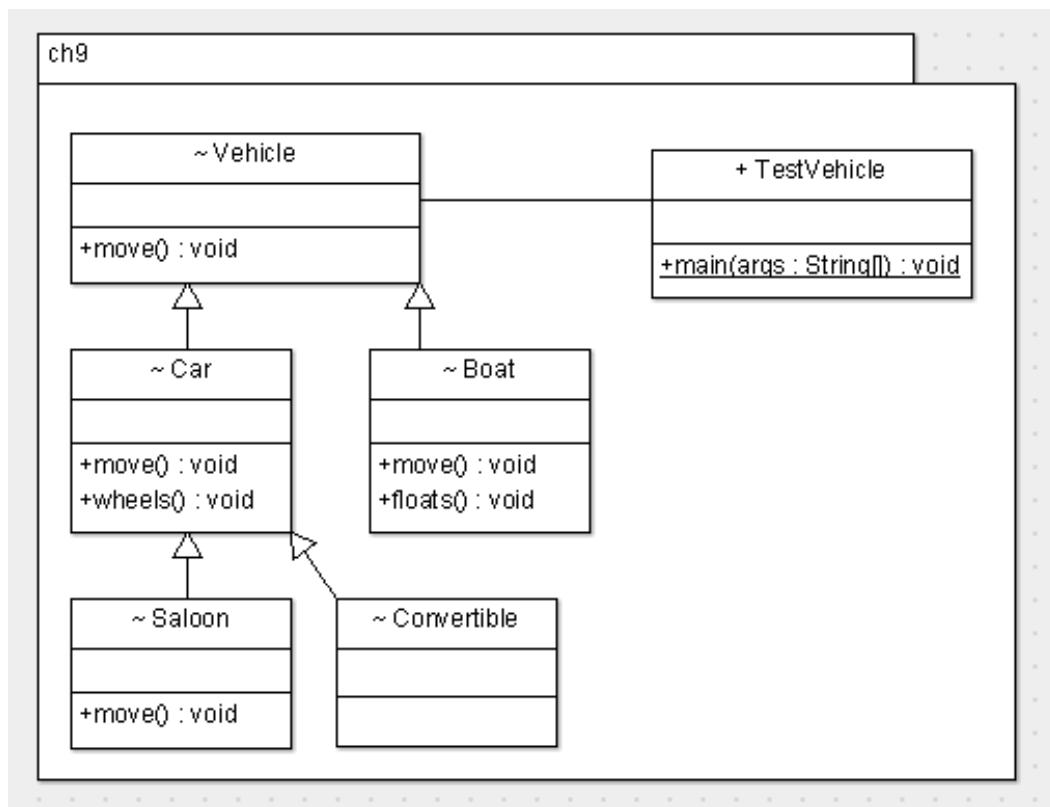


Figure 9.3 – UML for polymorphism example

In this figure, the `Vehicle` class has a `move()` method. It is an instance method, with a return type of `void`. Both `Car` and `Boat` extend `Vehicle` and override `move()`. `Car` adds a method called `wheels()` and `Boat` adds a method called `floats()`. Both `Saloon` and `Convertible` extend `Car`. `Saloon` overrides `move()` but `Convertible` does not.

*Figure 9.4* presents the code for this UML and demonstrates polymorphism in action:

#### @Override annotation

An annotation is a form of metadata that provides information about the program that is not part of the program itself. Annotations are preceded in Java with the `@` symbol and have several uses. For example, annotations are used by the compiler to detect errors or by the runtime to generate code.

When overriding a base class method, we can insert the `@Override` annotation just prior to the subclass method. While not mandatory, it is very useful, because, if we apply this annotation, the compiler will ensure that we override the method correctly.

```
3  class Vehicle{
4      public void move(){ System.out.println("Vehicle::move"); }
5  }
6  class Car extends Vehicle{
7      @Override public void move(){ System.out.println("Car::move()"); }
8      public void wheels(){ System.out.println("Car::wheels()"); }
9  }
10 class Boat extends Vehicle{
11     @Override public void move(){ System.out.println("Boat::move()"); }
12     public void floats(){ System.out.println("Boat::floats()"); }
13 }
14 class Saloon extends Car {
15     @Override public void move(){ System.out.println("Saloon::move()"); }
16 }
17 class Convertible extends Car {}
18
19 public class TestVehicle {
20     public static void main(String[] args) {
21         Vehicle v = new Car();
22         v.move();    // Car::move()
23         v = new Boat();
24         v.move();    // Boat::move()
25         v.floats(); // compiler error
26         v = new Saloon();
27         v.move();    // Saloon::move()
28         v = new Convertible();
29         v.move();    // Car::move()
30
31         Saloon s = (Saloon)new Vehicle(); // ClassCastException
32     }
33 }
```

Figure 9.4 – Polymorphism example

In this figure, Car and Boat both extend Vehicle; and Saloon and Convertible both extend Car. Note that the move () method in Vehicle (line 4) is a non-static/instance method and therefore polymorphic. In addition, as move () is non-private, it is inheritable. The move () method from Vehicle is overridden by Car (line 7), Boat (line 11), and Saloon (line 15). To highlight this fact, on each of those lines, we have used the @Override annotation. This means that the parent move () method is overridden by the respective subclass versions.

Line 21 creates a `Car` object and uses a `Vehicle` reference, namely `v`, to refer to it. It is worth repeating that this upcasting, from `Car` up to `Vehicle`, is only possible because, via inheritance, *every* `Car` “is-a” `Vehicle`. Therefore, any method available to the `Vehicle` reference will exist in `Car`. Consequently, as upcasting is never a risk, it is performed implicitly by the compiler; in other words, you do not need to explicitly state the (up)cast in code, as follows:

```
Vehicle v = (Vehicle) new Car();
```

## Compile time

Line 22 makes the polymorphic `v.move()` call. Every time a method call is in code, there are two perspectives to keep in mind: compile time and runtime. As we know, the compiler concerns itself with the reference type. So, in this case, the compiler checks the reference, `v`, and determines that it is of type `Vehicle`. The compiler then checks if there is a `move()` method, with that exact signature in the `Vehicle` class, either defined in `Vehicle` or inherited into `Vehicle` (from `Object` in this example). As there is a `move()` method defined in `Vehicle`, the compiler is happy.

## Polymorphism in action

At runtime, as `move()` is a non-static, polymorphic method, the object being referred to by the reference, `v`, applies. As `v` is referring to a `Car` object, the `Car` version of `move()` is executed. This is polymorphism in action! We have one method but many implementations of that method. The compiler ensures that the method exists and dynamically, at runtime, polymorphism kicks in and executes the version in the object being referred to. As `v` is referring to a `Car` object, the output from line 22 is `Car::move()`.

Line 23 reuses the `Vehicle` reference, `v` (which is perfectly valid), to refer to a `Boat` object. As `Boat` “is-a” `Vehicle`, this is fine. Line 24 makes the same polymorphic call to `v.move()` as was the case on line 22. However, this time, `v` is referring to a `Boat` object, and as `Boat` has overridden `move()`, the `Boat` version of `move()` is executed at runtime. Therefore the output is `Boat::move()`.

Line 25 demonstrates that the compiler looks at the reference type. As we know, `v` is of type `Vehicle`. However, `Vehicle` has no `floats()` method; this is a method specific to `Boat`. Therefore, the compiler complains about `v.floats()` on line 25, and hence, the line is commented out.

Line 26 reuses the `Vehicle` reference, `v`, to refer to a `Saloon` object. As `Saloon` “is-a” `Vehicle`, this is fine. Line 27 makes the same polymorphic call to `v.move()` as was the case on lines 22 and 24. As `v` is now referring to a `Saloon` object with an overridden `move()` method, the `Saloon` version of `move()` is executed polymorphically at runtime. Therefore, the output is `Saloon::move()`.

Line 28 creates a `Convertible` object and uses `v` to refer to it. This is not a problem as `Convertible` “is-a” `Vehicle` (indirectly, via `Car`). In other words, because `Convertible` is-a `Vehicle` and `Vehicle` is-a `Car`, `Convertible` is-a `Car` also. Line 29 makes the same polymorphic call, `v.move()`, as was the case on lines 22, 24, and 27. Note, however, that `Convertible` has not overridden `move()`. `Convertible` has an empty class body. Therefore, the methods in `Convertible` are

the `move()` and `wheels()` methods inherited from `Car` and the methods inherited from `Object`, such as `toString()`. So, at runtime, when `v.move()` is called, the JVM executes the version of `move()` in `Car`, resulting in `Car::move()`. You can also look at it this way: the runtime looks for `move()` in `Convertible`, and finds none; the JVM then checks the parent, `Car`, and finds one, which it executes. Note that if `Car` had not provided a `move()` method, its parent, `Vehicle`, would have been next in the search. So, there is an “up the hierarchy, one generation at a time” orderly search.

### Why do we get a `ClassCastException` error?

Line 31 demonstrates downcasting and a `ClassCastException` error. Exceptions will be discussed in *Chapter 11*, so we won’t go into detail here. Downcasting will be discussed in greater detail later in this chapter but this example is too good to pass up! Let’s examine line 31 in greater detail:

```
Saloon s = (Saloon) new Vehicle(); // ClassCastException
```

The first thing to note is that the cast (`Saloon`) is required. The compiler will not allow the following:

```
Saloon s = new Vehicle(); // Compiler error
```

This is a compiler error because every `Vehicle` is not a `Saloon` class; some are `Boats`. Indeed, even if the `Boat` class were not present, this line would still not compile. Why? Because, reading it right to left, you are going *down* the hierarchy from `Vehicle` to `Saloon`. As `Saloon` could (and indeed does) have extra methods not in the `Vehicle` class, this situation must be prevented. For example, the `Saloon` reference, `s`, has access to the `wheels()` method (inherited from `Car`), which is not present in `Vehicle`.

Now, we can override the compiler by using a (down)cast. This is what line 31 has done with the `(Saloon)` cast. In effect, by inserting the cast and overriding the compiler error, you are saying to the compiler: “Let me proceed, I know what I am doing.” So, the code compiles with the cast in place. However, at runtime, the JVM realizes that it has a `Saloon` reference referring *up* the inheritance tree to a `Vehicle` object. This is a big no-no because if the JVM allowed the `Saloon` reference `s` to refer to a `Vehicle` object, what would it do with a subsequent `s.wheels()` method call? Remember, we would be looking at a `Vehicle` object, which has no such method! Hence the JVM generates a `ClassCastException` error.

Let’s refactor this code to demonstrate polymorphism from another angle.

## Polymorphism in code – example 2

Figure 9.5 shows the refactored code from Figure 9.4:

```
3  class Vehicle{
4      public void move(){ System.out.println("Vehicle::move"); }
5  }
6  class Car extends Vehicle{
7      @Override public void move(){ System.out.println("Car::move()"); }
8      public void wheels(){ System.out.println("Car::wheels()"); }
9  }
10 class Boat extends Vehicle{
11     @Override public void move(){ System.out.println("Boat::move()"); }
12     public void floats(){ System.out.println("Boat::floats()"); }
13 }
14 class Saloon extends Car {
15     @Override public void move(){ System.out.println("Saloon::move()"); }
16 }
17 class Convertible extends Car {}
18
19 public class TestVehicle {
20     public static void doAction(Vehicle v){
21         v.move();
22     }
23     public static void main(String[] args) {
24         Vehicle v = new Car(); // Car::move()
25         doAction(v);
26         doAction(new Boat()); // Boat::move()
27         doAction(new Saloon()); // Saloon::move()
28         doAction(new Convertible()); // Car::move()
29     }
30 }
```

Figure 9.5 – Refactored polymorphism example

Note that, in this figure, the inheritance hierarchy remains untouched from *Figure 9.4*. The `TestVehicle` class (lines 19-30) has been refactored though. We have introduced a new method, namely `doAction()` (lines 20-22), that accepts a `Vehicle` reference. In the `doAction()` method, we simply call the `move()` method (line 21). As `Vehicle` has a `move()` method, this is fine.

Line 24 is as before; it creates a `Car` object and upcasts the reference to a `Vehicle` reference, `v`. Thus, `v` is referring to a `Car` object. Line 25 invokes the `doAction()` method, passing in the reference, `v`. This reference, `v`, which is declared on line 24, is copied into the separate (different scope) but similarly named reference, `v`, which is declared on line 20. Now, in `doAction()`, we have a local `v` reference referring to the same `Car` object created on line 24. Thus, when we invoke `v.move()` on line 21, polymorphism kicks in and we get the `Car` version of `move()`, resulting in `Car::move()`.

Line 26 does the same thing in one line of code as was done in the previous two lines of code (lines 24-25). On line 26, the `Boat` object is created, and the method call to `doAction()` results in the upcast to the `Vehicle` reference, `v` (line 20). After that, line 21 executes polymorphically and we get the `Boat` version of `move()`, resulting in `Boat::move()`.

Line 27 is the same as line 26 except we are creating a `Saloon` object. Thus, the `Vehicle` reference, `v`, in `doAction()` executes the `move()` method as `Saloon`, resulting in `Saloon::move()`.

Line 28 is the same as line 27 except we are creating a `Convertible` object. Thus, the `Vehicle` reference, `v`, in `doAction()` attempts to execute the `move()` method in `Convertible`. As there is none, the parent of `Convertible`, namely `Car`, is checked. `Car` does have a version of `move()`, resulting in `Car::move()`.

To be clear about when polymorphism applies and when it does not, we will revisit a callout box presented earlier.

## JVM – object type versus reference type usage

As discussed briefly in a previous callout, if you are dealing with any type of data (`static` or non-`static`), the reference type applies; when dealing with instance methods, the object type applies (polymorphism). *Figure 9.6* presents a code example:

```
3  ↪ class Vehicle{
4      double cost = 100.0;           // instance data
5      static int age = 1;            // class data
6  ↪     public void move(){        // instance method
7          System.out.println("Vehicle::move()");
8      }
9      public static void sm() {     // class method
10         System.out.println("Vehicle::sm()");
11     }
12 }
13 class Car extends Vehicle{
14     double cost = 20_000.0; // hiding
15     static int age = 2;     // hiding
16
17 ↪     @Override public void move(){ // overriding
18         System.out.println("Car::move()");
19     }
20     public static void sm() {       // hiding
21         System.out.println("Car::sm()");
22     }
23 }
24 ➤ public class TestVehicle {
25     public static void main(String[] args) {
26         Vehicle v = new Car();
27         System.out.println(v.cost); // 100.0
28         System.out.println(v.age); // 1
29         v.sm();                  // Vehicle::sm()
30         v.move();                // Car::move()
31     }
32 }
```

Figure 9.6 – When the JVM uses the reference type versus the object type

In this figure, the `Vehicle` class declares an instance variable, namely `cost` (line 4), and a class variable, namely `age` (line 5). In addition, `Vehicle` also declares an instance method called `move()` (lines 6-8) and a class method called `sm()` (lines 9-11). So, in `Vehicle`, we have both instance and static data and instance and static methods.

The `Car` class extends from `Vehicle` (lines 13-23) and simply replicates `Vehicle`. In other words, `Car` has the same data and methods as its parent, `Vehicle`.

In `Car`, we declare both instance and non-instance variables, namely `cost` and `age`, respectively (lines 14-15). These variables in `Car` have the same types and identifiers as their counterparts in the parent class, `Vehicle`. In other words, `Vehicle` has an instance variable called `cost`, which is a `double`; and `Car` also has an instance variable named `cost`, which is also a `double`. The same occurs with the `age` class variable in `Vehicle` – there is a class variable named `age` in the `Car` subclass also. This is known as *hiding* (or *shadowing*).

`Vehicle` defines the instance method, `move()` (lines 6-8), which is overridden by the version in `Car` (lines 17-19). As this is an instance method, polymorphism applies at runtime if `move()` is called.

`Vehicle` also defines a class method called `sm()` (lines 9-11), which is hidden (shadowed) by the version of `sm()` in `Car` (lines 20-22).

Line 26 creates a `Car` object and uses a `Vehicle` reference, `v`, to refer to it.

Line 27 outputs `v.cost`. As `cost` is data (an instance variable), the reference type applies. Consequently, we get `100.0`, which is the `cost` instance variable in `Vehicle` (as opposed to `20_000.0`, which is the `cost` instance variable in `Car`).

#### Using the class name when accessing a static member

Both lines 28 and 29 present syntax that you should *never* use: using a reference to access a `static` member. When accessing a `static` member, you should prefix the member with the class name. For example, line 28 should use `Vehicle.age` and line 29 should use `Vehicle.sm()` as this emphasizes the member's `static` nature. Using references here is confusing as it implies that the member is non-static. We accessed `static` members using the reference for demo purposes only!

Line 28 outputs `v.age`. As `age` is a `static` member, the compiler checks the type for `v` (namely `Vehicle`) and changes `v.age` to `Vehicle.age`. Therefore, `age` from `Vehicle` is used as opposed to `age` from `Car`. In other words, the output is 1, not 2.

Line 29 is the call to `v.sm()`. As `sm()` is also `static`, the compiler translates this into `Vehicle.sm()` and therefore the output is `Vehicle::sm()`.

Lastly, line 30 is the polymorphic call to `move()`, and as a result, the object type, `Car`, is used. This results in `Car::move()` being output.

Now that we understand polymorphism, let's ensure that we understand the difference between two terms that are often confused, namely method overriding and method overloading.

## Contrasting method overriding and method overloading

These two terms are often confused but in this section, we will compare and contrast both. We will show that concerning method overloading, the method signature must be different; whereas concerning method overriding, the method signature must be the same. Recall that the method signature consists of the method name and the parameter types, including their order. The return type and the parameter identifiers are *not* part of the method signature. So, for example, take the method from *Figure 9.5*:

```
public static void doAction(Vehicle v) {...}
```

The signature is `doAction(Vehicle)`.

With this in mind, we will initially discuss method overloading.

### Method overloading

Recall that the method signature consists of the method name and the parameter types. Method overloading is where you have the same method name but the parameters differ, either in type and/or order. This means that the method signatures are different even though the method names are the same. They have to be – how else will the compiler choose which method to bind to? Thus, method overloading is all about compile time.

#### *The rules*

Bearing in mind that the method signatures *must be different* (apart from the method name), the rules are quite straightforward:

- Overloaded methods must use *DIFFERENT* parameter lists; either the types used must be different or the order of the types must be different
- As the method signature only relates to the method name and the parameter list, overloaded methods are free to change the return type and the access modifier and use new or broader checked exceptions
- An overloaded method can be overloaded in the same type or a subtype

Now, let's look at an example of method overloading in code.

### Method overloading example

Figure 9.7 presents the example code:



```
2  class Animal{
3      public void eat(){}
4  }
5  class Cow extends Animal{
6      public void eat(){          // overriding, same signature
7          public void eat(String s){} // overloaded, different signature!
8  }
9  public class OverloadTest {
10     public void calc(int x, double y){} // calc(int, double)
11     public void calc(){}               // calc()
12     public void calc(int x){}         // calc(int)
13     public void calc(double y){}       // calc(double)
14     public void calc(double y, int x){} // calc(double, int)
15
16     // public void calc(int a, double b){}           // calc(int, double)
17     // public int calc(int a, double b){ return 1; } // calc(int, double)
18
19  public static void main(String[] args) {
20      Animal aa = new Animal();
21      aa.eat();
22      // aa.eat("Grass"); // compiler error
23
24      Animal ac = new Cow();
25      ac.eat();
26      // ac.eat("Grass"); // compiler error
27
28      Cow cc = new Cow();
29      cc.eat();           // inherited
30      cc.eat( s: "Grass");
```

Figure 9.7 – Method overloading

In this figure, we will first discuss the method overloading between lines 10-17. To help, the method signatures are in comments on each line. Line 10 defines a `calc` method that takes in an `int` and a `double`, in that order. Therefore, the signature is as follows:

```
calc(int, double)
```

We are not interested in the return type or the identifiers used for the `int` and `double` parameters. So long as we do not code another `calc(int, double)` method in the *same class*, we are okay. Note that if we coded a method with the same signature in a subtype, this is overriding! As the method signatures between lines 11-14 are different, they are fine.

Let's examine why lines 16 and 17 fail to compile. Line 16 attempts to just change the identifiers used for the parameters. This does not change the method signature. Consequently, this signature is an exact match for the method on line 10 and therefore, the compiler complains. Similarly, line 17 changes the return type (as well as the identifiers in the parameter list). Again, as this signature is a duplicate of the one on line 10, the compiler complains.

The inheritance hierarchy is interesting. We have a parent called `Animal` (lines 2-4) and a subclass class `Cow` (lines 5-8). On line 3, `Animal` defines an `eat()` method. On line 7, `Cow` overloads this method with an `eat(String)` method. The parent `Animal` version accepts no argument, whereas the subtype version accepts `String`. The compiler is happy.

But what about line 6, where `Cow` defines an `eat()` method that accepts no argument? This is overriding the parent version (polymorphism), so there is no conflict. The compiler will bind to the reference type used, be it `Animal` or `Cow`, as both have an `eat()` method. At runtime, depending on the object type, the JVM will execute the relevant code.

Let's examine this process to make sure it is clear. Line 20 creates an `Animal` object and uses an `Animal` reference, `aa`, to refer to it. Line 21 calls `aa.eat()`. At compile time, the compiler checks if there is an `eat()` method with that exact signature in `Animal`, as `Animal` is the type for `aa`. As there is, the compiler is happy. At runtime, as the method is an instance method, polymorphism applies and the JVM will execute the `Animal` version (as that is the object type).

Note how line 22 does not compile. This is because there is no `eat(String)` method in `Animal`. Remember, the compiler looks at the reference type only and as `aa` is of type `Animal`, it checks the `Animal` class.

Lines 24-26 take things one step further. Line 24 creates a `Cow` object and uses an `Animal` reference called `ac` to refer to it. Line 25 makes the polymorphic call to `eat()`, which will execute the `Cow` version at runtime. Line 26 is interesting and is there to prove that the compiler is looking at the reference type. Even though our object type is `Cow` and `Cow` has an `eat(String)` method, the `ac.eat("Grass")` class still does not compile (because `ac` is of type `Animal`).

So, how do we get access to the `eat(String)` method? We need a `Cow` reference. This is what lines 28-30 demonstrate. Line 30 successfully invokes `cc.eat("Grass")` using the `cc` reference declared on line 28.

What this code demonstrates is that an `Animal` reference only has access to the `eat()` method it defined. On the other hand, a `Cow` reference has access to both `eat()` and `eat(String)`. The `Cow` type inherited (and overrode) `eat()` and defined `eat(String)` itself. Note that the `Cow` class did not need to override `eat()` to have access to the inherited version.

## Method overriding

Method overriding occurs when you have the same method signatures in both a parent and subclass. Method overriding is critical for enabling (runtime) polymorphism. Remember, a method must first be inherited to be overridden. For example, methods that are defined as `private`, `static`, or `final` are not inherited because `private` methods are local to the class; `static` methods are not polymorphic and marking a method as `final` is stating that “this method is not to be overridden.”

To understand the rules, it is critical to remember that the compiler has compiled the code based on the reference. Therefore, the runtime polymorphic method *must not* behave differently from what the compiler verified. For example, the access modifier on the overriding method cannot be more restrictive.

Before we discuss the rules, we must first explain covariant returns.

### Covariant returns

When you are overriding a parent method in a subclass, if the return type is a primitive, then the overriding method’s return type must match. However, if the return type is a non-primitive, then there is one exception to the rule: covariant returns.

What a covariant return means is that if you return a type, X, in the parent method, then you can return X and any subtype of X in the overriding method. For example, if a parent method is returning `Animal`, then the overriding method can return `Animal`, (naturally) as well as any subtype of `Animal`; for example, `Cow`.

### The rules

As we discuss the rules, it is helpful to bear in mind that the compiler checks against the reference type. These overriding rules ensure that the runtime object cannot do something that the compiler (and thus your code) does not expect. The rules are as follows:

- The method signatures must match exactly in the parent and subclass; otherwise, you are just overloading the method.
- The return types must match also, except for covariant returns.
- The access modifier on the overriding method cannot be more restrictive. So, if the parent method defines a method as `public`, the subclass cannot override it with a `private` method. This makes sense, as your code, verified by the compiler, is expecting access to the method. *If*, however, you were allowed to reduce access when overriding, the compiler would have said “It is okay to access this method,” whereas the JVM would not! This rule helps keep the compiler and JVM in sync.

- Again, to keep the compiler and JVM in sync, an overriding method cannot throw (generate) new or broader checked exceptions (*Chapter 11*). Briefly, an exception is an error and checked exceptions must have code present to handle them. This is enforced by the compiler. If, at runtime, the overriding method threw/generated an exception for which there was no code to handle it, the JVM would be in trouble. So, the compiler steps in and prevents that from happening.

Now, let's look at an example of method overriding in code.

### Method overriding example

Figure 9.8 presents the example code:

```
3  import java.io.IOException;
4
5  class Dog{
6      public void walk(){System.out.println("Dog::walk()");}
7      public Dog run() { return new Dog(); }
8  }
9  class Terrier extends Dog{
10     //    public String walk(){ return "Walk the Dog"; } // return type should be void
11     //    private void walk(); // access rights cannot be weaker
12     //    public void walk() throws IOException {} // cannot throw new checked exceptions
13     public void walk(int metres){} // an overload, not an override
14     @Override public void walk(){System.out.println("Terrier::walk()");}
15
16     //    @Override public Dog run() {return new Dog();}           // ok
17     //    @Override public Terrier run() {return new Terrier();}   // ok
18     @Override public Dog run() {return new Terrier();}           // ok
19
20
21  public class OverridingTest {
22      public static void main(String[] args) {
23          Dog dt = new Terrier();
24          dt.walk(); // Terrier::walk()
25          Dog d = dt.run();
26          if(d instanceof Terrier){
27              System.out.println("Terrier object!"); // Terrier object!
28          }
29      }
30  }
```

Figure 9.8 – Method overriding



The code in this figure demonstrates what you can and cannot do when overriding a method. In the `Dog` class (lines 5-8), we have a `walk()` method that returns nothing (`void`). There is also a `run()` method that returns `Dog`.

The `Terrier` class subclasses from `Dog` (line 9). Therefore, any `Terrier` “is-a” `Dog`. As the two methods in `Dog` are public, `Terrier` automatically inherits them.

Let’s examine the lines in `Terrier` in turn.

Line 10 does not compile because, while the method signatures match (both are `walk()`), the return types are different. The parent return type is `void` and thus, the overriding return type must match; it does not, it is `String`, causing the compiler error.

Line 11 does not compile because you cannot weaken the access modifier when overriding. The `walk()` method in `Dog` is `public`, so `walk()` in `Terrier` cannot be `private`. If this was allowed, then when the JVM went to execute the `walk()` method in `Terrier`, using a `Dog` reference (as on line 24), there would be a serious problem. The compiler, looking at the `public Dog` version, said “All is well;” but the JVM would, polymorphically, encounter the `private` version in `Terrier`!

Line 12 fails to compile because the overridden method did not throw any exceptions but the overriding method is attempting to throw a new checked exception (`IOException`). This is similar to the previous access issue – the compiler will have checked the `walk()` version in `Dog` and as it throws no exceptions (errors), no code is present to handle (cater) for these exceptions. If the overriding method was allowed to throw new checked exceptions, what would the JVM do with them (as there is no code in place to handle them)?

Line 13 is simply an overload. `Dog` defines a `walk()` method; `Terrier` defines a `walk(int)` method. Two separate method signatures means two separate methods. As the methods have the same name, this is method overloading.

Line 14 is a correct method override. We used the `@Override` annotation to ensure that we have overridden properly (no typos, for example).

Line 16 is an exact duplication of the `run()` method defined on line 7. We just included it for demonstration purposes.

Line 17 demonstrates covariant returns because it defines a `Terrier` return type. This is a valid covariant return because `Terrier` is a subtype of the parent return type, `Dog` (line 7). The code for the overridden method (line 17) simply returns a `Terrier` object.

Line 18 is almost identical to line 17 except that the return type is now `Dog`. Thus, there is an upcast going on in the background. The code for `walk()` on line 18 is shorthand for the following:

```
Dog d = new Terrier():
    return d;
```

Now, let’s look at the `main()` method in `OverridingTest`.

Line 23 creates a `Terrier` object that can be accessed via a `Dog` reference, `dt`. Line 24 invokes the polymorphic `walk()` method in `Terrier`. As `Terrier` overrode the `walk()` method it inherited from `Dog`, the `Terrier` version is dynamically executed at runtime, resulting in `Terrier::walk()` being output.

Line 25 executes the `run()` method using the `dt` reference created on line 23. As `run()` is an instance method where `Terrier` overrode the version inherited from `Dog`, the version in `Terrier` is executed, resulting in the `d` reference (line 25) referring to a `Terrier` object (line 18). This is proven by the use of the `instanceof` operator (line 26). As the `Dog` reference, `d`, is indeed referring to a `Terrier` object, the `if` statement is `true`, resulting in `Terrier object` being output to the screen.

That concludes our discussion of method overloading and method overriding. Now, let's examine a keyword that is pivotal in inheritance: `super`.

## Exploring the super keyword

The `super` keyword is used in a subclass in two specific scenarios: to call a parent constructor and to access parent members (typically methods). When an object is constructed, the order of constructor calls is very important. Bearing in mind that we now have the possibility of having many classes in an inheritance hierarchy, *the order of constructor calls is from the top down*. This means that, the parent constructor is *always* called before the subclass constructor. If you have a hierarchy where `Toyota` “is-a” `Car` and `Car` “is-a” `Vehicle`, then when you go to create a `Toyota` object, the order of constructor calls is as follows: `Vehicle` is first, `Car` is second, and `Toyota` is last.

There is a good reason for this. Firstly, remember that the constructor’s role is to initialize the instance members of the class. Now, given that the subclass constructor *may use inherited members* from its parent when initializing its own members, it stands to reason that the parent must first get a chance to initialize those members.

Let’s discuss the situations where the `super` keyword is very often used. We will then present code, supported by a UML diagram, where both contexts are demonstrated.

### `super()`

When you use the parentheses after `super`, as in `super()`, you are invoking the parent constructor. If required, you can pass in arguments inside the parentheses as constructors are just (special) methods. There are two rules for the use of `super()`:

- The call to `super()` can only appear inside a constructor and not a regular method
- If present, the call to `super()` must be the very first line in the constructor (there is one exception – see the callout)



We have coded several constructors so far and none of them had a call to `super()` present. How did that work? Well, if you *do not* provide any constructor at all, the default constructor will be synthesized by the compiler for you and its first line of code is `super();`. Please refer back to *Figure 8.1* and *Figure 8.2* for examples of this. If you *do* provide a constructor, then the compiler will also insert `super();` as the first line (unless the first line is already a call to `super()` or `this()`).

### The first line of any constructor

The very first line of any constructor is `this()` or `super()`. You cannot have both. A call to `this()` is a call to another constructor in the same class. From the inheritance hierarchy perspective, this is a sideways call. Remember that the parent constructor must be called before the subclass constructor. Regardless of whether `this()` is present or not, the order of constructor calls is from the top down. Now, if the subclass constructor has a `this()` call present, it is only delaying the call to `super()`. At some point, either explicitly or implicitly, the call to `super()` will execute. Note that, as with `super()`, the call to `this()` can contain arguments.

So, `super()` relates only to constructors and must be the first line of code (assuming `this()` is not there already). Now, let's examine the other scenario.

### **super.**

To access a parent member (not the constructor), you can use the `super.` dot notation syntax. As with the `this` keyword, the `super` keyword relates to instances and thus cannot be used from within a `static` context (`static` methods or `static` blocks). This can be very useful when you want to piggyback on parent functionality. For example, the subclass method can invoke its parent version first and then execute its own version. This is what we will demonstrate in the example.

So, rather than call a parent constructor from a subclass constructor (which is what `super()` is for), `super.` gives us access to the other (non-constructor) members.

### An example of using `super`

Let's examine both `super()` and `super.` in code. *Figure 9.9* presents the UML inheritance diagram:

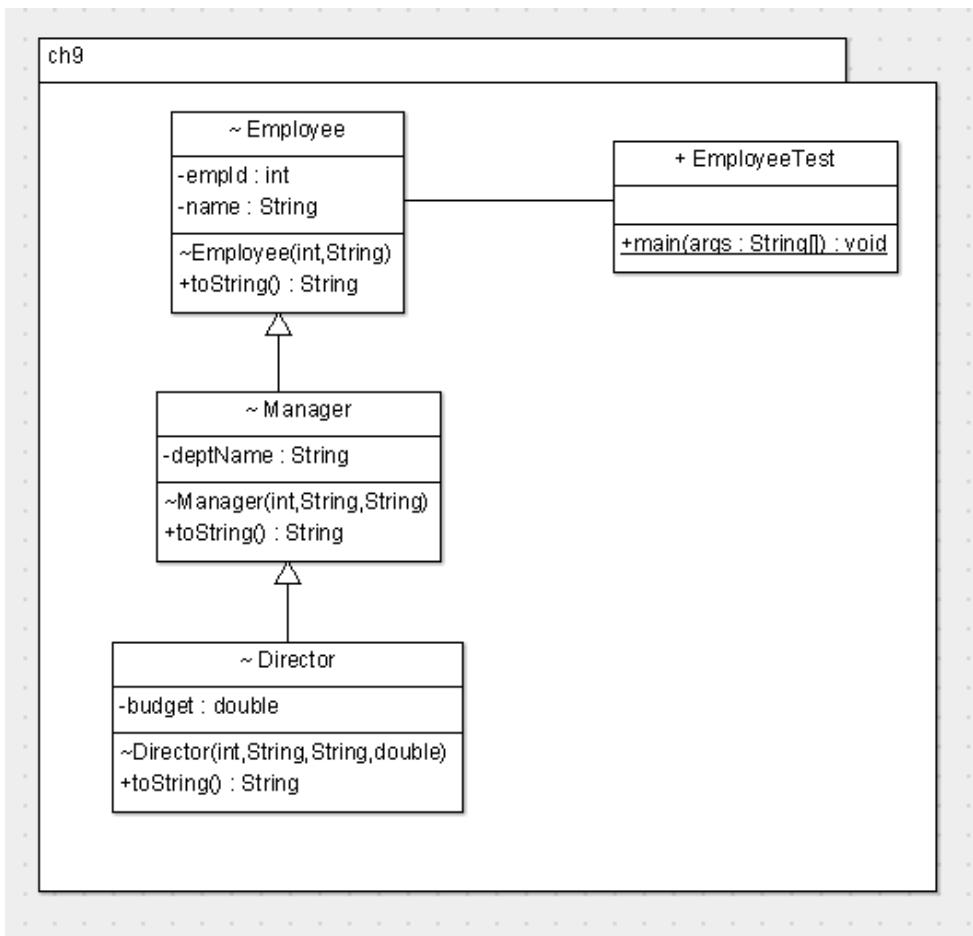


Figure 9.9 – UML for demonstrating super() and super.

In this figure, we have three classes representing a class inheritance hierarchy. `Employee` is at the top of the hierarchy. `Manager` “is-a” `Employee` and `Director` “is-a” `Manager`. Indirectly, `Director` “is-a” `Employee` also. Each of the classes has private instance variables that its respective constructors will initialize, based on the arguments passed into the respective constructor. For example, the `Employee` constructor takes in two parameters, `int` followed by `String`; these parameters will be used to initialize the `Employee` instance variables, namely `empId` (`int`) and `name` (`String`).

`EmployeeTest` is simply the driver to ensure the code is working as it should. Let's examine the code. *Figure 9.10* is the code for the UML in *Figure 9.9*:

```

3  class Employee {
4      private int empId;
5      private String name;
6
7      Employee(int empId, String name) {
8          this.empId = empId;
9          this.name = name;
10     }
11     @Override public String toString() { return "ID: " + empId + ", " + "Name: " + name + ","; }
12 }
13 class Manager extends Employee // a Manager "IS-A" Employee
14     private String deptName; // a Manager "HAS-A" department
15
16     Manager(int empId, String name, String deptName) {
17         super(empId, name); // call parent constructor
18         this.deptName = deptName;
19     }
20     @Override
21     public String toString() {
22         // call the parent toString()
23         return super.toString() + " Department: " + deptName + ",";
24     }
25 }
26 class Director extends Manager {
27     private double budget;
28
29     Director(int empId, String name, String department, double budget) {
30         super(empId, name, department);
31         this.budget = budget;
32     }
33     @Override public String toString() { return super.toString() + " Budget: " + budget; }
34 }
35 public class EmployeeTest {
36     public static void main(String[] args) {
37         Employee emplDir = new Director( 754, name: "Joe Bloggs", department: "Marketing", budget: 10_000.00);
38         System.out.println(emplDir); // ID: 754, Name: Joe Bloggs, Department: Marketing, Budget: 10000.0
39     }
40 }
41
42 }
```

Figure 9.10 – Code demonstrating super

In this figure, the `Employee` class initializes its instance variables (lines 8-9). The `toString()` method for `Employee` (line 11) returns a `String` outlining the values in the `empId` and `name` instance variables. Line 11 also uses the `@Override` annotation because it is overriding the `toString()` method inherited from `Object`.

The `Manager` class “is-a” `Employee` (line 13). `Manager` contains (is composed of) a `String` instance variable, namely `deptName`. This is known as composition.

### Composition versus inheritance

Composition defines a “*has-a*” relationship whereas, inheritance defines an “*is-a*” relationship. Composition is where an object is “composed” of other objects. For example, Car has Engine. In *Figure 9.10*, Manager “*is-a*” Employee (line 13), but Manager “*has-a*” department, which is represented by the String instance variable deptName (line 14).

The Manager constructor (lines 16-19) is where things get interesting. Line 17, `super (empId, name)`, is the call to the parent constructor in Employee passing up the employee ID (`empId`) and employee name (`name`) that are required by the Employee constructor. That is why the Manager constructor requires those parameters in the first place – it needs the employee ID and employee name so it can invoke its parent Employee constructor. The Manager constructor also requires the department name so that it can initialize its own instance variable, `deptName`. Thus, when executing a Manager constructor, the Employee constructor is executed first and then the Manager constructor executes.

Note that if line 17 is commented out, the code will not compile. Why? Because the compiler will now insert `super ()`; which is attempting to call the Employee constructor with no-arguments (the no-args constructor, namely `Employee ()`). There is no such constructor in Employee. Additionally, as Employee has already defined a constructor, the compiler will not insert the default (no-args) constructor.

The Manager classes’ `toString()` method (lines 21-24), overrides the version inherited from Employee. However, Manager can still access the Employee version, which it does by using `super.toString()` on line 23. Thus, the `toString()` method in Manager first executes the `toString()` method in Employee, which returns the employee ID and employee name. The Manager classes’ `toString()` method then appends its own instance variable, `deptName`, to the overall String to be returned.

The Director class behaves similarly to Manager. The constructor “supers up” (line 30) the required data for the Manager constructor; in turn, the Manager constructor supers up the required data for the Employee constructor. So, when creating a Director object, the order of constructor calls is as follows: Employee is first; Manager is second; Director is last. On line 31, Director initializes its own instance data.

The Director version of `toString()`, on line 33, first calls the Manager version of `toString()` using `super.toString()`. The Manager version (line 23) then calls the Employee classes’ `toString()` method, which is on line 11. So, the employee’s ID and name are the first employee details in the string. Next, the manager data (`deptName`) is appended (after the call to the Employee classes’ `toString()` method returns). Lastly, the Director data (`budget`) is appended to the string (after the call to the Manager classes’ `toString()` method returns). Note that you cannot bypass a level in the hierarchy; meaning that, `super.super.` is not allowed.

EmployeeTest is the driver class. In main() on line 39, we create a Director object that can be accessed via an Employee reference of emp1Dir (implicit upcasting). Using super() as outlined, this results in the Employee constructor being executed first, followed by the Manager constructor, and lastly the Director constructor being executed.

Line 40 passes the emp1Dir reference to System.out.println(), resulting in a polymorphic call to the Director classes' toString() method. Using super.toString(), Director invokes the Manager classes' toString() method, which also has a super.toString() method resulting in Employee toString() being executed first. Then, the Manager classes' toString() method finishes, and lastly, the Director classes' toString() method finishes. The output shows this:

```
ID: 754, Name: Joe Bloggs, Department: Marketing, Budget:  
10000.0
```

Regarding the output, ID: 754, Name: Joe Bloggs is output from Employee toString(), Department: Marketing is output from the Department toString(), and Budget: 10000.0 is output from Director toString().

That concludes our discussion on super. Now that we understand inheritance, as promised in Chapter 8, let's return to the protected access modifier.

## Revisiting the protected access modifier

Recall that a protected member is accessible from within its own package and any subclasses outside of the package: *protected = package + children*. On the face of it, this seems very straightforward. However, some nuances lead to confusion. The subclasses that access the protected member (via inheritance), can only do so *in a very specific way*. A subclass from outside the package cannot use a superclass reference to access the protected member! In addition, an unrelated class from outside the package cannot use a reference to the subclass outside the package either to access the protected member. In effect, once the subclass that's outside the package inherits the protected member, that member becomes private to the subclass (and subclasses of the subclass). This is quite tricky and definitely needs an example.

### The UML diagram

Figure 9.11 shows the UML diagram for this example:

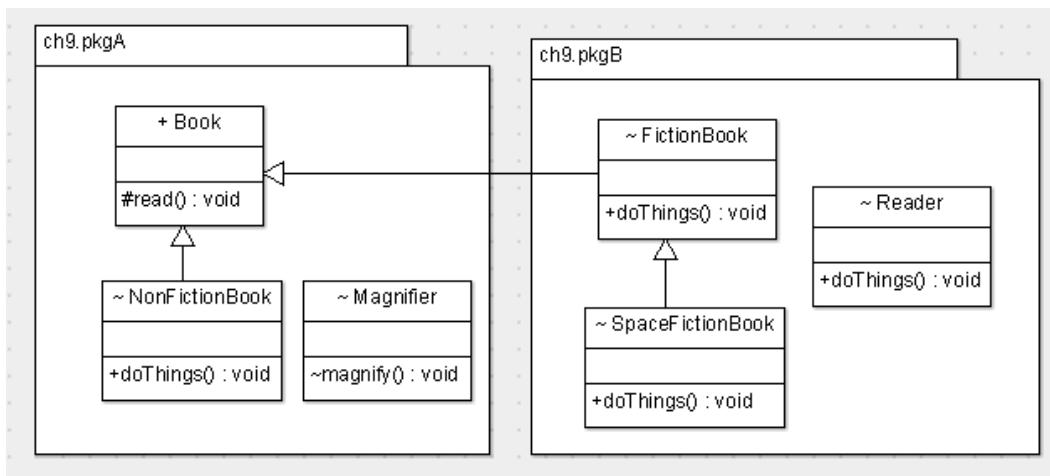


Figure 9.11 – UML for “protected” code

In this figure, we have two packages, namely `ch9.pkgA` and `ch9.pkgB`. In `ch9.pkgA`, we have a `Book` class and its subclass, `NonFictionBook`. The `read()` method in `Book` is marked with the `#` symbol, which means it is `protected`. The `Magnifier` class is not related to `Book` and is simply another class in the same package.

In `ch9.pkgB`, `FictionBook` subclasses `Book` from `ch9.pkgA` and provides a `doThings()` method, which we will use to demonstrate what is/is not allowed. In addition, `SpaceFictionBook` subclasses `FictionBook` and overrides the `doThings()` method inherited from `FictionBook`. Lastly, `Reader` is a completely separate class from the `Book` hierarchy; its `doThings()` method is also a sample method for demo purposes.

Recall from the previous chapter that we had not fully completed the access modifiers table (as we had not covered inheritance at that point). *Table 9.1* represents the completed access modifiers table. Bear in mind that the table represents annotating a member in the `Book` class.

Access Modifier	<code>Book</code>	<code>NonFictionBook</code>	<code>Magnifier</code>	<code>FictionBook</code>	<code>SpaceFictionBook</code>	<code>Reader</code>
<code>private</code>	Y	N	N	N	N	N
<code>package-private</code>	Y	Y	Y	N	N	N
<b><code>protected</code></b>	Y	Y	Y	Y	Y	N
<code>public</code>	Y	Y	Y	Y	Y	Y

Table 9.1 - Access modifiers table with ‘protected’ row fully filled out

Examining the `protected` row, we can now see that subclasses, regardless of the package, can access inherited `protected` members.

Now let us examine the code for each of the packages in turn. Firstly, we will examine the package that defines the `protected` member.

## The package with the protected member

Figure 9.12 shows the code for the first package, `ch9.pkgA`, from Figure 9.11:

```
1 package ch9.pkgA;
2
3 public class Book {
4     protected void read(){}
5 }
6 class NonFictionBook extends Book{
7     public void doThings(){
8         read(); // same package; no problem
9     }
10}
11 class Magnifier{
12     void magnify(){
13         Book b = new Book();
14         b.read(); // same package; no problem
15     }
16}
```

Figure 9.12 – Code for “ch9.pkgA” from UML

In this figure, we have a class called `Book` (lines 3-5) that defines a `protected read()` method (line 4). `NonFictionBook` is a subclass of `Book` and has its own `doThings()` method (lines 7-9). In addition, there is a completely unrelated class to `Book`, namely `Magnifier`.

The first thing to note is that, as the `read()` method is `protected`, other code in the same package can access it, even if the code is *not* a subclass. This is demonstrated by line 14, where the `read()` method in `Book` is accessed from `Magnifier`, a completely unrelated class.

Of course, regardless of the package, subclasses can access the `protected` member. This is shown on line 8, where the `NonFictionBook` subclass invokes `read()`. Remember that line 8 is essentially `this.read()`. So, whichever `NonFictionBook` object is used to invoke `doThings()` on line 7 will be used to invoke the inherited (and `protected`) `read()` method on line 4.

The interesting code is in the other package, namely `ch9.pkgB`. Let’s examine that now.

## The other package

Figure 9.13 presents the code:

```
1  package ch9.pkgB;
2
3  import ch9.pkgA.Book;
4
5  class FictionBook extends Book{
6      public void doThings(){
7          read(); // different package, via inheritance; no problem
8          this.read(); // different package, via inheritance; no problem
9          FictionBook fb = new FictionBook(); // default ctor created for us
10         fb.read(); // no problem
11
12         // Here, I create an instance of the superclass that has the protected
13         // member. Note that even though FictionBook has access via inheritance
14         // to read(), FictionBook must access it properly.
15         Book b = new Book();
16         b.read(); // not public!
17     }
18 }
19 class SpaceFictionBook extends FictionBook{
20     public void doThings(){
21         read(); // different package, via inheritance; no problem
22         new Book().read();
23         new FictionBook().read();
24         new SpaceFictionBook().read(); // ok
25     }
26 }
27 class Reader{
28     public void doThings(){
29         Book b = new Book();
30         b.read();
31
32         // can I access the protected member via the subclass that inherits it?
33         FictionBook fb = new FictionBook();
34         fb.read();
35     }
36 }
```

Figure 9.13 – Code for “ch9.pkgB” from UML

In this figure, we can see that `FictionBook` “is-a” `Book` (line 5) and `SpaceFictionBook` “is-a” `FictionBook` (line 19). For this hierarchy to be possible, the `Book` class needed to be imported from the other package (line 3). We were only able to import `Book` from another package because `Book` is a `public` class. In addition, we have a completely unrelated class called `Reader` (lines 27-36).

Now, for the fun! Let's examine the `dothings()` method in `FictionBook` (lines 6-17). Lines 7 and 8 are essentially equivalent and show that subclasses outside the package, when using *inheritance directly*, can access **protected** members.

Lines 9-10 also show that when inside the subclass outside the package, if you create an instance of that particular subclass (`FictionBook`, in this instance), then all is ok. This makes sense because the two references used to invoke `read()` without issue, namely `this` and `fb`, are both of type `FictionBook`, where the code resides.

Note that line 15, where we instantiate a `Book` object, compiles because the `Book` class (*Figure 9.12*, line 3) is **public**. The `Book` class did not define a constructor, so the default constructor was created for us. This default constructor takes on the same access as the class, namely **public**, and as a result, we can invoke the constructor from a different package.

Line 16, which does not compile, is very interesting. When inside the subclass outside the package, you cannot access the **protected** member using the superclass reference – even though the **protected** member resides in that superclass! Remember that, once outside the package, the **protected** member becomes **private** to subclasses (and their subclasses). In other words, you must use inheritance in a very specific way.

`SpaceFictionBook` (lines 19-26) shows that access is available to subclasses of the subclass outside the package. Line 21 is the same as line 7, except that they are in two separate classes. As this line compiles, it demonstrates that subclasses of the subclass outside the package have access to the **protected** member in the base class.

Lines 22 and 23 both fail to compile. Line 22 tries to access the **protected** member via a `Book` reference and line 23 tries to access it via a `FictionBook` reference. Both fail. Contrast this with line 24, which uses an instance of the current class, namely `SpaceFictionBook`, which works. Note that line 24 is similar to line 21 in that a `SpaceFictionBook` reference is used in both instances (as line 21 is equivalent to `this.read()`). In addition, line 24 is very similar to lines 9-10. Therefore, when in a subclass outside the package, access the **protected** member directly, as on lines 7, 21; or use a reference to the current subclass, as on lines 10, 24.

The `Reader` class (lines 27-36) is a completely separate class from the `Book` hierarchy. Line 30 attempts to access the **protected** member using a reference to the class that defines the **protected** member, namely `Book`, and fails. Line 34 attempts to access the **protected** member using a reference to the subclass outside the package that inherits the **protected** member, namely `FictionBook`, and also fails.

So **protected** is somewhat tricky. While we are revisiting previous topics, it is an ideal opportunity to revisit `switch`. To be more specific, to discuss pattern matching for `switch`.

## Pattern matching for switch

As promised from *Chapter 4*, now that we understand inheritance and polymorphism, we are going to revisit the `switch`. Given the following code:

```
public static void patternMatchingSwitch(Vehicle v) {  
    System.out.println(  
        switch(v){  
            case Boat b -> "It's a Boat";  
            case Train t -> "It's a Train";  
            case Car c when c.getNumDoors() == 4 ->  
                "Saloon "+ c.onRoad(); // custom Car method  
            case Car c when c.getNumDoors() == 2 ->  
                "Convertible: " + c.onRoad();  
            case null, default -> "Invalid type";  
        }  
    );  
}
```

Assume that `Car`, `Boat`, and `Train` all extend from `Vehicle` and that `Car` has a custom method `onRoad()`. As you can see, in this `switch` expression, the selector expression, `v`, can be any reference type (`Boat`, `Train`, `Car`, and so forth). The `case` labels demonstrate *type patterns and pattern matching*; for example, `Boat b`.

In addition, both `case` labels for `Car` are known as *guarded patterns*. A guarded pattern is a `case` label protected by a “guard” on the right-hand side of a `when` clause. A guard is a conditional expression, evaluating to true or false. Note the use of the custom `Car` method `onRoad()` and the fact that no cast is required, as the cast is done for us in the background (provided we are dealing with a `Car`).

The last `case` label, containing `default`, ensures exhaustiveness is catered for, thereby keeping the compiler happy. In other words, all possible `Vehicles` are catered for. Note also the use of `null` as a valid label and the fact that `null` and `default` can be comma separated.

Now, let’s examine the effect on inheritance of two particular keywords, namely `abstract` and `final`.

## Explaining the abstract and final keywords

As we know, when coding methods, we can apply the access modifier keywords, namely `private`, `protected`, `public`, and package-private (no keyword). Two other keywords have special significance regarding inheritance: `abstract` and `final`. Both are opposites of each other, which is why both cannot be applied to a method at the same time. Let’s discuss them now, starting with `abstract`.

## The `abstract` keyword

The `abstract` keyword is applied to classes and methods. While `abstract` classes will be discussed more fully in *Chapter 10*, we will be discussing them here also (for reasons that will soon become obvious). An `abstract` method has no implementation (code). In other words, the method signature, rather than following it with curly braces, `{ }`, which represents the implementation, an `abstract` method signature is simply followed by a semi-colon. Marking a method as `abstract` implies the following:

- The class must be `abstract` also
- The first concrete (non-`abstract`) subclass must provide an implementation for the `abstract` method

Let's discuss this in more detail. When you mark a method (or methods) as `abstract`, you are saying that this method has no implementation code. As there is something "missing," the class itself must be marked as `abstract` also. This tells the compiler that the class is incomplete and as a result, you cannot instantiate (create) an object based on an `abstract` class. In other words, you cannot execute `new` on an `abstract` class (although a reference is perfectly ok). The whole rationale for `abstract` methods (and thus `abstract` classes) is for them to be overridden by subclasses, where the "missing" implementation code is provided. Now, if the direct subclass does not provide the implementation code for the inherited `abstract` method, that subclass must also be `abstract`. Therefore, the first non-`abstract` (concrete) subclass of an `abstract` class must provide the implementation code for the `abstract` method. *Figure 9.14* demonstrates these principles:

```
3  abstract class Pencil{  
4      abstract void write(); // no {}  
5  }  
6  class CharcoalPencil extends Pencil{}  
7  abstract class WaterColorPencil extends Pencil{}  
8  class GraphitePencil extends Pencil{  
9      @Override  
10     void write(){  
11         System.out.println("GraphitePencil::write()");  
12     }  
13 }  
14  
15 public class PencilsExample{  
16     public static void main(String[] args) {  
17         Pencil pp    = new Pencil(); // cannot "new" a Pencil (abstract)  
18         Pencil pdp   = new GraphitePencil();  
19         pdp.write(); // GraphitePencil::write()  
20     }  
21 }
```

Figure 9.14 – The “`abstract`” keyword in action

In this figure, we have an `abstract` method, namely `write()`, on line 4. Notice how there are no curly braces for the method; we just have the semi-colon immediately after the parentheses. As the `Pencil` class (lines 3-5) contains an `abstract` method, the class itself must be `abstract`; which it is (line 3).

On line 6, `CharcoalPencil` attempts to subclass `Pencil`. But because (a) it does not provide an implementation for the `abstract` method `write()`, which it inherited from `Pencil`, and (b) `CharcoalPencil` *itself* is not `abstract`, `CharcoalPencil` fails to compile.

Contrast line 6 with line 7. As we saw, line 6 does not compile. However, line 7, `WaterColorPencil`, does compile. Why? Because `WaterColorPencil` is `abstract`; the fact that it does not provide an implementation for the `abstract` method `write()` is no problem.

#### Abstract classes do not have to have abstract methods

As we know, if you have 1 (or more) `abstract` methods, then the class must be `abstract`. However, the opposite is not true. In other words, an `abstract` class does not have to have any `abstract` methods at all! Note that `WaterColorPencil` (line 7 in *Figure 9.14*) is an example of such a class. It is `abstract` and yet has no methods at all. This is fine. This could be a design decision whereby, even if the class contains only concrete methods, you simply want this class to be used as a reference type and not as an object type (as you cannot `new` it).

The `GraphitePencil` class (lines 8-13) is a concrete, non-`abstract` class. As it extends the `abstract` class, `Pencil`, it must provide an implementation for the `abstract` method `write()`. This is done on lines 10 to 12 and we use the `@Override` annotation to emphasize this.

Line 17 demonstrates that you cannot instantiate an object of an `abstract` class. The reference part of the `Pencil pp` statement is fine. The issue is with the `new Pencil()` part.

Line 18 shows what is allowed. Again, we are using a `Pencil` reference but this time, we are referring to a `GraphitePencil` object. `GraphitePencil` is a concrete class (line 8). Line 19 polymorphically calls the `write()` method provided by `GraphitePencil` (lines 10-12). Assuming lines 6 and 17 are commented out (so the code will compile), line 19 outputs `GraphitePencil::write()`.

Now that we understand `abstract` methods and classes, let's examine the `final` keyword.

## The `final` keyword

The `final` keyword can be applied in various contexts. Inheritance is the main focus here, but we will examine other situations also. We will examine each in turn and then look at code that demonstrates them. We will start with `final` methods.

### ***final methods***

A `final` method cannot be overridden in a subclass. This prevents any unwanted changes by subclasses. We can take this a stage further with `final` classes.

### ***final classes***

A class that is marked `final` cannot be used as a base type. This means you cannot extend from a `final` class. All the methods in the class are implicitly `final`. Java uses this in its API to guarantee behavior. For example, the `String` class is `final` so that nobody can extend it and provide a custom implementation. Therefore, Java always knows how strings behave. Now, we will examine `final` method parameters.

### ***final method parameters***

A `final` method parameter is a parameter that cannot be changed. However, be aware that the semantics are subtly different depending on the parameter type. If the parameter type is a primitive, such as `int`, then you cannot change the value of the `int` parameter.

However, if the parameter in question is a reference (as opposed to a primitive), `final` applies to the reference and therefore, it is the reference that cannot be changed. In other words, the object the reference is pointing to is modifiable, but the reference itself is not. What this means is that, for example, if the method accepts a `Dog` reference, namely `dog`, then using the `dog` reference, you can change the properties of the object, such as `dog.setAge(10)`. You cannot, however, change `dog` to refer to a different object, such as `dog = new Dog()`.

### ***final (constants)***

A constant is a value that cannot change. It is customary and good practice to use capital letters as the identifiers for constants, with each word separated by an underscore. This makes them stand out and developers know they cannot change them. One example from the Java API is the `PI` constant from the `Math` class (in the auto-imported `java.lang` package). It is `final` so that it cannot be changed. To provide easy access, `PI` is also `public` and `static`.

Now, let's look at a code example to re-enforce the use of `final`. *Figure 9.15* presents the code:

```
3   final class Earth{}  
4   // cannot extend a 'final' class  
5   class SubEarth extends Earth{}  
6  
7   ↪ class Pen{  
8       final void write(){}
9       // 'final' and 'abstract' not allowed together
10      // as they have opposite meanings
11      // final abstract scribble();
12  }  
13  ↪ class FountainPen extends Pen{
14      // cannot override a 'final' method
15  ↗ @Override void write(){}
16  }  
17  ↪ public class DemoOfFinal {
18      final int ONE_YEAR = 1;
19  ↗ @ void print(final String name, final int age){
20      // primitives
21      age = age + ONE_YEAR;
22      // references - ok to access the object
23      System.out.println(name.toUpperCase());
24      // references - cannot modify the reference
25      name = "Alexander";
26
27      ONE_YEAR = 2; // cannot change a constant
28  }  
29 }
```

Figure 9.15 – The “final” keyword in action

In this figure, we have a `final` class called `Earth` (line 3). Line 5 demonstrates, via a compiler error, that you cannot extend from a `final` class.

Line 8 defines a `final` method called `write()` in the `Pen` class. Consequently, the `FountainPen` class encounters a compiler error (line 15) when attempting to override `write()`.

Line 11 shows that you cannot annotate a method as both `abstract` and `final - abstract` implies that this method is to be overridden in a subclass; `final` means that this method must not be overridden.

Line 18 declares a constant called `ONE_YEAR` and sets it to 1. Line 27 attempts to change the constant value – as this is not allowed, the compiler complains.

The `print()` method (lines 19-28) outlines what `final` means for method parameters. The method parameters (line 19) are `final String name` and `final int age`, respectively. `String` is a non-primitive type and therefore `name` is a reference. In other words, the value inside `name` is a memory location (reference) of where the object is on the heap. On the other hand, `age` is simply a primitive `int`, whose value is simply a whole number, such as 1. It is easy to understand what you can and cannot do with `final` parameters when you view the *value* as `final`. Thus, if 1 is in `age`, it cannot be changed and neither can the reference (address) in `name`. However, the object referred to by `name` can be modified.

Line 21 is a compiler error and demonstrates that `final` primitives cannot be changed.

Line 23 shows us that the object that the reference is referring to can be accessed (and changed if required). Note that, in this particular example, as `Strings` are immutable objects, the `toUpperCase()` method returns the new, uppercase `String`, as opposed to changing the original. We will talk more about `Strings` in *Chapter 12*. The important thing to note is that the compiler had no issue with line 23.

Line 25 attempts to change the `String` reference `name` to refer to a different `String`. As the reference is `final`, the compiler complains. Once again, the separation of reference and object makes things much easier to understand.

At this point, we know how to create (unlimited) inheritance hierarchies (using `extends`). We also know that `final` disables inheritance. What if we wanted a “middle ground,” where we could customize our hierarchy to certain types? This is what sealed classes enable. Let’s discuss them now.

## Applying sealed classes

Sealed classes were introduced in Java 17. What we are going to cover here relates to classes but the same logic applies to interfaces (*Chapter 10*). With inheritance, you can extend from any class (or interface) using the `extends` keyword, unless the class is `final` of course.

**Note**

Interfaces cannot be `final` because their whole rationale is to be implemented.

Consider the following scenario: what if you wanted your class to be available for inheritance, but only for certain classes? In other words, you want to scope the subclasses allowed. So far, inheritance, using `extends`, enables every class to become a subclass, whereas `final` prevents a class from having subclasses.

This is where sealed classes are useful – they enable you to specify what subclasses are allowed. Just to reiterate, this also applies to interfaces, where we can specify what classes are allowed to implement the interface.

Before we look at an example, there are some new keywords that we need to understand.

## **sealed and permits**

These keywords work together. To state that a class is sealed, you can simply specify that it is just that, `sealed`. Once you do that, however, you must specify which classes can extend from this class. To do that, you use the `permits` keyword, followed by the comma-separated list of classes.

## **non-sealed**

When you start to scope/restrict a hierarchy, you must use certain keywords when specifying the subclasses. A subclass involved in a sealed hierarchy must state one of the following:

- It is also sealed. This means, we have further scoping to perform and therefore we must use the `sealed/permits` pairing on this subclass to specify the subclasses allowed.
- It is the `final` class in the hierarchy (no more subclasses allowed).
- It ends the scoping. In effect, you want to open up the hierarchy again for extension. To do this, we use the `non-sealed` keyword as `non-sealed` classes can be subclassed.

Now, let's look at an example.

## **Example using sealed, permits, and non-sealed**

*Figure 9.16* presents a UML diagram for the code example we will use:

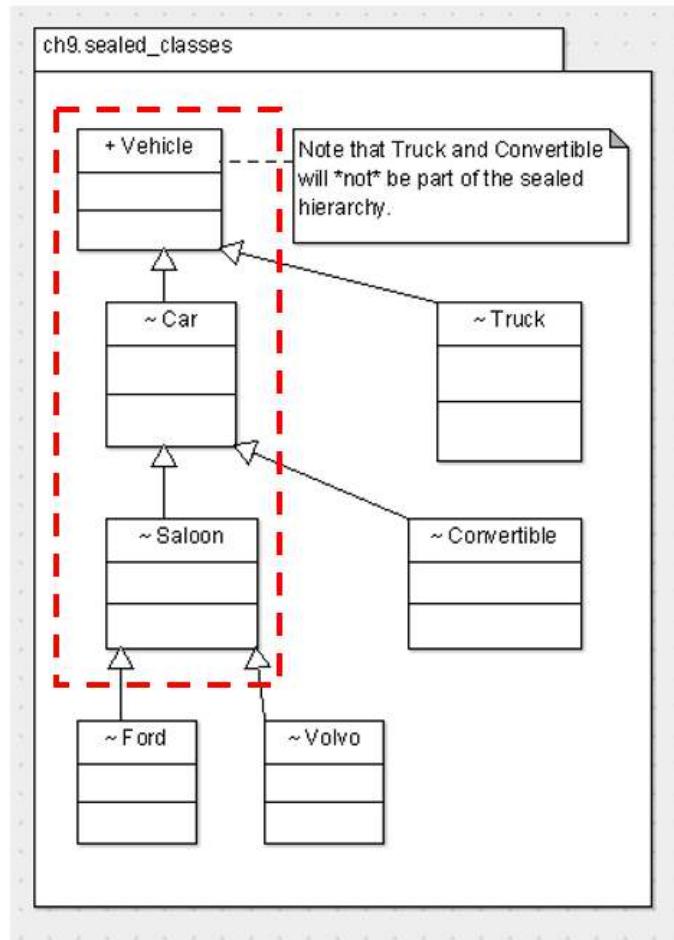


Figure 9.16 – UML diagram for “sealed” classes

In this figure, we have a `Vehicle` hierarchy. The parts we are going to restrict (seal) are the `Vehicle`, `Car`, and `Saloon` classes. Therefore, the only class that can subclass `Vehicle` is `Car`; and the only class that can subclass `Car` is `Saloon`. Note that even though the diagram implies `Truck` “is-a” `Vehicle` and `Convertible` “is-a” `Car`, for this example, we will prevent that in code.

The goal of the code is to ensure that the only `Vehicles` we are interested in are `Cars` and the only `Cars` we are interested in are `Saloons`. In addition, all `Saloons` (`Ford` and `Volvo`) are of interest. *Figure 9.17* presents the code.

```
3 ❸ public sealed class Vehicle permits Car{ } // scoping hierarchy
4 ❸ sealed class Car extends Vehicle permits Saloon {}
5 //sealed class Truck extends Vehicle {}      // compiler error
6 ❸ non-sealed class Saloon extends Car{}      // opening up hierarchy again
7 class Volvo extends Saloon{}
8 class Ford extends Saloon{}
9 //class Convertible extends Car{}           // compiler error
```

Figure 9.17 – “sealed” code

In the preceding figure, line 3 states that we have a sealed class called `Vehicle` and that the only subclass allowed (permitted) is `Car`. At this point, the `Car` class must exist; otherwise, the compiler will complain.

Line 4 defines a sealed class called `Car` as a subclass of `Vehicle` (which it must do due to line 3) and that the only subclass permitted is `Saloon`. Note that when we were defining `Car`, we had to specify that `Car` was either sealed, non-sealed, or final.

Line 5 is the `Truck` class attempting to subclass `Vehicle`. However, as we have sealed `Vehicle` to only allow `Car` as a subclass, this generates a compiler error.

Line 6 defines `Saloon` as a subclass of `Car` (as expected from line 4). In this instance, we have chosen to open up the hierarchy for further extension (by any class) by stating that `Saloon` is non-sealed. Lines 7 and 8 demonstrate that `Saloon` is a non-sealed class by allowing `Volvo` and `Ford` to extend from it, respectively.

Lastly, on line 9, `Convertible` attempts to subclass `Car`. This is not allowed as line 4 states that the only subclass of `Car` allowed is `Saloon`.

Let's move on now and discuss both instance and static blocks.

## Understanding instance and static blocks

As we know, in Java, a block is delimited by curly braces, { }, and these code blocks are no different. What is different about instance and static code blocks is *where* these blocks appear – in other words, their scope. Both of these code blocks appear outside every method but inside the class.

We will discuss each in turn and then present a code example to demonstrate them in operation. We will start with instance blocks.



## Instance blocks

An instance block is a set of braces that appear outside of any method but inside the class. Assuming an instance block is present in a class, every time an object is created (using new), the instance block is executed. Note that the instance block executes *before* the constructor. To be technically accurate, `super()` is executed first so that the parent constructor has a chance to execute; this is followed by the instance block, after which the rest of the constructor executes. Use the “sic” (super, instance block, constructor) acronym to help remember the order. You can think of the compiler inserting the instance block into the constructor code just after the call to `super()`. If more than one instance block exists in a class, they are executed in order of appearance, from top to bottom.

As instance blocks execute as part of every constructor, they are an ideal location for inserting code that you want every constructor to have. In other words, code that is common across all constructors should go into an instance block. This saves you from duplicating code across constructors.

As we know, the parent constructor must execute before the child constructor. The same occurs with instance blocks. In other words, the parent instance blocks must execute before the child instance blocks. We will see this in the code example.

## static blocks

A `static` block is a set of braces, preceded by the `static` keyword, that appears outside of any method but inside the class. The `static` block is only executed once, the very first time the class is loaded. This could occur when the first object of the class is created or the first time a `static` member is accessed. Static blocks execute before instance blocks (as we have to load the class file/bytocode before we can execute a constructor). Once executed, given that the class file is now loaded into memory, the `static` block is never executed again.

As with instance blocks, if more than one `static` block exists in a class, they are executed in order of appearance, from top to bottom. Similarly, as with instance blocks, if inheritance is involved, then the parent `static` blocks execute before the child `static` blocks.

This will all make a lot more sense with a code example, where we will be able to compare and contrast both types of code blocks in an inheritance hierarchy.

*Figure 9.18* presents the code:

```
3  class Parent{
4      // instance initialization block
5      { System.out.println("6. Parent instance init block 1"); }
6      // static initialization block
7      static {System.out.println("2. Parent static init block 1");}
8      Parent() { System.out.println("8. Parent constructor"); }
9      { System.out.println("7. Parent instance init block 2"); }
10     static {System.out.println("3. Parent static init block 2");}
11 }
12 }
13 }
14 class Child extends Parent{
15     { System.out.println("9. Child instance init block 1"); }
16     static {System.out.println("4. Child static init block 1");}
17     Child() { System.out.println("11. Child constructor"); }
18     { System.out.println("10. Child instance init block 2"); }
19     static {System.out.println("5. Child static init block 2\n");}
20 }
21 }
22 }
23 public class InitializationBlocks {
24     static {System.out.println("1. InitializationBlocks static init block");}
25     { System.out.println("InitializationBlocks instance init block"); }
26     public static void main(String[] args) {
27         System.out.println("----> Creating first Child object...");
28         new Child();
29         System.out.println("\n---->Creating second Child object...");
30         new Child();
31     }
32 }
```

Figure 9.18 – Instance and “static” code blocks example

In this figure, we have a parent class called `Parent` and a child class called `Child` (it took a while to come up with those names!). Both classes have two instance initialization blocks, two static initialization blocks, and a constructor. Notice that the `static` initialization blocks (lines 7, 12, 16, and 21) are all simply blocks of code preceded by the `static` keyword. Also, note their location/scope – outside the methods but inside the class. The same is true for the instance initialization blocks (lines 5, 11, 15, and 20), except that the instance blocks have no keyword preceding them.

The main driver class, `InitializationBlocks`, also has one `static` and one instance initialization block (lines 24 and 25, respectively).

Each of these blocks simply outputs a tracer message so that we know which block of code is currently executing. The tracer messages are annotated with ascending numbers so we can follow the order of execution more easily. *Figure 9.19*, presents the output from the code in *Figure 9.18*:

```
1. InitializationBlocks static init block
---> Creating first Child object...
2. Parent static init block 1
3. Parent static init block 2
4. Child static init block 1
5. Child static init block 2

6. Parent instance init block 1
7. Parent instance init block 2
8. Parent constructor
9. Child instance init block 1
10. Child instance init block 2
11. Child constructor

---> Creating second Child object...
6. Parent instance init block 1
7. Parent instance init block 2
8. Parent constructor
9. Child instance init block 1
10. Child instance init block 2
11. Child constructor
```

Figure 9.19 – Output from the code in Figure 9.18

#### Note

To avoid confusion between numbers representing output in *Figure 9.19* with line numbers in *Figure 9.18*, all numbers mentioned here refer to output numbers in *Figure 9.19*. Any line numbers relating to *Figure 9.18* will be explicitly annotated as “line....”

All Java programs start with the `main()` method. Therefore, the JVM has to find, using the `CLASSPATH` environment variable, the `.class` file containing `main()`, namely `InitializationBlocks.class`. As the JVM loads the class, if the class has a parent, it loads the parent first. In this example, as `InitializationBlocks` is not a subclass, this does not apply. However, there is a `static` block and this gives us our first line of output. Note that the instance block for `InitializationBlocks` is never executed. This is because no instance of `InitializationBlocks` was ever created. In other words, there is no new `InitializationBlocks()` in the code.

Line 27 simply outputs "---> Creating first Child object...". What is interesting to note is that it is not the first line output to the screen – the output from the `static` block is first.

Line 28 creates a `Child` object. Its output is represented by numbers 2-11. As this is the first time a `Child` object has been created (as no `static` member in `Child` has been accessed before this), the class file for `Child` is loaded. During this process, the JVM realizes that `Child` is a subclass of `Parent`, so it loads the `Parent` class first. Therefore, the `static` blocks in `Parent` are executed first, in order of appearance (2 and 3); followed by the `Child` static blocks, also in order of appearance (4 and 5).

Now that the `static` blocks are done, the instance blocks and constructors are executed. First, the superclass `Parent` instance blocks are executed in order of appearance (6 and 7), followed by the `Parent` constructor (8). Then, the subclass `Child` instance blocks are executed in order of appearance (9 and 10), followed by the `Child` constructor (11). That is a lot of processing from a simple `new Child()` line of code.

Line 29 simply outputs "---> Creating second Child object...".

Line 30 creates another `Child` object. As the class was already loaded previously, the `static` blocks will already have run for both `Child` and its superclass, `Parent`. Therefore, they do not run again. So, we run the `Parent` instance blocks (6 and 7), followed by the `Parent` constructor (8). Then, we run the `Child` instance blocks (9 and 10), followed by the `Child` constructor (11).

Note the repetition of line numbers 6-11 when creating a `Child` object. The `Parent` instance blocks are executed in order; followed by the `Parent` constructor. The `Child` instance blocks and constructor follow in a similar fashion.

That covers `static` and instance initialization blocks. Before we conclude this chapter on inheritance, we would just like to delve a little deeper into one of the topics we touched on earlier: upcasting and downcasting.

## Mastering upcasting and downcasting

Earlier, we touched upon why we get `ClassCastException` errors. The rule is that a reference can refer to objects of its own type or objects of subclasses. In effect, a reference can point across and down the inheritance hierarchy, but never up. If a reference does point up the hierarchy, you will get a `ClassCastException` error. Recall that the reason this occurs is that the subclass reference could have extra methods that any superclass object would have no code for. Whether that is the case or not is immaterial, *could have* is enough.

Keep in mind that assignment works from right to left; so, when reading code involving upcasting/downcasting, the direction in the hierarchy is from right to left as well. In addition, remember that the compiler is always looking at the reference type.

Now, let's discuss, with the aid of code examples, both upcasting and downcasting. Let's start with upcasting.

## Upcasting

With upcasting, you are going from a more specific type “up to” a more general type. For example, let's look at the following line of code:

```
Vehicle vc = new Car()
```

Here, we are going from `Car` *up* to `Vehicle`. The more specific type (`Car`) is further down the hierarchy and potentially has extra methods. Due to inheritance, whatever methods the parent reference has access to, the subclass will also have. So, any methods available to the `Vehicle` reference, `vc`, will exist in the `Car` object! Therefore, upcasting is never an issue, and an explicit cast is not required.

*Figure 9.20* presents upcasting in code:

```
3  class Machine {  
4      void on(){ System.out.println("Machine::on()"); }  
5  }  
6  class Tractor extends Machine {  
7      @Override void on(){ System.out.println("Tractor::on()"); }  
8      void drive() { System.out.println("Tractor::drive()"); }  
9  }  
10 public class UpcastingAndDowncasting {  
11     public static void doAction(Machine machine){  
12         machine.on();  
13     }  
14     public static void main(String[] args) {  
15         Machine mt = new Tractor(); // upcasting  
16         doAction(mt);           // polymorphism, Tractor::on()  
17         doAction(new Tractor()); // polymorphism, Tractor::on()  
18     }  
19 }
```

The code shows a `Machine` class with an `on()` method. A `Tractor` class extends `Machine` and overrides the `on()` method. The `doAction` method takes a `Machine` reference and calls its `on()` method. In the `main` method, a `Tractor` object is created and passed to `doAction`, demonstrating upcasting and polymorphism.

Figure 9.20 – Upcasting in action

In this figure, we have a class called `Machine` (lines 3-5) and a subclass called `Tractor` (lines 6-9). The `on()` method in `Tractor` (line 7) overrides the `on()` method in `Machine` (line 4).

Line 15 involves an implicit upcast. Reading it right to left (as assignment is right to left), we are going from `Tractor` “up to” `Machine`. This is possible because every `Tractor` “is-a” `Machine`. Thus, line 15 results in a `Machine` reference referring to a `Tractor` object.

Line 16 invokes the `doAction()` method while passing in the reference created on line 15, namely `mt`. This `mt` reference is copied (remember Java is call by value) into the `Machine` reference, namely `machine`, on line 11. Thus, the `mt` reference in the `main()` method and the `machine` reference in the `doAction()` method are pointing at the one and same object, which was created on line 15.

Inside the `doAction()` method, we invoke the `on()` method using the `machine` reference (line 12). As the `machine` references type, namely `Machine`, has an `on()` method, the compiler is happy. At runtime, the object that `machine` is referring to, namely `Tractor`, is used. In other words, the `on()` method from `Tractor` is dynamically executed (polymorphically).

Line 17 is just accomplishing in one line what was coded over lines 15 and 16. With the invocation of `doAction()` on line 17, the upcasting is as follows:

```
Machine machine = new Tractor()
```

The `Machine` reference, namely `machine`, is provided by the `doAction()` signature (line 11), and the `Tractor` instance creation comes from line 17.

Both lines 16 and 17 result in the same output: `Tractor::on()`. Now, let’s discuss the trickier of the two: downcasting.

## Downcasting

With downcasting, you are going from a more general type “down to” a more specific type. For example, let’s look at the following line of code:

```
Car cv = (Car) new Vehicle(),
```

Reading it from right to left, we are going from `Vehicle` *down* to `Car`. Again, the more specific type (`Car`) is further down the hierarchy and potentially has extra methods. The compiler spots this and complains. We can overrule the compiler by inserting a (down)cast, `(Car)`. This is what we have done here. However, at runtime, this line of code results in a `ClassCastException` error. This is because, on the right-hand side of the assignment statement, we are attempting to create a `Car` reference that will point up the inheritance tree at a `Vehicle` object!

Figure 9.21 presents downcasting in code:

```

3  class Machine {
4      void on(){ System.out.println("Machine::on()"); }
5  }
6  class Tractor extends Machine {
7      @Override void on(){ System.out.println("Tractor::on()"); }
8      void drive() { System.out.println("Tractor::drive()"); }
9  }
10 public class UpcastingAndDowncasting {
11     public static void doAction(Machine machine){
12         // machine.on();
13
14         // Let us try and call the Tractor-specific method 'drive()'
15         // machine.drive(); // compiler error
16         ((Tractor)machine).drive(); // possible ClassCastException
17         if(machine instanceof Tractor t){
18             t.drive(); // safe
19         }
20     }
21     public static void main(String[] args) {
22         doAction(new Machine());    // outputs nothing
23         doAction(new Tractor());   // Tractor::drive()
24     }
25 }
```

Figure 9.21 – Downcasting in action

The code in this figure is very similar to the code in Figure 9.20. The inheritance hierarchy is the same. The changes are in the `doAction()` and `main()` methods. Line 12 works normally and we have commented it out to focus on downcasting.

Our goal, as stated on line 14, is to *safely* invoke the `Tractor` object's `drive()` method. Note that this method is specific to `Tractor`. Let's look at the changes in baby steps.

Firstly, as `drive()` is specific to `Tractor` (and not `Machine`), this means that we need a `Tractor` reference to get the code to compile. The fact that line 15 does not compile demonstrates this – the `machine` reference is of the `Machine` type and `Machine` does not have a `drive()` method.

Line 16 addresses the compiler error from line 15. Line 16 compiles because it (down)casts the `machine` reference to a `Tractor` reference *before* it calls the `drive()` method. That is why the extra set of parentheses are needed – method invocation has higher precedence than casting, so we change the order of precedence by using parentheses. Without the extra set of parentheses, we have `(Tractor)`

---

`machine.drive()`, and this does not compile (for the same reason as line 15 does not compile). However, the extra set of parentheses forces the cast from `Machine` to `Tractor` to be performed first, and thus the compiler looks for `drive()` in `Tractor`.

However, we are still not “out of the woods.” Yes, the compiler is happy, but the JVM is vulnerable to `ClassCastException` errors at runtime. If line 16 were uncommented, then line 22 would cause a `ClassCastException` error at runtime. This is because line 22 passes in a `Machine` object, so inside the `doAction()` method, the `machine` reference is referring to a `Machine` object. Therefore, on line 16, we would be trying to create a `Tractor` reference to point *up* to the `Machine` object, which is a `ClassCastException` all day long.

Line 17 uses the `instanceof` keyword, in conjunction with a type pattern and pattern matching. Line 17 is only true if the reference `machine` refers to a `Tractor` object; when it is, the cast is done for us in the background and `t` is initialized to refer to the `Tractor` object. This is why line 22 outputs nothing – the `Machine` object passed in, fails the `instanceof` test and therefore line 18 is *not* executed. However, as line 23 passes in a `Tractor` object, it passes the `instanceof` test. This means that line 18 is executed and outputs `Tractor::drive()`.

That completes another hugely important chapter. Now, let’s apply what we have learned!

## Exercises

Our park is full of diversity, not just in the species of dinosaurs but also in the roles of our employees. To model this diversity, we will be incorporating the concept of inheritance into our applications:

1. Not all dinosaurs are the same. Some are small, others big. Some are herbivores, others carnivores. Create at least three subclasses for different types of dinosaurs that inherit from the base `Dinosaur` class.

If you need inspiration, you can create a `FlyingDinosaur` subclass and an `AquaticDinosaur` subclass from the `Dinosaur` class, each with its unique properties. (This is not the most optimal way to model this, but don’t worry about that now.)

2. Just like our dinosaurs, our employees also have diverse roles. Some are park managers, while others are security officers or veterinarians. Create subclasses for these employee roles that inherit from the `Employee` base class. Come up with at least three subclasses.
3. Inheritance doesn’t just stop at properties and methods. Even the behavior of some methods can be customized in subclasses. Provide a custom implementation of the `toString()` method in the `Dinosaur` and `Employee` classes (from exercises 1 and 2) and their subclasses to display detailed information about each object.
4. Also, override the `equals()` method in the `Dinosaur` and `Employee` classes to compare objects of these classes.

5. Create a class called App with a main method. In there, add functionality to check if an employee is qualified to work in a specific enclosure, considering the employee's role and the enclosure's safety level.
6. The park offers regular tickets and season tickets. Create a SeasonTicket class that extends the Ticket class and add properties such as start date and end date.

## Project

You will be developing a more advanced version of the Mesozoic Eden park manager console application. Your task is to implement the concept of polymorphism to handle different types of dinosaurs and employees. By incorporating polymorphism, the application can accommodate an even wider range of dinosaur species and employee roles. The key features of the system should now include the following:

- The ability to manage diverse types of dinosaur profiles, representing a variety of species
- The ability to manage different types of park employee profiles that represent a variety of roles, such as park rangers, janitors, veterinarians, and more
- All previous features, such as editing and removing profiles, real-time dinosaur tracking, employee scheduling, guest admissions, and handling special events, should now accommodate these new varieties

Here's what you need to do, broken down into smaller steps if you need it:

1. **Extend the data structures:** Extend your Dinosaur and Employee classes into different subclasses to represent different types of dinosaurs and employee roles. Make sure these subclasses demonstrate the principle of polymorphism.
2. **Enhance initialization:** Upgrade your data initialization so that it supports different types of dinosaurs and employees. This could involve creating arrays or lists of Dinosaur and Employee objects, where each object could be an instance of any subclass.
3. **Update the interaction:** Modify your interactive console-based interface to handle the new types of dinosaurs and employees. You might need to add more options or submenus.
4. **Enhance menu creation:** Your menu should now handle different types of dinosaurs and employees. Make sure each option corresponds to a particular function in the program.
5. **Handle actions:** Each menu item should trigger a function that is now able to handle different types of dinosaurs and employees. For example, the "Manage Dinosaurs" option could now trigger a function to add, remove, or edit a profile of any dinosaur species.
6. **Exit the program:** Ensure your program continues to provide an option for the user to exit the program.

The starting code snippet will remain mostly the same as the previous one. However, when implementing `manageDinosaurs()`, `manageEmployees()`, and other similar functions, you'll need to handle different types of dinosaurs and employees:

```
public void handleMenuChoice(int choice) {  
    switch (choice) {  
        case 1:  
            manageDinosaurs(); // This function now needs  
                               to handle different types of dinosaurs  
            break;  
        case 2:  
            manageEmployees(); // This function now needs  
                               to handle different types of employees  
            break;  
        case 3:  
            // manageTickets();  
            break;  
        case 4:  
            // checkParkStatus();  
            break;  
        case 5:  
            // handleSpecialEvents();  
            break;  
        case 6:  
            System.out.println("Exiting...");  
            System.exit(0);  
    }  
}
```

The `manageDinosaurs()`, `manageEmployees()`, `manageTickets()`, `checkParkStatus()`, and `handleSpecialEvents()` methods now need to be updated to be able to handle the increased complexity.

## Summary

In this chapter, we examined one of the cornerstones of OOP, namely inheritance. Inheritance defines an “is-a” relationship between the sub- and parent classes – for example, Fox “is-a” Animal, and Train “is-a” Vehicle. Inheritance promotes code reuse as inheritable base class members are automatically available to subclasses. Class inheritance is enabled via the `extends` keyword and interface inheritance is enabled via the `implements` keyword.

Regarding methods, the subclasses are free to override (replace) the base class implementation. This is how we enable another cornerstone of OOP, namely polymorphism.

Polymorphism is a feature where the instance method from the object is only selected at runtime. Hence, other terms for polymorphism are “late binding,” “runtime binding,” and “dynamic binding.” For polymorphism to work, the signature of the instance method in the subtype must match that of the parent method. The only caveat to that rule is covariant returns, where, in the overriding method, a subtype of the parent return type is allowed. The overriding method, when comparing it with its parent version, must not reduce the access privileges or add extra checked exceptions.

Method overloading, on the other hand, is where the method signatures must be different (apart from the matching method name). Thus, the number of parameters, their types and/or their order, must be different. The return type and parameter names do not matter (as they are not part of the method signature). Method overloading can occur at any level in the hierarchy.

With inheritance, the reference type and object types are often different. As assignment works right to left, when we discuss upcasting and downcasting, we refer to going “up” or “down” the inheritance tree. Upcasting is always safe as the subtype will always have the methods accessible via the supertype reference. Downcasting, however, is not safe and requires a cast for the compiler to be happy. Even at that, if you end up creating a reference that is pointing up the hierarchy tree, you will get a `ClassCastException` error at runtime. Pointing up the hierarchy is not allowed because the subclass reference type could have methods that the parent type object has no code for.

The `super` keyword is used in two situations. The first is to access the parent constructor using `super()`. This call is only allowed as the very first line in any constructor. If not coded explicitly, `super()` will be inserted by the compiler to ensure that the parent constructor executes before the subtype constructor. Construction occurs from the base down because a subtype may rely on parent members and thus, the parent must have a chance to initialize them first. The second scenario is accessing a parent member from subtype code, using `super.parentMember`.

We already know from *Chapter 8*, that the `protected` access modifier ensures members are available within the package and also to any subclasses, regardless of the package. We revisited this and demonstrated that, when accessing a `protected` member from a subclass in a different package, you have to do so, via inheritance, in a very specific way.

An `abstract` method is a method with no code (implementation). Even though a class does not need to have any `abstract` methods to be `abstract` itself; once the class has even one `abstract` method, the class must be `abstract`. Any subclass of an `abstract` class must provide the implementation code for the `abstract` method(s) inherited, or the subclass must also be `abstract`.

Concerning inheritance, a `final` class cannot be inherited from. A `final` method cannot be overridden. Other uses for `final` are for defining constants and ensuring that (the values of) method parameters are constant.

---

The use of sealed classes enables us to restrict parts of a hierarchy to certain types. Rather than the general `extends`, which allows a class to subclass any base class it wants; and without turning off inheritance altogether using `final`; sealed classes achieve a custom restriction using the `sealed`, `non-sealed`, and `permits` keywords.

Instance and `static` initialization blocks are coded outside the methods but inside the class. The `static` block precedes the block with the `static` keyword. The instance uses no keyword (instance semantics are implied). Both enable initialization at various points. Static initialization occurs just once – the first time a class is loaded. Instance initialization occurs every time a constructor is called. Consequently, instance blocks are perfect locations for inserting code that is common across all constructors.

Lastly, we took a deeper dive into upcasting and downcasting. This helped deepen our understanding as to why upcasting is not an issue, why downcasting needs a cast, and why we get `ClassCastException` errors. In addition, using the `instanceof` operator ensures that we prevent `ClassCastException` errors from occurring.

That completes our discussion on inheritance – this was a big chapter! We will now move on to interfaces and `abstract` classes.

Trial Version



Wondershare  
**PDFelement**

# 10

## Interfaces and Abstract Classes

In *Chapter 9*, we learned about another core pillar of OOP, namely inheritance. We saw that Java uses the `extends` keyword to define an “is-a” inheritance relationship between the child and the parent class. The subclass inherits functionality from its parent that enables code reuse, a core benefit of inheritance. Java prevents multiple class inheritance by ensuring you can only extend from one class at a time.

We also took a deep dive into the other remaining pillar of OOP, polymorphism. Polymorphism is enabled by subclasses overriding the parent class instance methods. We saw that, regarding the hierarchy, references can point (across) to objects of their own type and (down) to subclass objects. An exception occurs if a reference attempts to point (up) to parent objects in the hierarchy.

Next, we compared and contrasted method overloading and method overriding. In method overriding, the method signatures must match (except for covariant returns). In method overloading, while the method names are the same, the method signatures must be different.

We also discovered that the order of constructor calls is from the top (base class) down. This is facilitated by the `super()` keyword. To access a parent (non-constructor) member, we can use the `super.` syntax.

We then revisited the `protected` access modifier and demonstrated that, for subclasses outside the package to access the `protected` member, they must do so via inheritance in a very specific manner. In effect, once outside the package, the `protected` member becomes private to subclasses (of the class containing the `protected` member).

We then covered two keywords that have an impact on inheritance: `abstract` and `final`. As an `abstract` method has no implementation code, it is intended to be overridden. The first non-`abstract` (concrete) subclass must provide implementation code for any inherited `abstract` methods. The `final` keyword can be applied in several scenarios. Concerning inheritance, a `final` method cannot be overridden and a `final` class cannot be subclassed.

Next, we discussed `sealed` classes, which enable us to scope parts of the inheritance tree. Using the `sealed` and `permits` keywords, we can state that a class can only be subclassed by certain other named classes. The `non-sealed` keyword ends the scoping task and thus enables us to subclass as normal.

We examined both `instance` and `static` blocks in an inheritance hierarchy. A `static` block is only executed once when a class is first loaded. An `instance` block, on the other hand, is executed every time an object instance is created, making it an ideal place to insert code common to all constructors.

Lastly, we examined upcasting and downcasting. Whereas upcasting is never an issue, downcasting can lead to an exception. Use of the `instanceof` keyword helps prevent this exception.

In this chapter, we will cover `abstract` classes and interfaces. We will compare and contrast them. Interfaces have had several changes over the years. With the aid of examples, we will examine these changes. Java 8 introduced both `static` and `default` methods for interfaces, thereby enabling code to be present in an interface for the first time. In Java 9, to reduce code duplication and improve encapsulation, `private` methods were introduced to interfaces. Finally, Java 17 introduced `sealed` interfaces, which enable us to customize what classes can implement our interface.

This chapter covers the following main topics:

- Understanding `abstract` classes
- Mastering interfaces
- Examining `default` and `static` interface methods
- Explaining `private` interface methods
- Exploring `sealed` interfaces

## Technical requirements

The code for this chapter can be found on GitHub at <https://github.com/PacktPublishing/Learn-Java-with-Projects/tree/main/ch10>.

## Understanding abstract classes

In *Chapter 9*, we covered the `abstract` keyword. Let's review some key points that we discussed. An `abstract` method is exactly that – it is abstract. It has no code. It doesn't even have curly braces – `{ }`. This is typically a design decision. The class containing the `abstract` method wants subclasses to provide the code. This means that the class itself is “incomplete” and therefore any class defining an `abstract` method must itself be `abstract`. Any subclass of the `abstract` class must either override the `abstract` method or declare that it too is `abstract`. The compiler will complain otherwise.

However, the inverse is not the case – an abstract class need not have any abstract methods at all. Again, this is a design decision. Since the class is marked as `abstract`, it is considered “incomplete” (even though it may contain code for all the methods). This prevents objects based on abstract classes from being instantiated. In other words, you cannot new an object based on an abstract class. You can, however, have a reference based on an abstract type.

Please refer to *Figure 9.14* for a code example of abstract methods and classes.

## Mastering interfaces

By default, an interface is an abstract construct. Before Java 8, all the methods in an interface were `abstract`. In general, when you create an interface, you are defining a contract for *what* a class can do without saying anything about *how* the class will do it. A class signs the contract when it implements an interface. A class implementing an interface is agreeing to “obey” the contract defined in the interface. “Obeying” here means that, if a concrete (non-abstract) class is implementing an interface, the compiler will ensure that the class has implementation code for each abstract method in the interface. As the Oracle tutorials state, “*Implementing an interface allows a class to become more formal about the behavior it promises to provide.*”

In contrast to classes, where you can (directly) inherit from only one other class, a class can implement many interfaces. Thus, interfaces enable multiple inheritance. Let’s look at an example:

```
class Dog extends Animal implements Moveable, Loveable {}
```

This line of code states that `Dog` “is-a” `Animal`, `Moveable`, and `Loveable`. Interface names are often adjectives as they often describe a quality of a noun. Thus, interface names often end in “able.” For example, `Iterable` and `Callable` are interface names in the Java API.

In the previous line of code, we are limited to extending from one class but we can implement as many interfaces as we like. This flexibility is very powerful as we can link into hierarchies without forcing artificial class relationships. This is one of the core reasons for interfaces – *to be able to cast to more than one base type.*

As with abstract classes, given that interfaces are also abstract, you cannot new an interface type. In addition, similarly to abstract classes, you can (and often do) have references that are interface types.

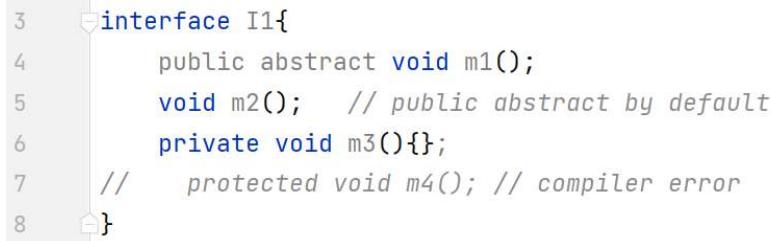
In later sections, we will discuss the `static`, `default`, and `private` methods, all of which have implementation code. Before that, we will deal with the other type of methods we can use in an interface: abstract methods. Additionally, we will discuss interface constants.

## Abstract methods in interfaces

Prior to Java 8, all of the methods in an interface were implicitly `public` and `abstract` by default. Back then, you could state that an interface was a “purely abstract class.”

Concerning the `public` access modifier, this is still the case, even though Java 9 introduced `private` methods. This means that, you can explicitly mark a method in an interface as `public` or `private`. However, if you do *not* specify any access modifier, `public` is the default.

What about their abstract nature? Well, any method that is *not* denoted as `static`, `default`, or `private` is still `abstract` by default. *Figure 10.1* encapsulates this:



```
3  interface I1{
4      public abstract void m1();
5      void m2(); // public abstract by default
6      private void m3(){}
7      // protected void m4(); // compiler error
8  }
```

The screenshot shows a code editor with the following Java interface definition. Line 3 starts with 'interface I1{'. Lines 4 through 8 define methods: 'm1()' is marked as 'public abstract', 'm2()' is marked as 'void' (implying it's abstract), 'm3()' is marked as 'private', and 'm4()' is preceded by a comment indicating it's a compiler error.

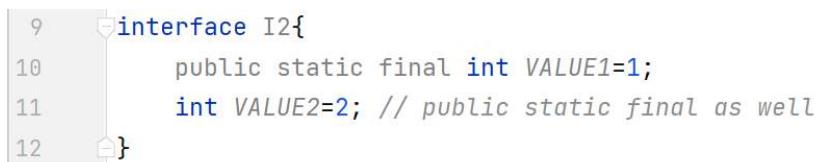
Figure 10.1 – Abstract methods in an interface

In this figure, we can see that the `m2()` method is `public` and `abstract`, even though none of those keywords are explicitly coded. The only other valid access modifier is `private`, as shown when declaring `m3()` on line 6. The fact that `m4()` does not compile (line 7) demonstrates that `protected` is not a valid access modifier on interface methods.

Can we declare variables in an interface? Yes, we can. Let’s discuss them now.

## Interface constants

Any variables specified in an interface are `public`, `static`, and `final` by default. In effect, they are constants, and thus, their initial values cannot be changed. By placing these constants in the interface, any class implementing the interface has access to them (via inheritance), but they are read-only. *Figure 10.2* shows some interface constants:



```
9  interface I2{
10     public static final int VALUE1=1;
11     int VALUE2=2; // public static final as well
12 }
```

The screenshot shows a code editor with the following Java interface definition. Line 9 starts with 'interface I2{'. Lines 10 and 11 define constants: 'VALUE1' is marked as 'public static final' and assigned the value 1, and 'VALUE2' is marked as 'int' and assigned the value 2, with a note indicating it's also a 'public static final' variable. Line 12 ends the interface with a closing brace '}'.

Figure 10.2 – Interface constants

In the preceding figure, we have two variables, namely VALUE1 and VALUE2. Both are constants. VALUE1 states explicitly that it is `public`, `static`, and `final`, whereas VALUE2 does the same implicitly (no keywords are used).

Now, let's look at an example where a class implements an interface.

Figure 10.3 represents a class implementing an interface:

```
3  ① interface Moveable{
4      String HOW="walk"; // constant - public static final
5      ② void move();    // public abstract by default
6  }
7  ➤ public class Dog implements Moveable{
8      // MUST be public - cannot assign weaker privileges
9      // void move(){}
10     @Override
11     ③ public void move(){// MUST be public
12         System.out.println("Dog::move()");
13     }
14    ➤ public static void main(String[] args) {
15        // HOW = "walk"; // cannot change a final variable
16        System.out.println(Moveable.HOW); // walk
17        System.out.println(HOW);          // walk
18        // cannot refer to an instance member from a static context
19        // move();
20        new Dog().move();              // Dog::move()
21    }
22 }
```

Figure 10.3 – A class implementing an interface

In this figure, lines 3-6 represent an interface called `Moveable` that declares a constant, `HOW`, and a method, `move()`. The `Dog` class on line 7 declares that it implements `Moveable`. Therefore, since `Dog` is a concrete, non-abstract class, it must provide an implementation for `move()`.

As we know, interface methods are `public` by default. However, this is not the case for classes. In classes, methods are package-private by default; which means, if you do not provide an access modifier on a method in a class, the method is package-private. Therefore, when overriding an interface method in a class, ensure that the method is `public`. As `package-private` (line 9) is weaker than `public` (line 5), we get a compiler error – hence this line is commented out. Line 11 shows that `move()` must be explicitly declared `public` in `Dog`.

Line 15 shows that `HOW`, declared on line 4, is a constant. If uncommented, line 15 gives a compiler error as constants, once assigned a value, cannot change.

Lines 16 and 17 demonstrate both ways we can access the `HOW` constant – either by prepending it with the interface name (line 16) or directly (line 17).

Line 19 shows that once inside a `static` method, which `main()` is, you cannot directly access an instance method, which `move()` is. This is because instance methods are secretly passed a reference to the (object) instance responsible for calling it, namely the `this` reference. Since `static` methods relate to the class and not a specific instance of the class, there is no `this` reference available in `static` methods. Thus, as per line 20, we need to create an instance and then use that instance to invoke `move()`.

When we run this program, lines 16 and 17 both output the value of the `walk` constant. Line 20 outputs `Dog : move()`, the output from the `Dog` implementation of `move()` (line 12).

#### Note

Since Java 8, code is allowed in `default` methods. As `default` methods are inheritable, the compiler must step in to prevent multiple inheritance in interfaces from causing an issue. We will return to this when we discuss `default` methods in interfaces.

Now, let's look at multiple interface inheritance.

## Multiple interface inheritance

Unlike classes, where multiple inheritance is prohibited in Java, multiple inheritance is allowed in interfaces. Note that the issue with multiple class inheritance is that *if* multiple *class* inheritance was allowed, you could potentially inherit two distinct implementations for the same method.